

Introduction to Computer Vision

2. Basic Image Processing

UCLA – CS 188 – Fall 2019

Fabien Scalzo, Ph.D.

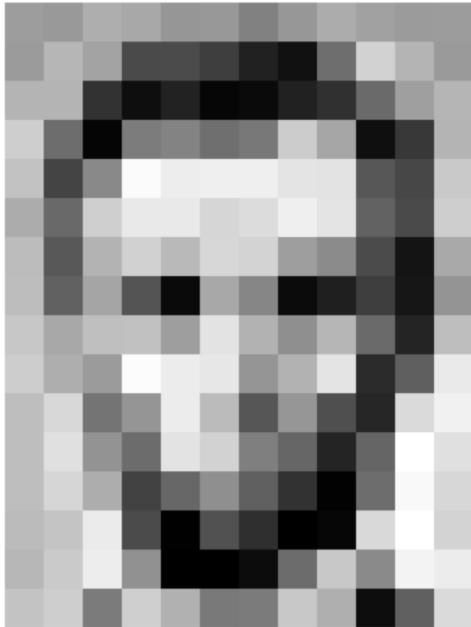
Today: Basic Image Processing

Week 1			26-Sep	Introduction
Week 2	1-Oct	Basic Image Processing	3-Oct	Feature Extraction and Classification
Week 3	8-Oct	Feature Tracking/Optical Flow	10-Oct	SVD, 2D camera model, projective plane
Week 4	15-Oct	2D Image transformations, RANSAC	17-Oct	Euclidean geometry, rigid body motion
Week 5	22-Oct	Epipolar Geometry	24-Oct	3D Cameras and processing
Week 6	29-Oct	Midterm	31-Oct	Statistical decision theory/Pattern Recognition
Week 7	5-Nov	Deep Learning	7-Nov	Deep Learning for Image Classification
Week 8	12-Nov	Object Detection	14-Nov	Generative models
Week 9	19-Nov	Medical Imaging	21-Nov	Autonomous Navigation
Week 10	26-Nov	Guest Lecture (Nikhil Naik)	28-Nov	
Week 11	3-Dec	Recursive 3D reconstruction and pose estimation	5-Dec	Recap
Week 12	10-Dec		12-Dec	Final



Progressive Growing of GANs

Image representation (2D)

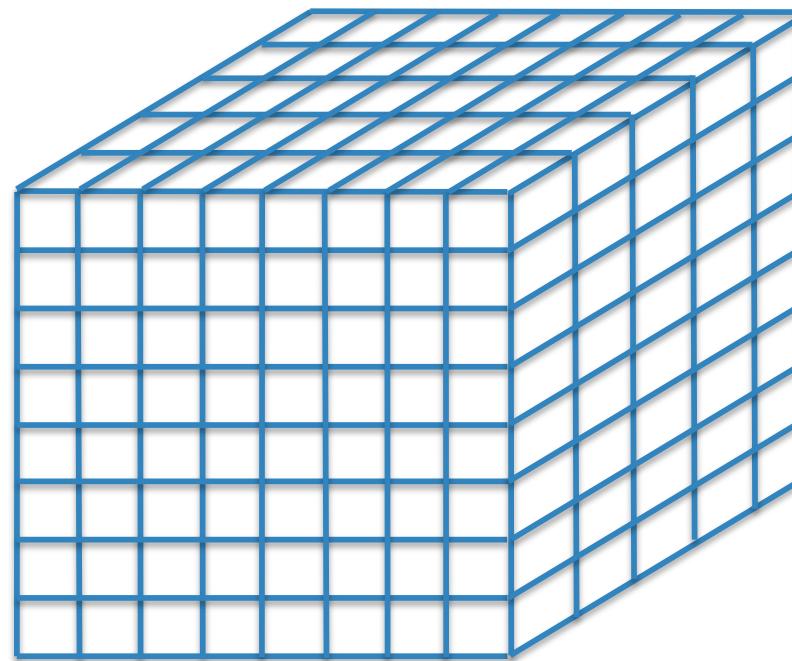


157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	83	17	110	210	180	154
180	180	50	14	84	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	83	17	110	210	180	154
180	180	50	14	84	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	209	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

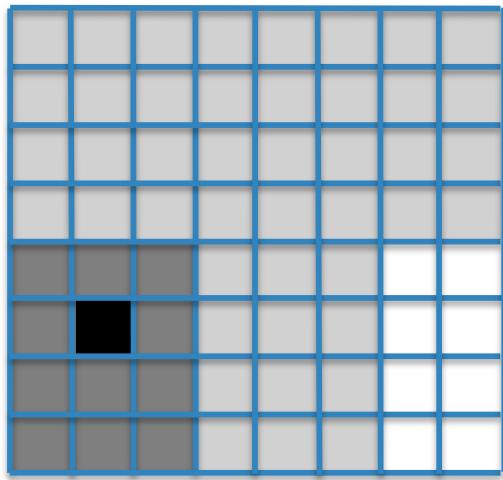
Array of pixels

Image representation (3D)



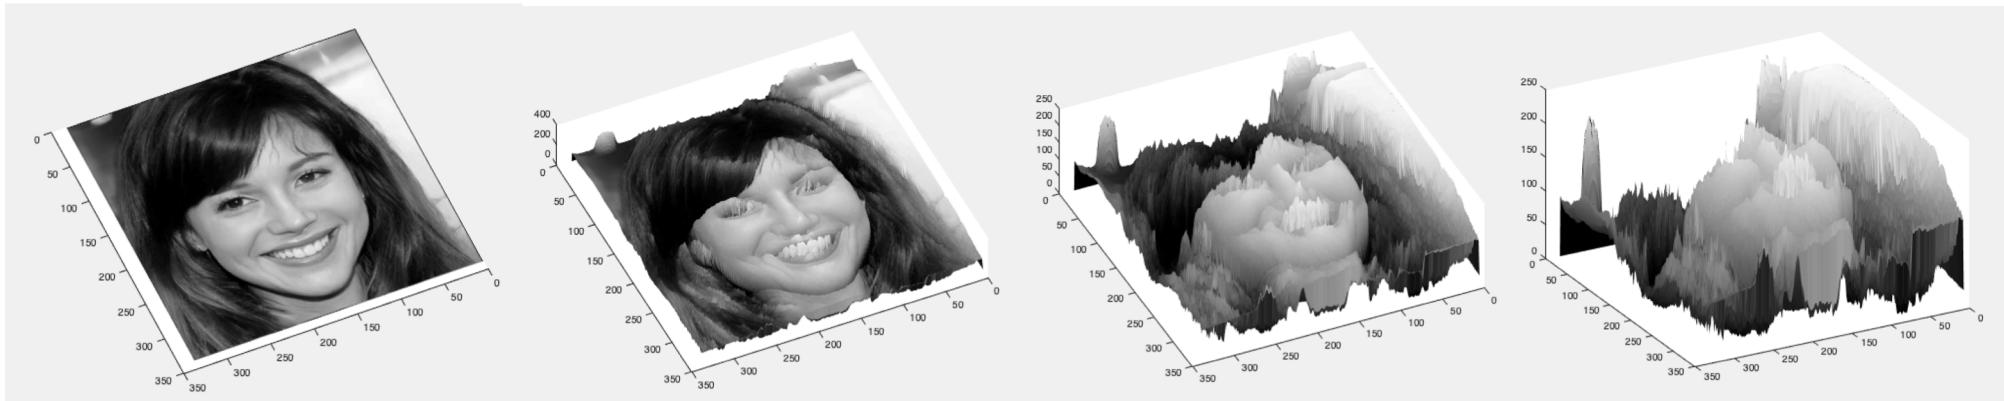
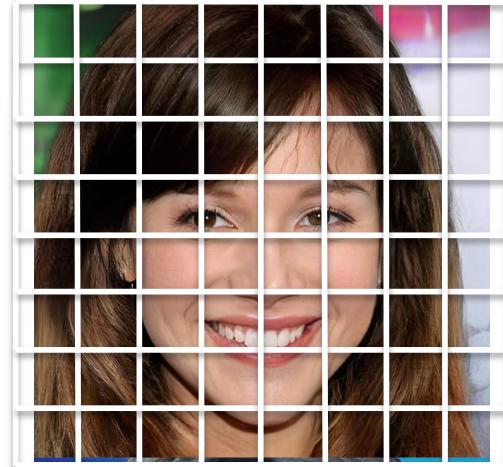
Array of voxels

Image representation



Each pixel is described by a position (x,y) in the image and an intensity value $I(x,y)$

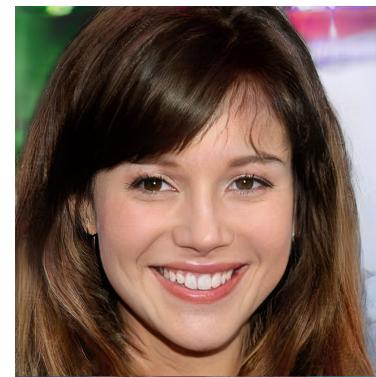
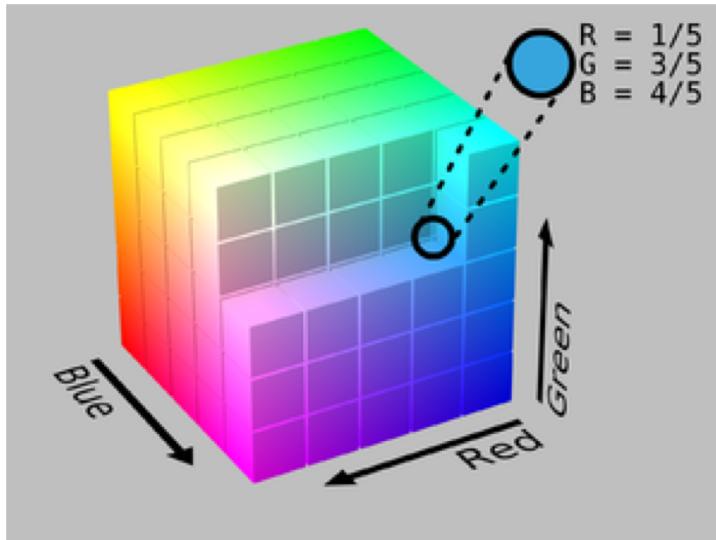
Image representation



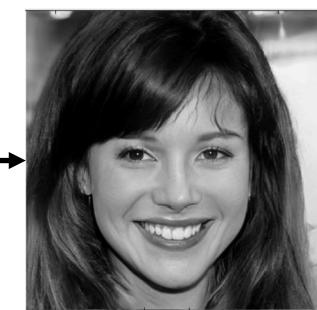


RGB Image representation

Each pixel is described by a position (x,y) in the image and its RGB intensity values $R(x,y)$, $G(x,y)$, $B(x,y)$



red



green

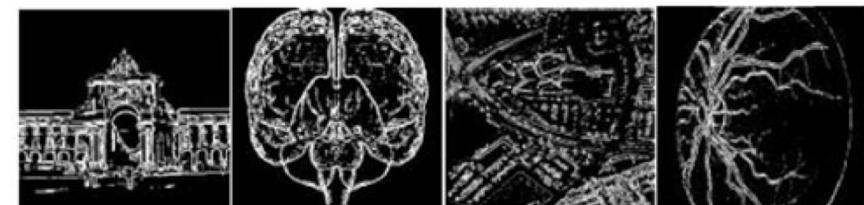
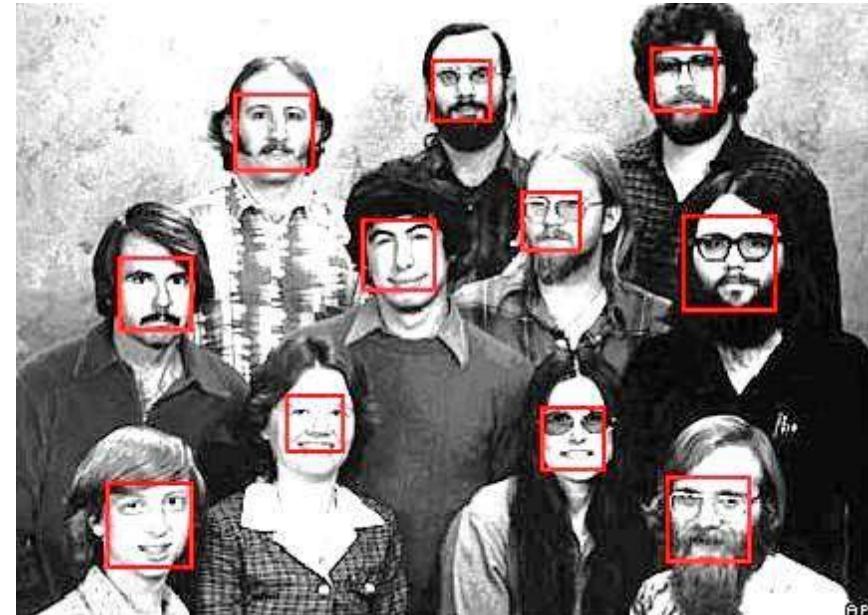


blue

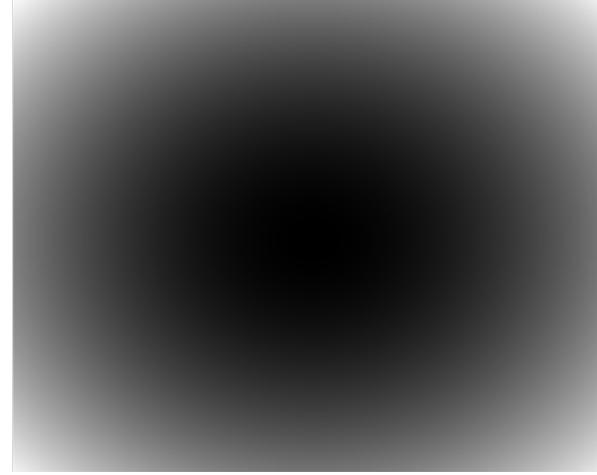
Image filtering

- Purpose: image quality improvement, eliminate unwanted structures, emphasize structures of interest.
- In most of computer vision and imaging applications.

Feature correspondence, image registration, retrieval, categorization, face detection/recognition, ...



We don't “see” most information in image

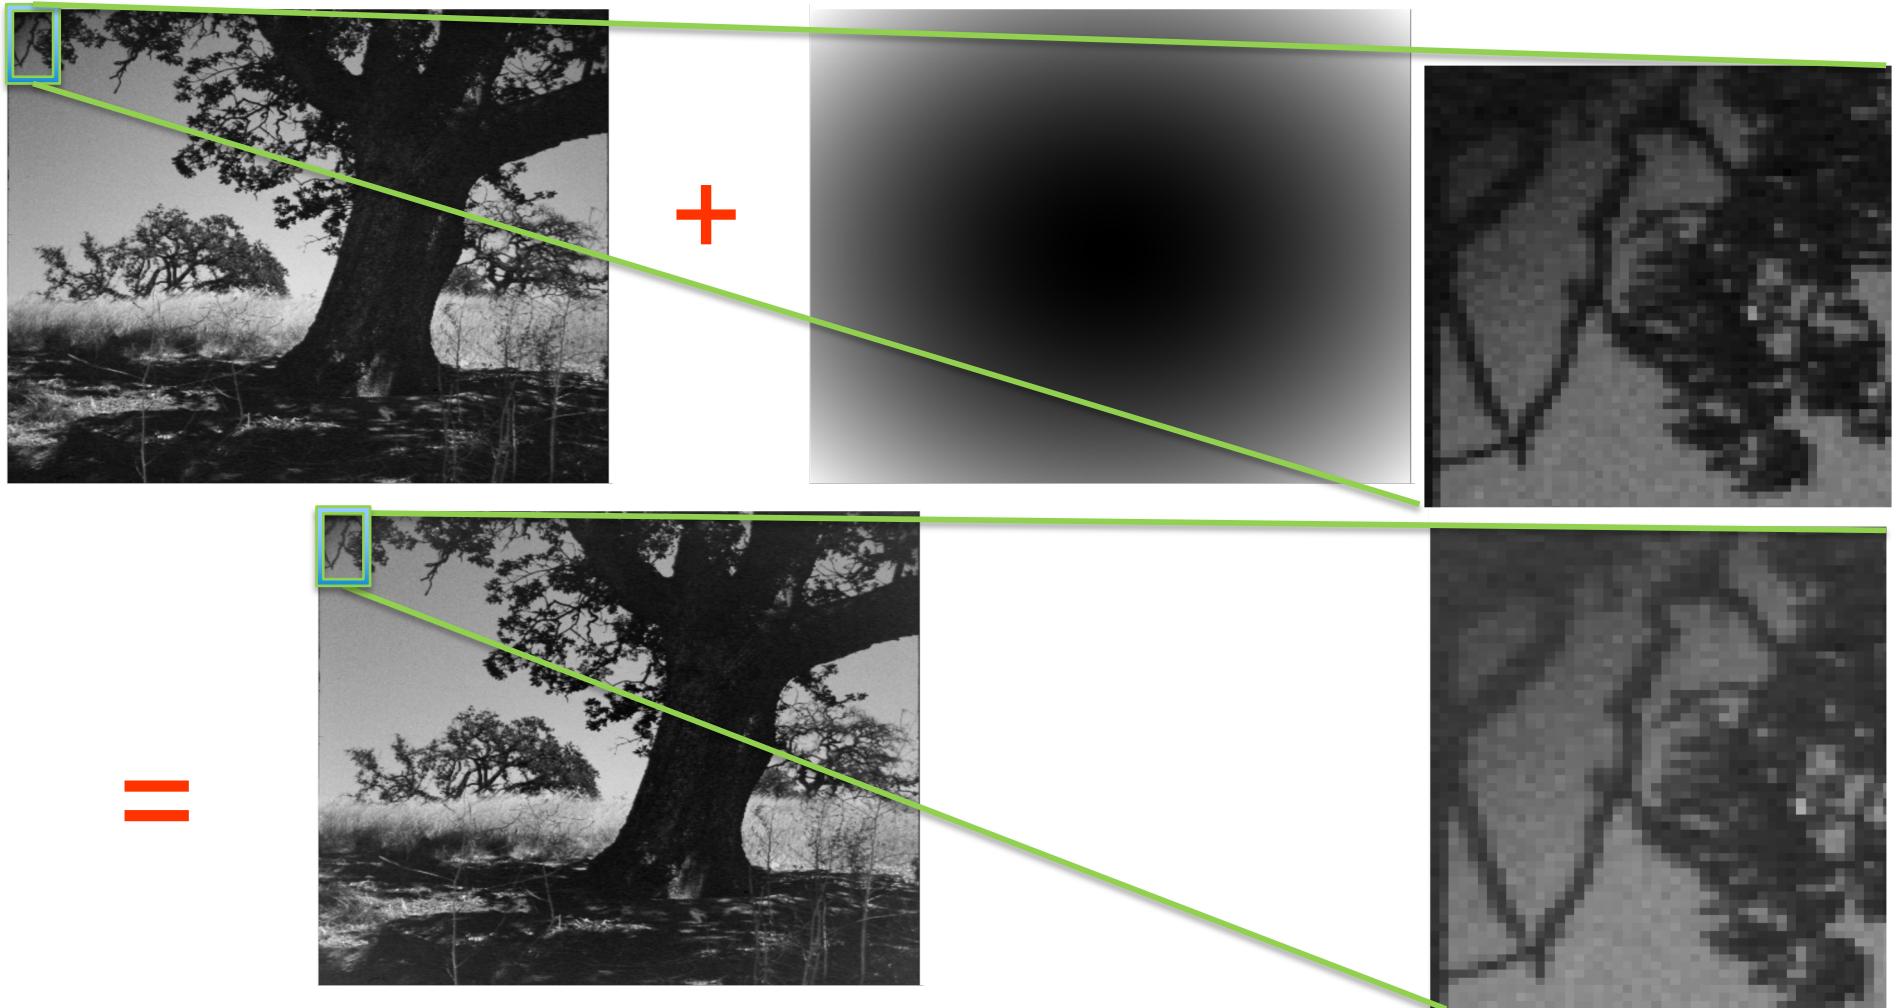


—



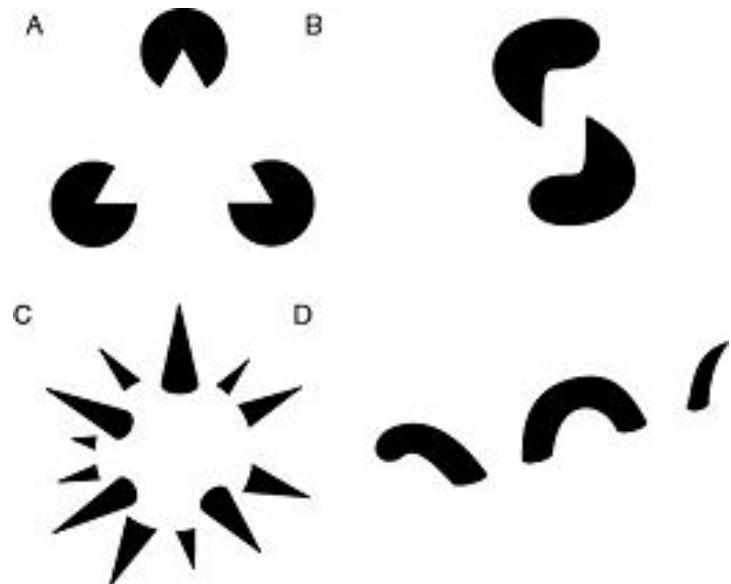
Can add invisible change

We don't “see” most information in image



Images are complex to process

Human vision uses various tricks, such as **Proximity**, **Similarity**, **Continuity** to make processing more efficient.



The brain “filters” what matters,
and what matters is purpose dependent

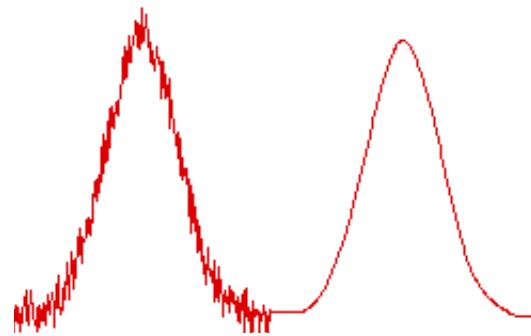
see **Gestalt psychology**

Filtering: the basics

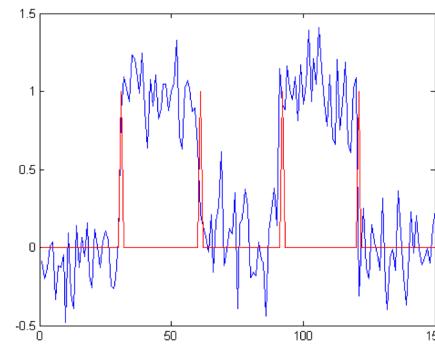
Performs mathematical operations to reduce or enhance certain aspects of a signal (1D, 2D, ...)

A **low-pass** filter passes signals with a frequency lower than a certain cutoff frequency and attenuates signals with frequencies higher than the cutoff frequency.

A **high-pass** does the opposite and its output is proportional to rate of change of the input signal.

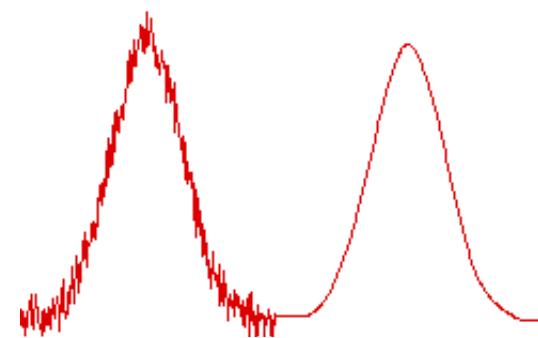


noise removal (low-pass)

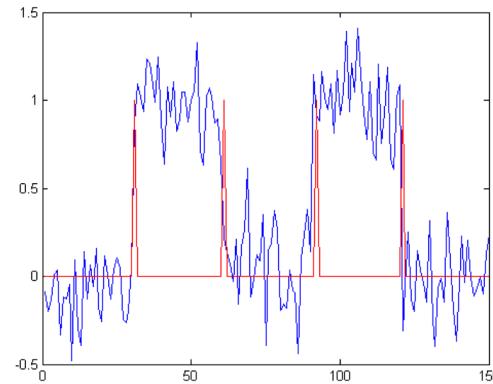


change detection (high-pass)

Filtering: the basics



noise removal (low-pass)



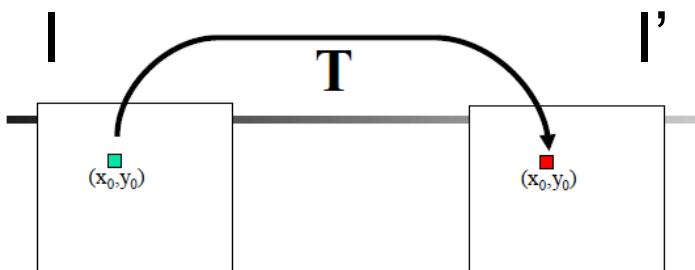
change detection (high-pass)

Filters can be applied either:

- In the frequency domain: apply Fourier transform and its inverse after application of the filter.
- **In the spatial domain:** convolve the image with the filter.

Image Filtering

Modify the pixels in an image based on some function T



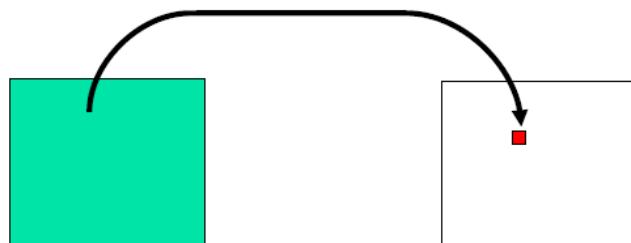
Point-wise

$$I'(x_0, y_0) = T[I(x_0, y_0)]$$



Local

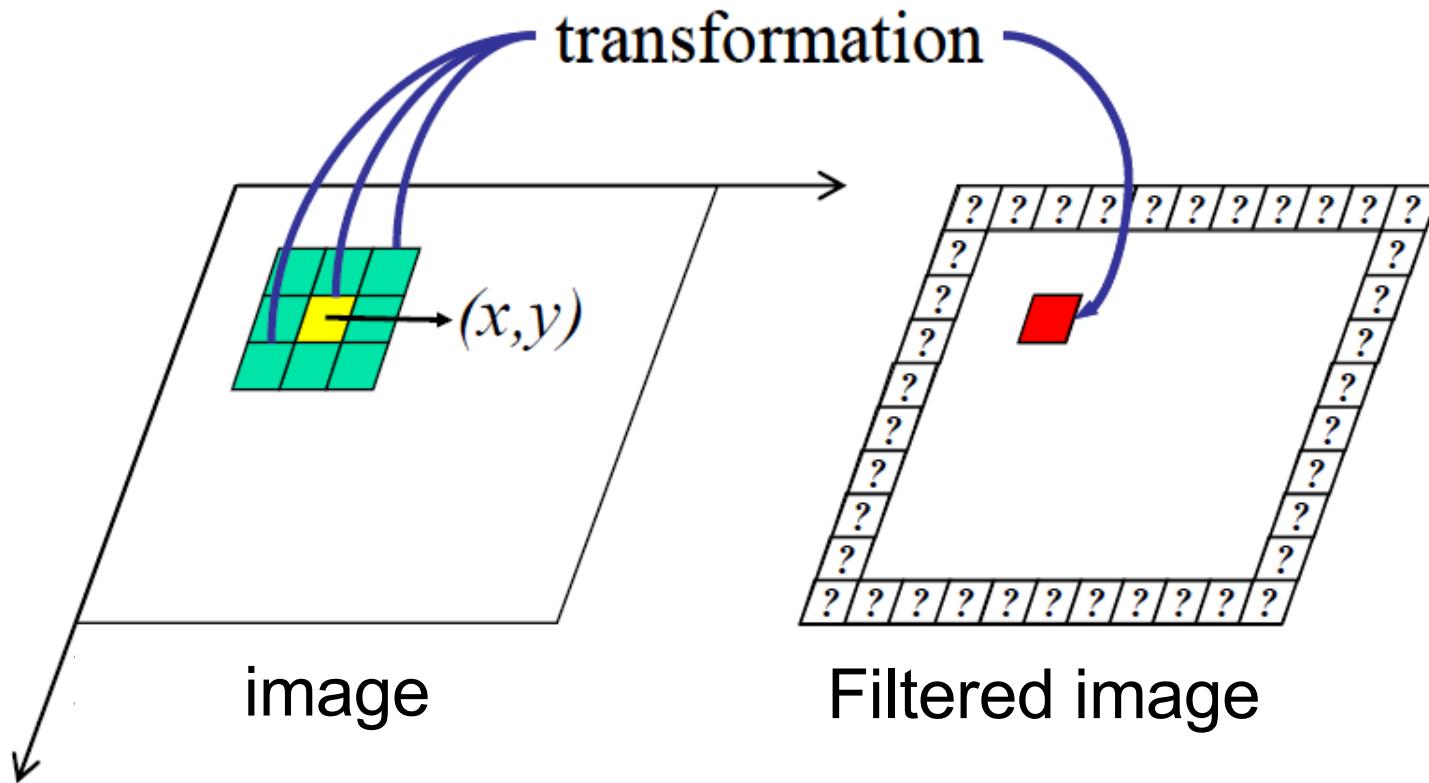
$$I'(x_0, y_0) = T[I(V)]$$



Global

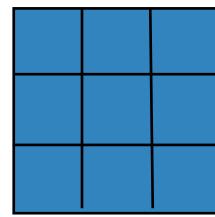
$$I'(x_0, y_0) = T[I(x, y)]$$

2D Local filter

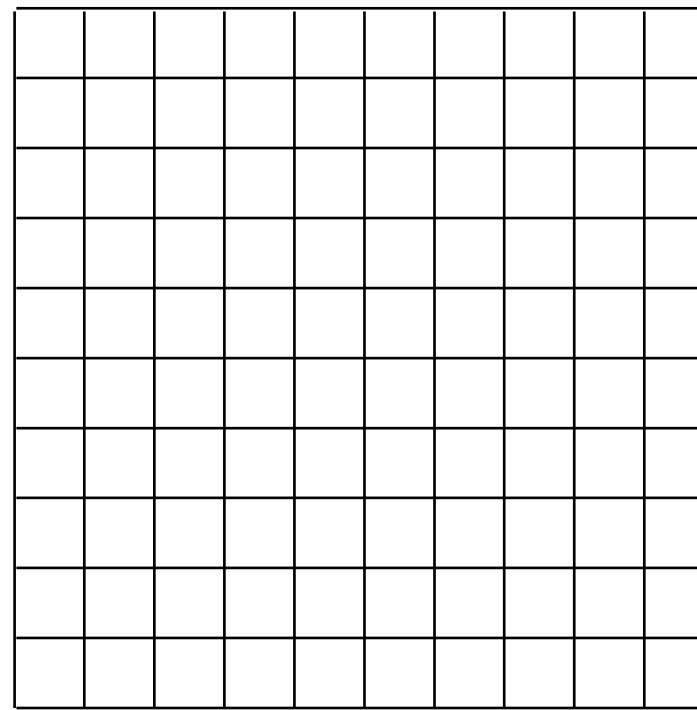


Examples: minimum, maximum, average

2D Local Filtering



Filter



Image

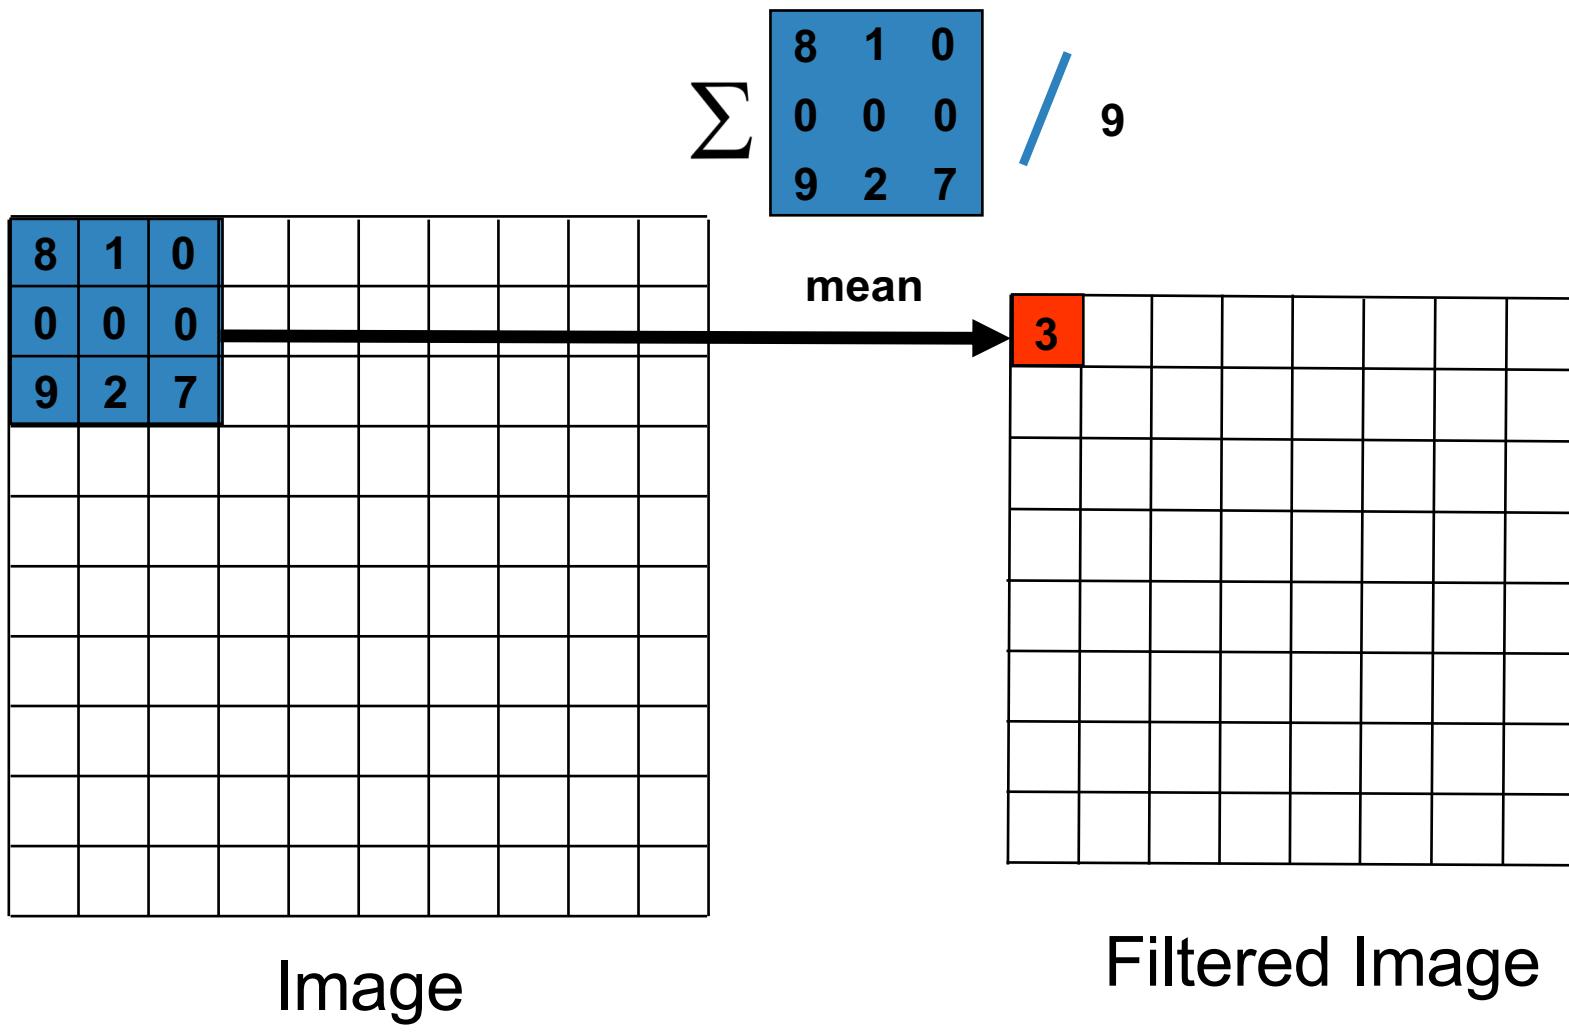
8	1	0						
0	0	0						
9	2	7						

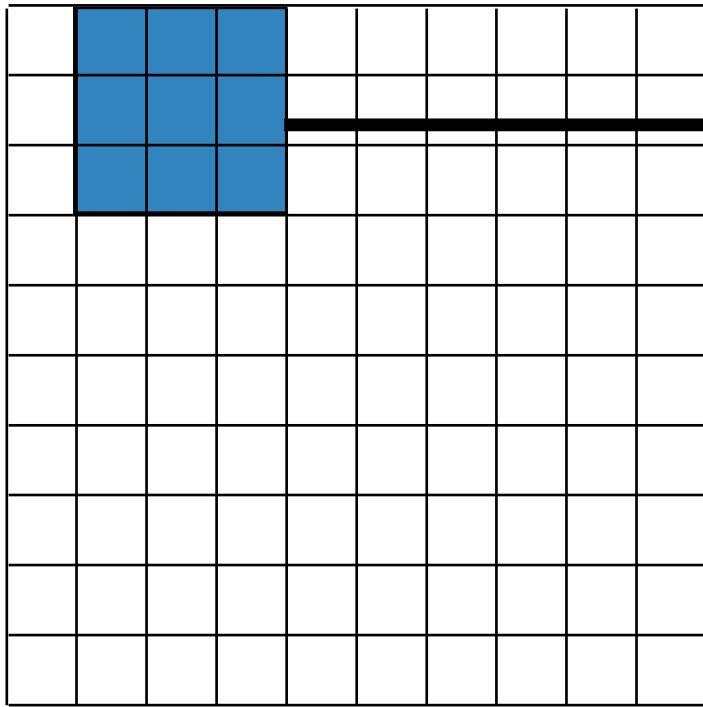
max

9								

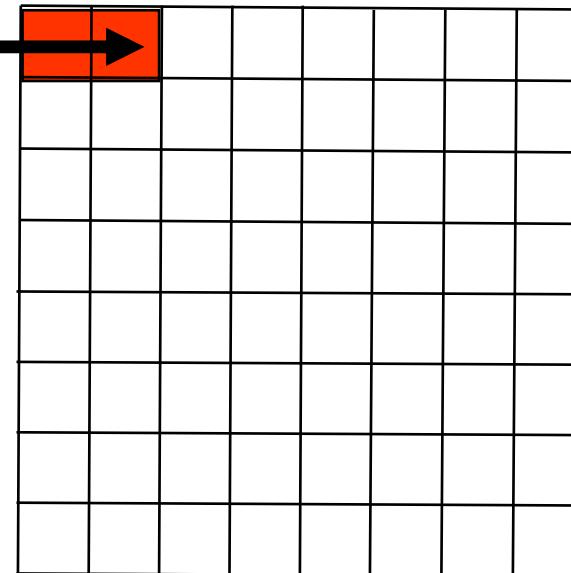
Image

Filtered Image

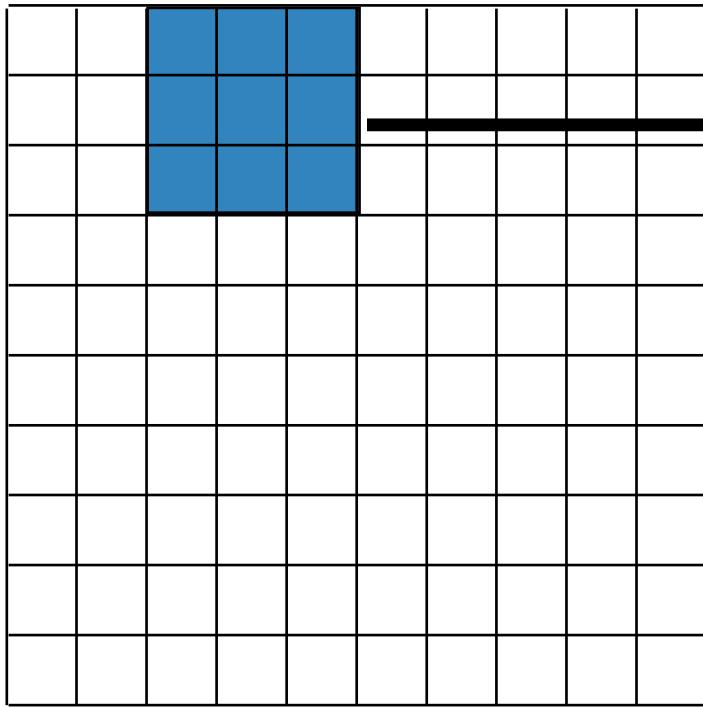




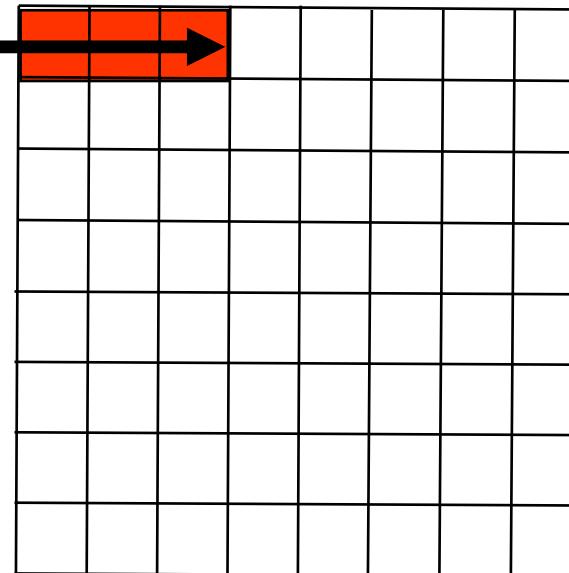
Image



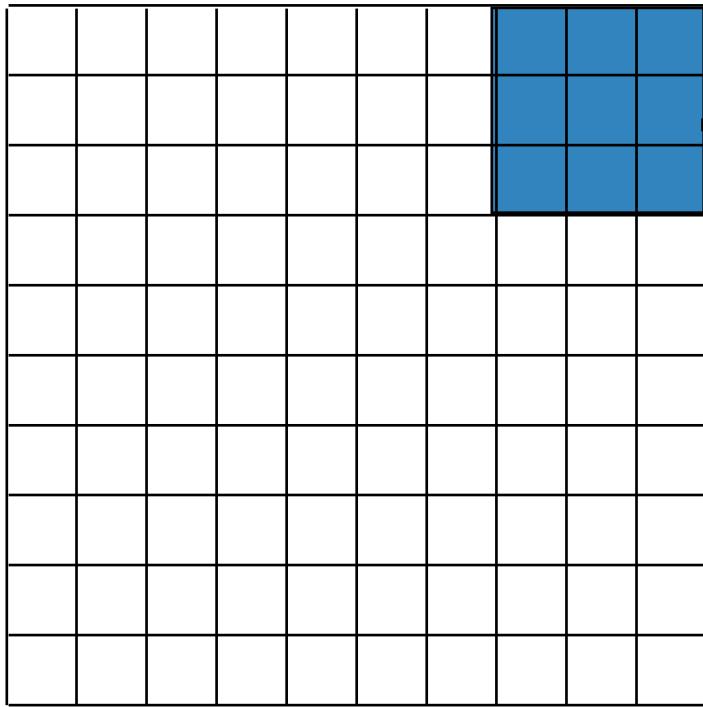
Filtered Image



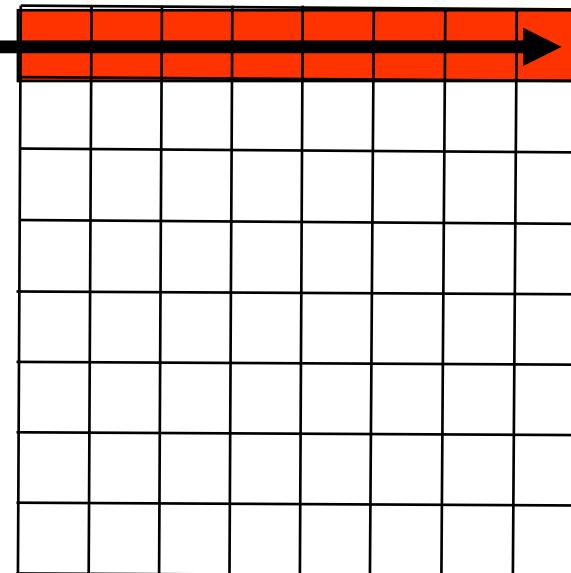
Image



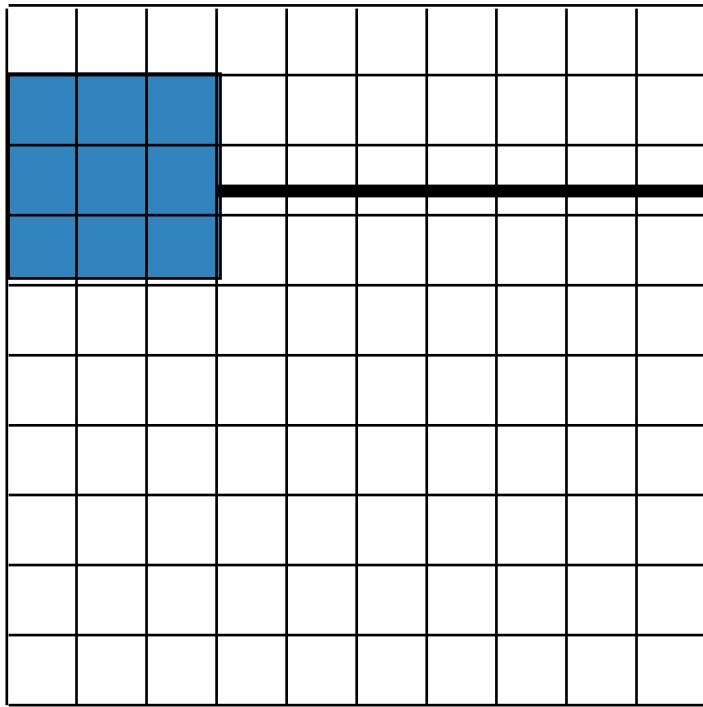
Filtered Image



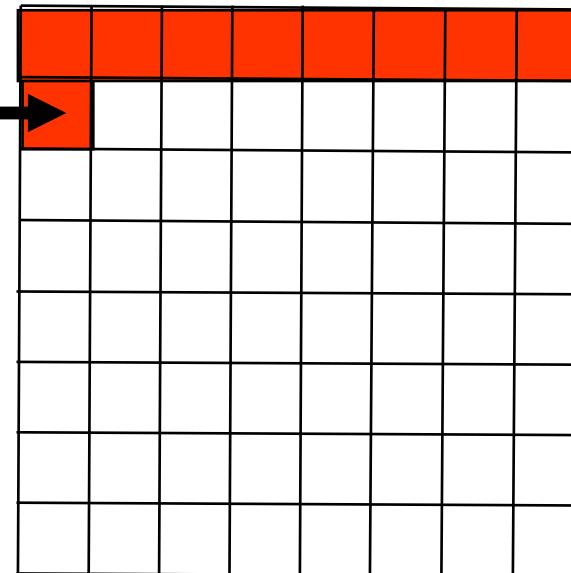
Image



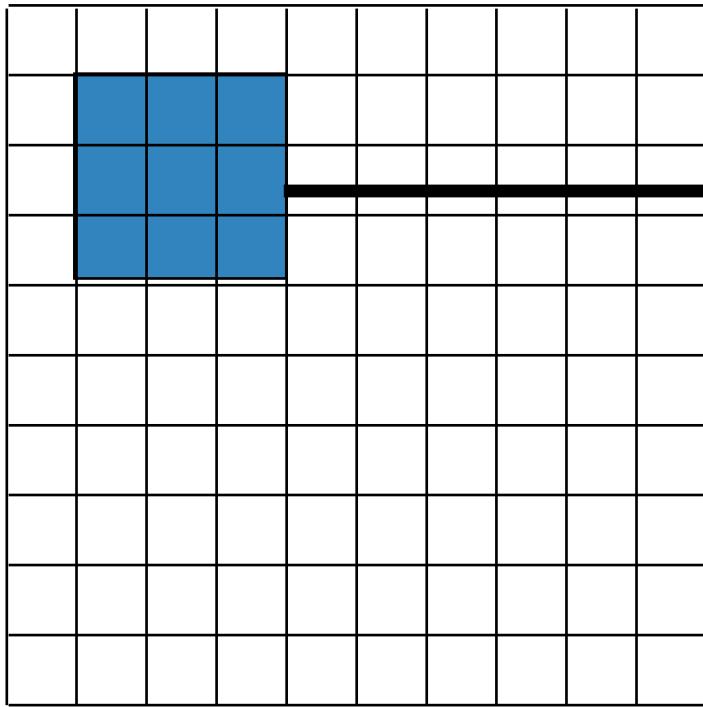
Filtered Image



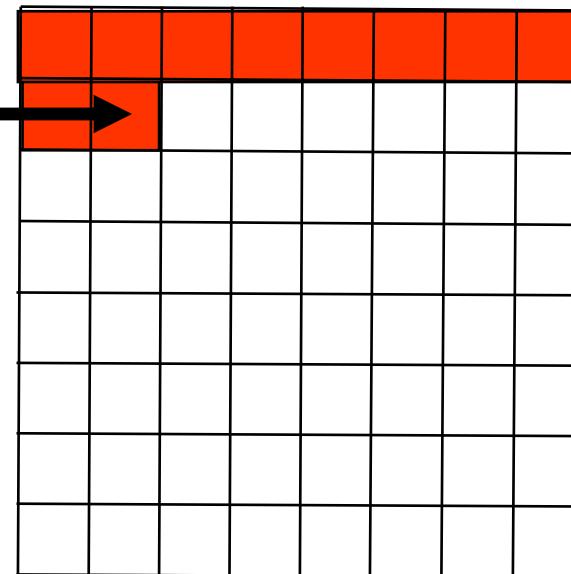
Image



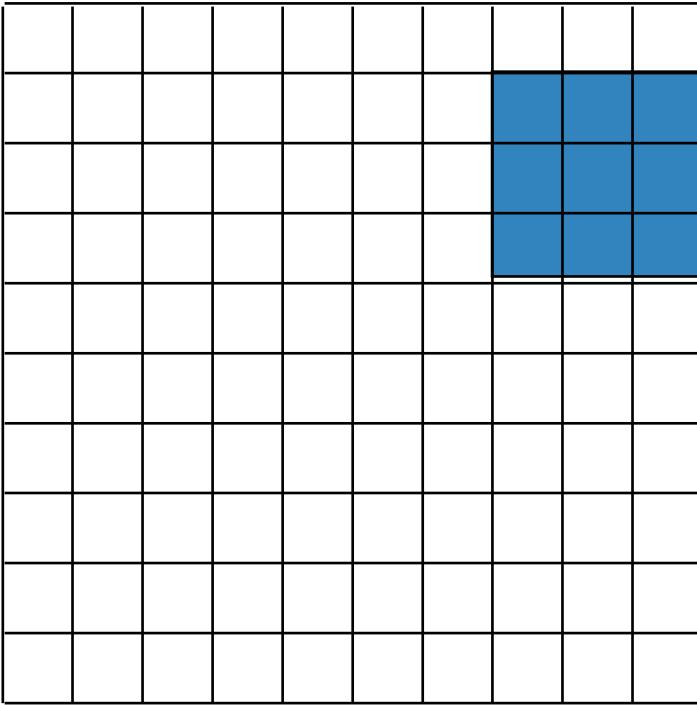
Filtered Image



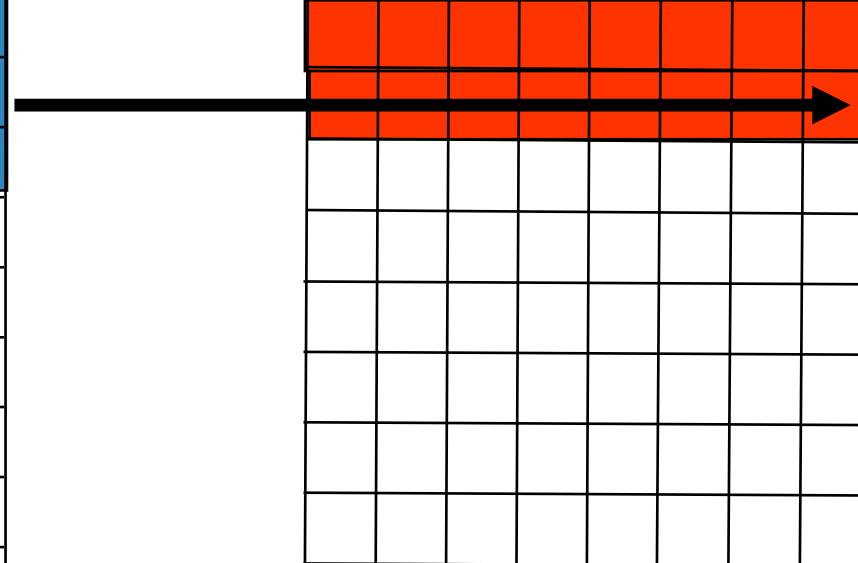
Image



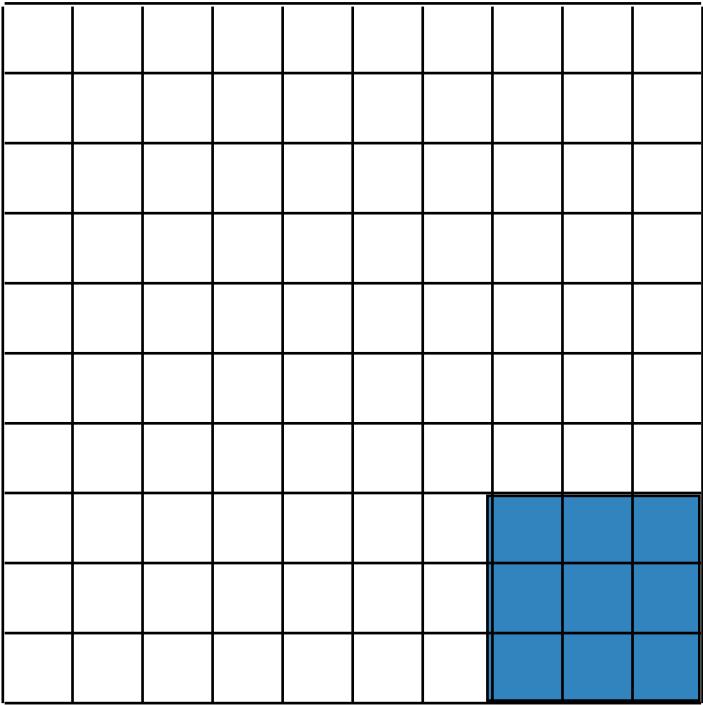
Filtered Image



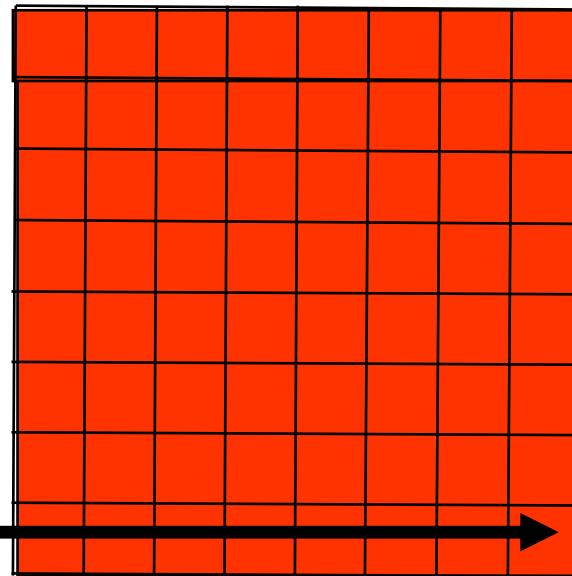
Image



Filtered Image

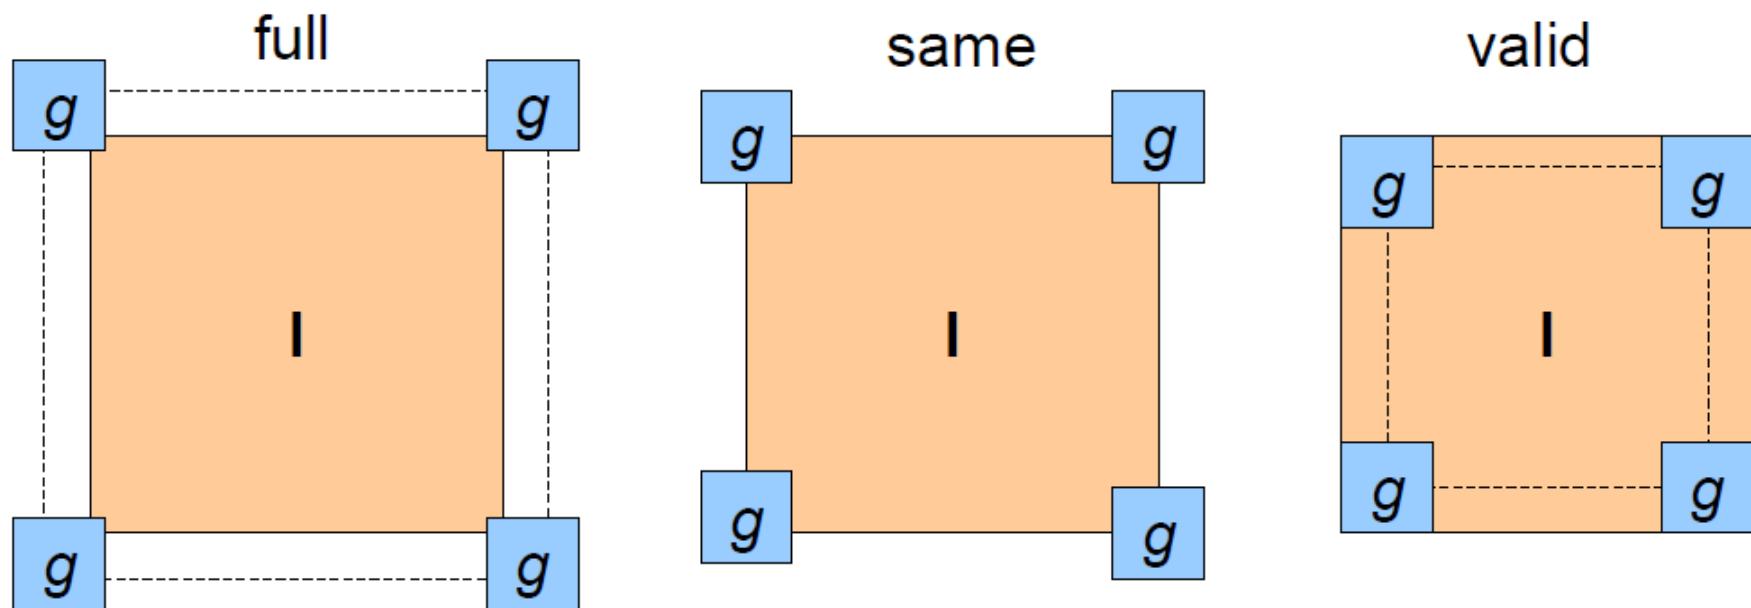


Image



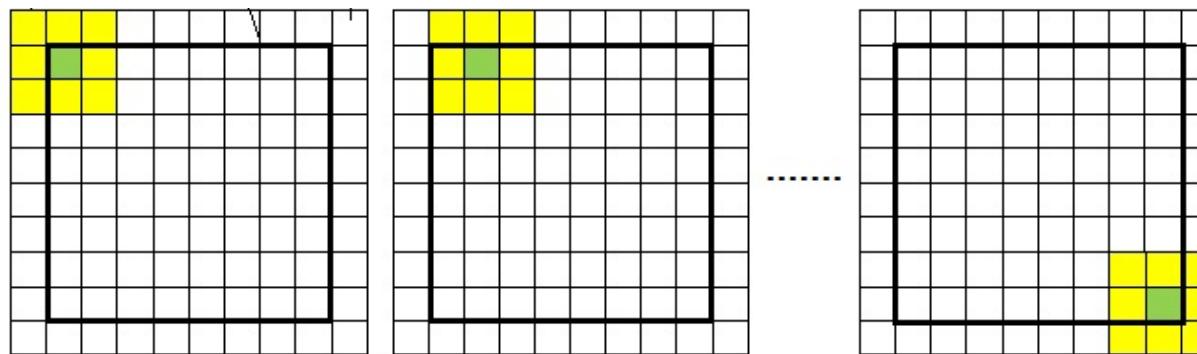
Filtered Image

Output size



Dealing with Borders

- ❑ the filter window falls off the edge of the image
- ❑ need to extrapolate
- ❑ methods
 - clip filter (black): Pad with zeros
 - wrap around: Pad with circular repetition of elements within the dimension,
Assume periodicity of the image
 - ■ copy edge: Pad by repeating border elements of image
 - reflect across edge: Pad with mirror reflection of image along dimension



Examples



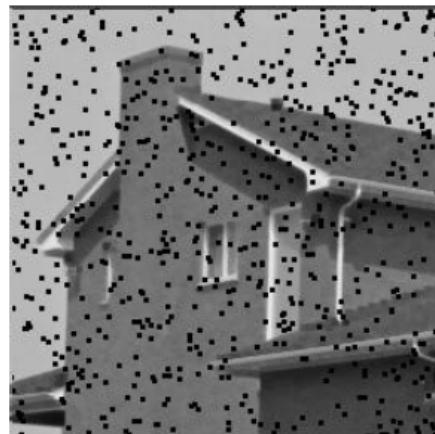
Original image



Noise



Average 3 x 3



Min 3 x 3



Max 3 x 3



Median 3 x 3

Problems



Median, 5pt



Min, 2pt

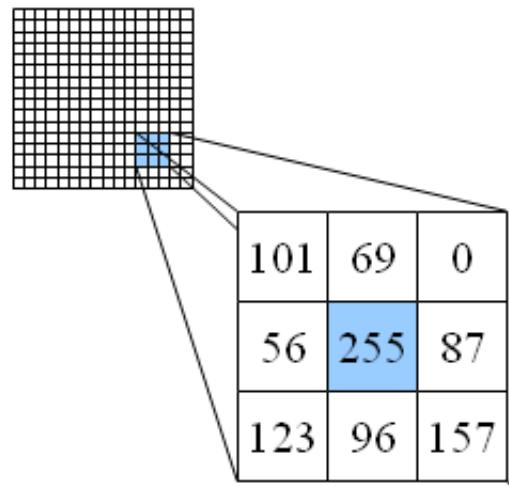


Mean, 5pt

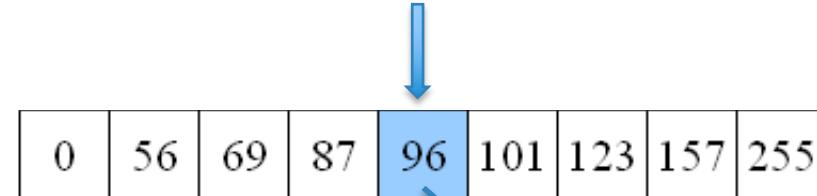


Max, 2pt

Median Filter



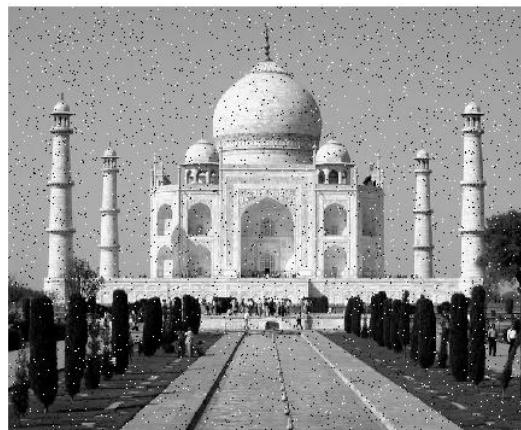
sort →



new pixel value

A 3x3 output window with the new pixel value 96 highlighted in blue:

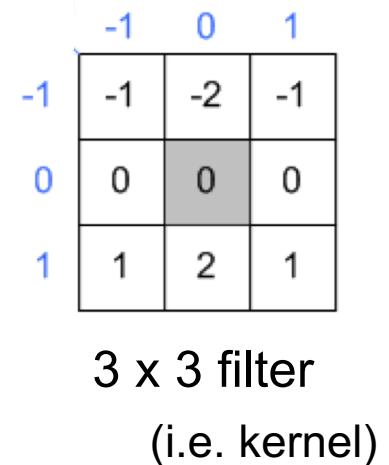
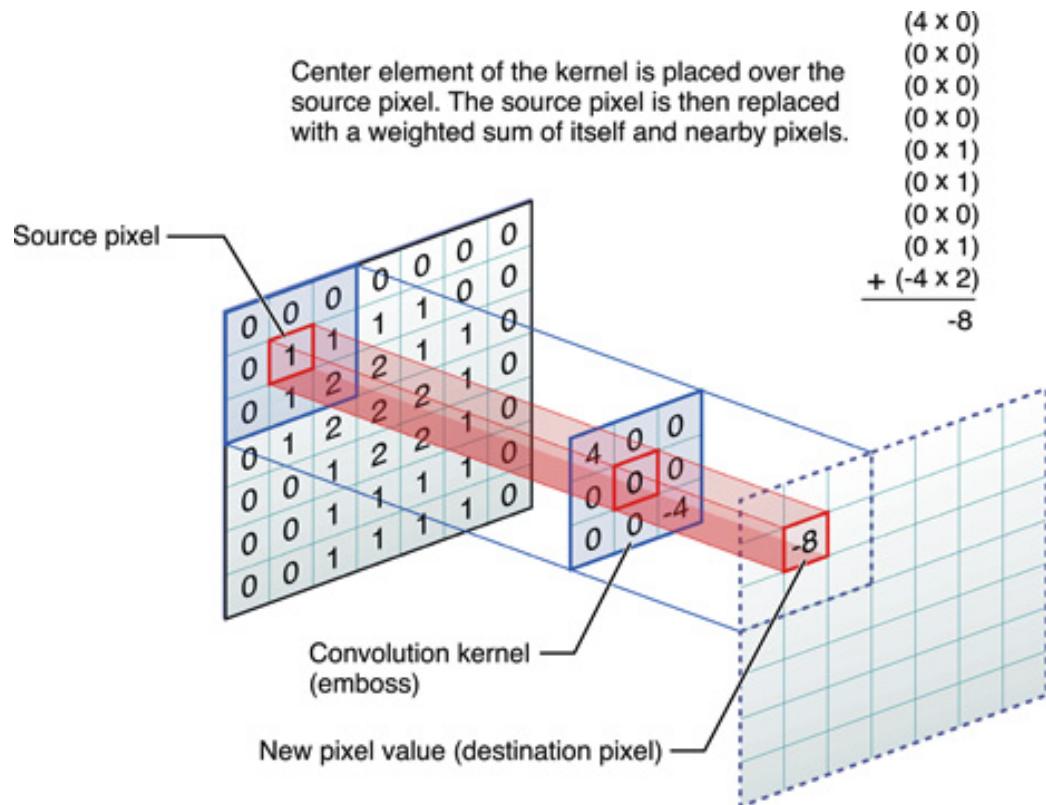
101	69	0
56	96	87
123	96	157



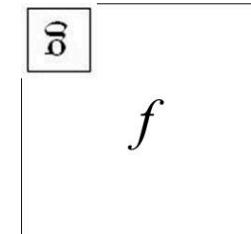
Convolution

$$I'(x, y) = \sum_{i=-1}^1 \sum_{j=-1}^1 I(x+i, y+j) \cdot \text{filter}^{\text{flipped}}(i, j)$$

Use *odd sized filters*, symmetric index range $[-N, N]$ so filter center at $(0,0)$



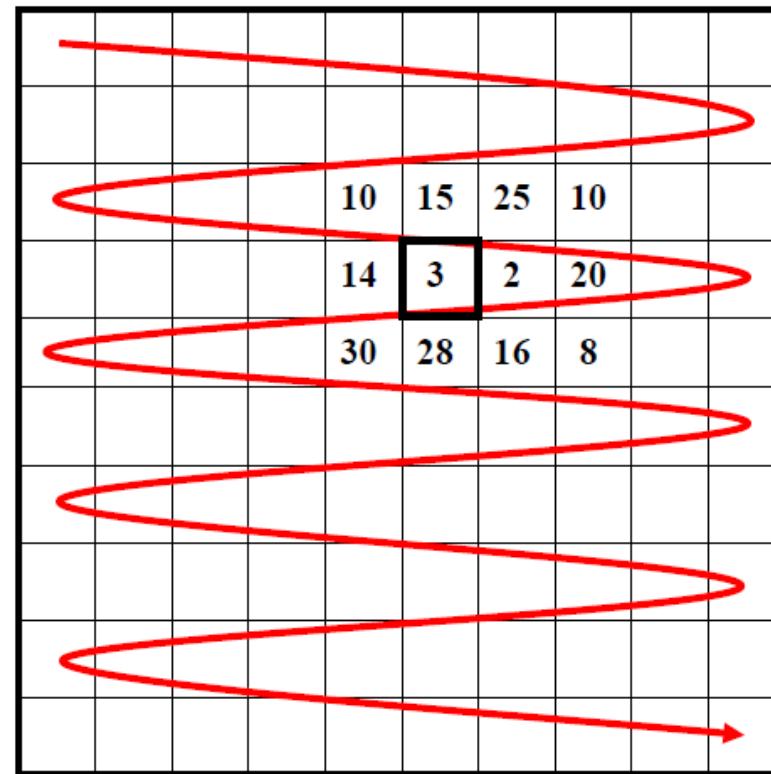
Convention:
kernel is “flipped”



Average

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Average 3 x 3 kernel

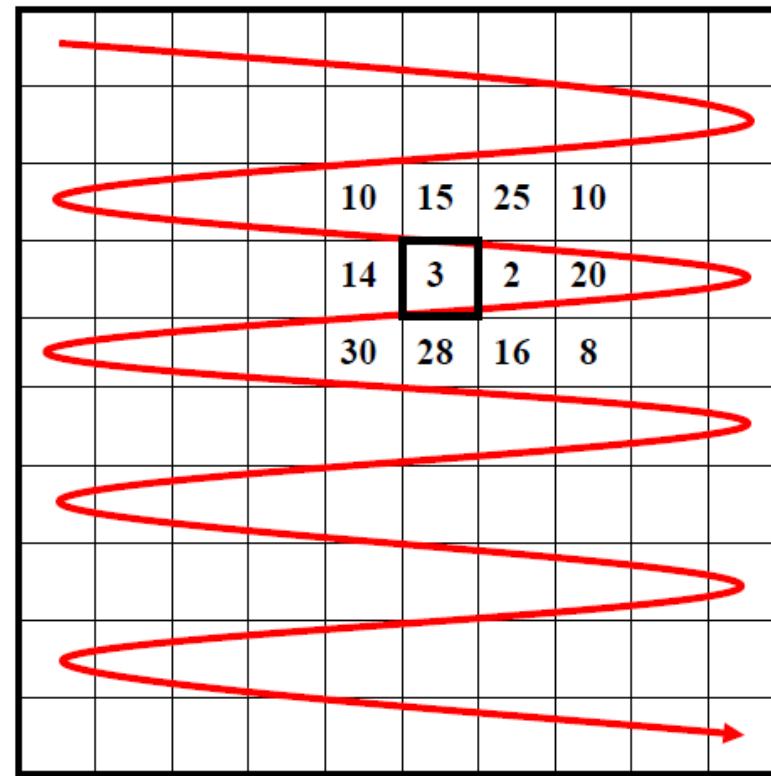


image

Average

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Average 3 x 3 kernel



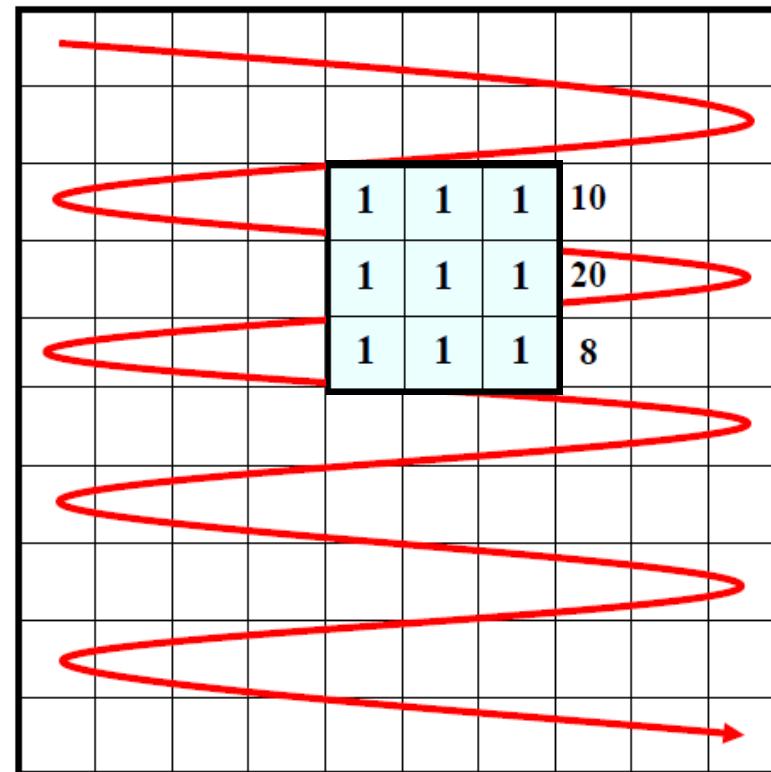
image

Average

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Average 3 x 3 kernel

$$(10*1+15*1+25*1+14*1+3*1 +2*1+30*1+28*1+16*1) / 9 = 143/9 \sim 16$$



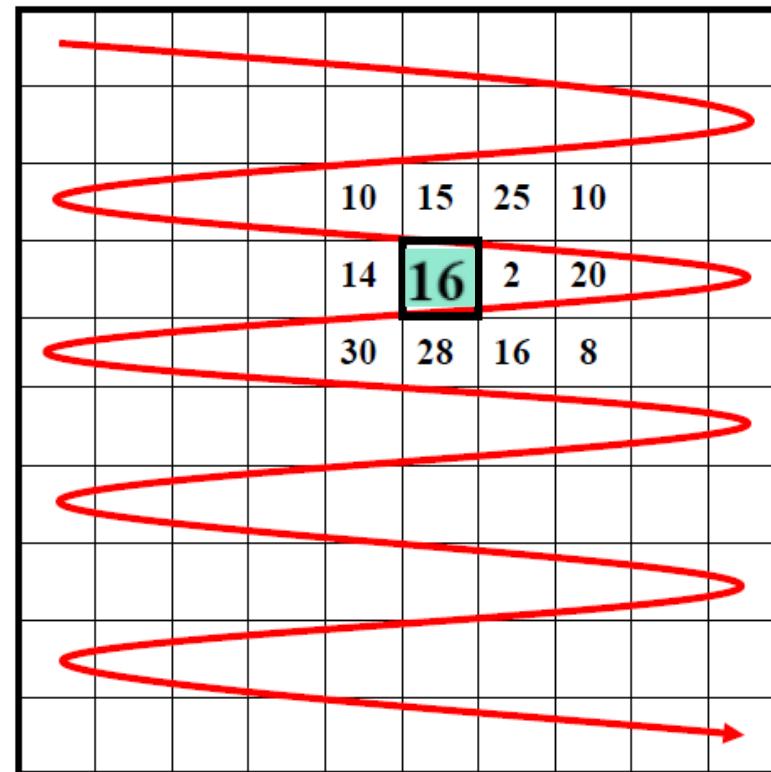
image

Average

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Average 3 x 3 kernel

$$(10*1+15*1+25*1+14*1+3*1 +2*1+30*1+28*1+16*1) / 9 = 143/9 \sim 16$$



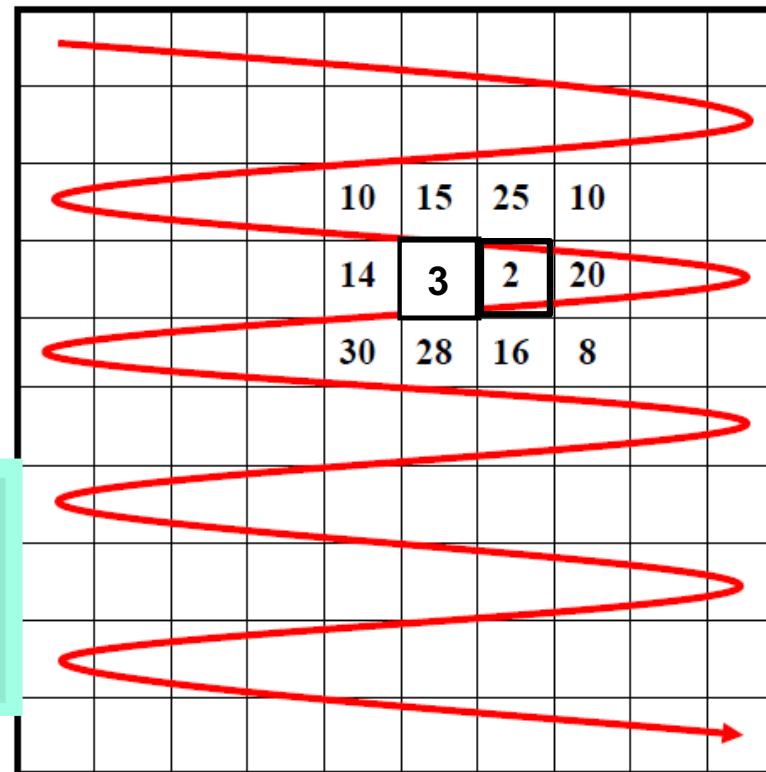
image

Average

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Average 3 x 3 kernel

$$(10*1+15*1+25*1+14*1+3*1 +2*1+30*1+28*1+16*1) / 9 = 143/9 \sim 16$$



image

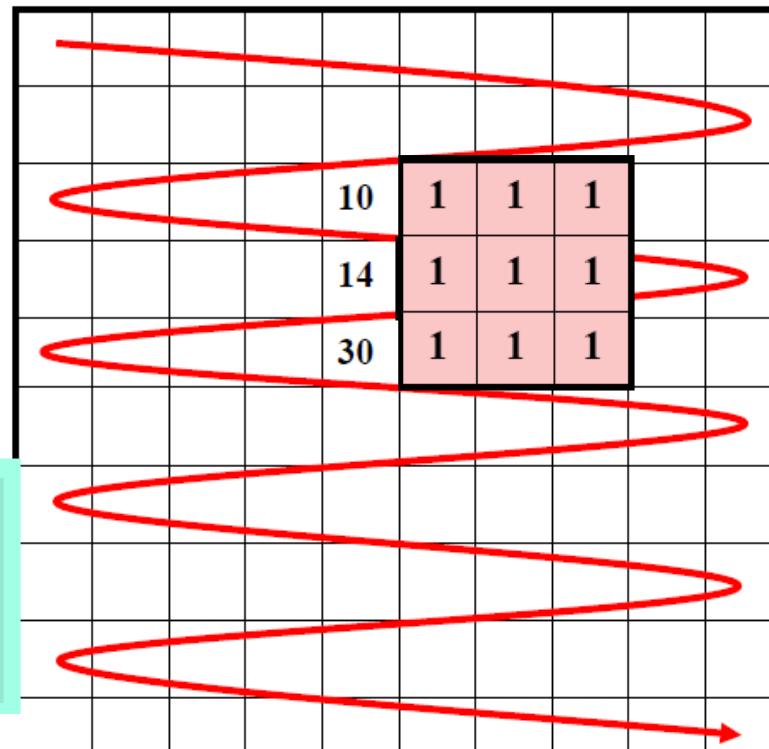
Average

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

Average 3 x 3 kernel

$$(10*1+15*1+25*1+14*1+3*1 +2*1+30*1+28*1+16*1) / 9 = 143/9 \sim 16$$

$$(15*1+25*1+10*1+3*1+2*1+ 20*1+28*1+16*1+8*1) / 9 = 127/9 \sim 14$$



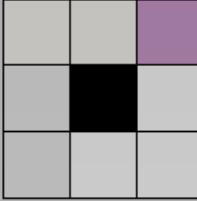
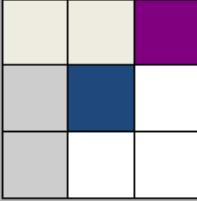
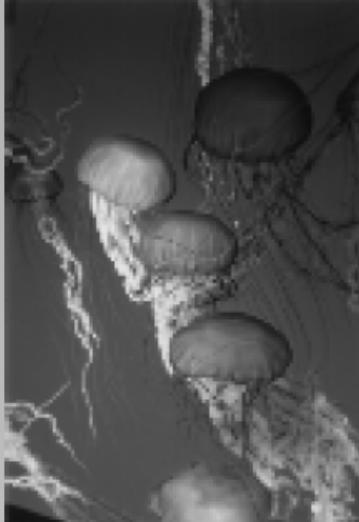
image

Main properties of convolution

- **Commutative:** $a * b = b * a$
 - no difference between filter and signal
- **Associative:** $a * (b * c) = (a * b) * c$
 - applying several filters one after another:
$$(((a * b1) * b2) * b3)$$
is equivalent to applying one filter: $a * (b1 * b2 * b3)$
- **Distributive:** $a * (b + c) = (a * b) + (a * c)$

Enhancing vs Smoothing

Sharpen / Blur

$\begin{matrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{matrix}$ 	 $\begin{matrix} 1 & 1 & 1 \\ 1 & 2 & 1 \\ 1 & 1 & 1 \end{matrix}$	
		

← sharpen blurring →

Low-pass filters

$$\frac{1}{9} \times \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\frac{1}{81} \times \begin{bmatrix} 1 & 2 & 3 & 2 & 1 \\ 2 & 4 & 6 & 4 & 2 \\ 3 & 6 & 9 & 6 & 3 \\ 2 & 4 & 6 & 4 & 2 \\ 1 & 2 & 3 & 2 & 1 \end{bmatrix}$$

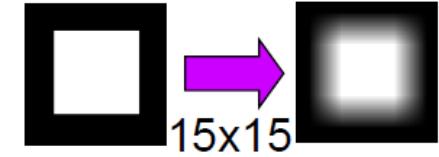
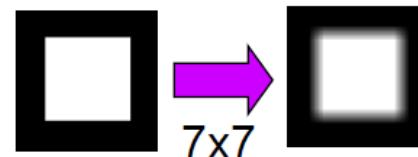
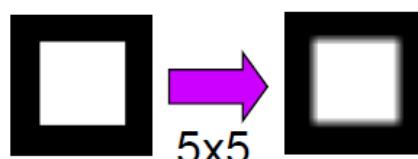
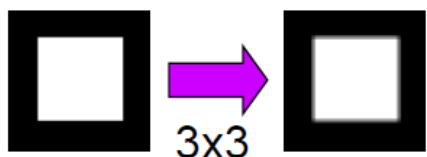
$$\frac{1}{25} \times \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 2 & 2 & 2 & 0 \\ 1 & 2 & 5 & 2 & 1 \\ 0 & 2 & 2 & 2 & 0 \\ 0 & 0 & 1 & 0 & 0 \end{bmatrix}$$

$$\frac{1}{273} \times \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{bmatrix}$$


pyramidal

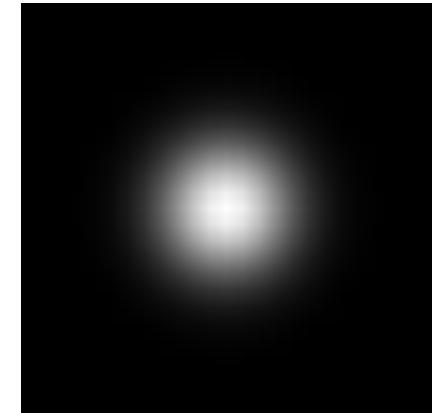
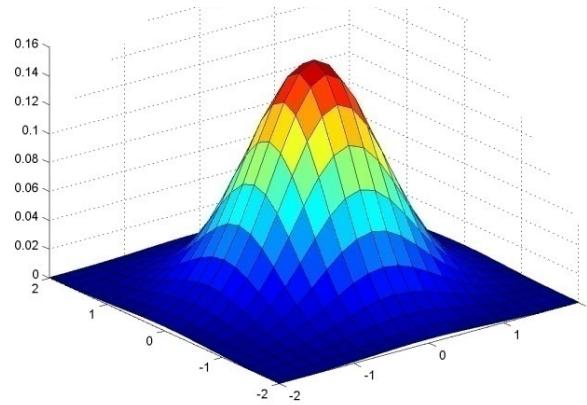
conical

Gaussian

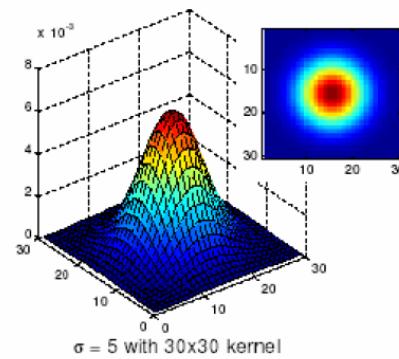
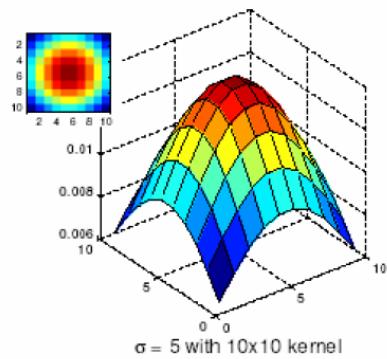


Gaussian Kernel

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



- Gaussian filters have infinite support, but discrete filters use finite kernels



Rule of thumb: Use kernel with each side 5 or 6σ

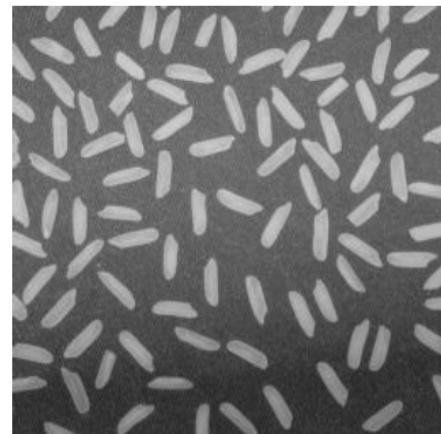
Gaussian Blur



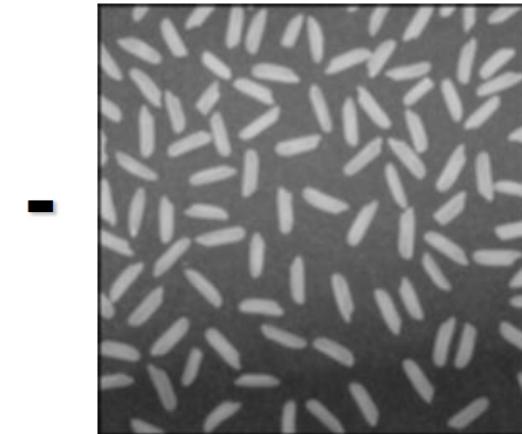
High-pass filters

Example of edge detection using low-pass filter:

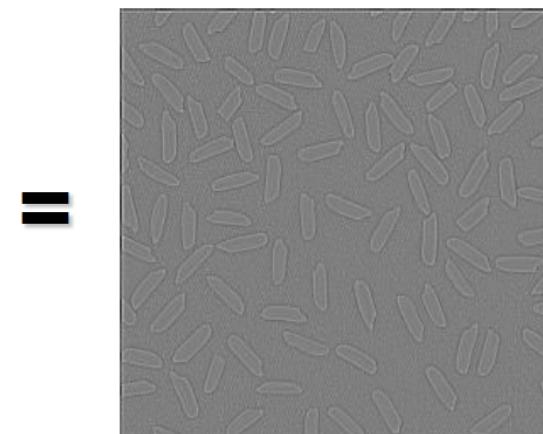
$$\text{HighPassImage} = \text{SourceImage} - \text{LowPassImage}$$



Source Image



Low-pass

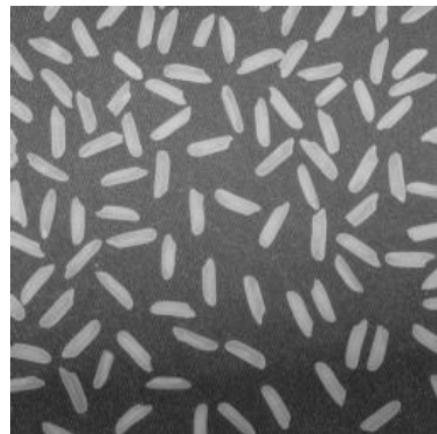


High-pass

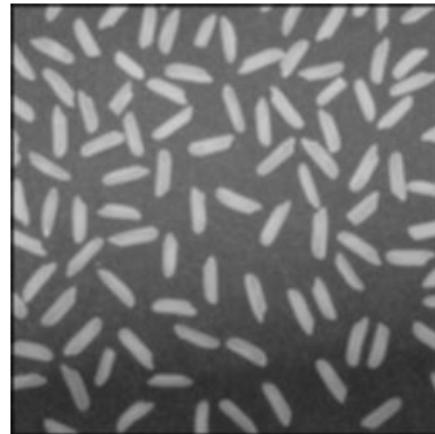
High-pass filters

Example of edge detection using low-pass filter:

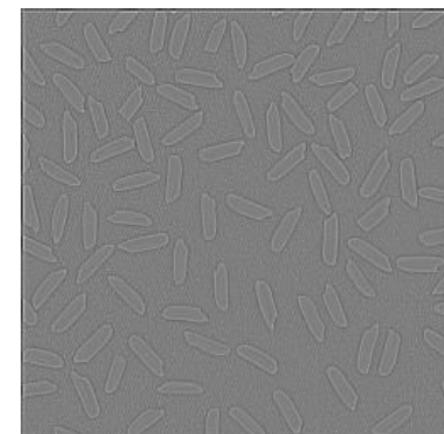
$$\text{HighPassImage} = \text{SourceImage} - \text{LowPassImage}$$



-



=



0	0	0
0	1	0
0	0	0

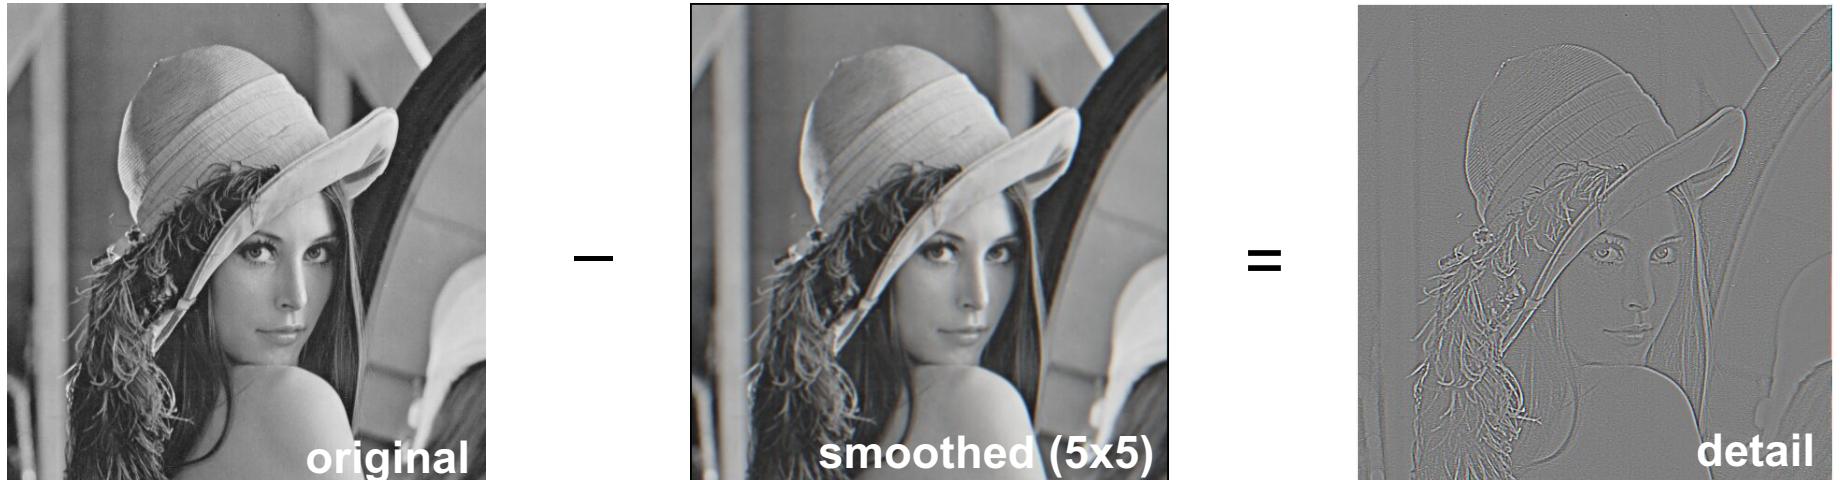
Source Image

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

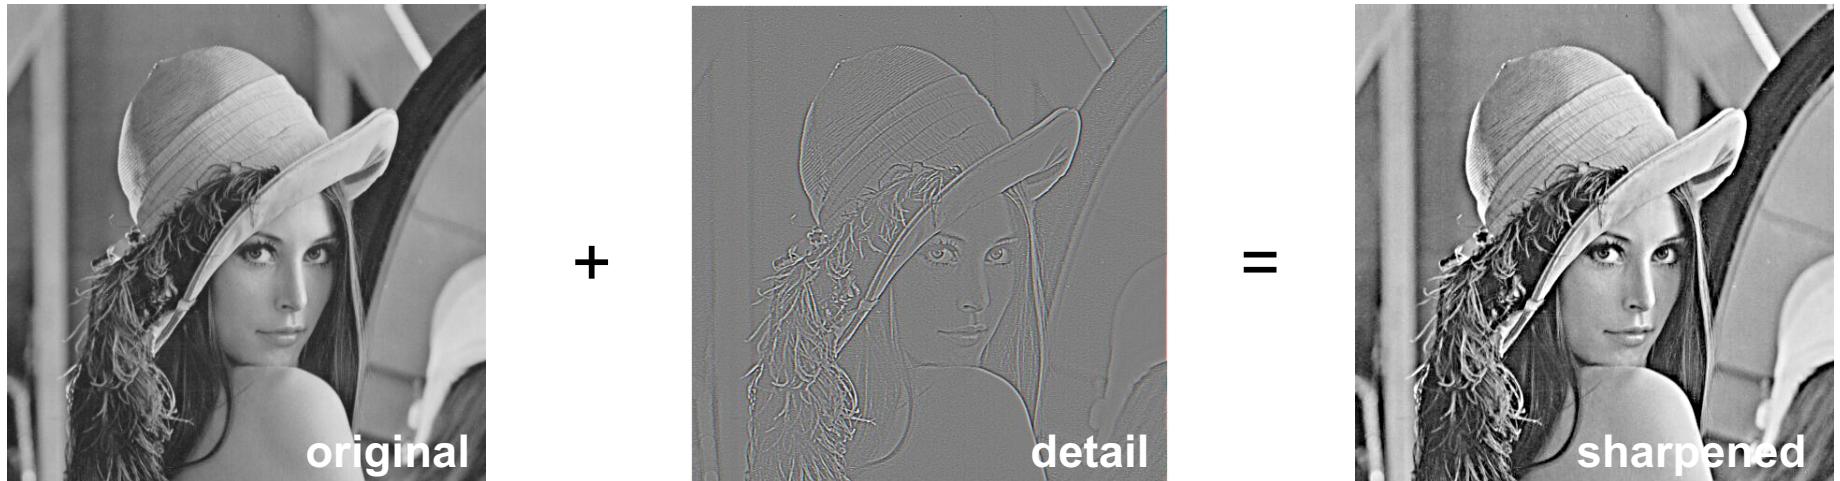
Low-pass

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$

High-pass



Let's add it back:

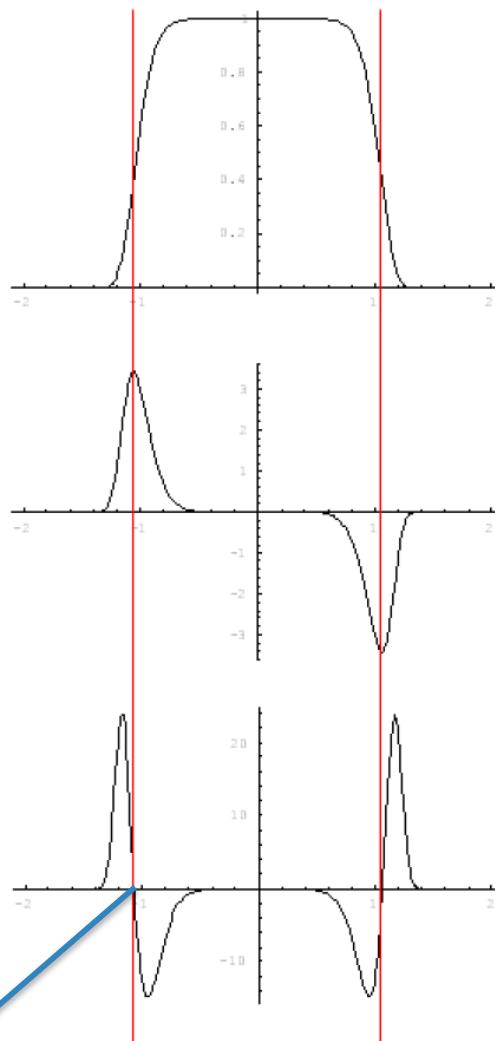


Signal Derivatives

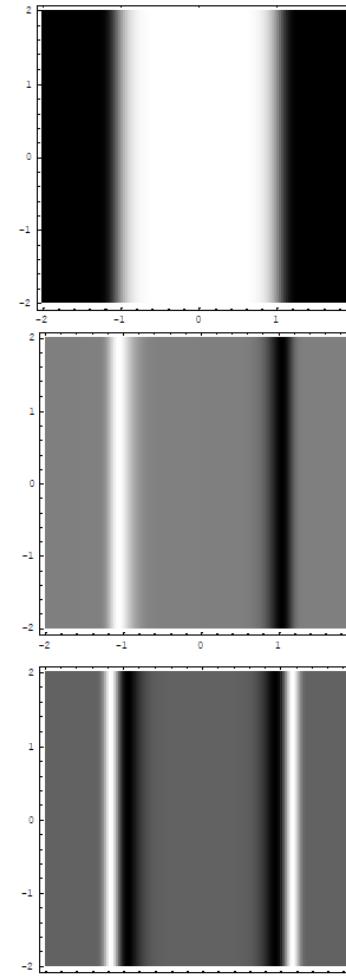
$$I(x, y) = e^{-x^{10}/2}$$

$$\frac{\partial I}{\partial x}$$

$$\frac{\partial^2 I}{\partial x^2}$$



Zero crossing



1st

2nd

Gradient

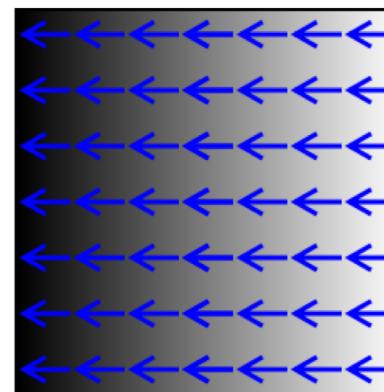
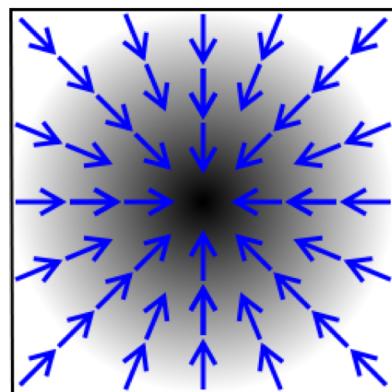
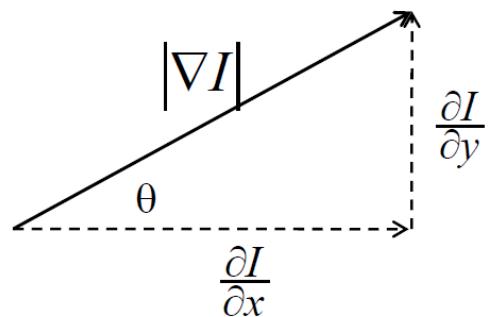
$$\nabla I = \begin{bmatrix} \frac{\partial I}{\partial x} \\ \frac{\partial I}{\partial y} \end{bmatrix}$$

Magnitude

$$|\nabla I| = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2} \approx \left|\frac{\partial I}{\partial x}\right| + \left|\frac{\partial I}{\partial y}\right|$$

Direction

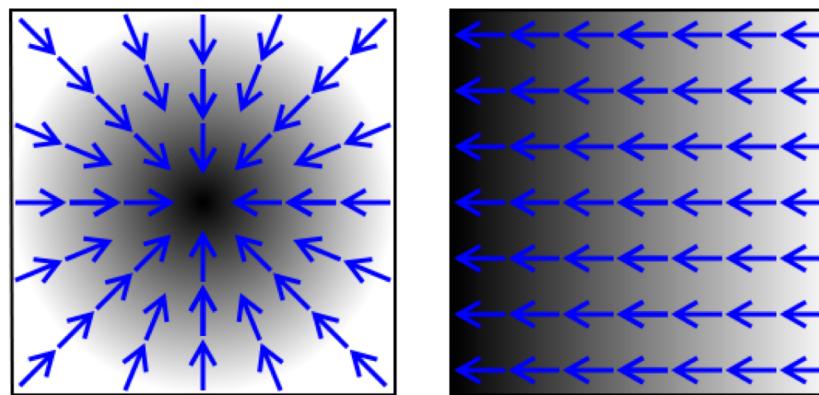
$$\theta = \tan^{-1} \frac{\left(\frac{\partial I}{\partial y}\right)}{\left(\frac{\partial I}{\partial x}\right)}$$



Gradient

Discrete approximation

$$\frac{\partial I}{\partial x} \approx \frac{I(x+1, y) - I(x-1, y)}{2}$$



Gradient

Image

I_1	I_2	I_3
I_4	I_5	I_6
I_7	I_8	I_9

Filter

0	-1	0
0	0	0
0	1	0

0	0	0
-1	0	1
0	0	0

$$\frac{\partial I}{\partial y} = [I_8 - I_2]$$

$$\frac{\partial I}{\partial x} = [I_6 - I_4]$$

Filters

Prewitt:

-1	-1	-1
0	0	0
1	1	1

-1	0	1
-1	0	1
-1	0	1

Sobel:

-1	-2	-1
0	0	0
1	2	1

-1	0	1
-2	0	2
-1	0	1

Prof. Judith M.S. Prewitt

Mathematical Methods Applied to Image Processing in Medicine

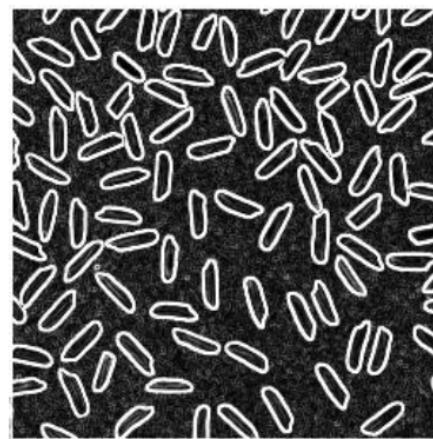
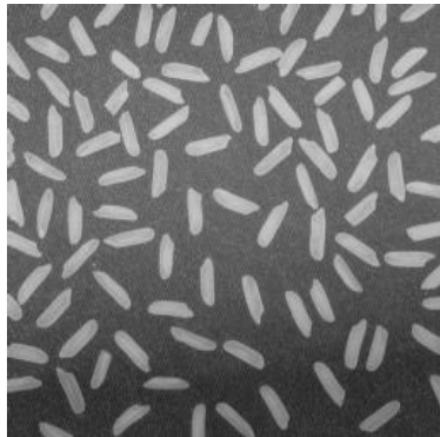
Chapter Dec 1980

 J. M. S. Prewitt

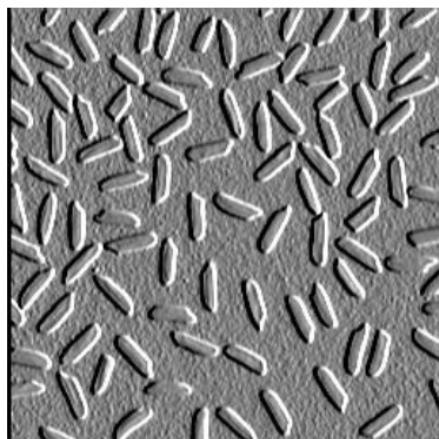
Image processing is a mathematical and engineering discipline which has as its domain digitized pictures, that is, pictures which have been spatially sampled and photometrically quantized by an electro-optical scanning device and have been recorded on a medium such as magnetic tape, magnetic disc or computer memory core. Each spatially discrete image element is called a pixel and with each pixel is associated a photometric quantity such as brightness, transmittance, extinction or optical density. Processing refers to mathematical and logical operations performed on these image pixels in order to extract desired information. Medical images of interest arise from blood smears showing white cells and showing red cells, Papanicolaou cervical smears, radiographs, electron micrographs, scintigrams, etc.



Sobel

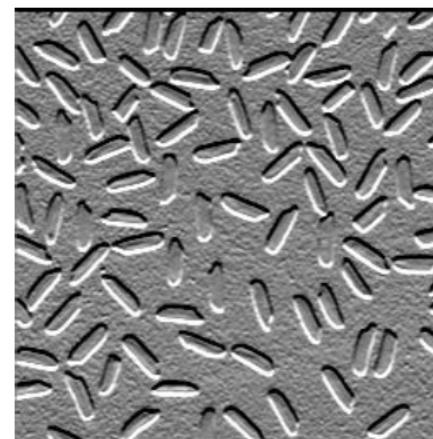


$$|\nabla I| = \left| \frac{\partial I}{\partial x} \right| + \left| \frac{\partial I}{\partial y} \right|$$



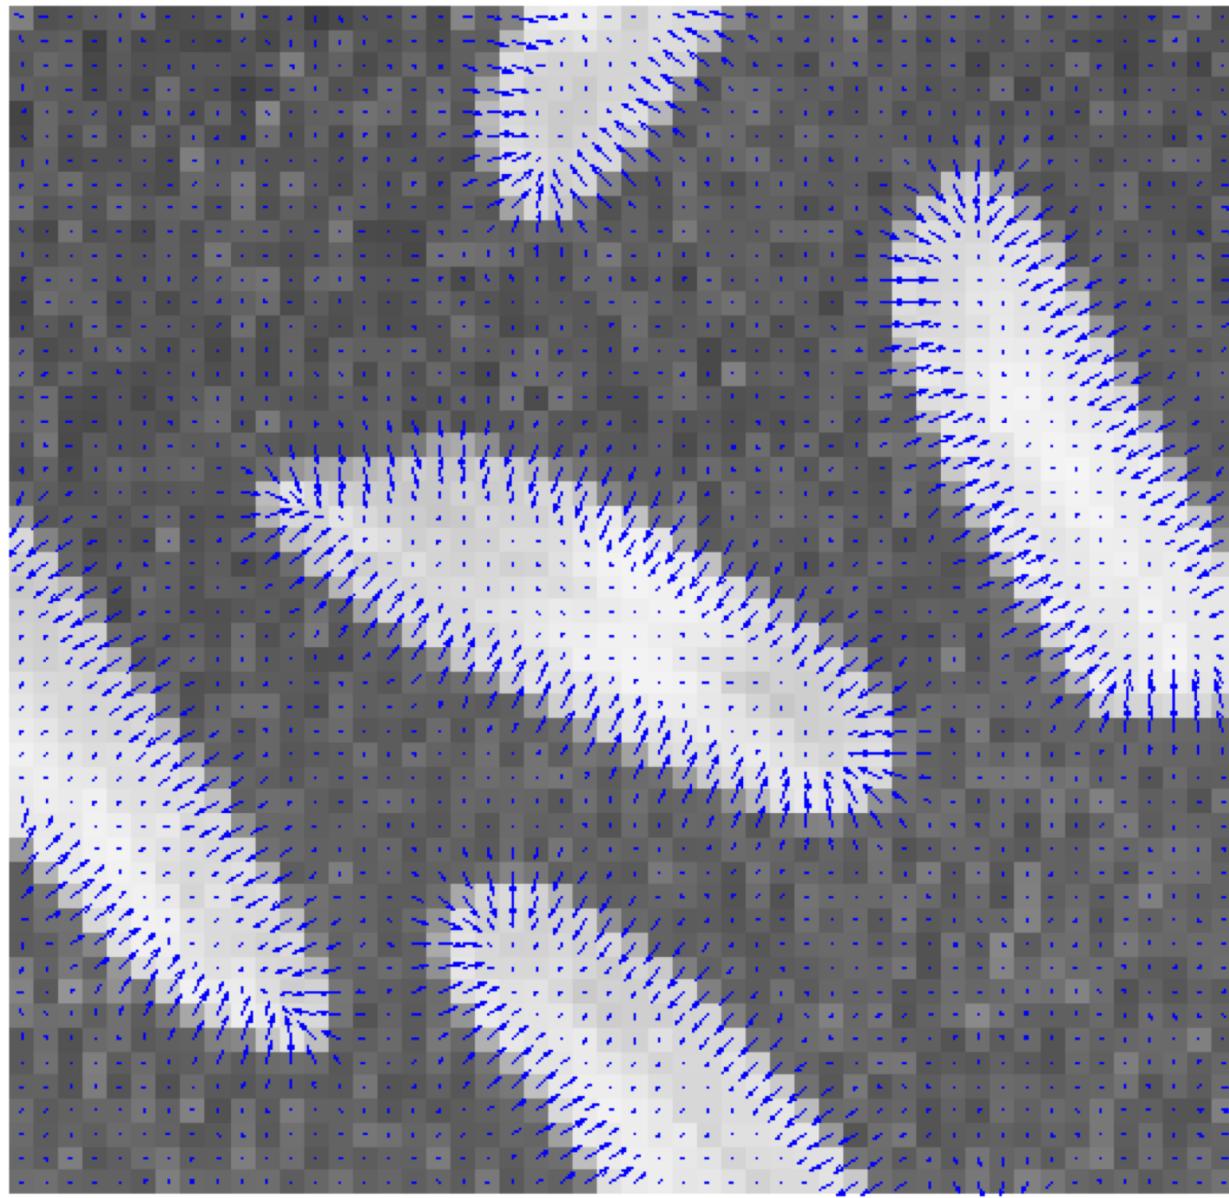
$$\frac{\partial I}{\partial x}$$

-1	0	1
-2	0	2
-1	0	1



$$\frac{\partial I}{\partial y}$$

-1	-2	-1
0	0	0
1	2	1



Laplacian: 2nd derivative

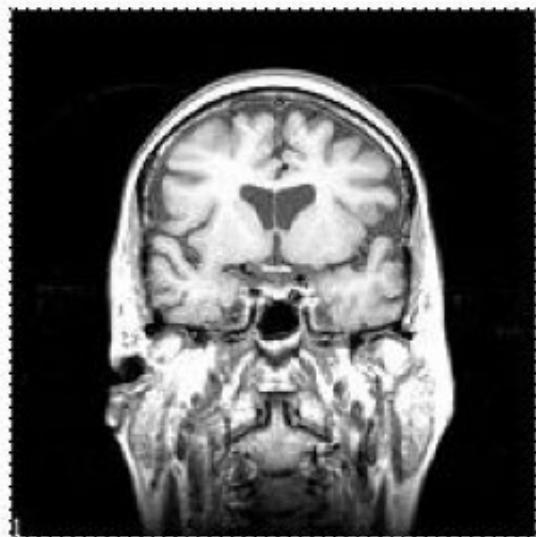
$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Discrete:

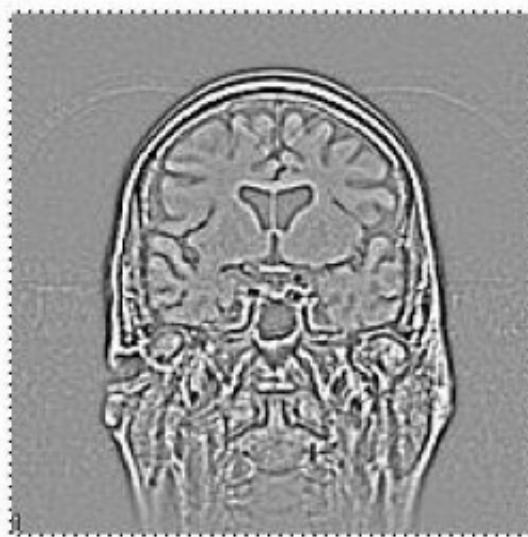
0	-1	0
-1	4	-1
0	-1	0

Laplacian

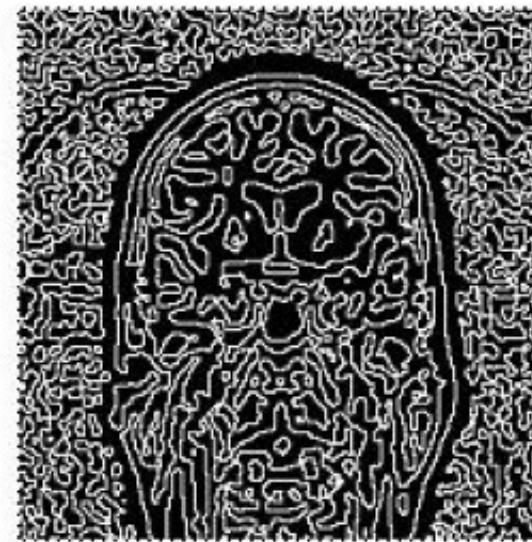
Laplacian: 2nd derivative



(A) Original MR image

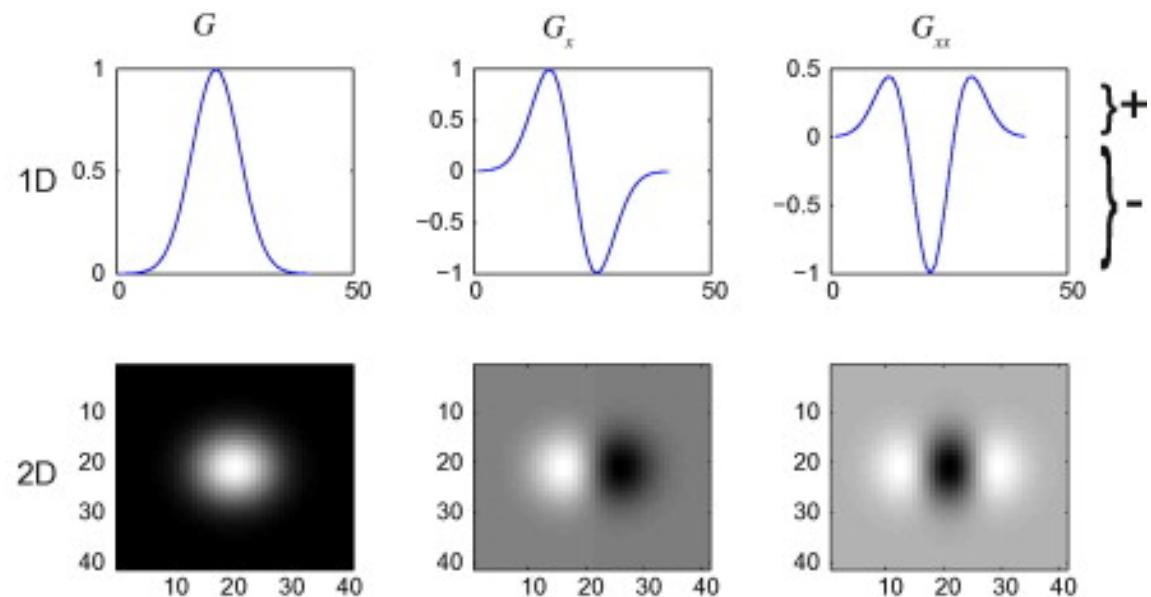


(B) Laplacian results



(C) Extraction of the zero crossing of the Laplacian (object edges)

Gaussian Derivatives: 1D vs 2D





prewitt



laplacian

Separability

- A two-dimensional filter kernel is **separable** if it can be expressed as the outer product of two vectors.
- Example: Sobel kernel

$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} [+1 \ 0 \ -1]$$

Separable Convolution

- **Associativity of convolution:**

$$f * (g * h) = (f * g) * h$$

$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} +1 & 0 & -1 \end{bmatrix}$$

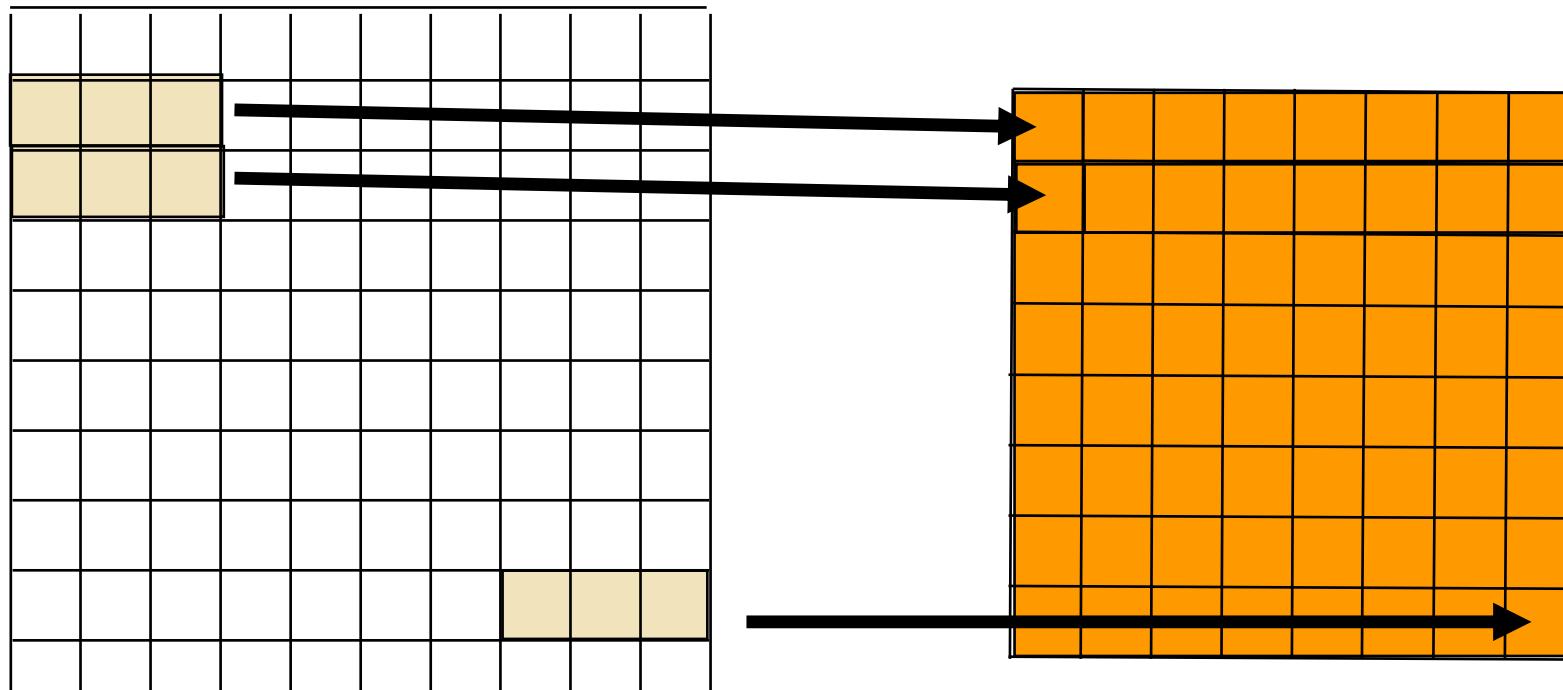
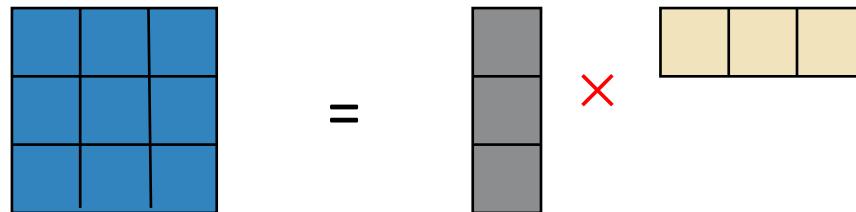
$$K = \quad g \quad * \quad h$$

$$Dest[i, j] = \sum_{l=0}^{Kh} \sum_{k=0}^{Kw} Kernel[k, l] Source[i+k, j+l]$$

$$Dest[i, j] = \sum_{l=0}^{Kh} Vert[l] \sum_{k=0}^{Kw} Horiz[k] Source[i+k, l+j]$$

The Gaussian kernel is separable

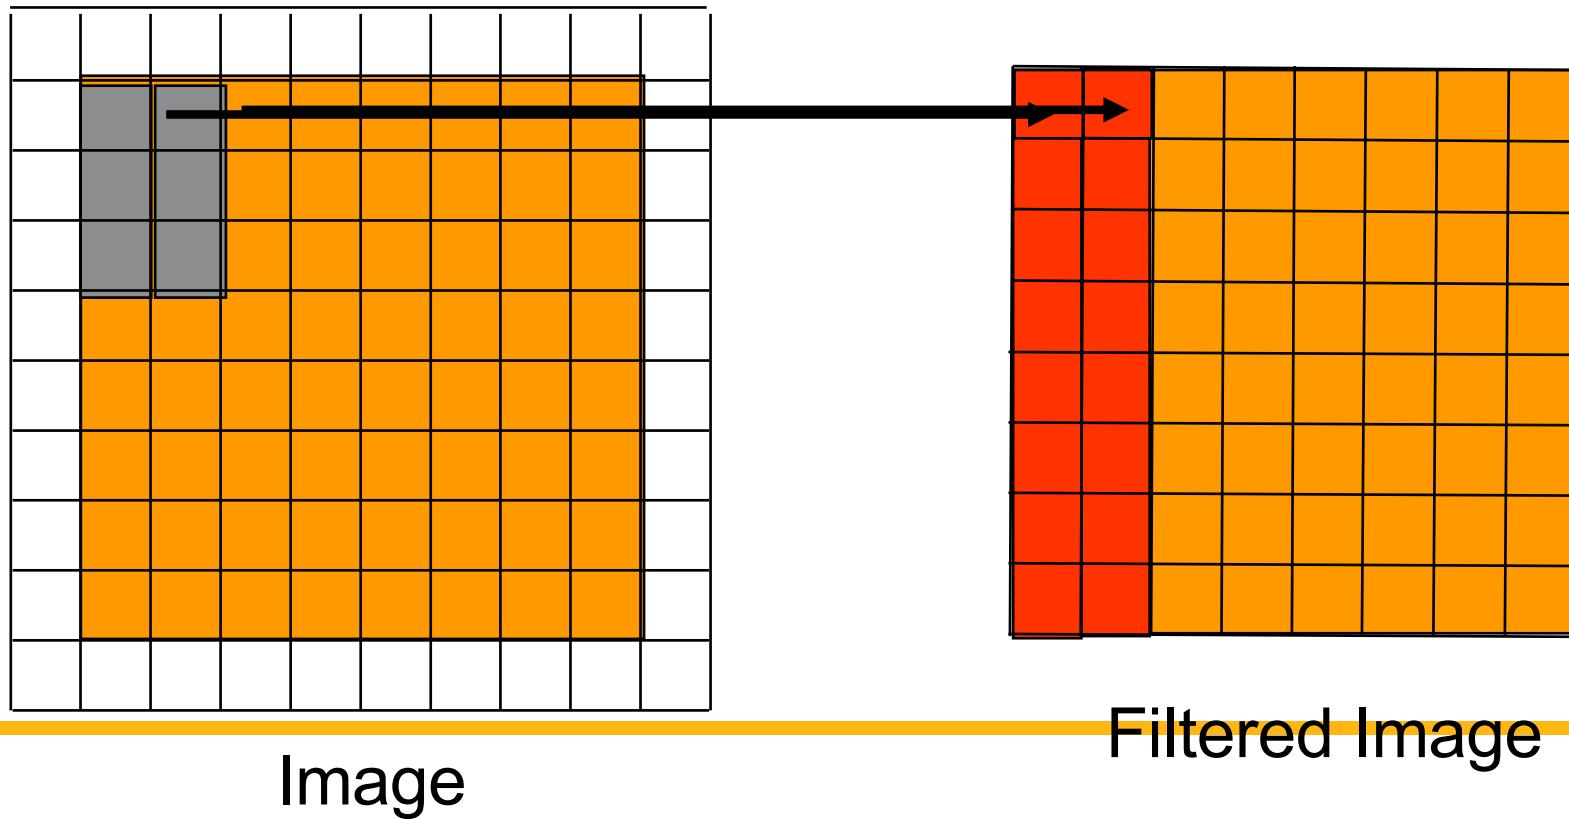
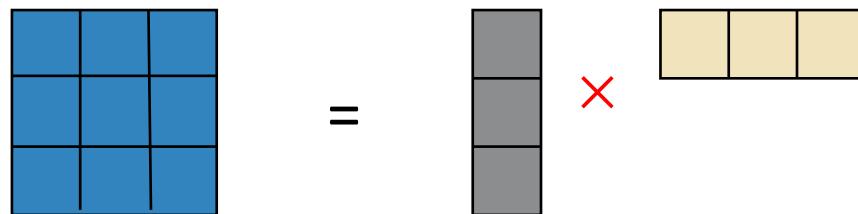
Separable Filtering



Image

Filtered Image

Separable Filtering



Separable Filtering

What is the number of operations for the 3×5 kernel H ?

Answer: $15wh$

What is the number of operations for H_x followed by H_y ?

Answer: $3wh + 5wh = 8wh$

What about the case of a $M \times M$ kernel?

Answer:

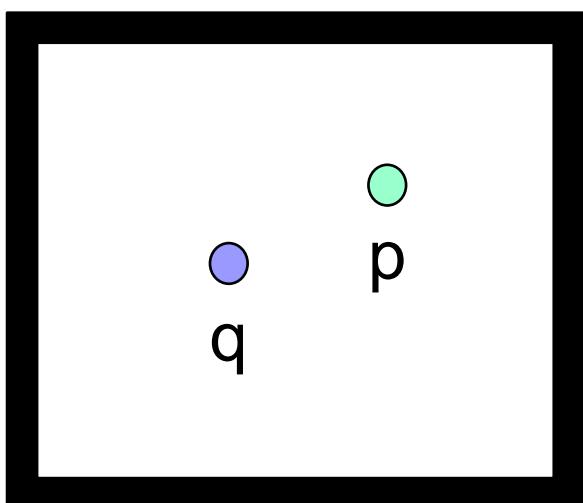
$O(M^2)$ – no separability (M^2wh operations)

$O(M)$ – with separability ($2Mwh$ operations)

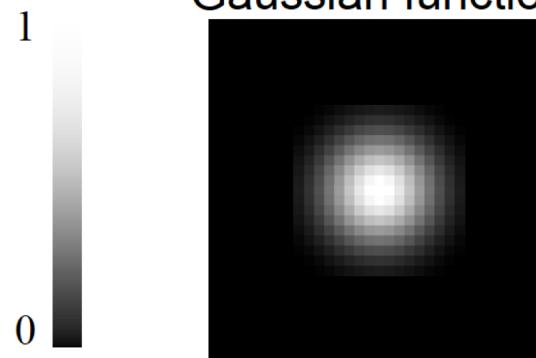
Properties of Gaussian Filter

Same idea: **weighted average of pixels.**

$$GB[I]_p = \sum_{q \in S} G_\sigma(\|p-q\|) I_q$$



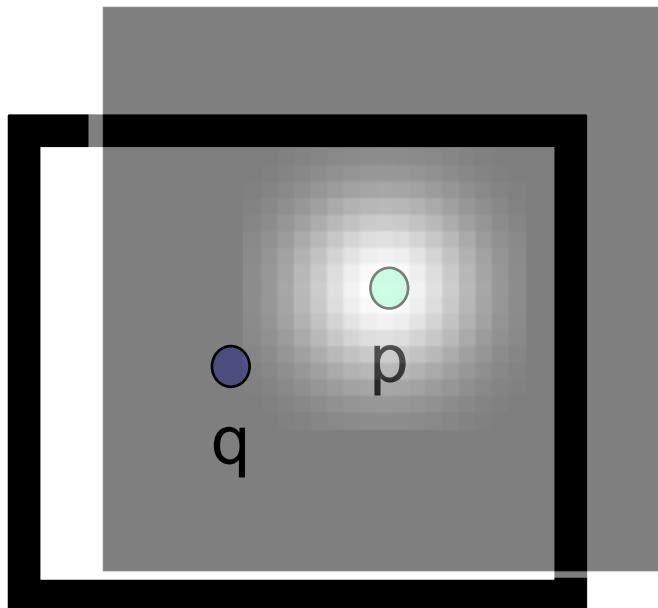
normalized
Gaussian function



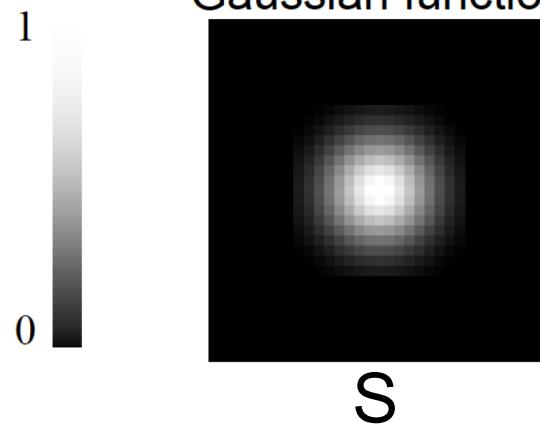
Properties of Gaussian Filter

Same idea: **weighted average of pixels.**

$$GB[I]_p = \sum_{q \in S} G_\sigma(\|p-q\|) I_q$$



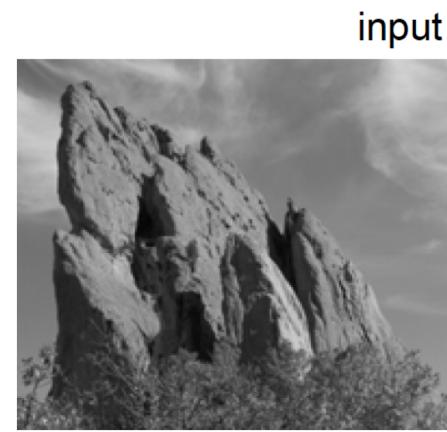
normalized
Gaussian function



Properties of Gaussian Filter

- Does smooth images
- But smoothes too much:
edges are blurred.
 - Only spatial distance matters
 - No edge term

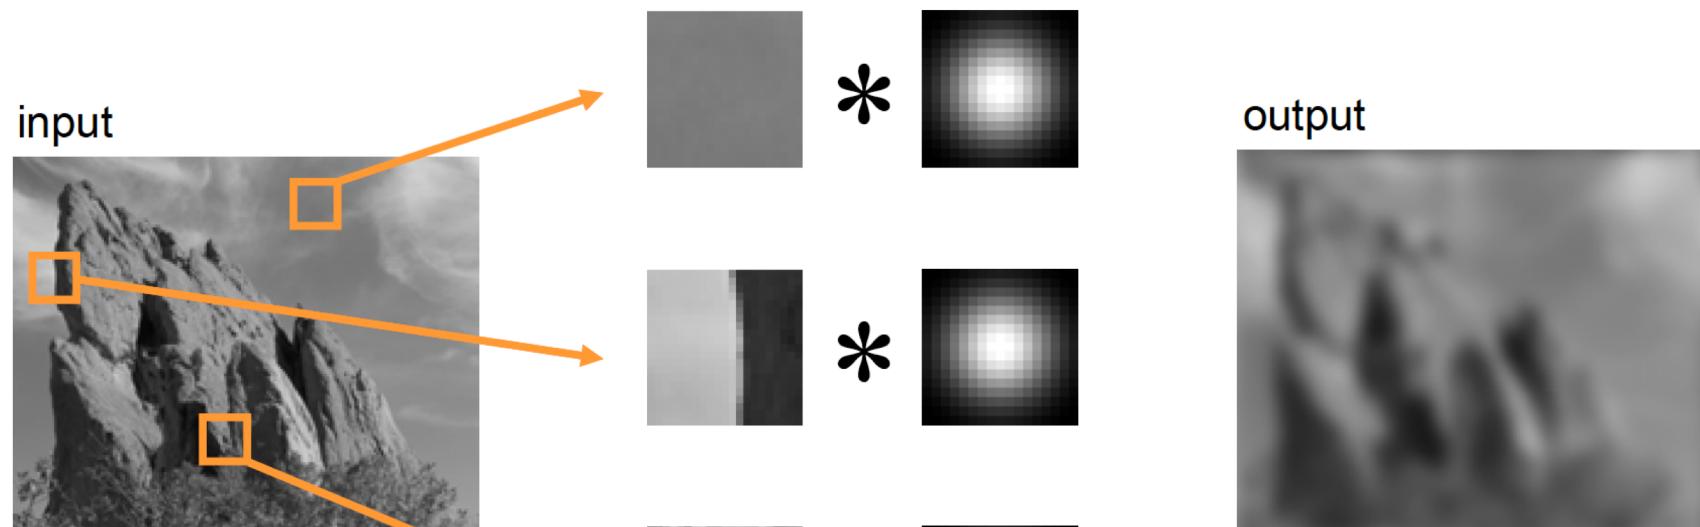
$$GB[I]_p = \sum_{q \in S} G_\sigma(\| p - q \|) I_q$$



input

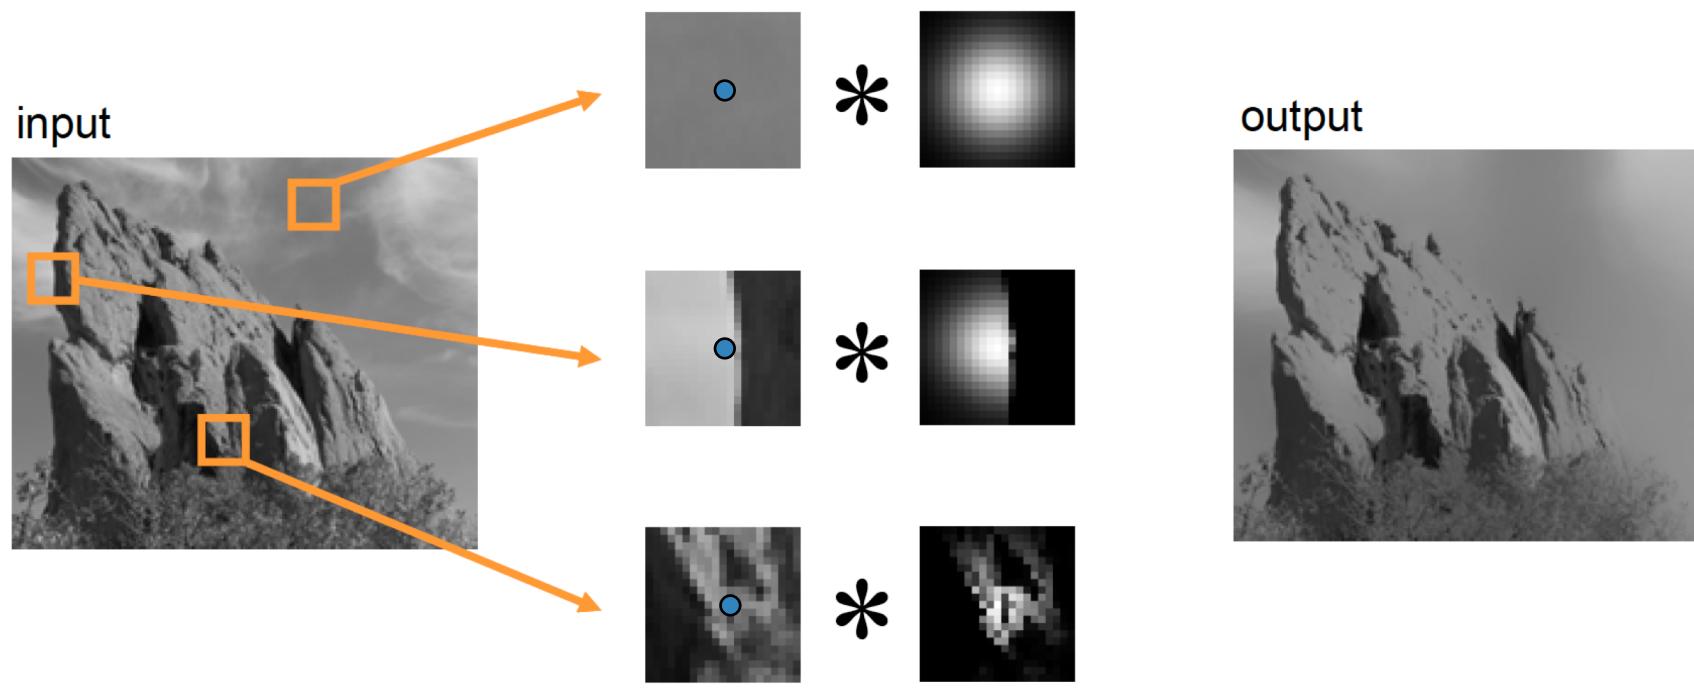
output

Properties of Gaussian Filter



Same Gaussian kernel everywhere.

Bilateral Filter



The kernel shape depends on the image content.

Bilateral Filtering



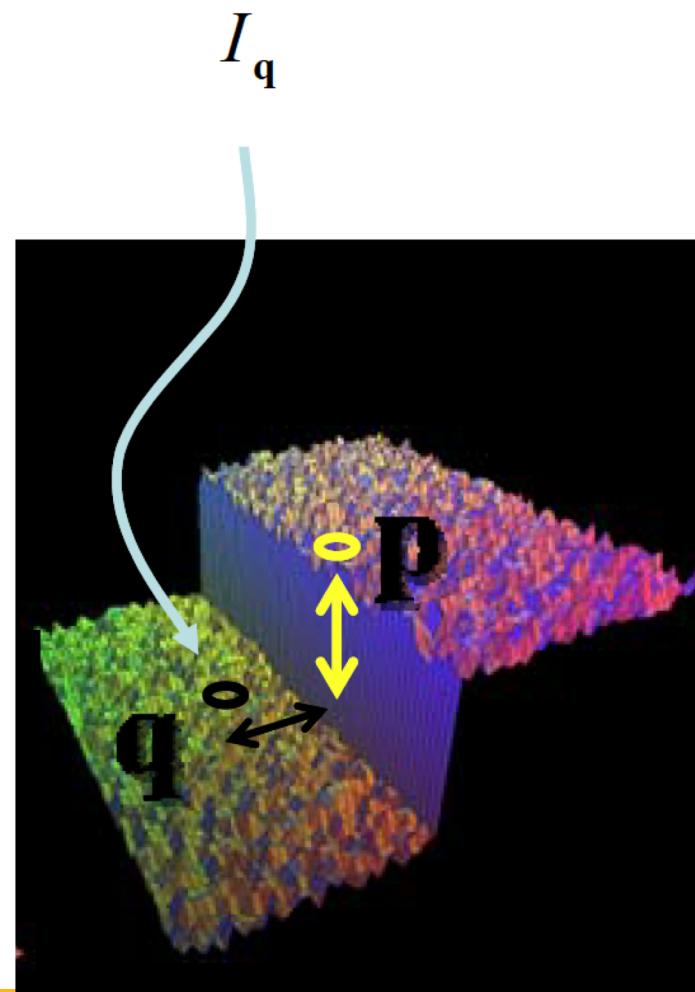
Bilateral Filtering

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s} (\| p - q \|) G_{\sigma_r} (\| I_p - I_q \|) I_q$$

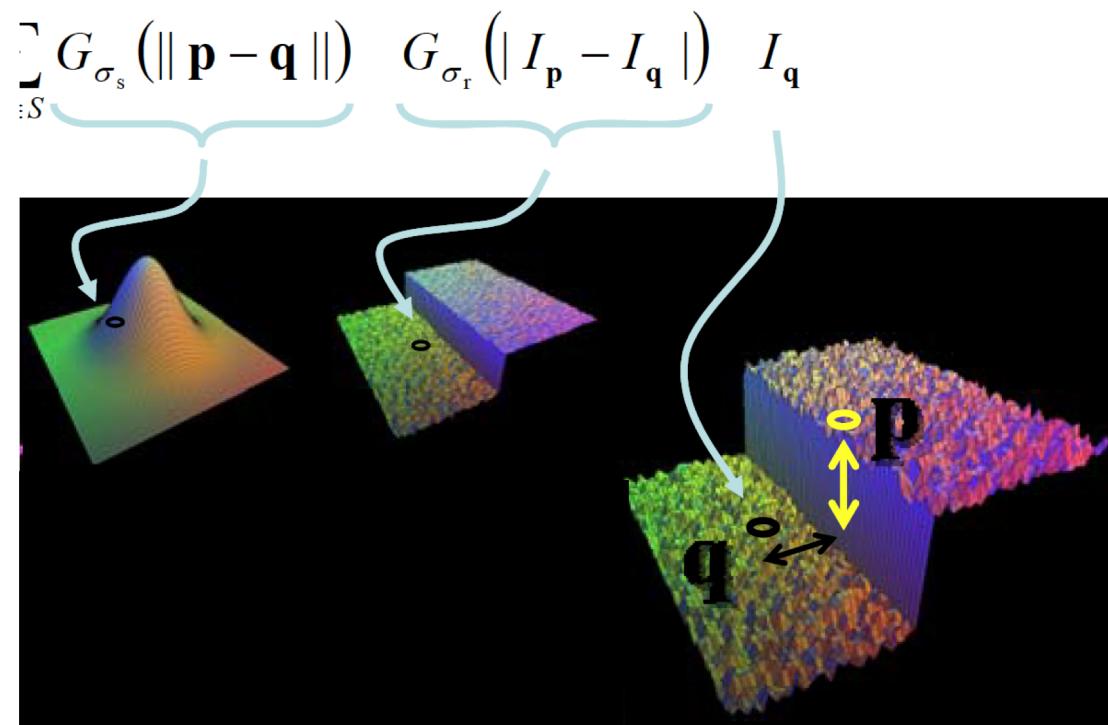
new
not new
new

normalization factor space weight range weight

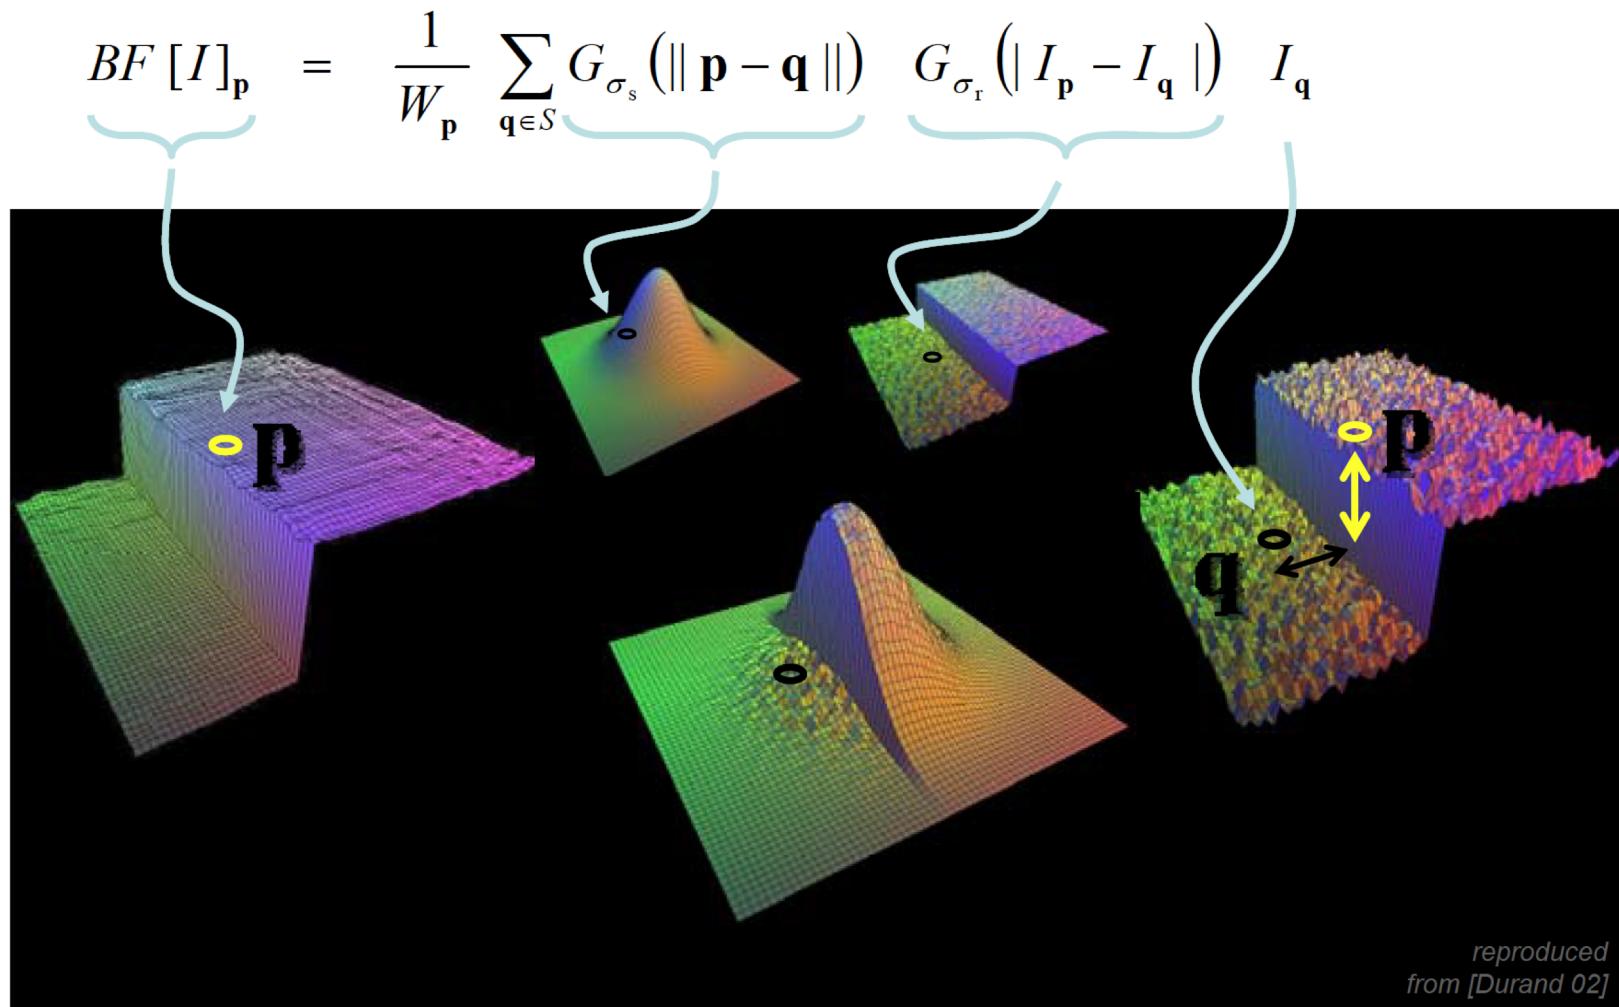
Bilateral Filtering



Bilateral Filtering



Bilateral Filtering



Bilateral filtering: parameters

$$BF [I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s}(\| p - q \|) G_{\sigma_r}(|I_p - I_q|) I_q$$


- space σ_s : spatial extent of the kernel, size of the considered neighborhood.
- range σ_r : “minimum” amplitude of an edge

Bilateral filtering: parameters

input

Exploring the Parameter Space

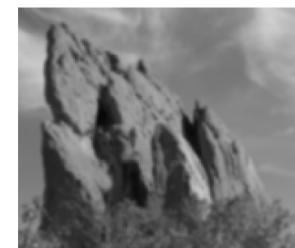
$$\sigma_r = 0.1$$



$$\sigma_r = 0.25$$



$$\sigma_r = \infty \\ (\text{Gaussian blur})$$



$$\sigma_s = 2$$

$$\sigma_s = 6$$

$$\sigma_s = 18$$