

Introduction to Computer Vision

3. Feature Extraction and Classification

UCLA – CS 188 – Fall 2019

Fabien Scalzo, Ph.D.

Part I: Part II:

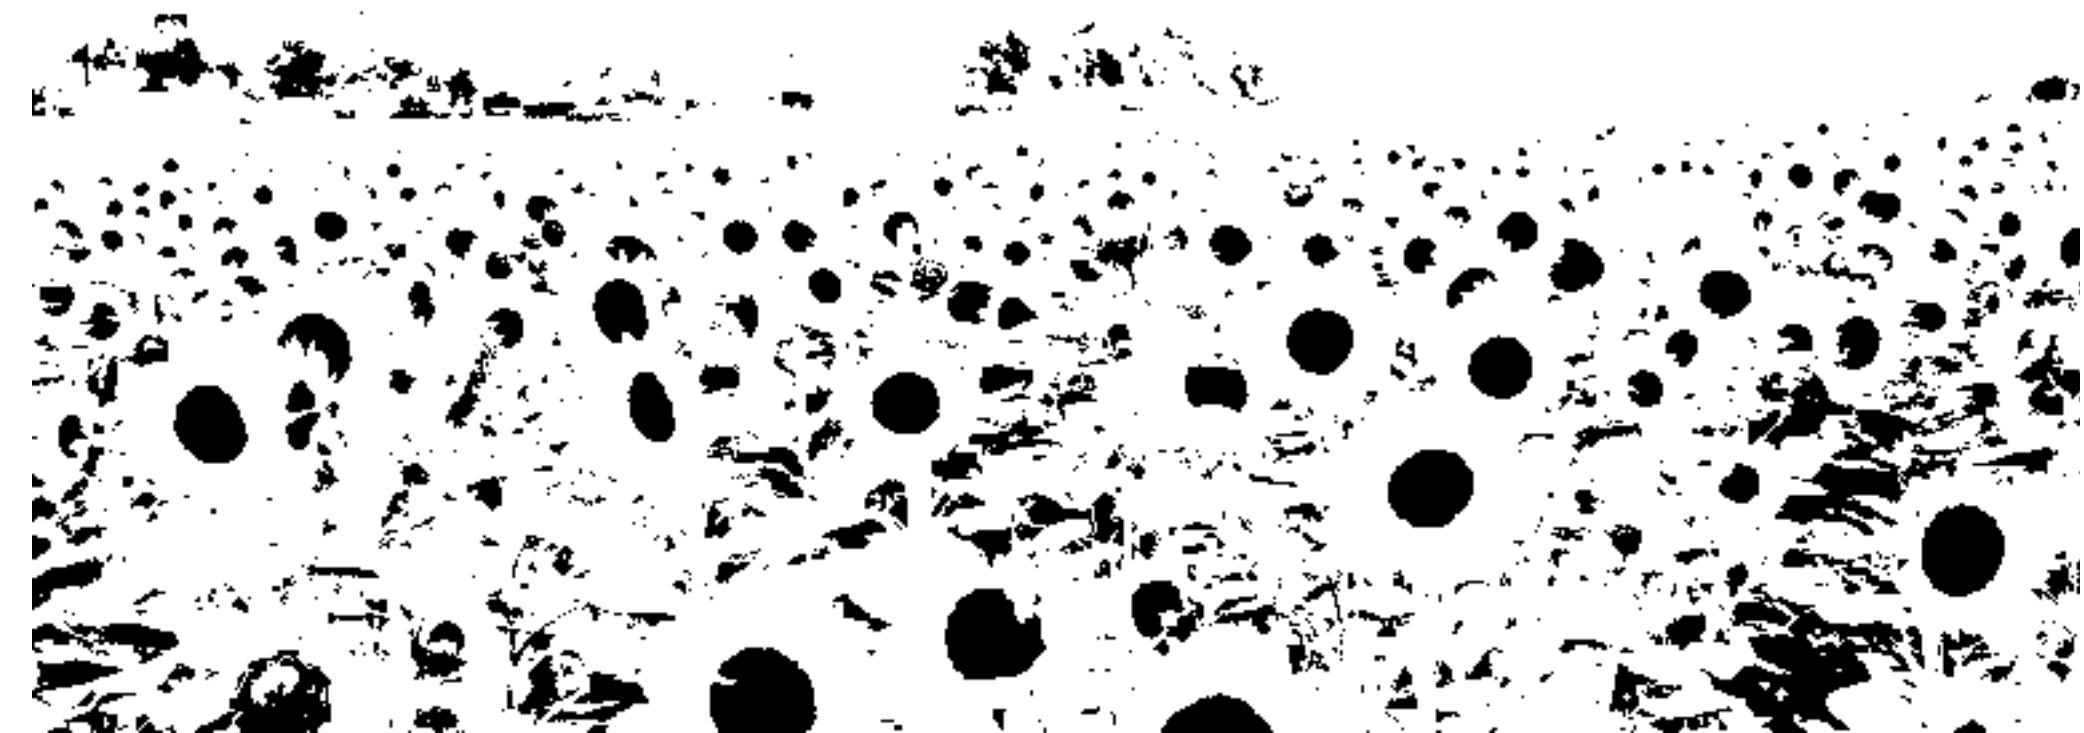
Feature Extraction Clustering and Classification

Week 1			26-Sep	Introduction
Week 2	1-Oct	Basic Image Processing	3-Oct	Feature Extraction and Classification
Week 3	8-Oct	Feature Tracking/Optical Flow	10-Oct	2D Image transformations, RANSAC
Week 4	15-Oct	SVD, 2D camera model, projective plane	17-Oct	Euclidean geometry, rigid body motion
Week 5	22-Oct	Epipolar Geometry	24-Oct	3D Cameras and processing
Week 6	29-Oct	Midterm	31-Oct	Statistical decision theory/Pattern Recognition
Week 7	5-Nov	Deep Learning	7-Nov	Deep Learning for Image Classification
Week 8	12-Nov	Object Detection	14-Nov	Generative models
Week 9	19-Nov	Medical Imaging	21-Nov	Autonomous Navigation
Week 10	26-Nov	Guest Lecture (Nikhil Naik)	28-Nov	
Week 11	3-Dec	Recursive 3D reconstruction and pose estimation	5-Dec	Recap
Week 12	10-Dec		12-Dec	Final

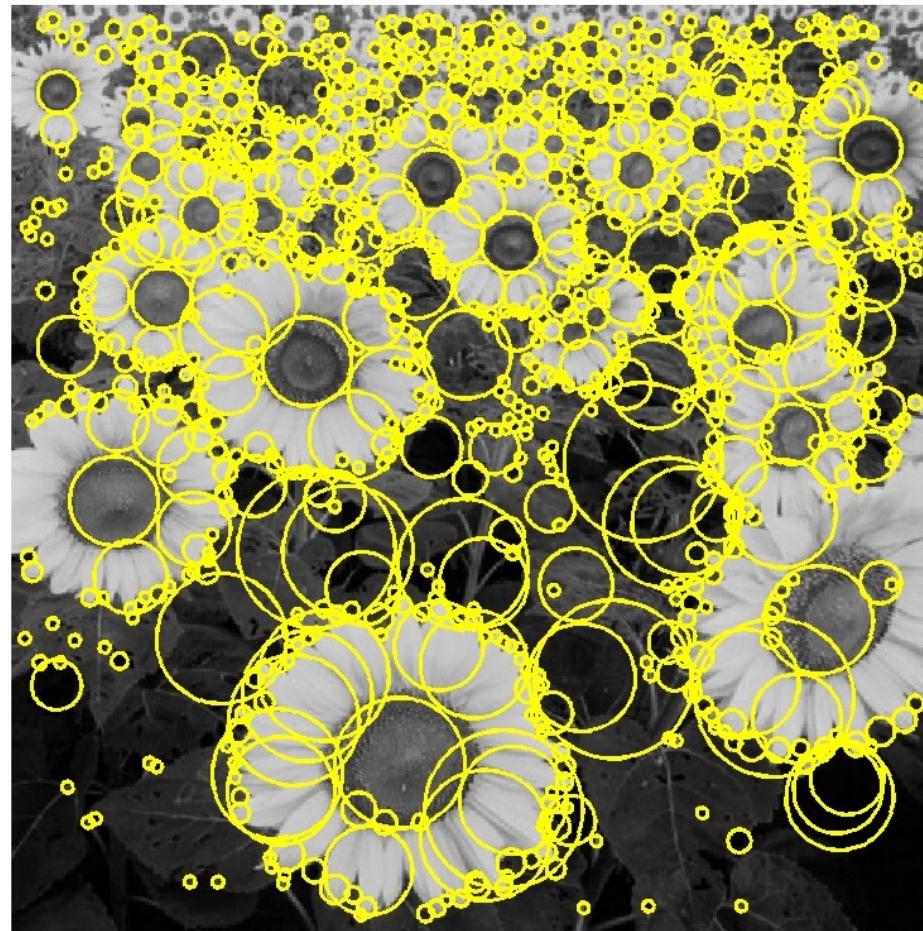
Which features best describe this image?







Blob detection



Automatic Feature Extraction from an image

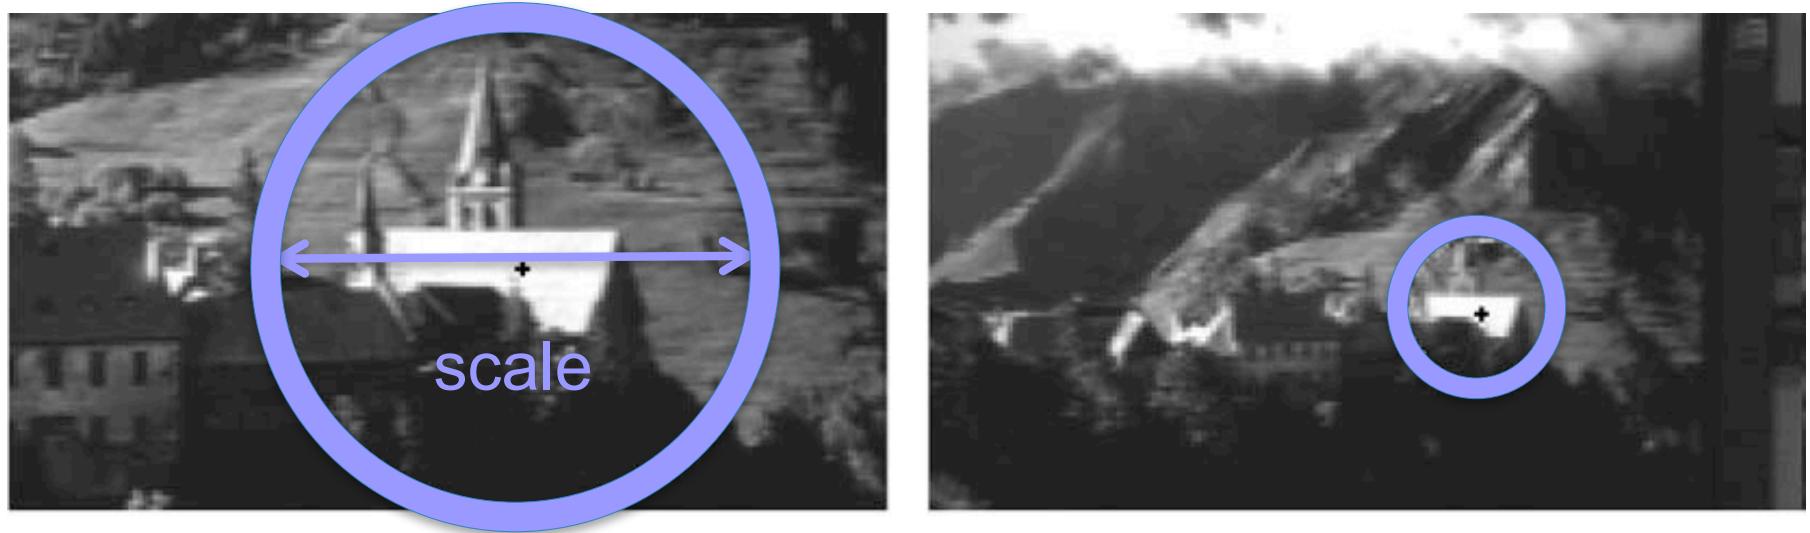
- Detection: Where is it?
- Description: What does it look like?

should be robust to common variations

- ✓ Intensity
- ✓ Scaling
- ✓ Viewpoint

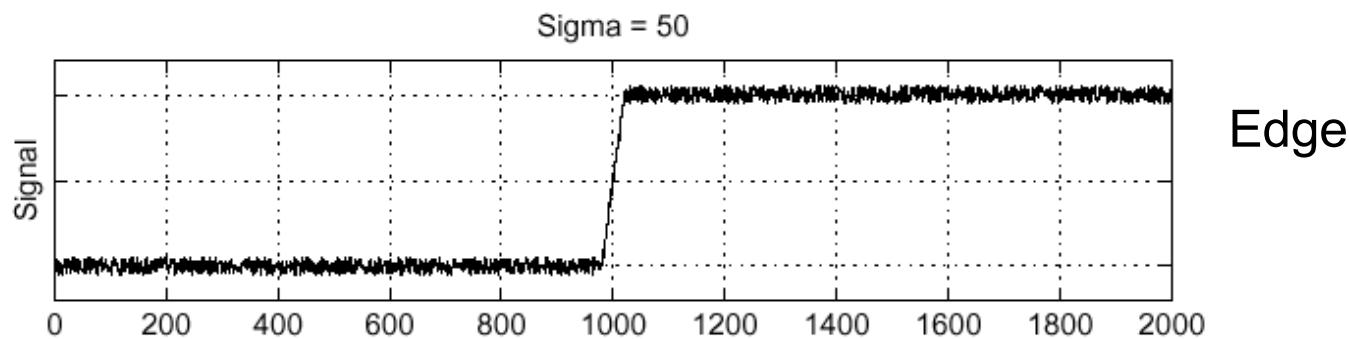
Blob detection and Scale invariance

- Goal: independently detect corresponding regions in scaled versions of the same image
- Characteristic region size that is **covariant** with the image.
Content description is **invariant** to transformation.



Edge detection, version 1

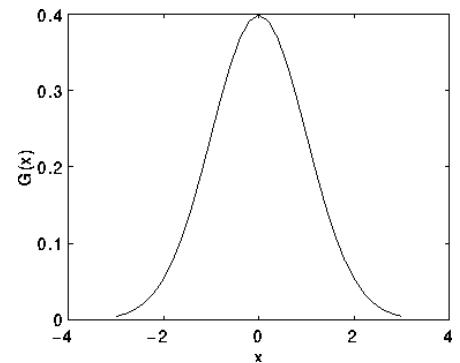
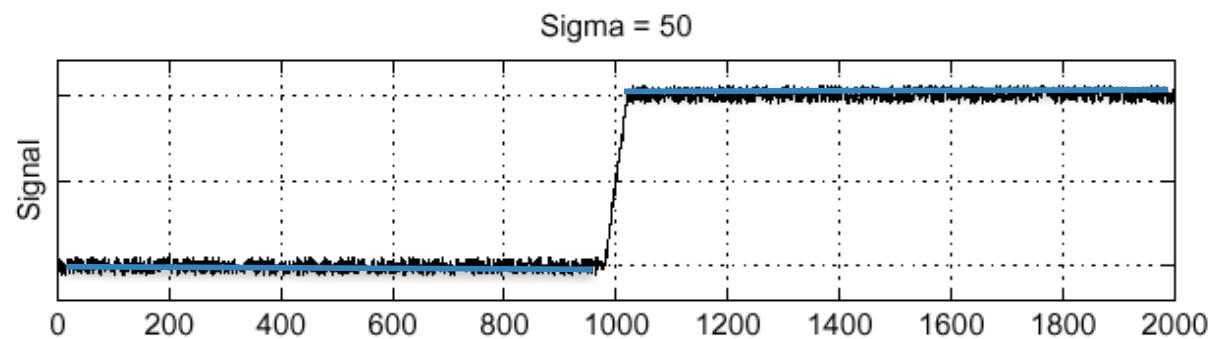
f



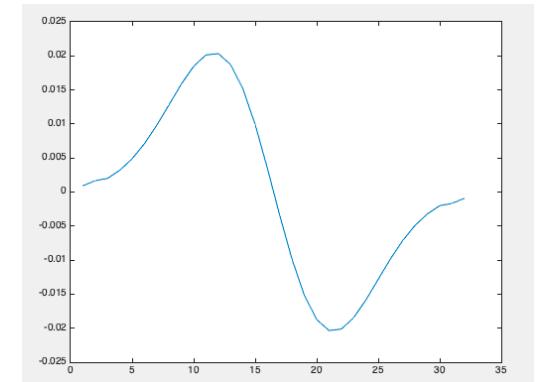
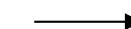
1	0	-1
----------	----------	-----------

Edge detection, version 1

f



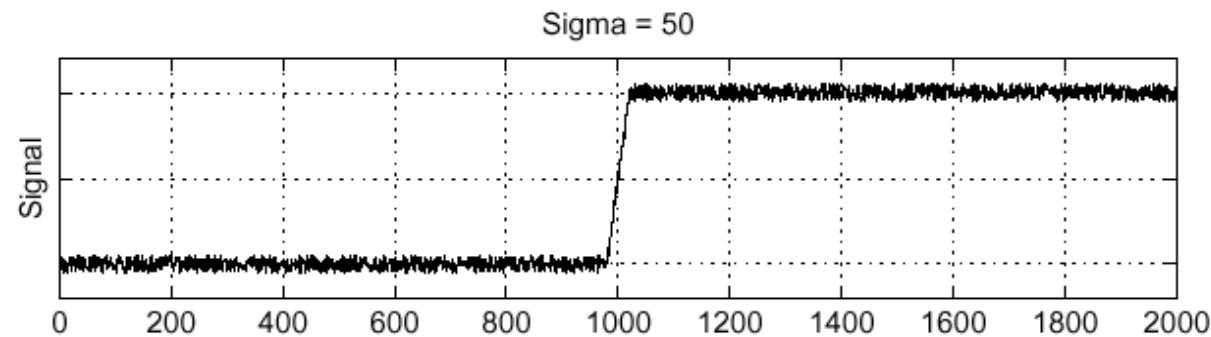
$$\ast \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array}$$



Gaussian filter

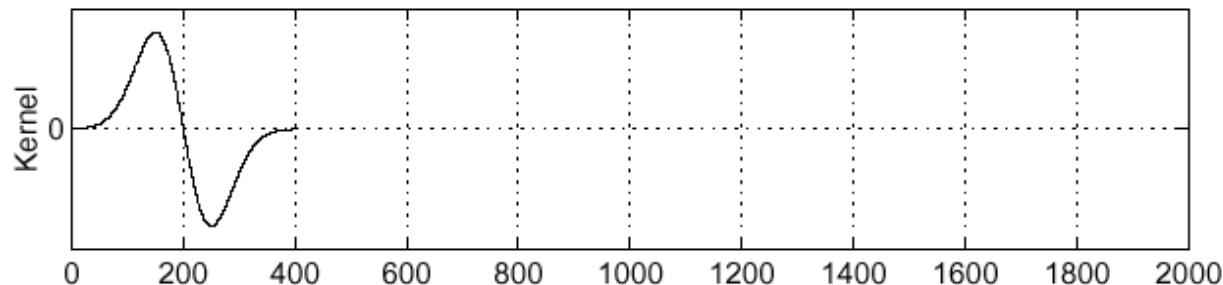
Edge detection, version 1

f



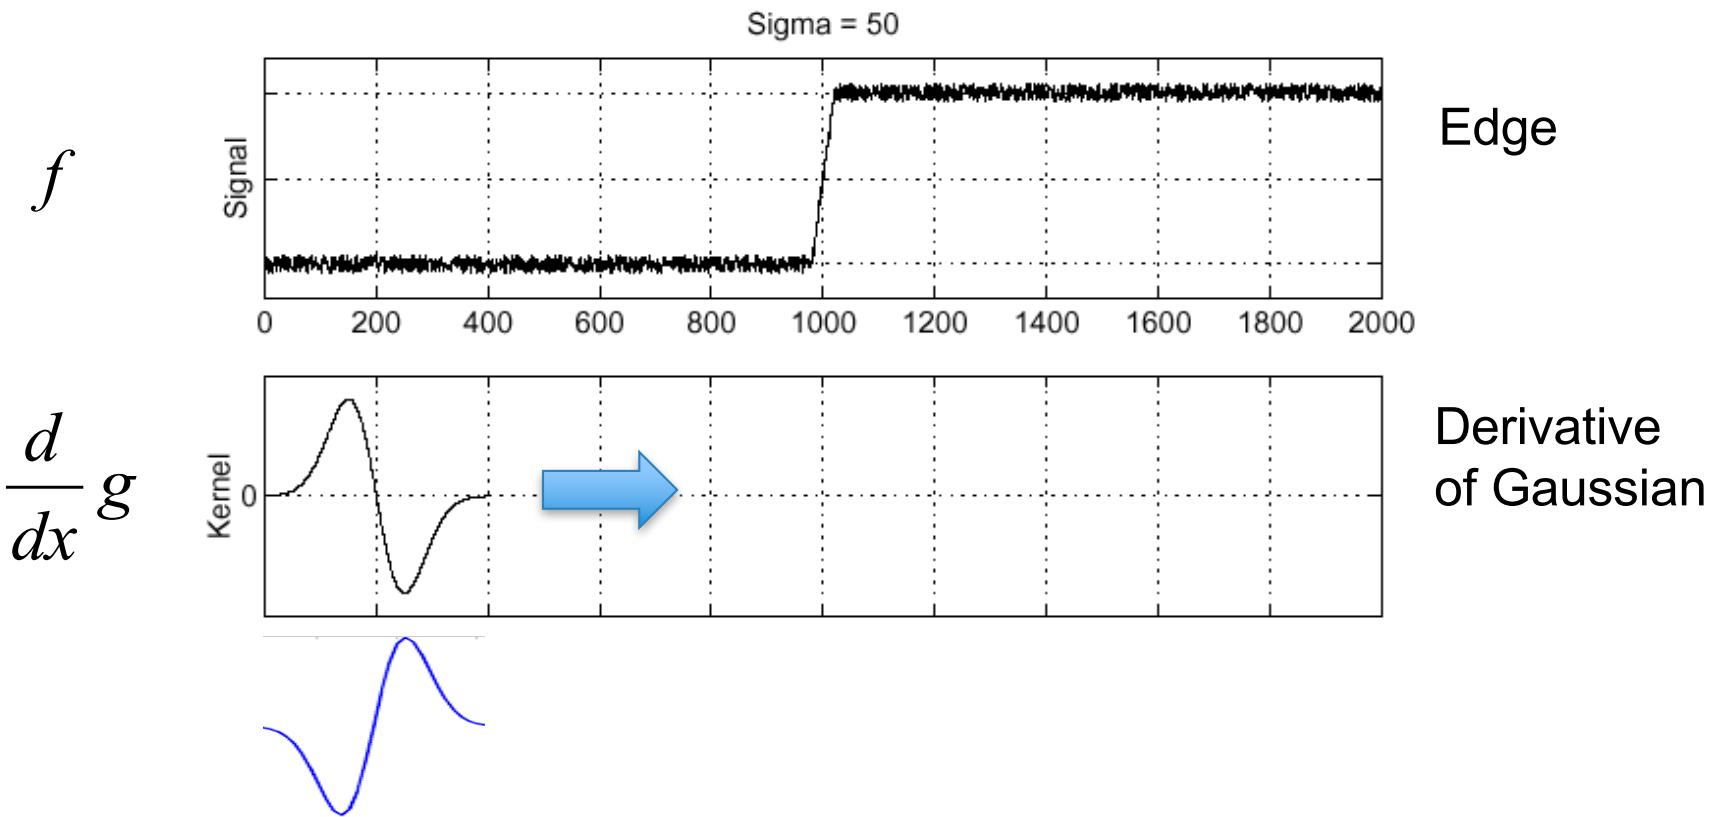
Edge

$\frac{d}{dx} g$

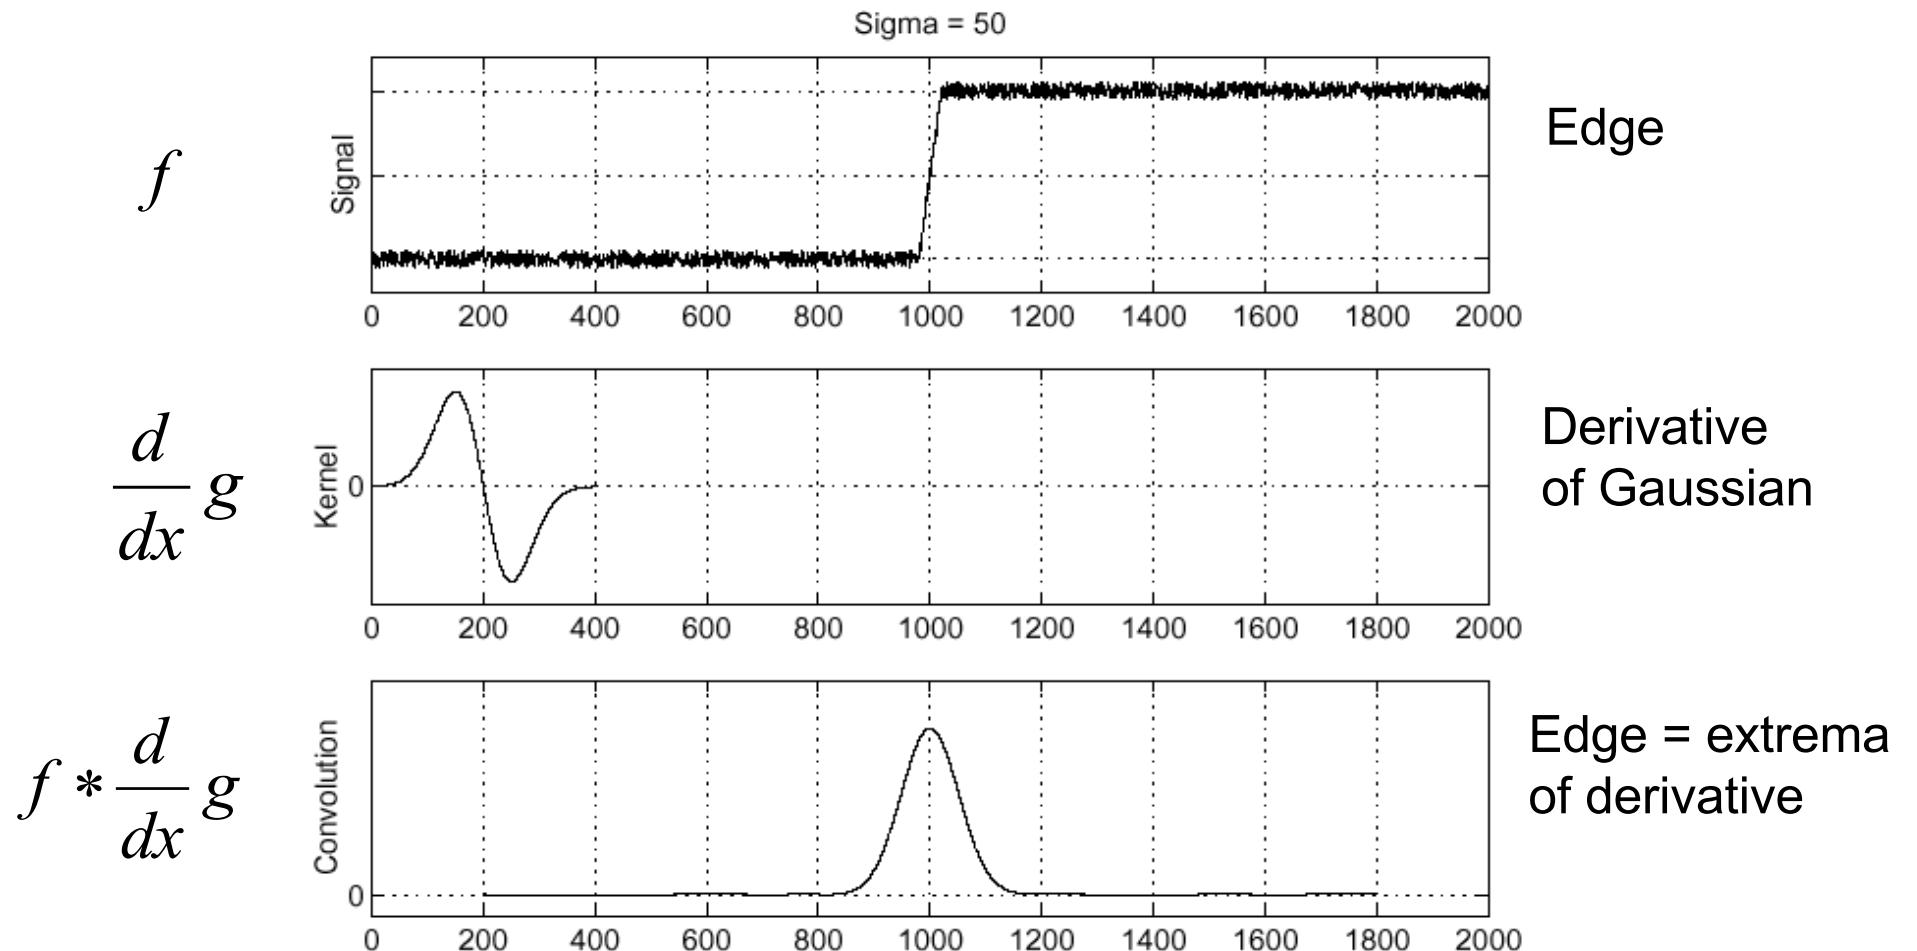


Derivative
of Gaussian

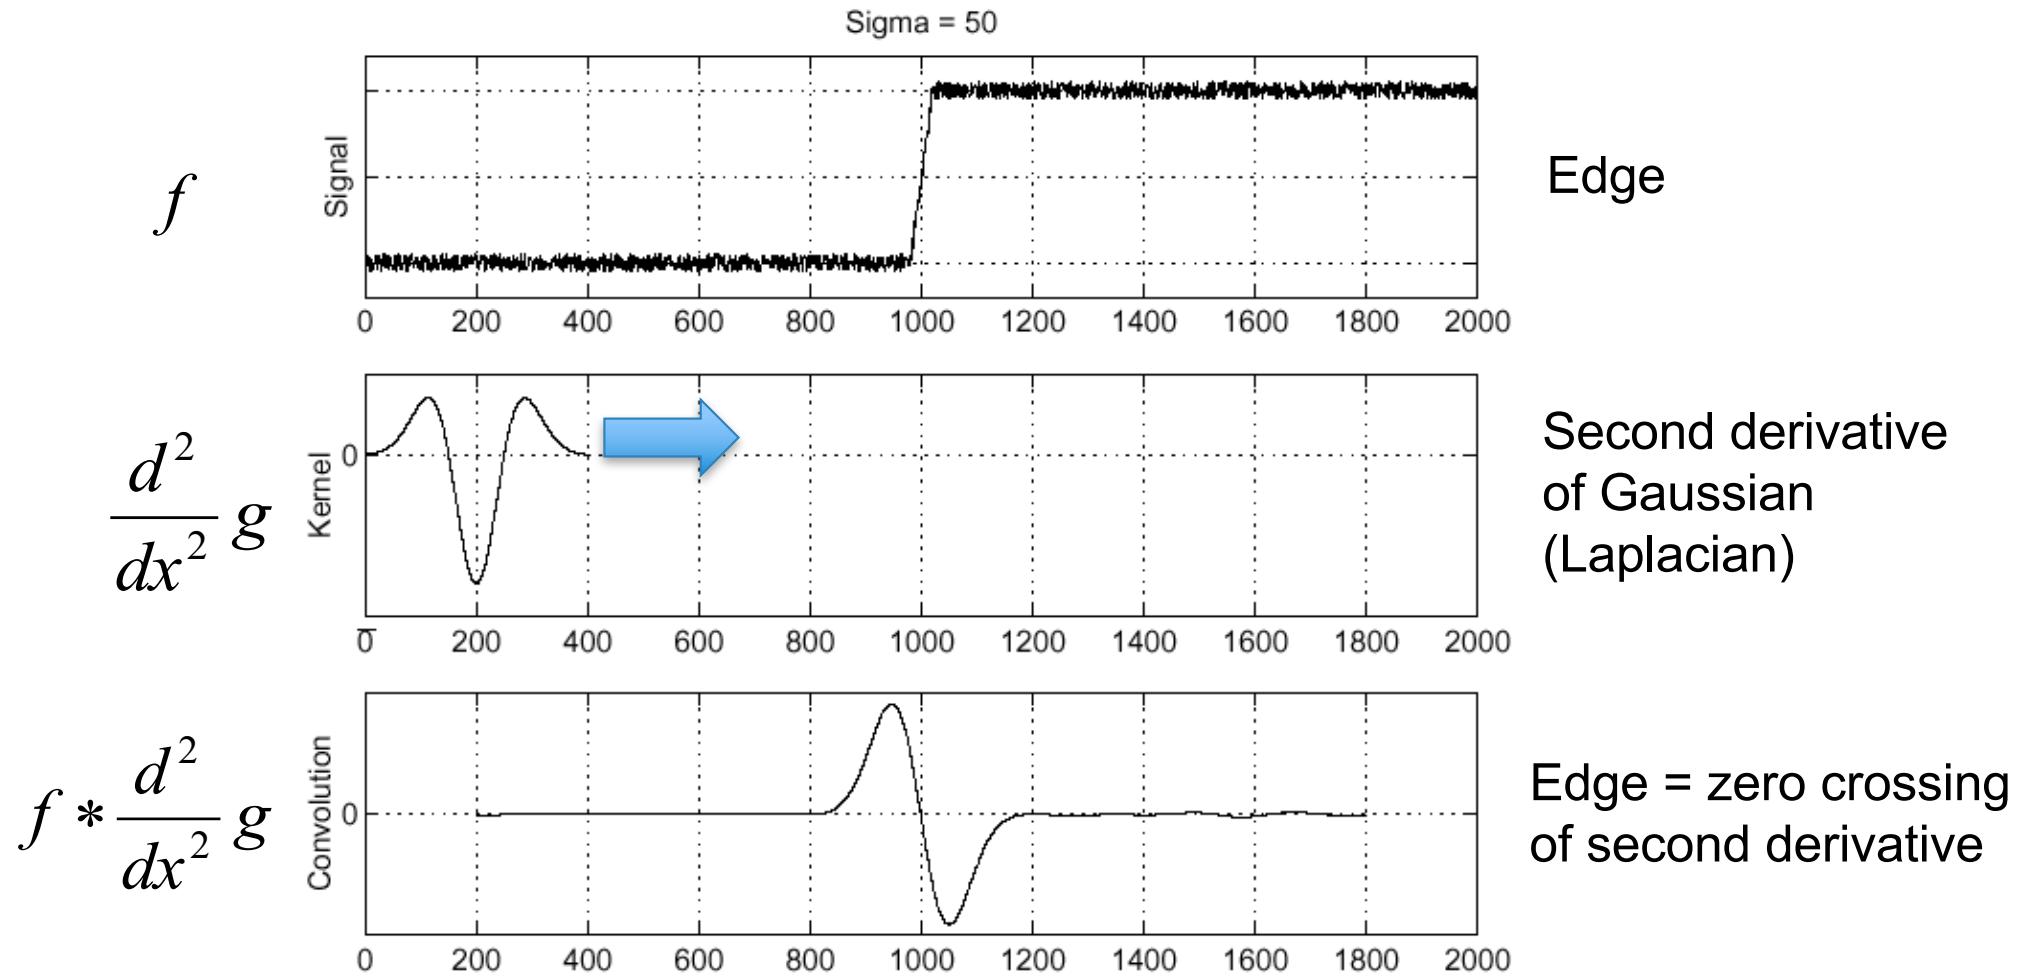
Edge detection, version 1



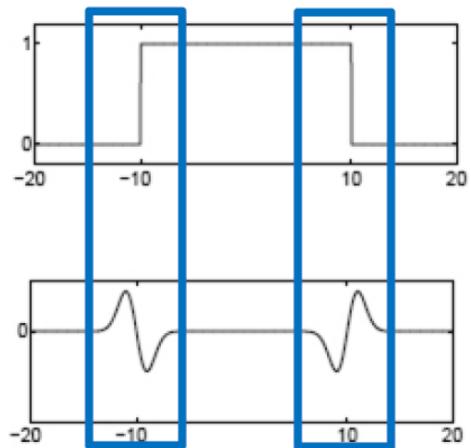
Edge detection, version 1



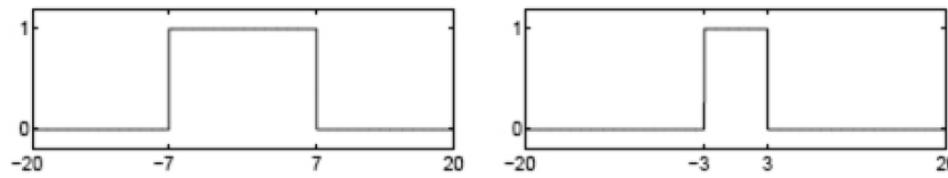
Edge detection, version 2



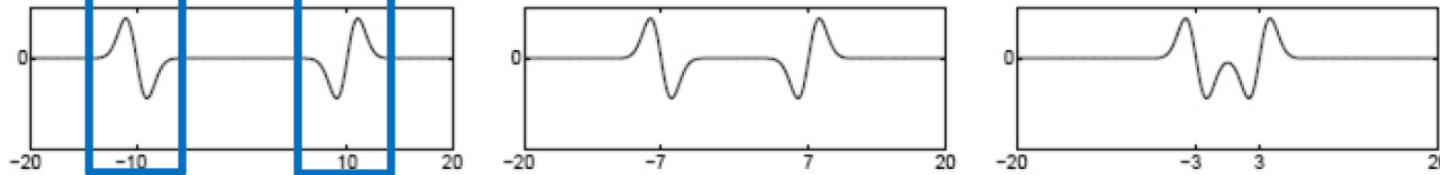
Ripple at edges

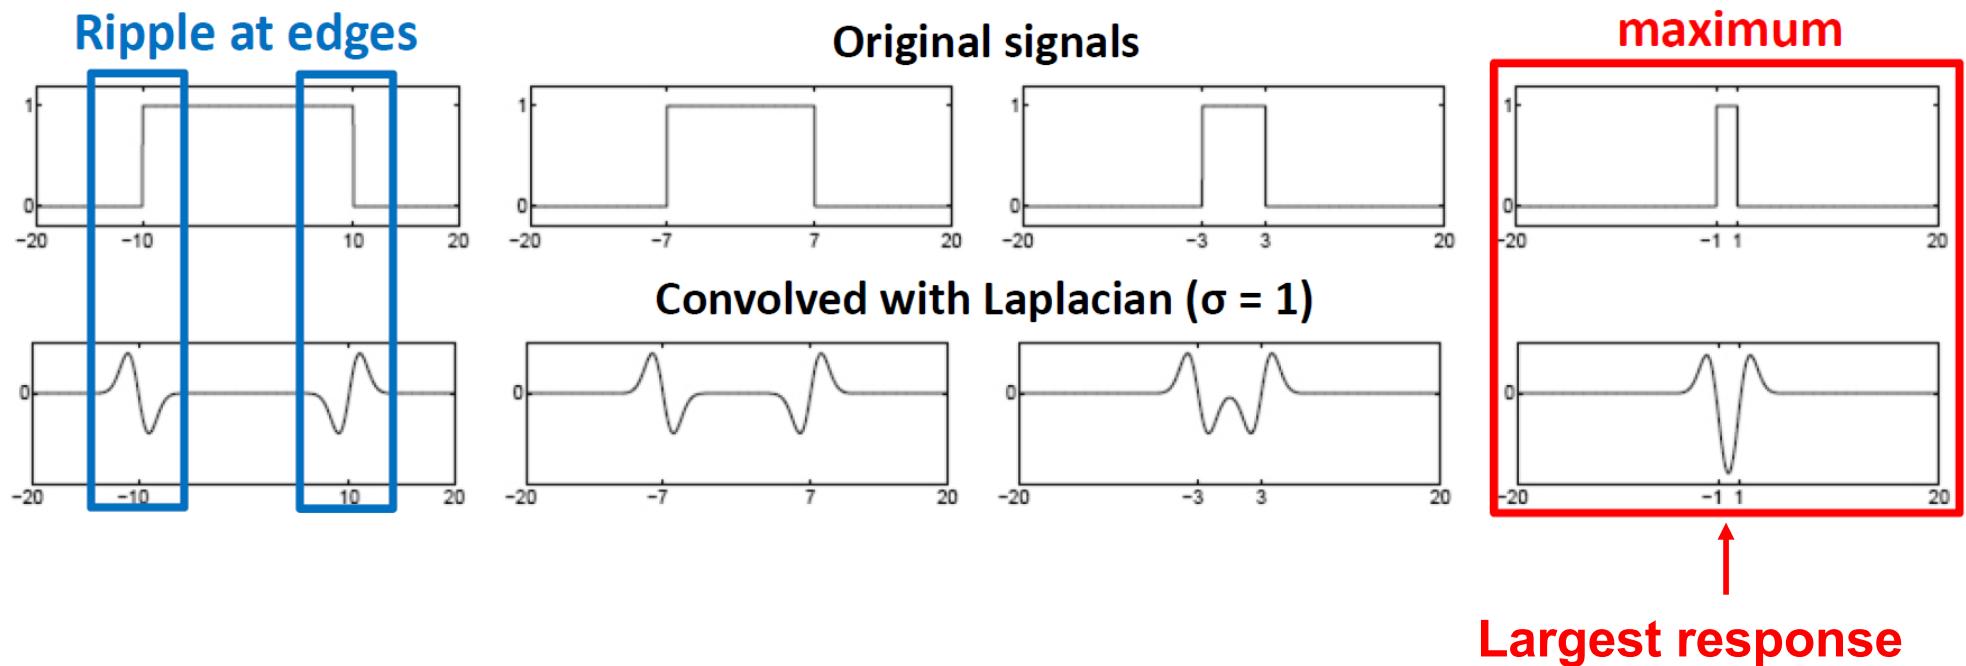


Original signals



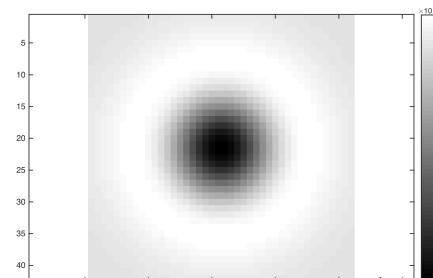
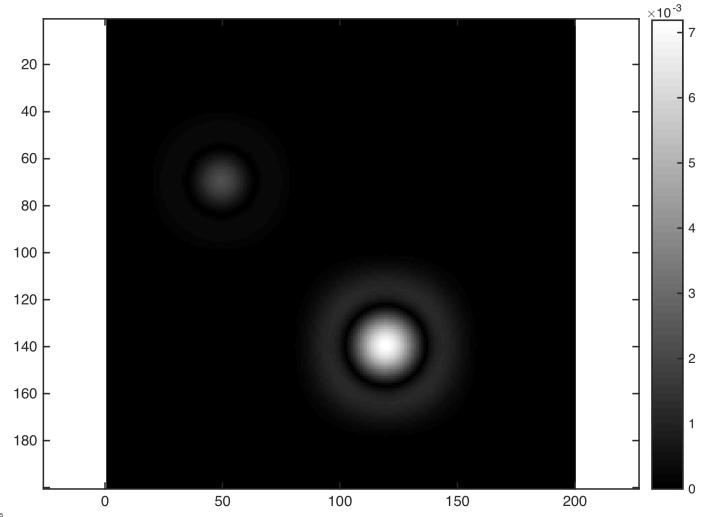
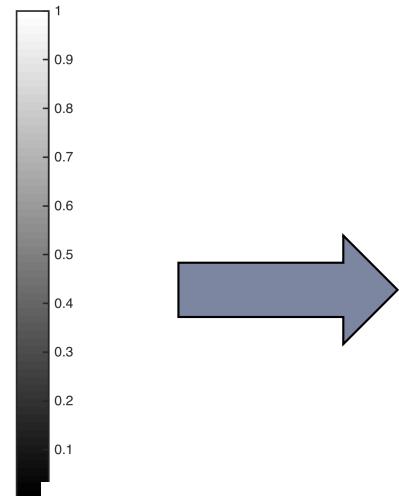
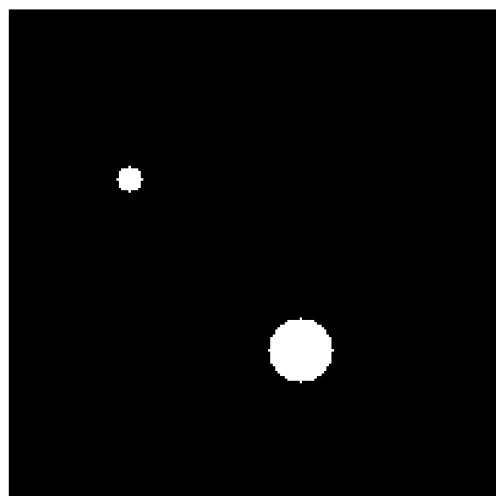
Convolved with Laplacian ($\sigma = 1$)



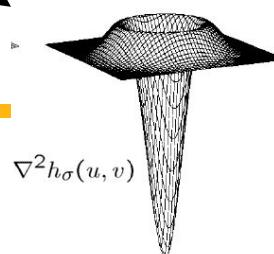


The largest response will be at the center of the blob;
provided the scale of the Laplacian is “matched” to the scale of the blob

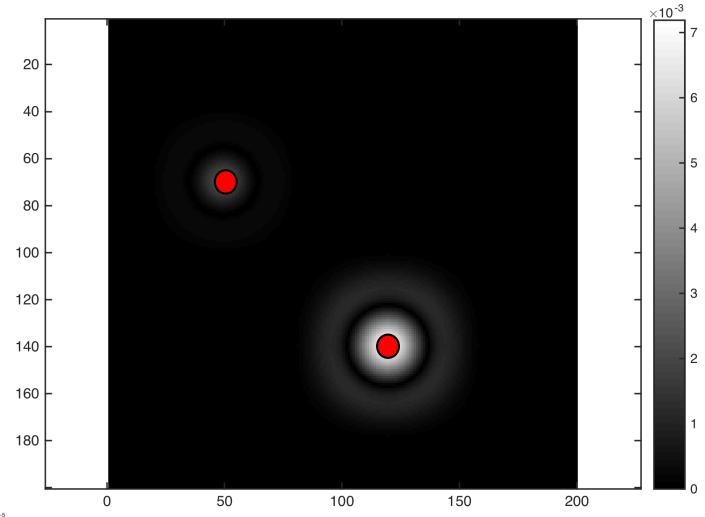
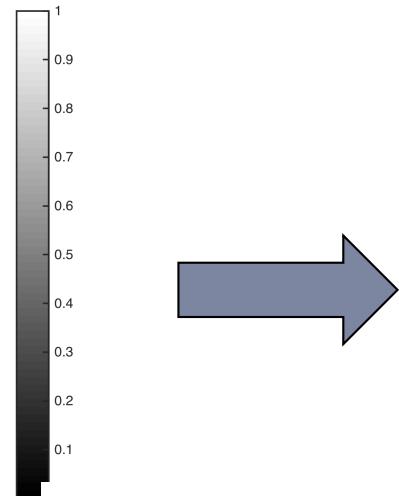
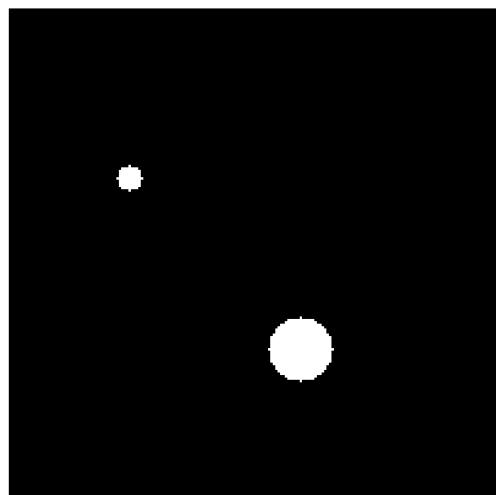
Blob Detector



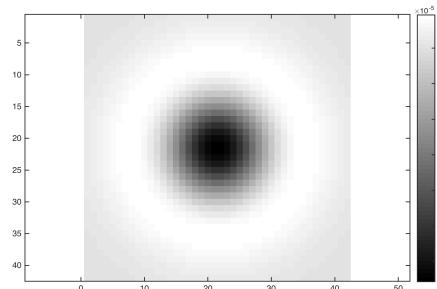
Laplacian of Gaussian



Blob Detector

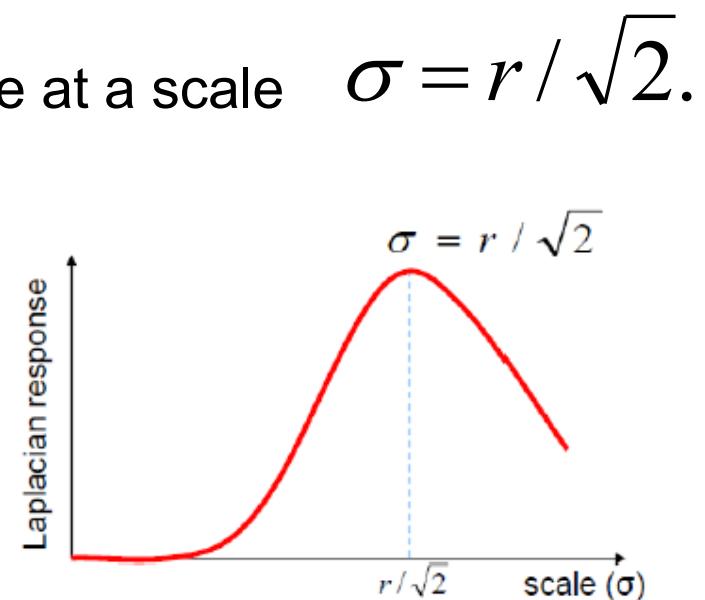
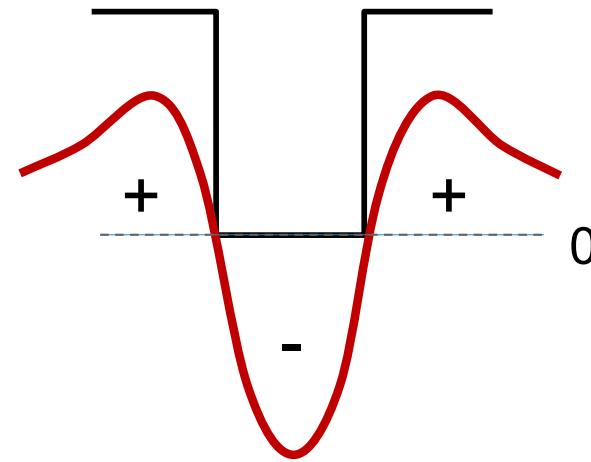
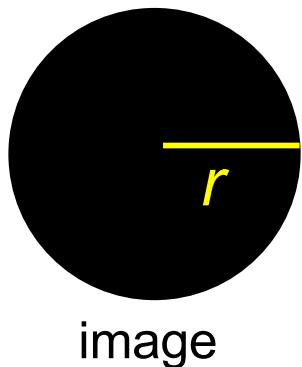


Local maxima



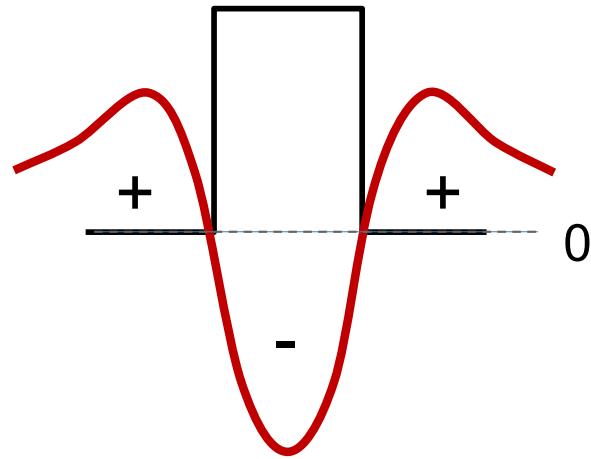
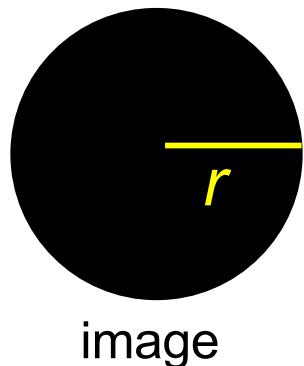
Scale selection

- To get maximum response, the zeros of the Laplacian have to be aligned with the circle.
- The Laplacian achieves a maximum response at a scale $\sigma = r / \sqrt{2}$.



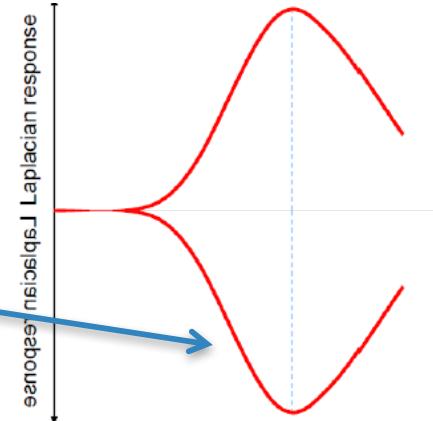
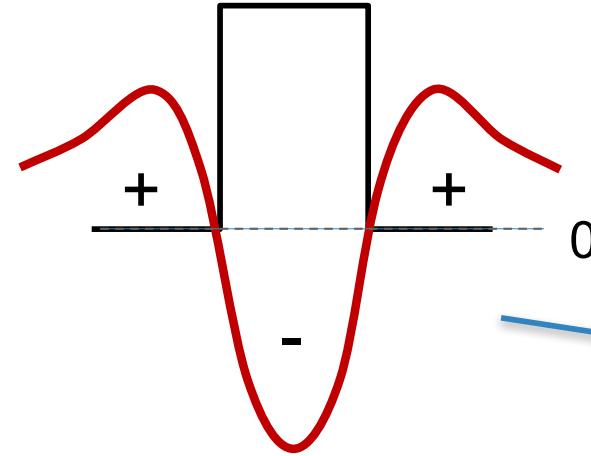
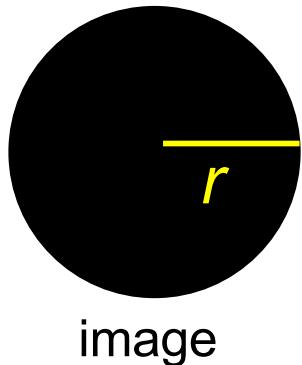
Scale selection

- To get maximum response, the zeros of the Laplacian have to be aligned with the circle.
- The Laplacian achieves a maximum response at a scale $\sigma = r / \sqrt{2}$.



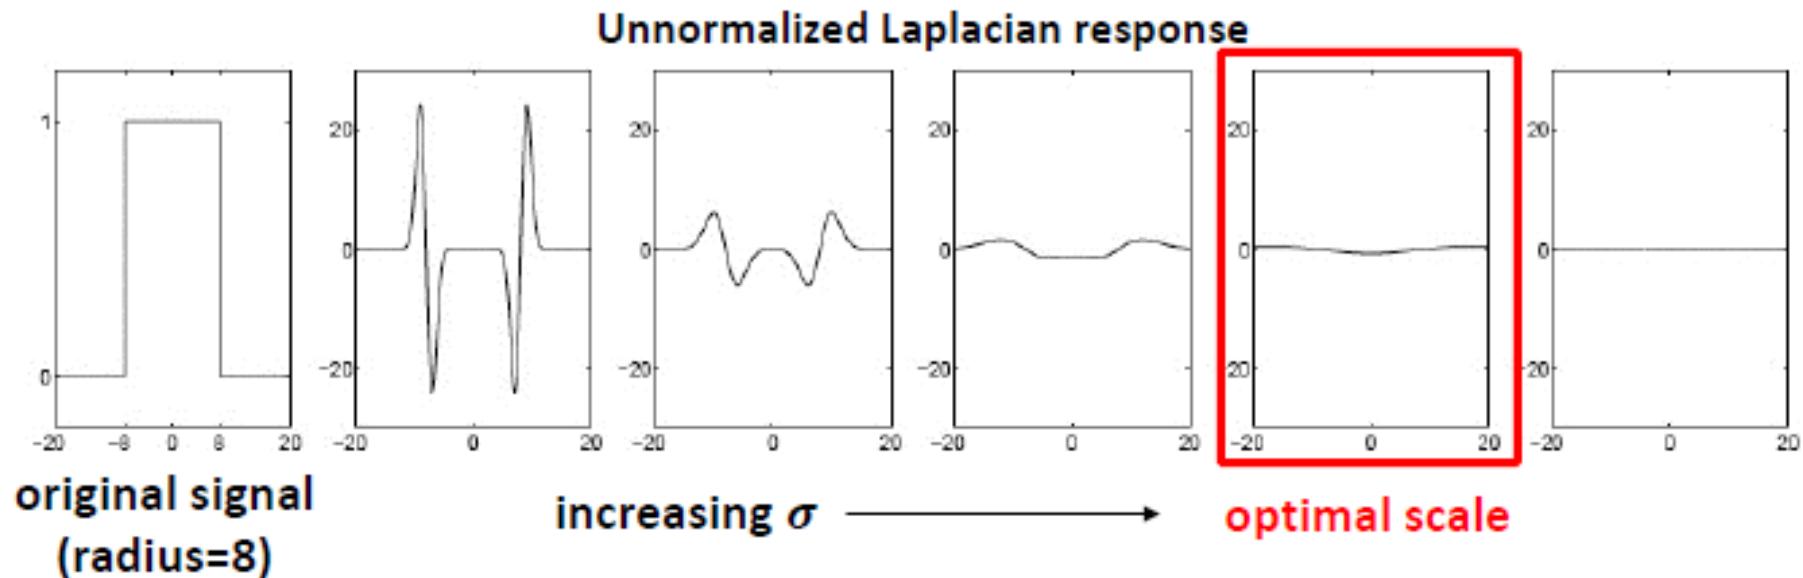
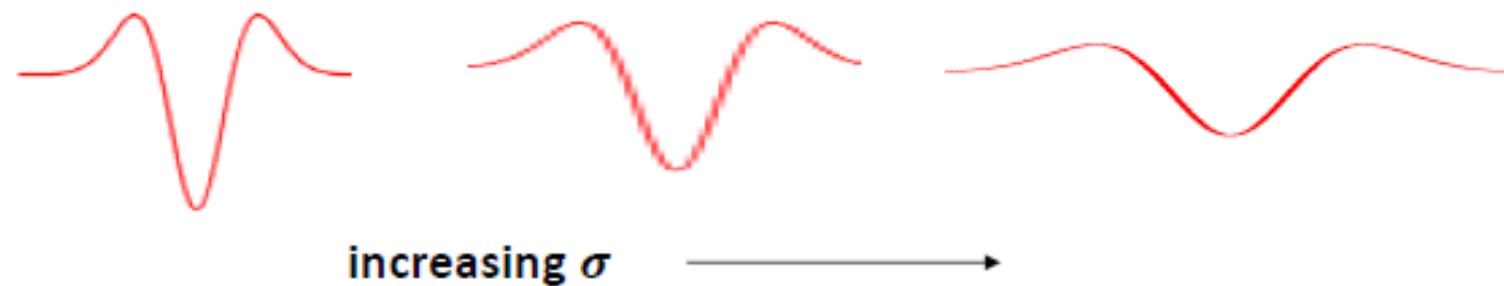
Scale selection

- To get maximum response, the zeros of the Laplacian have to be aligned with the circle.
- The Laplacian achieves a maximum response at a scale $\sigma = r / \sqrt{2}$.



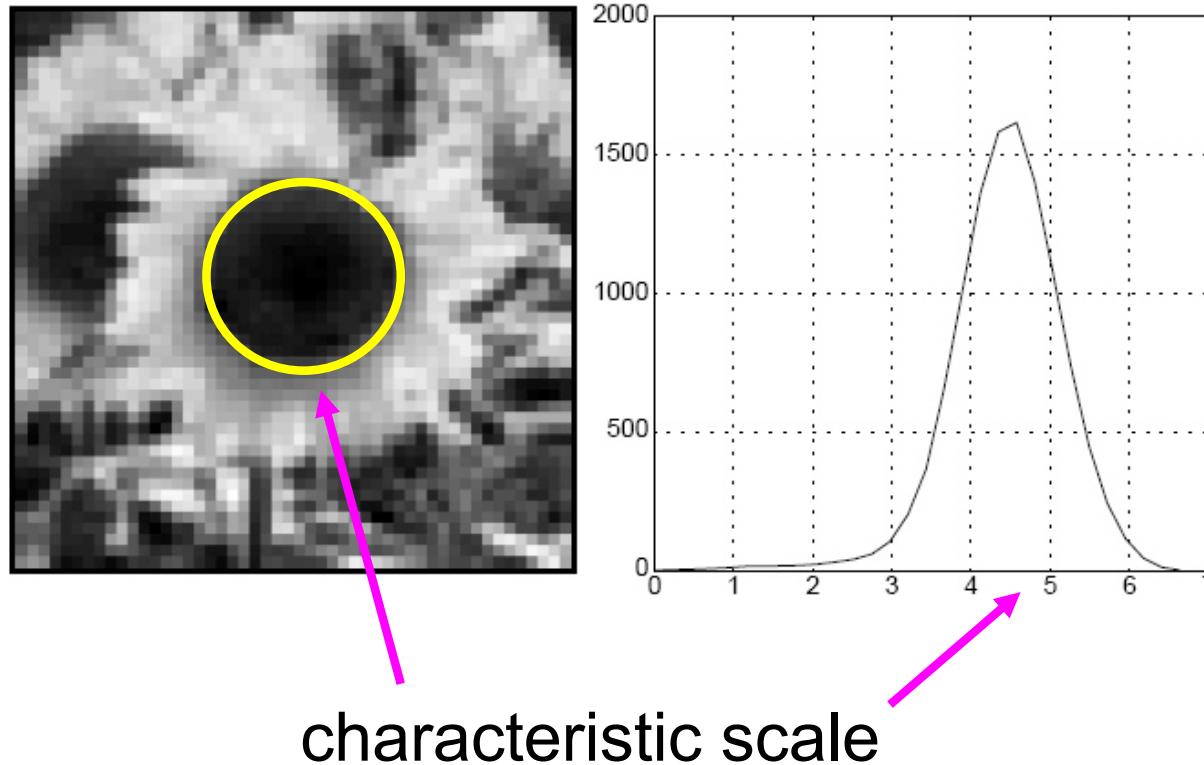
Effect of scale normalization

- However, the magnitude of the Laplacians depends on scale.



Characteristic scale

$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$



Scale-space blob detector: Example



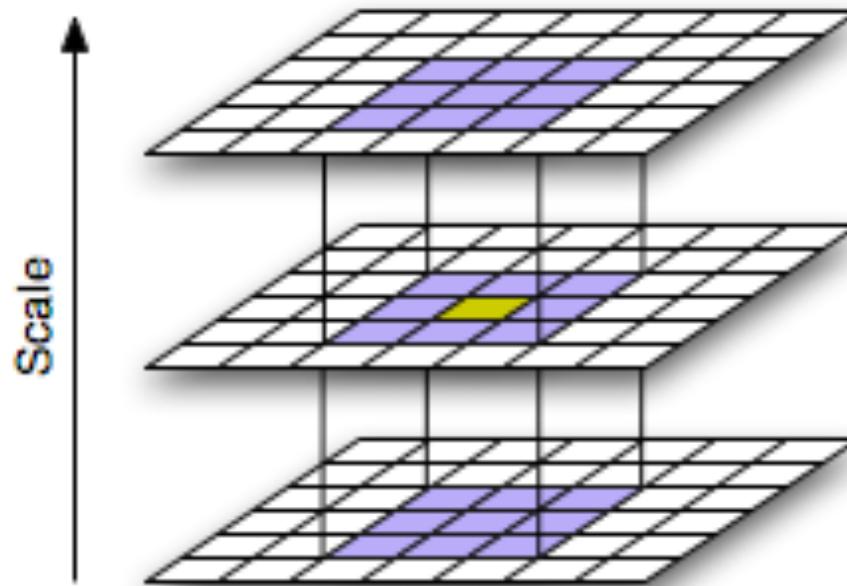
Scale-space blob detector. Example



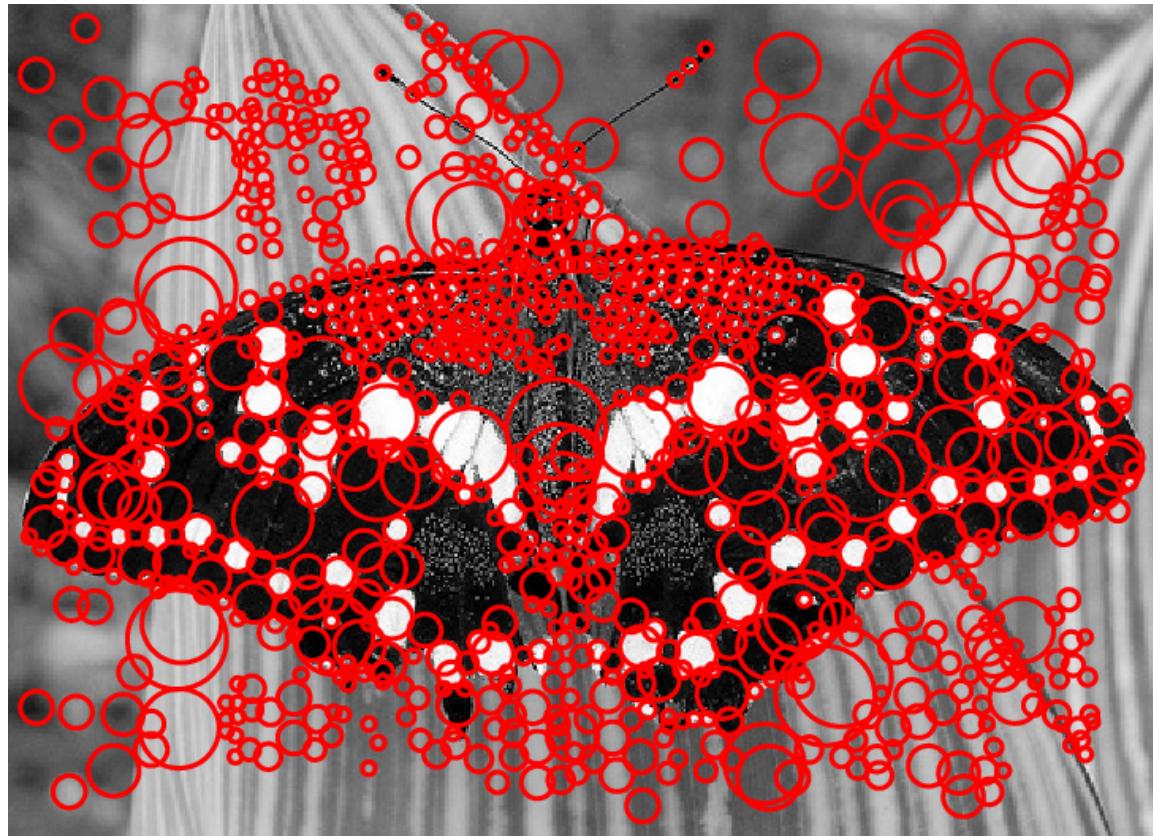
sigma = 11.9912

Scale-space blob detector

1. Convolve image with normalized Laplacian at several scales
2. Find maxima of Laplacian response in scale-space

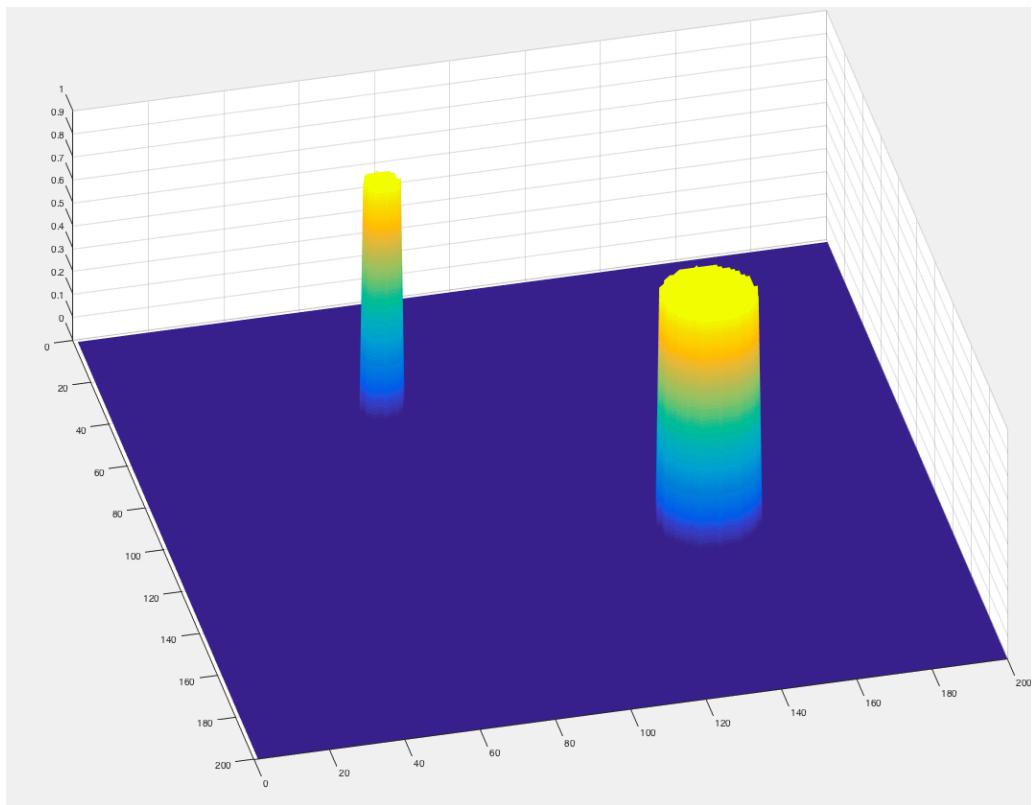
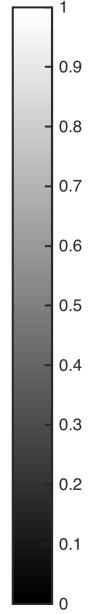
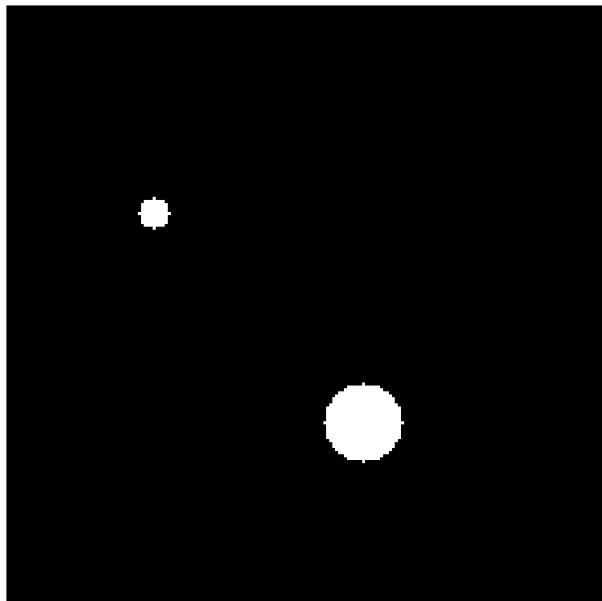


Scale-space blob detector: Example



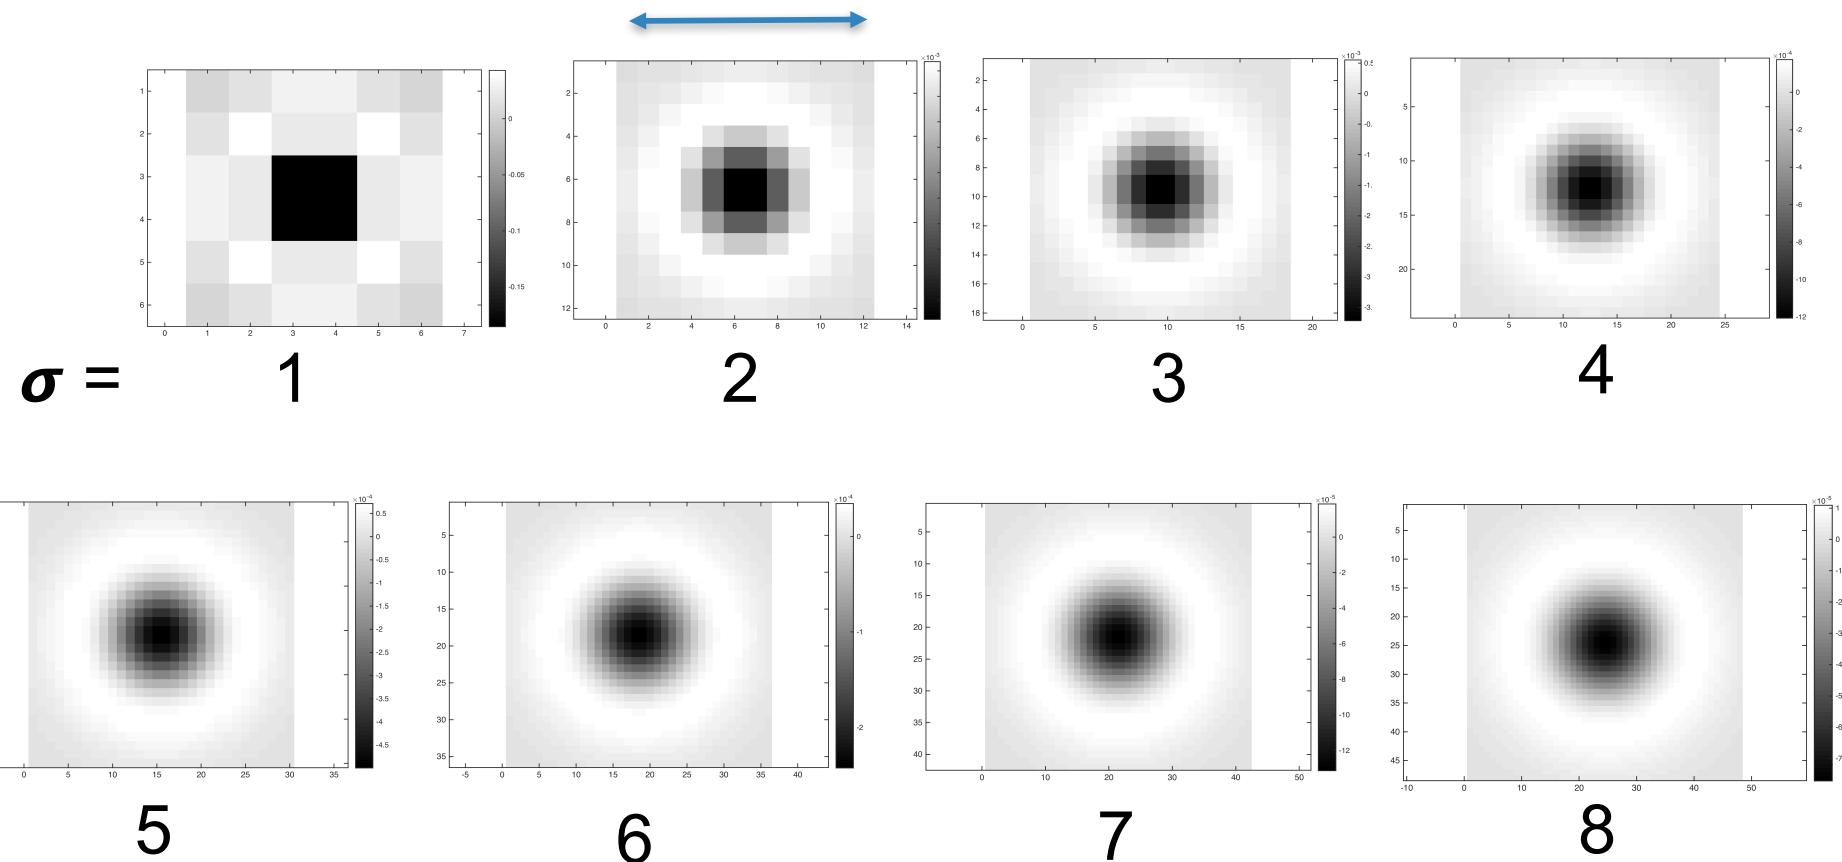
Spatial + Scale Selection

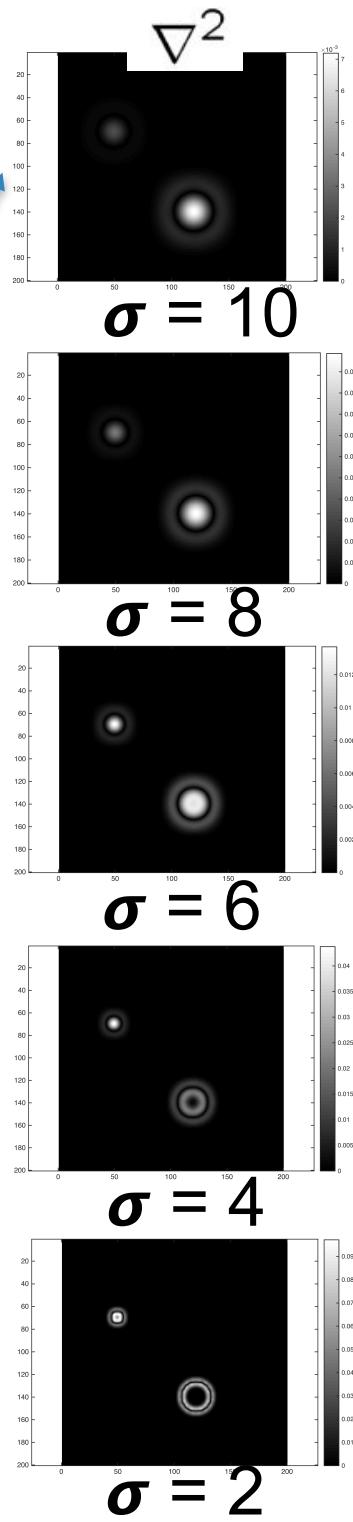
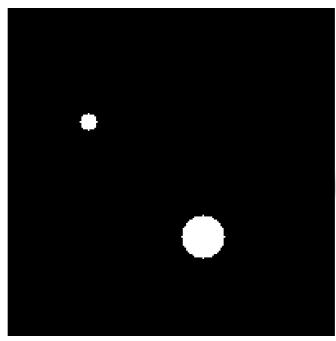
- Purpose: to identify the location and scale of objects

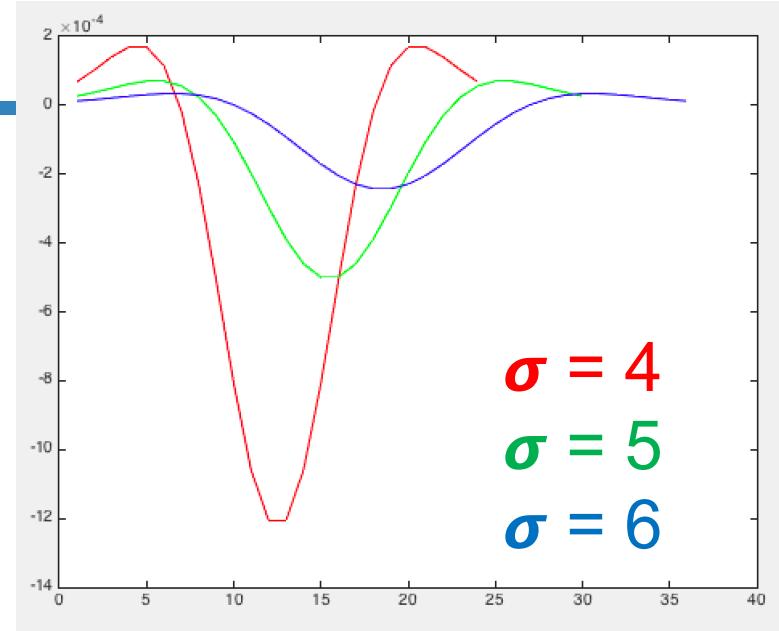
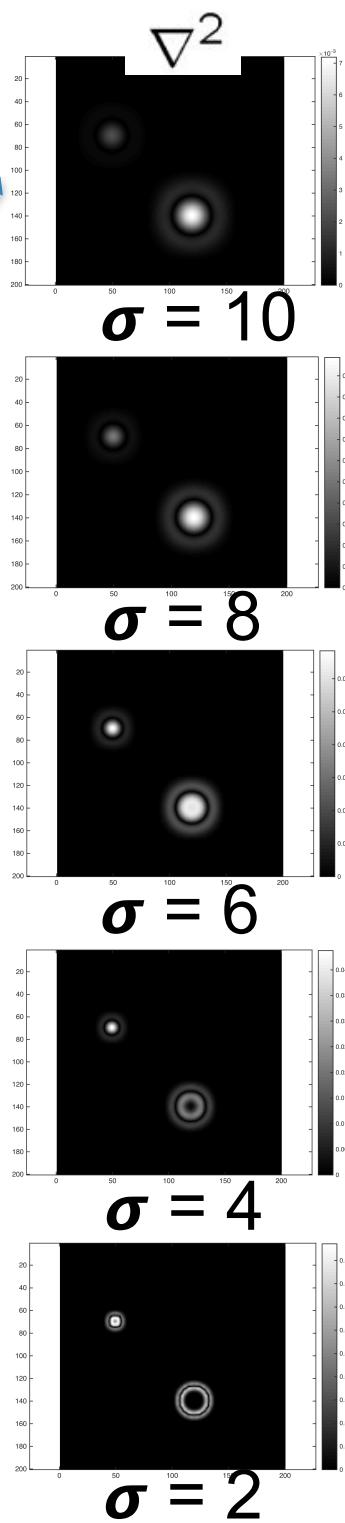
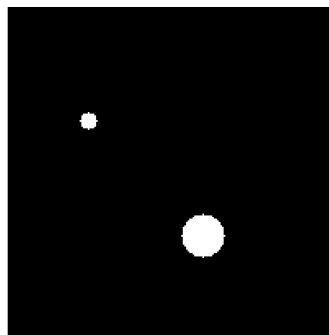


2D Laplacian filter

Filter size = $6 * \sigma$



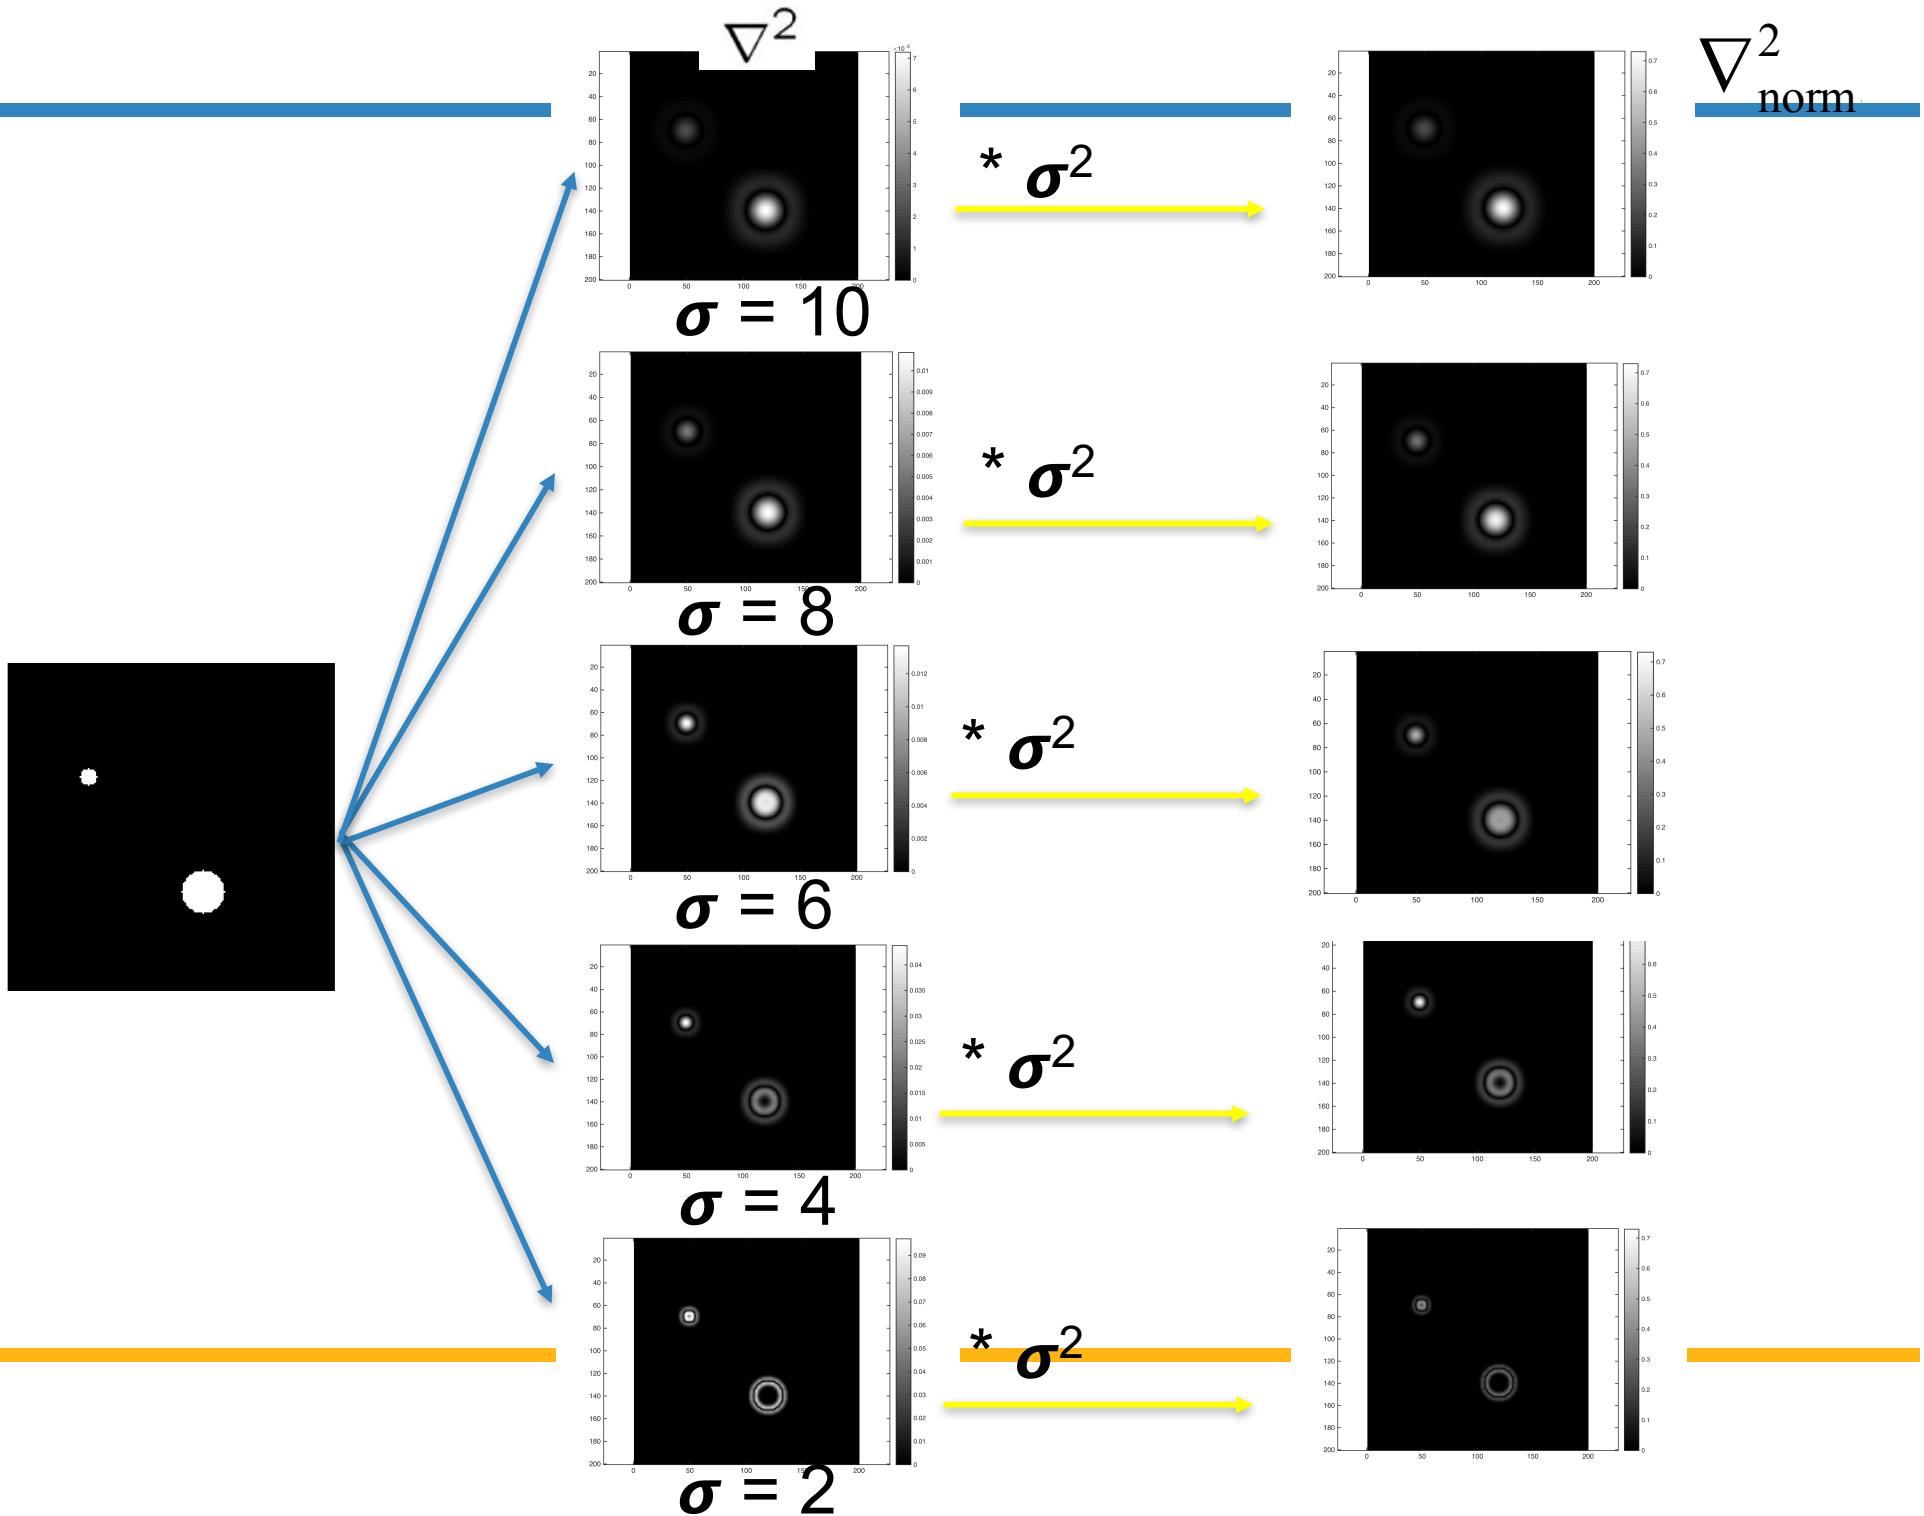


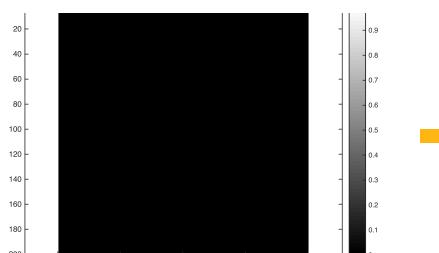
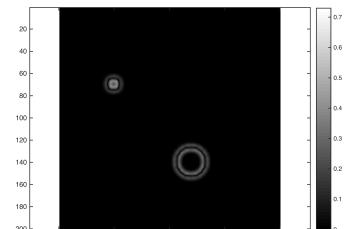
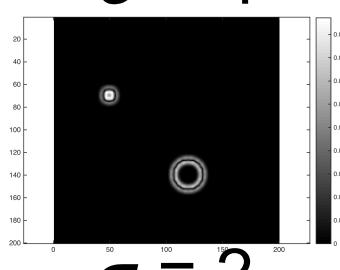
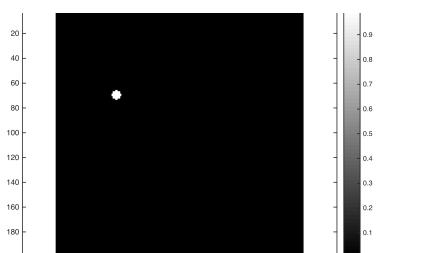
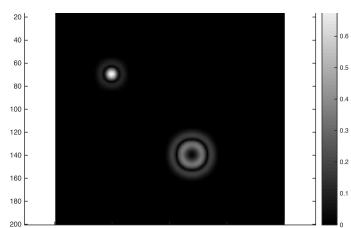
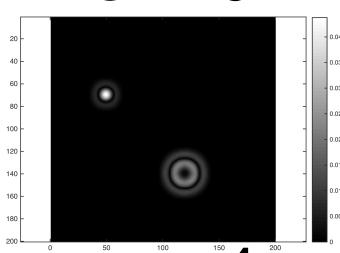
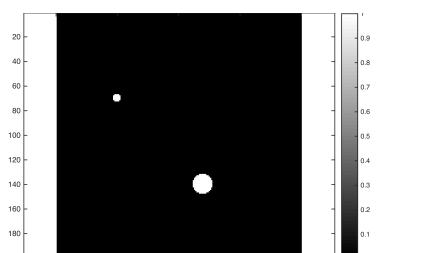
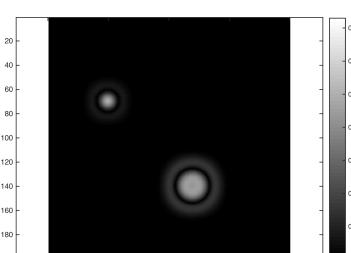
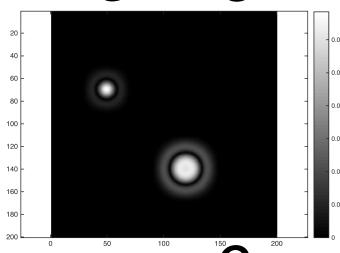
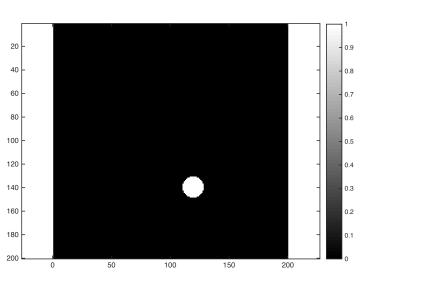
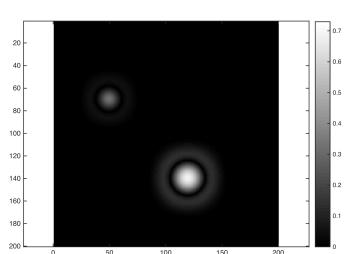
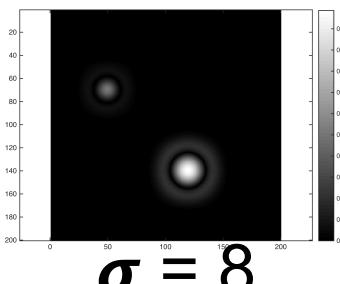
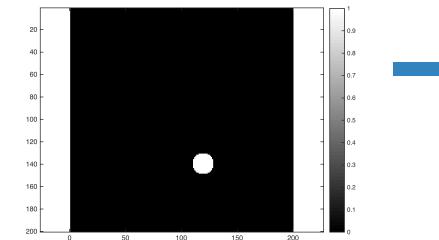
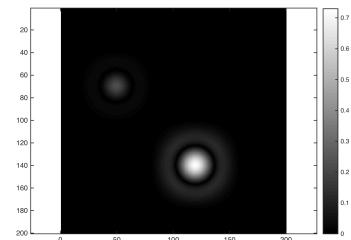
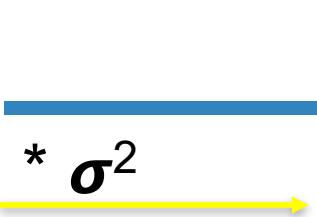
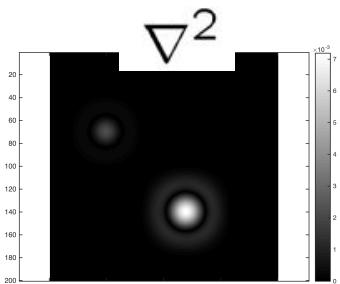


Range is different
(not directly comparable)

$$\nabla^2_{\text{norm}} g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

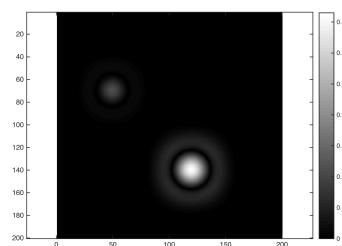
We adjust for the scale



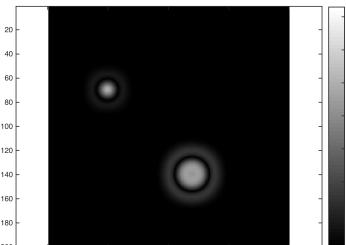
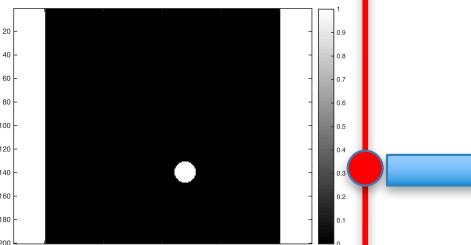
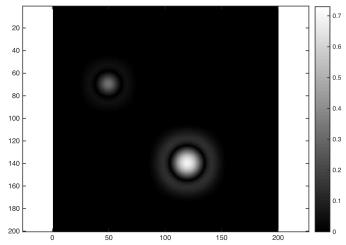
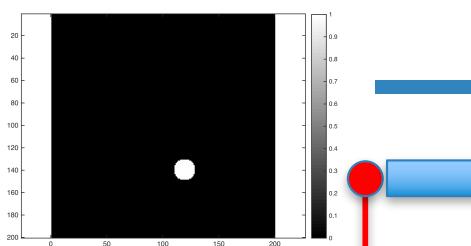


Threshold

> 0.4

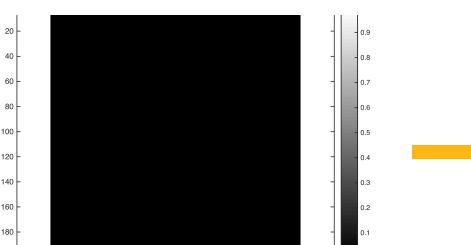
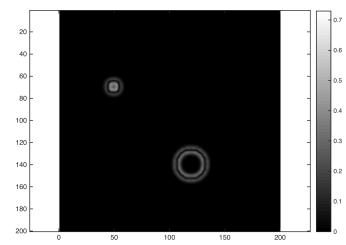
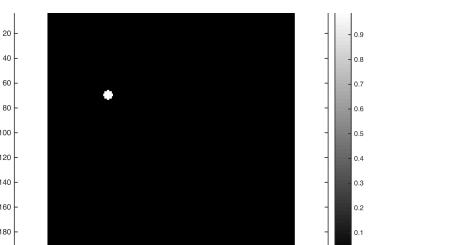
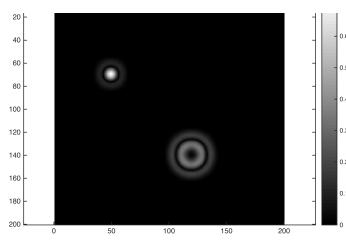


∇^2
norm

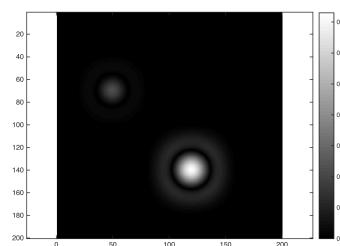


Threshold

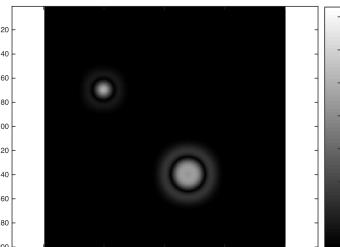
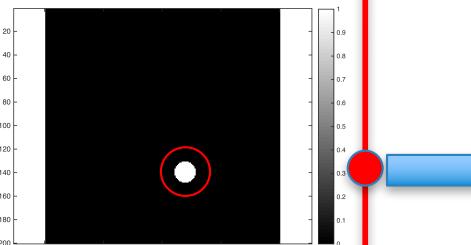
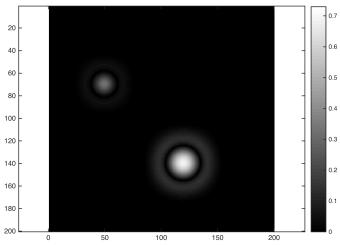
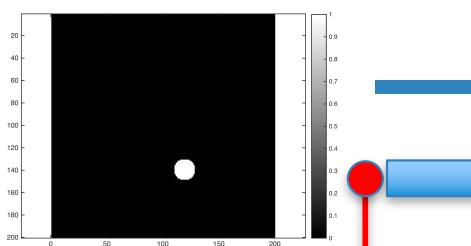
> 0.4



Local maxima #1
scale = 8
radius = $8 * \sqrt{2}$
= 11.3 pixels

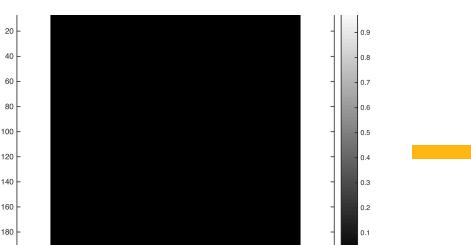
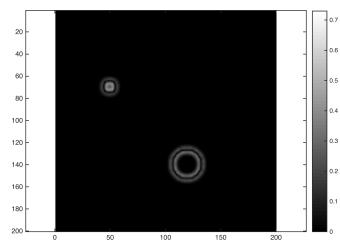
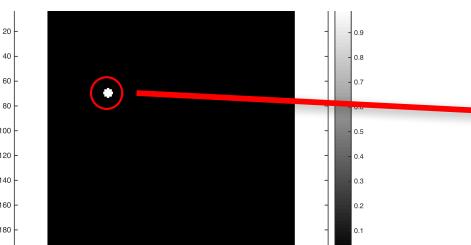
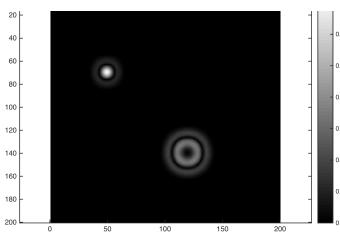


∇^2
norm



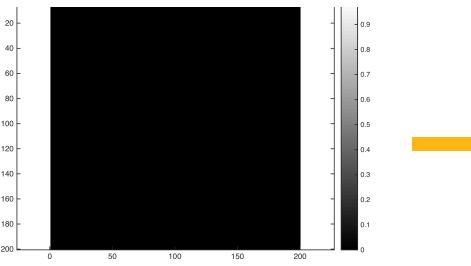
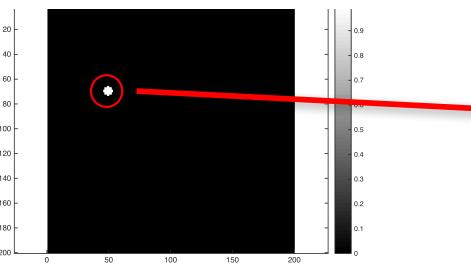
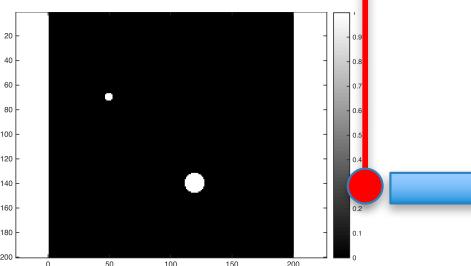
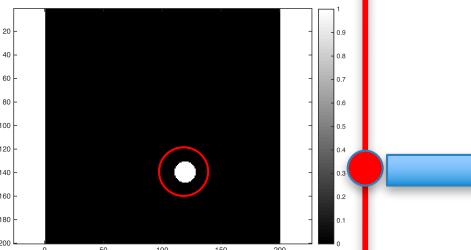
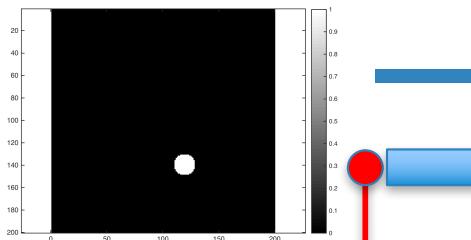
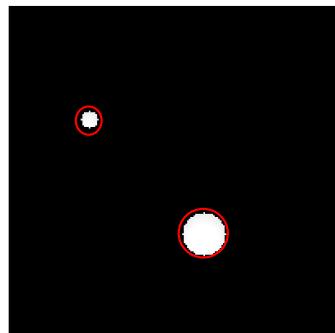
Threshold

> 0.4



Local maxima #1
scale = 8
radius = $8 * \sqrt{2}$
= 11.3 pixels

Local maxima #2
scale = 4
radius = $4 * \sqrt{2}$
= 5.6 pixels



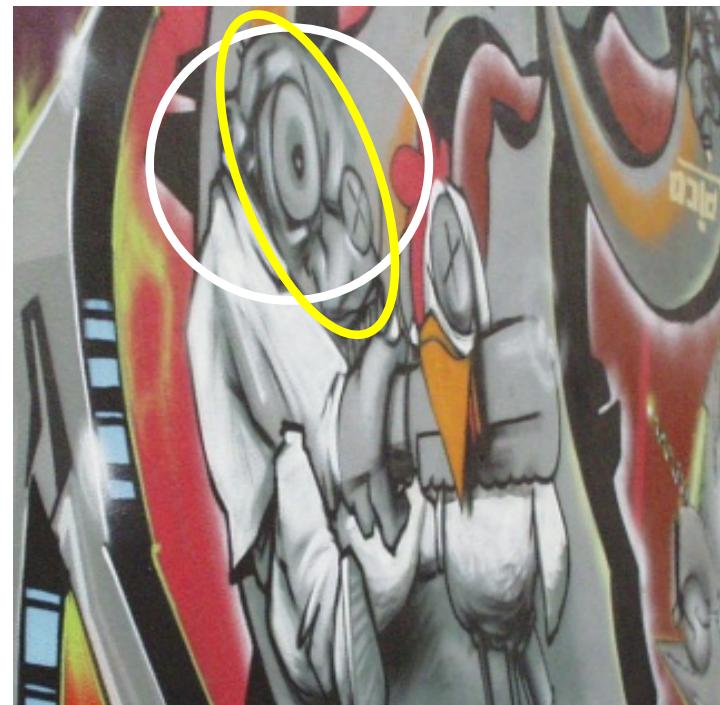
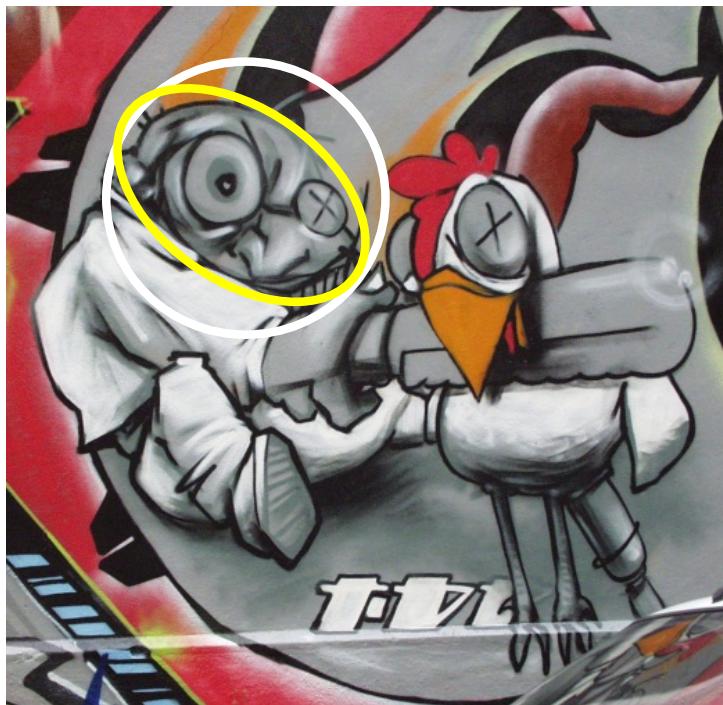
scale = 4
radius = $4 * \sqrt{2}$
= 5.6 pixels

Not always circular



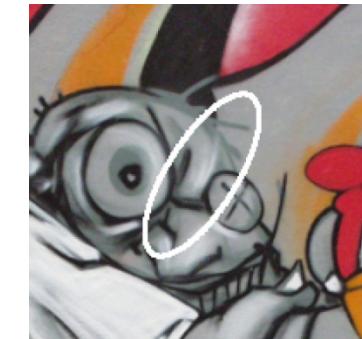
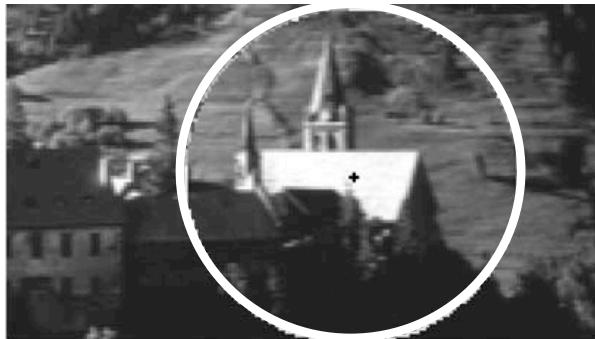
Affine covariance

- Affine transformation approximates viewpoint changes
(for planar objects)



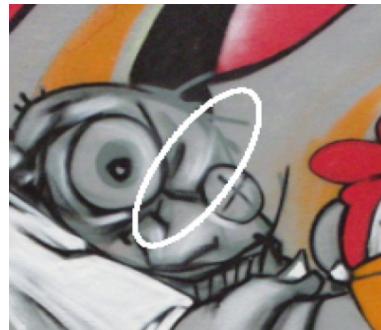
From covariant detection to invariant description

- How to compare the appearance of these image regions?
 - *Normalization*: transform these regions into same-size circles



Affine normalization

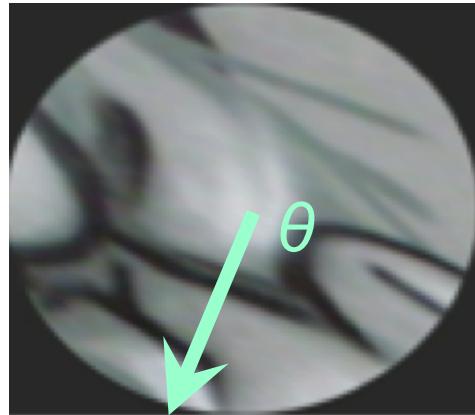
Problem: There is no unique transformation from an ellipse to a unit circle. (rotated or flipped versions remain a circle)



-
- 5-min Break
-

Eliminating rotation ambiguity

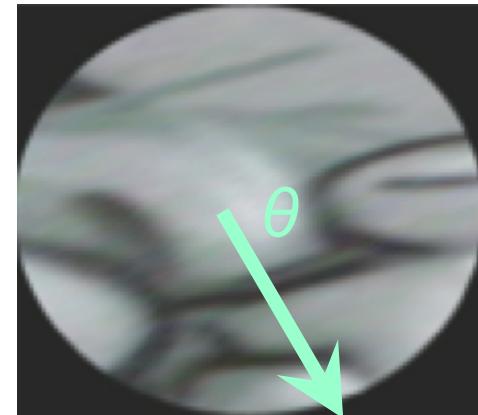
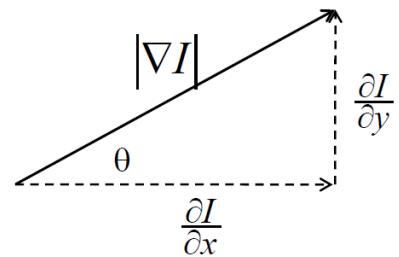
Assign unique orientation to image patch using gradient orientation



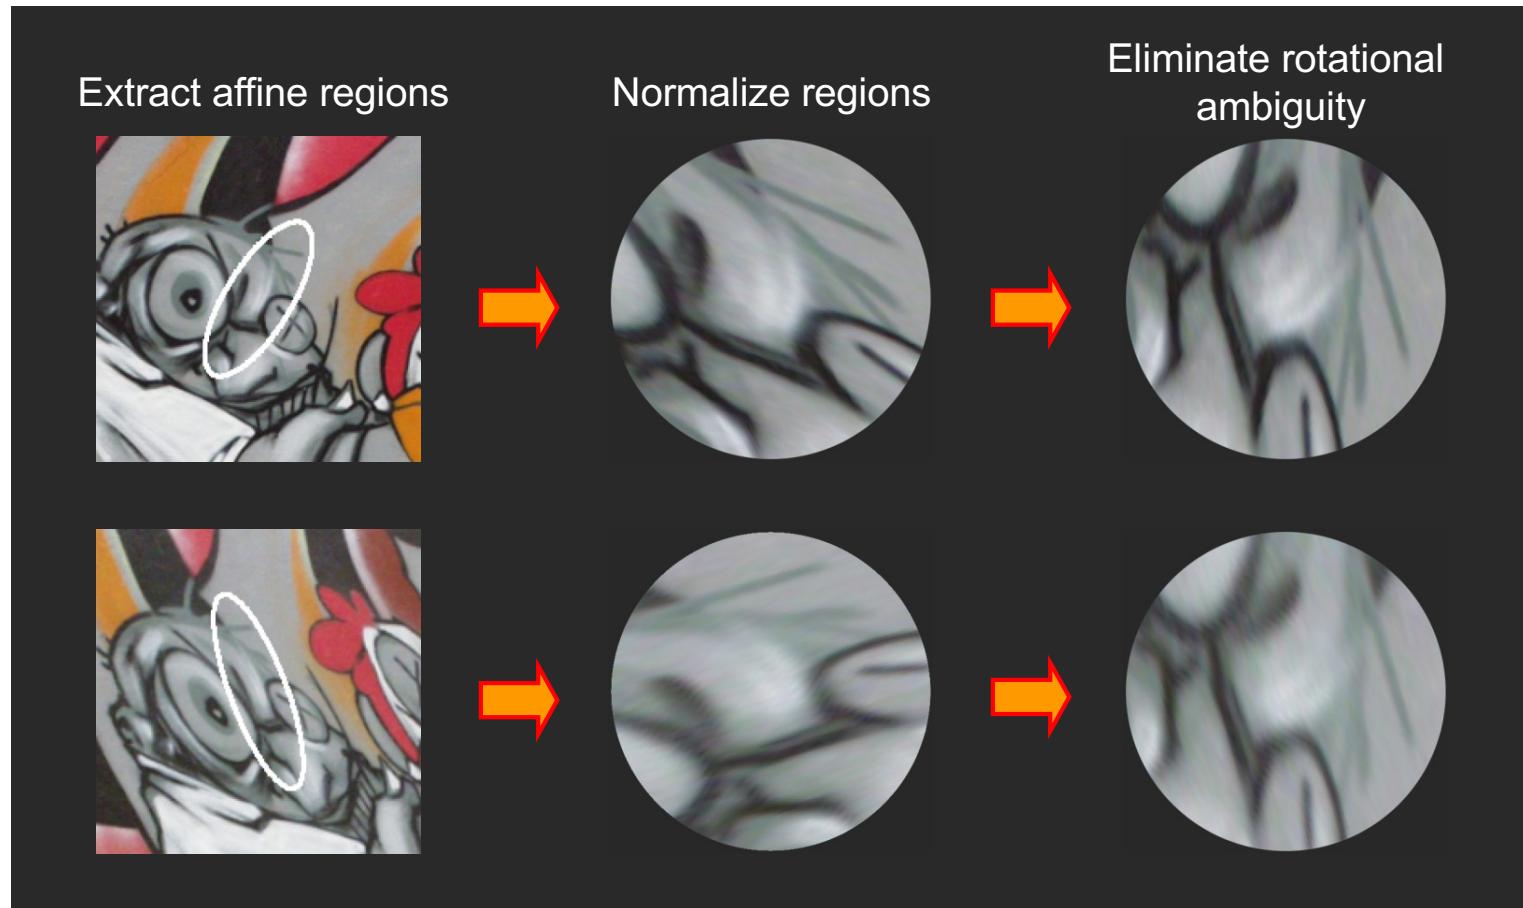
$$\nabla I = \begin{bmatrix} \frac{\partial I}{\partial x} \\ \frac{\partial I}{\partial y} \end{bmatrix}$$

$$|\nabla I| = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2} \approx \left|\frac{\partial I}{\partial x}\right| + \left|\frac{\partial I}{\partial y}\right|$$

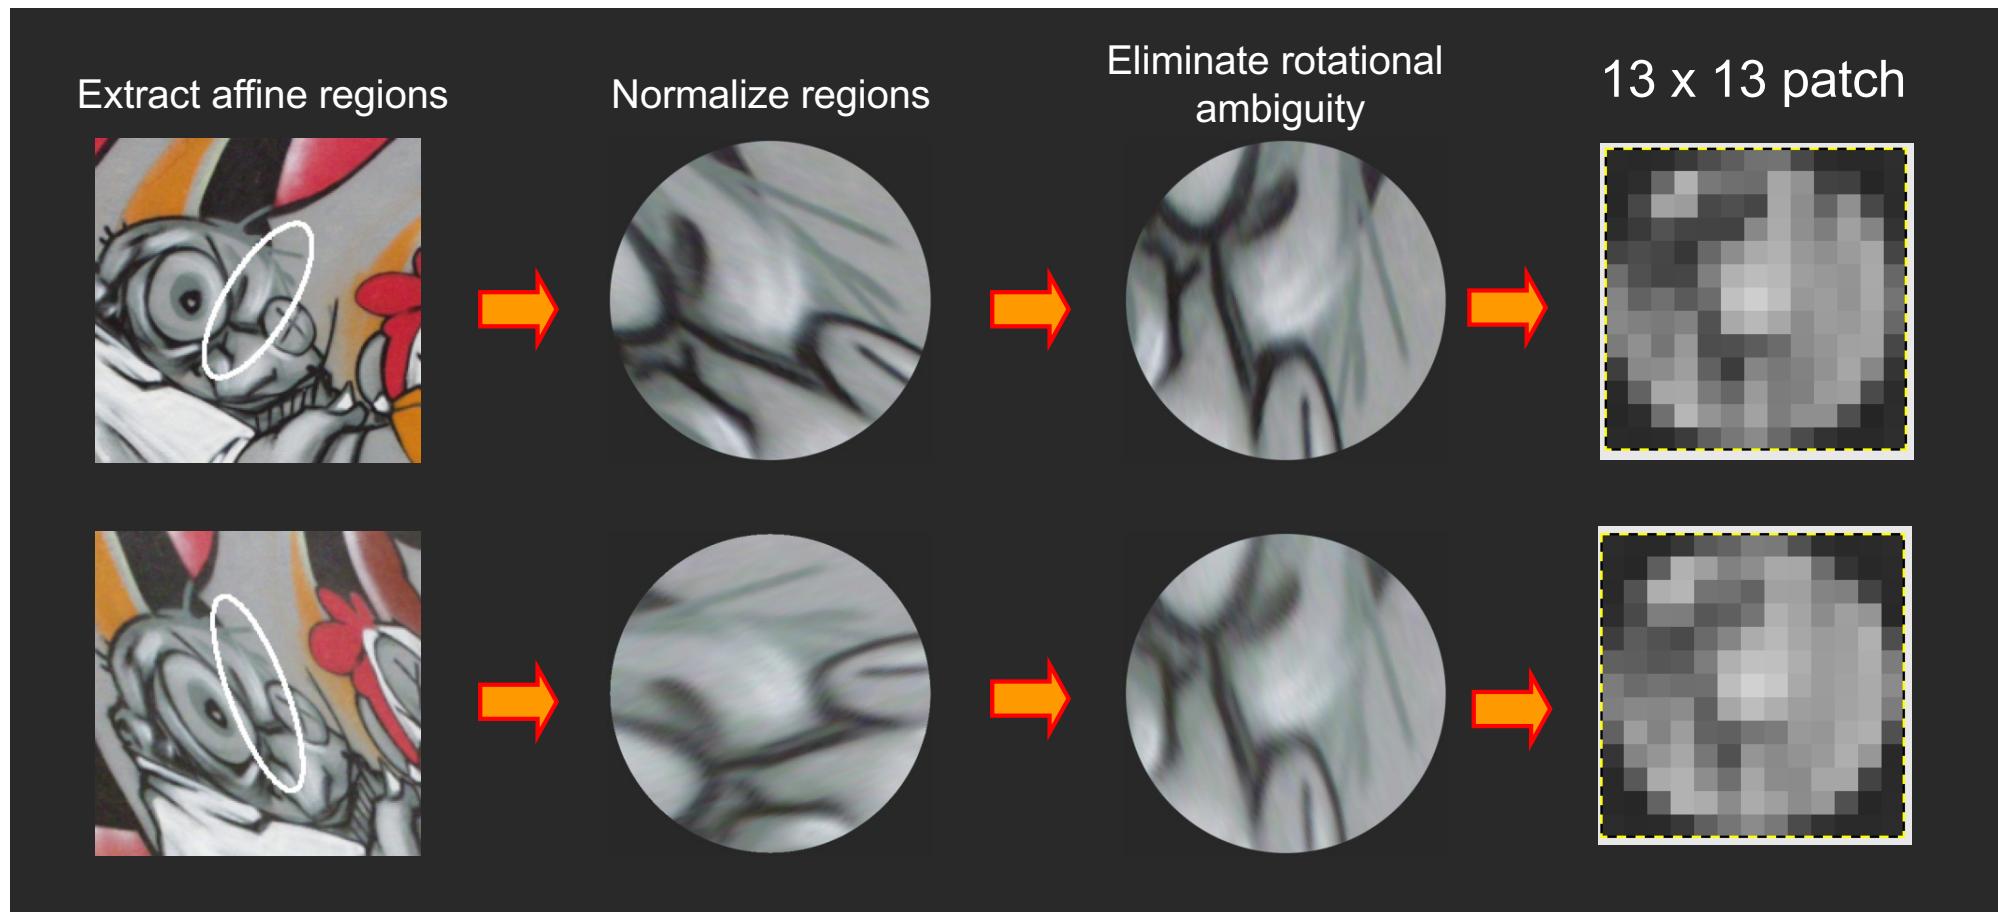
$$\theta = \tan^{-1} \frac{\left(\frac{\partial I}{\partial y}\right)}{\left(\frac{\partial I}{\partial x}\right)}$$



From covariant regions to invariant features

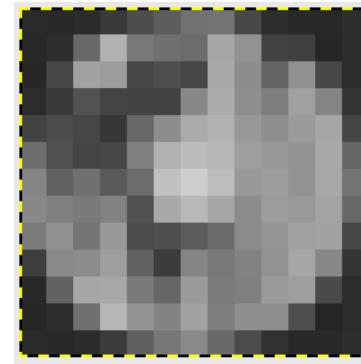


Local Descriptor <option 1>: Raw pixels

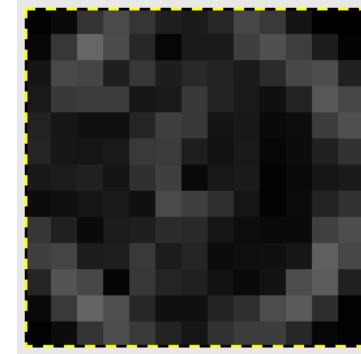
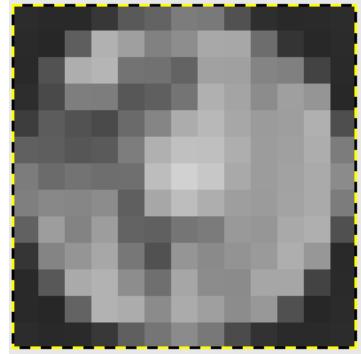
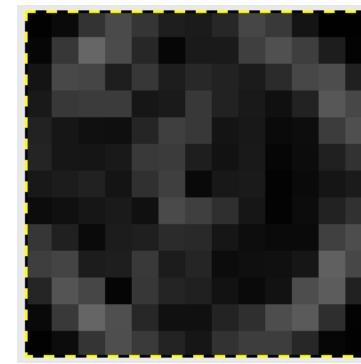


Local Descriptor <option 2>: Derivatives

Eliminate rotational
ambiguity

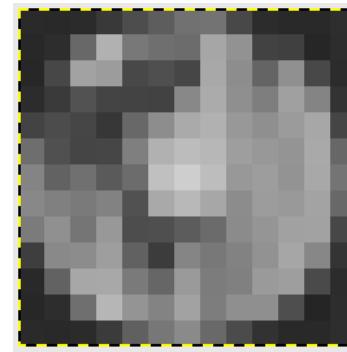
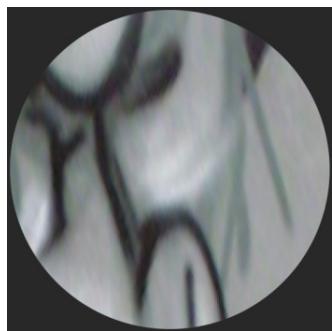


gradient



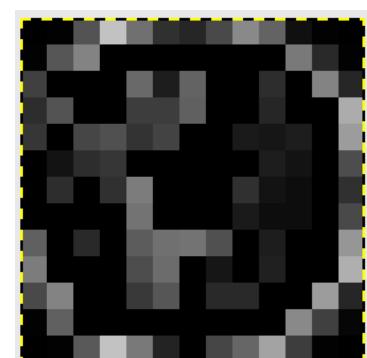
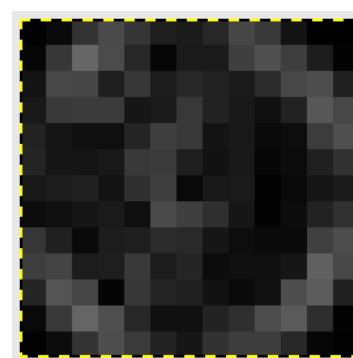
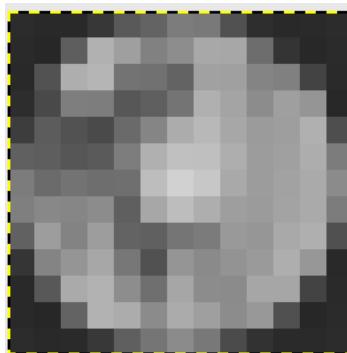
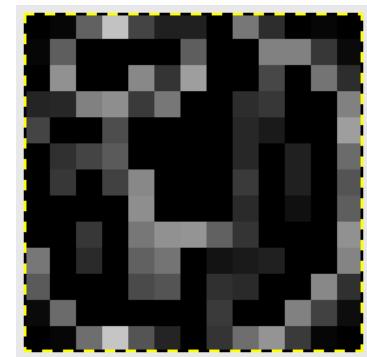
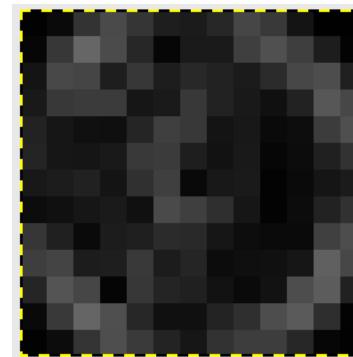
Local Descriptor <option 2>: Derivatives

Eliminate rotational
ambiguity



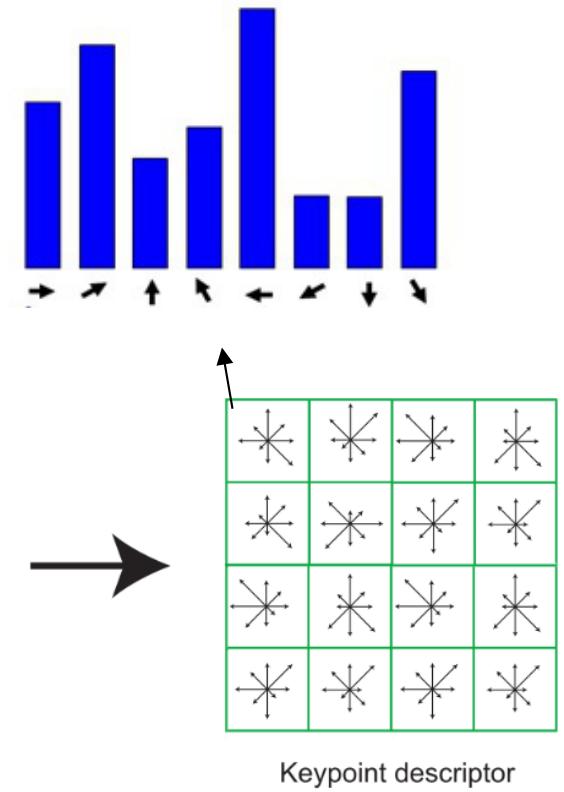
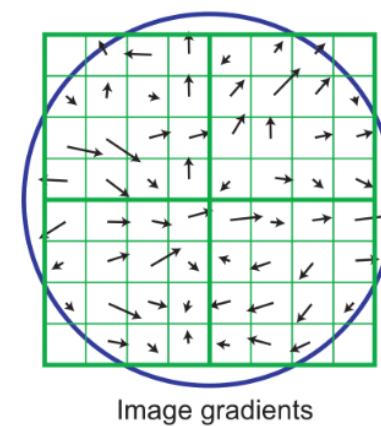
1st
Gradient

2nd
Laplacian



Descriptors: <option 3>: Local Distributions of Derivatives

- SIFT feature descriptor
 - 128 dimension: 4x4 quadrant, 8 bin histogram



From covariant regions to invariant features

SIFT (Lowe '04)

Extract affine regions



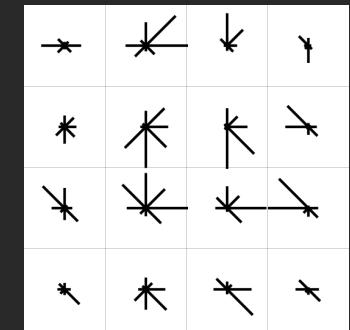
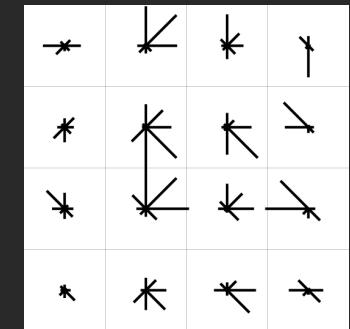
Normalize regions



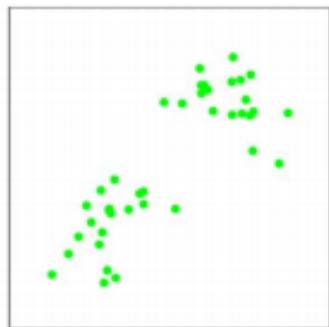
Eliminate rotational ambiguity



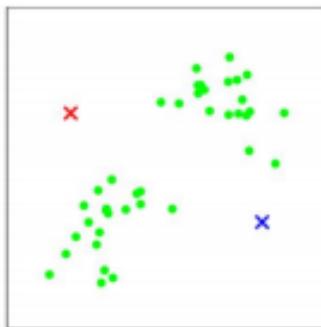
Compute appearance descriptors



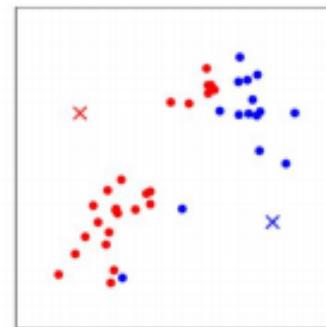
K-means clustering



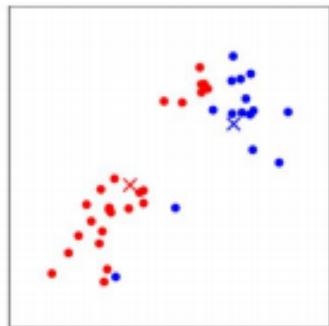
(a)



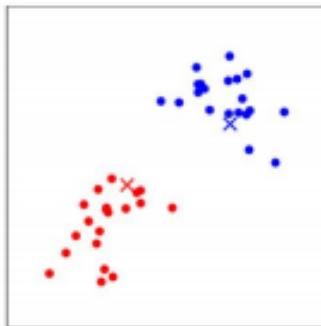
(b)



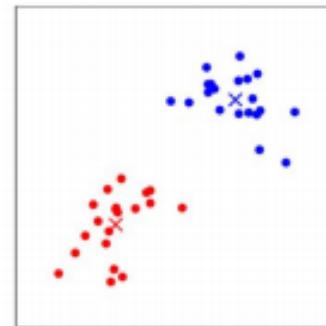
(c)



(d)



(e)



(f)

K-means algorithm.

Training examples are shown as dots, and cluster centroids are shown as crosses. (a) Original dataset. (b) Random initial cluster centroids. (c-f) Illustration of running two iterations of k-means. In each iteration, we assign each training example to the closest cluster centroid; then we move each cluster centroid to the mean of the points assigned to it.

K-means clustering

In the clustering problem, we are given a training set $x^{(1)}, \dots, x^{(m)}$, and want to group the data into a few cohesive "clusters." Here, we are given feature vectors for each data point $x^{(i)} \in \mathbb{R}^n$ as usual; but no labels $y^{(i)}$ (making this an unsupervised learning problem). Our goal is to predict k centroids **and** a label $c^{(i)}$ for each datapoint. The k-means clustering algorithm is as follows:

1. Initialize **cluster centroids** $\mu_1, \mu_2, \dots, \mu_k \in \mathbb{R}^n$ randomly.
2. Repeat until convergence: {

For every i , set

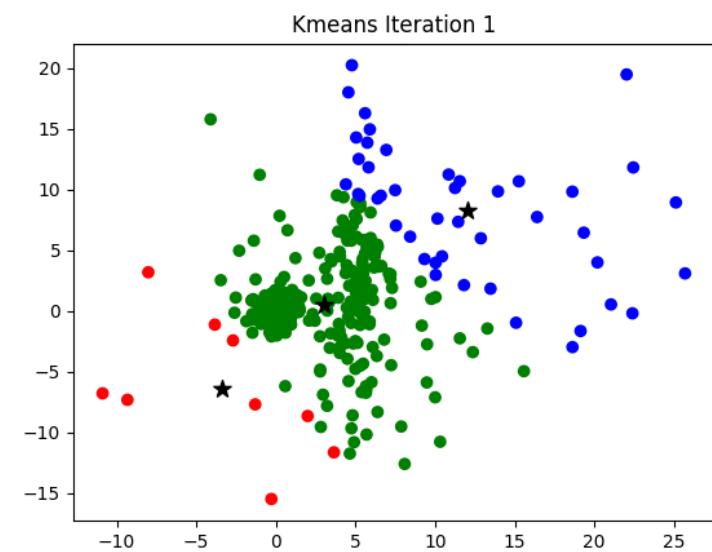
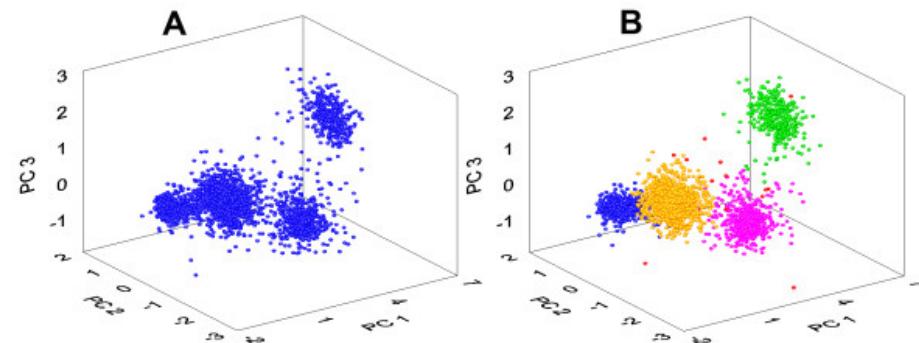
$$c^{(i)} := \arg \min_j \|x^{(i)} - \mu_j\|^2.$$

For each j , set

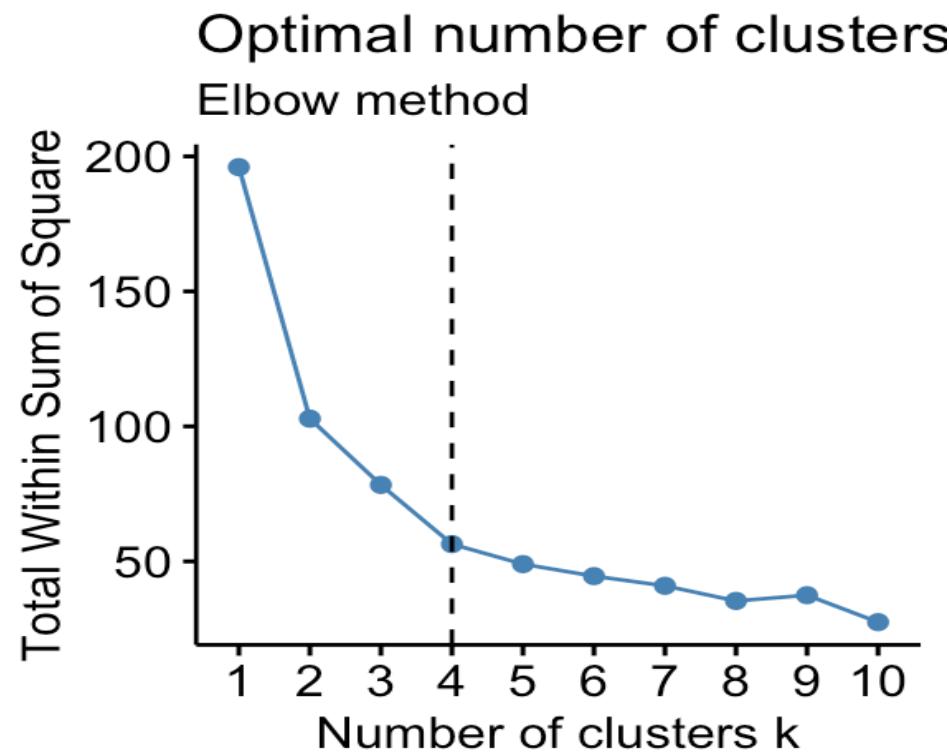
$$\mu_j := \frac{\sum_{i=1}^m \mathbf{1}\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m \mathbf{1}\{c^{(i)} = j\}}.$$

}

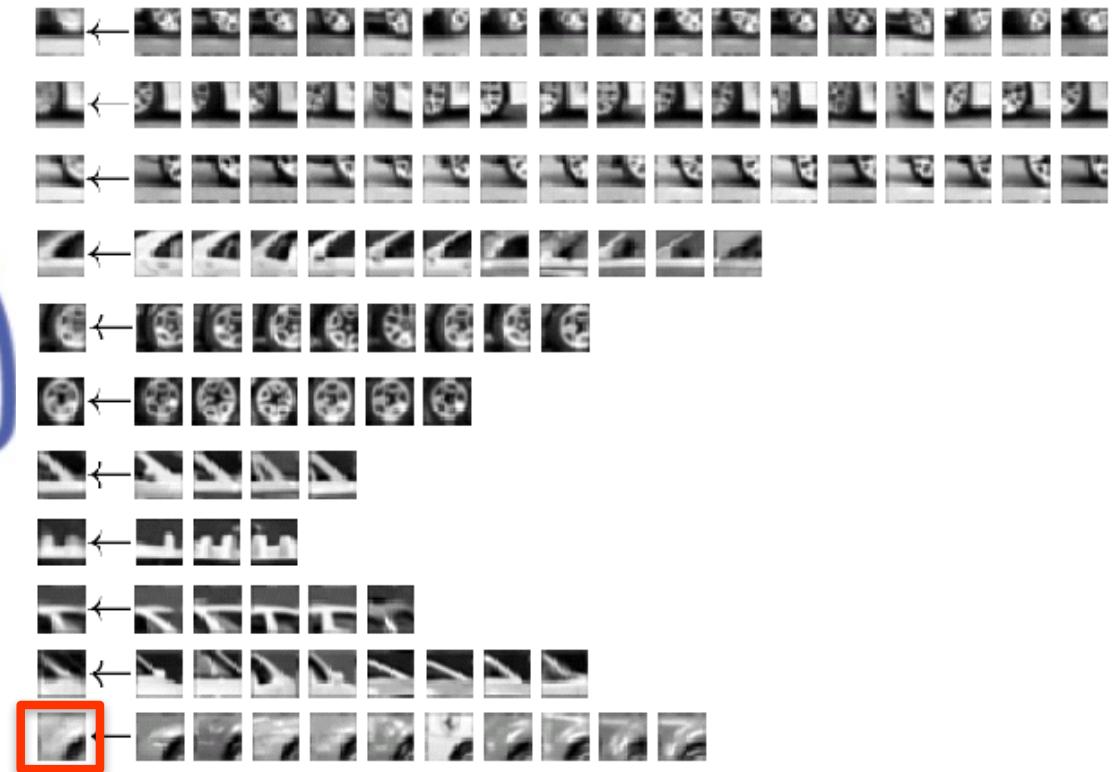
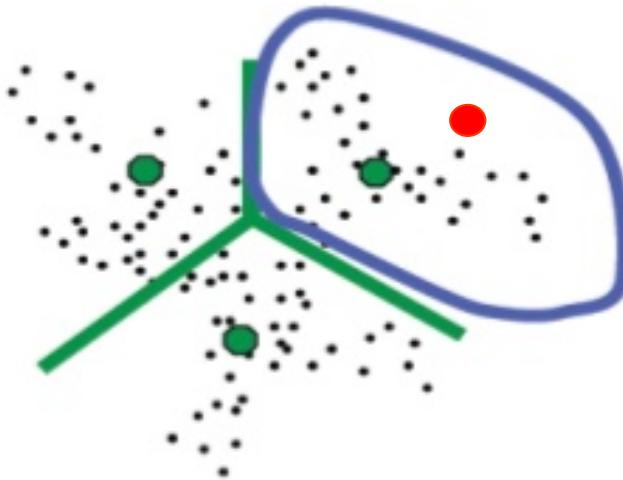
K-means clustering



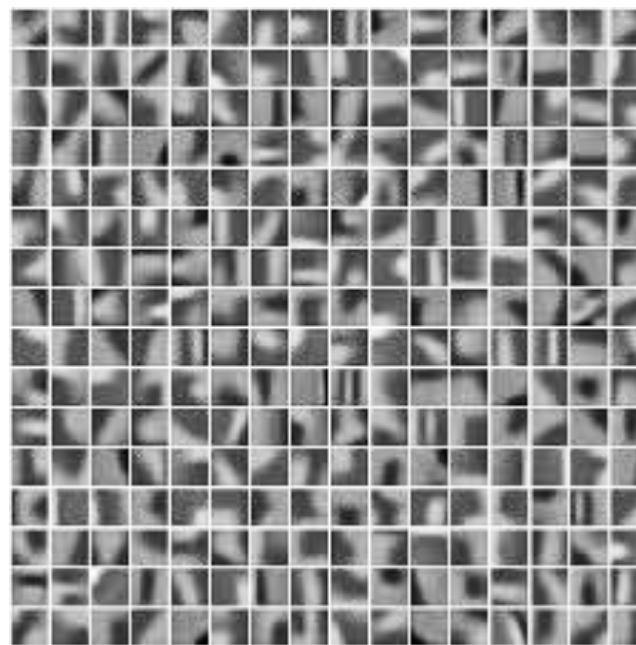
How many clusters?



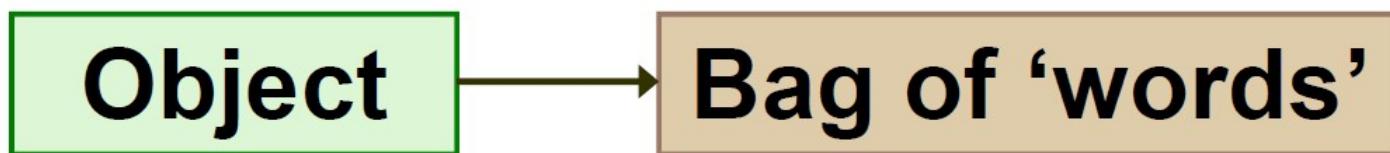
Visual words per object category



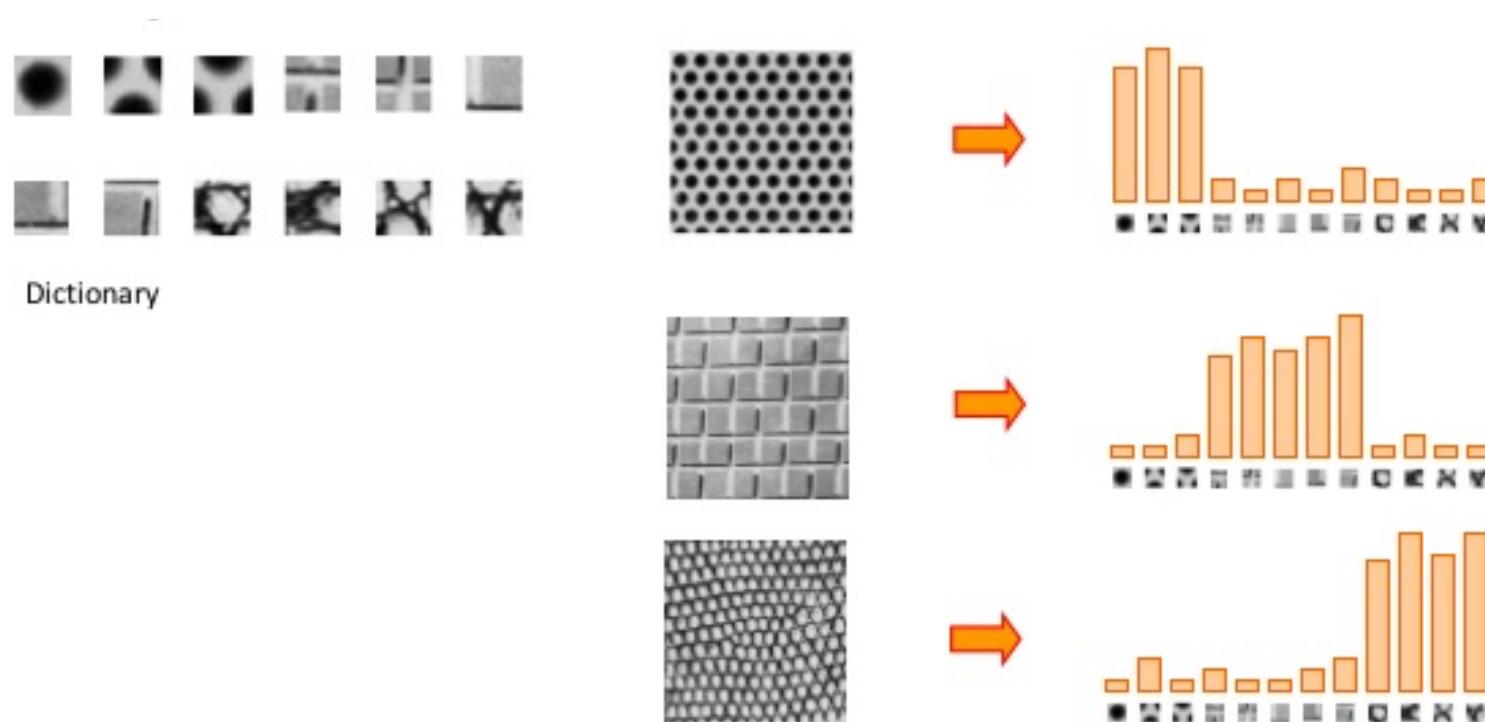
Example of learned visual dictionary



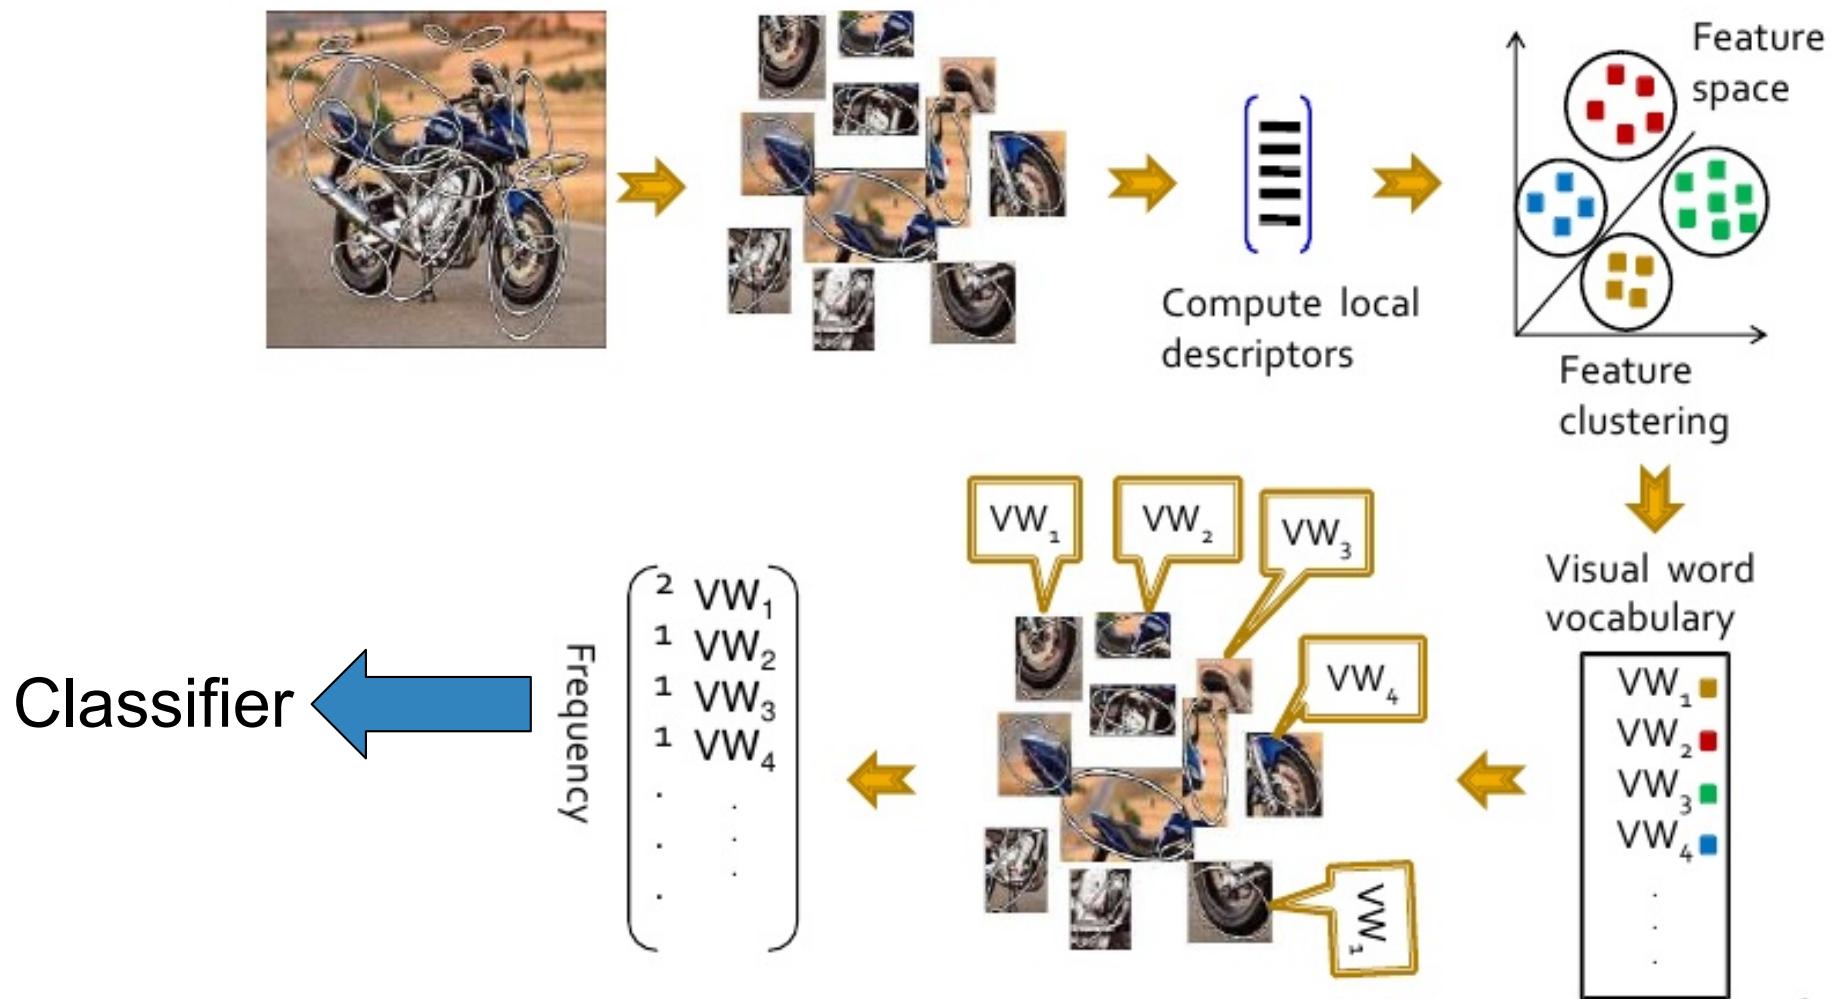
“Bag of words” classification



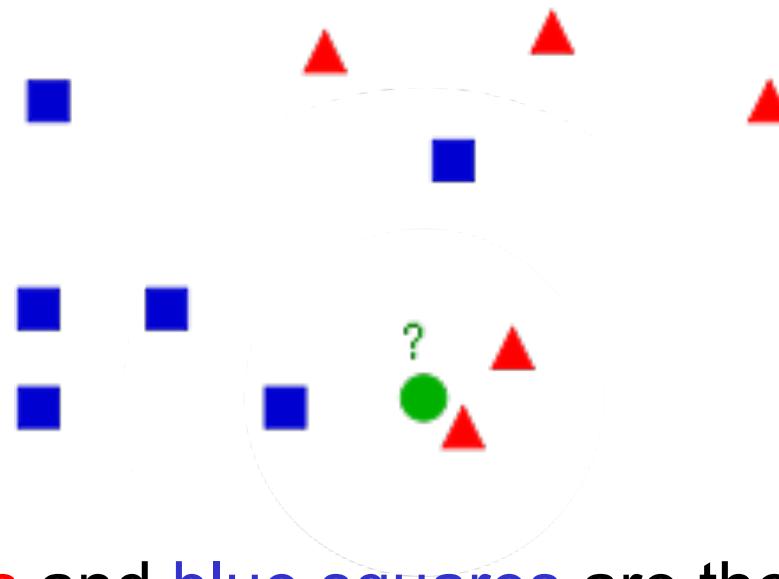
Bag of words



“Bag of words” classification



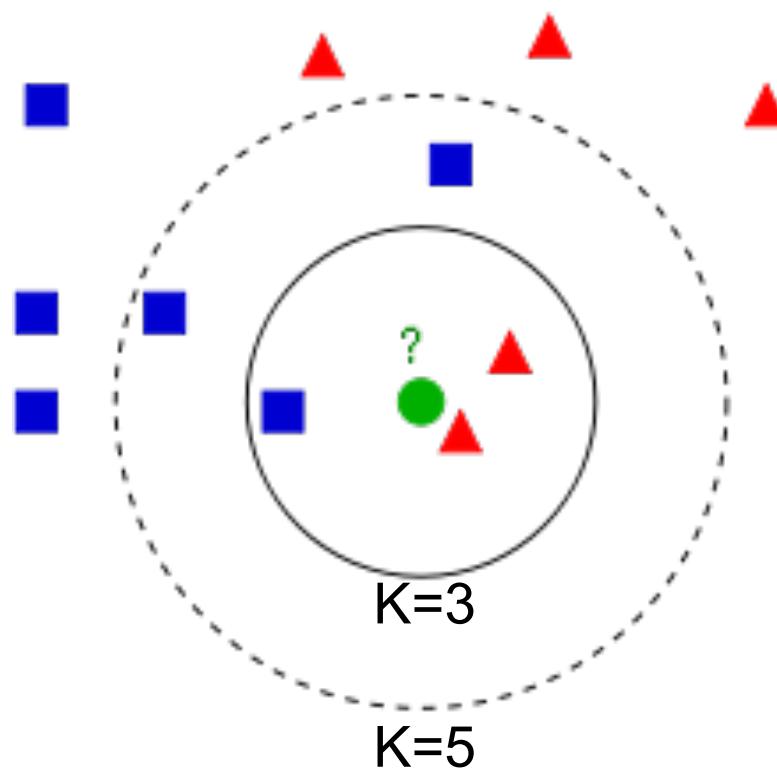
K-nearest Neighbor classifier



Red triangles and blue squares are the points in the training dataset.

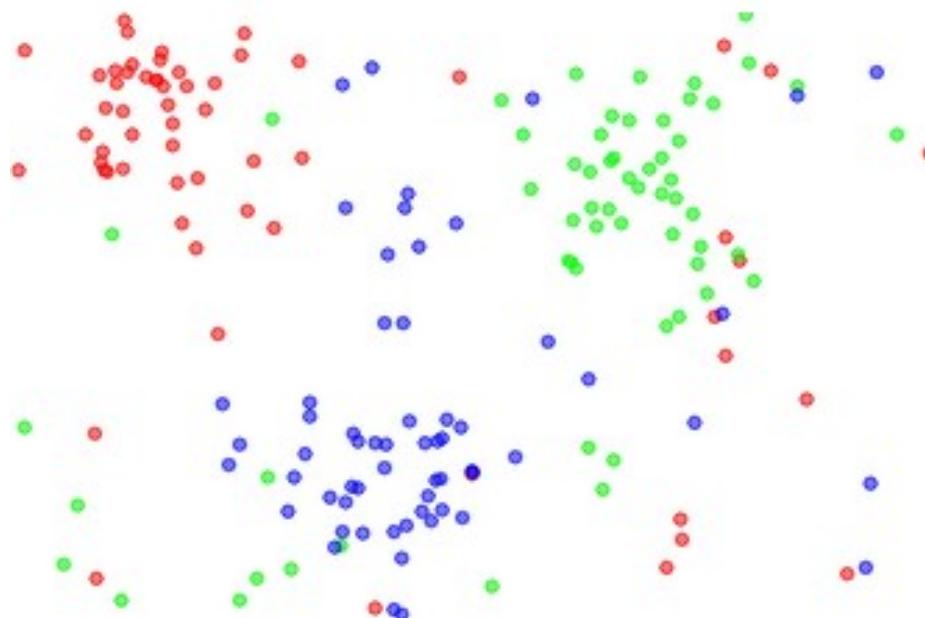
To which class should ● (test point) be assigned to?

K-nearest Neighbor classifier



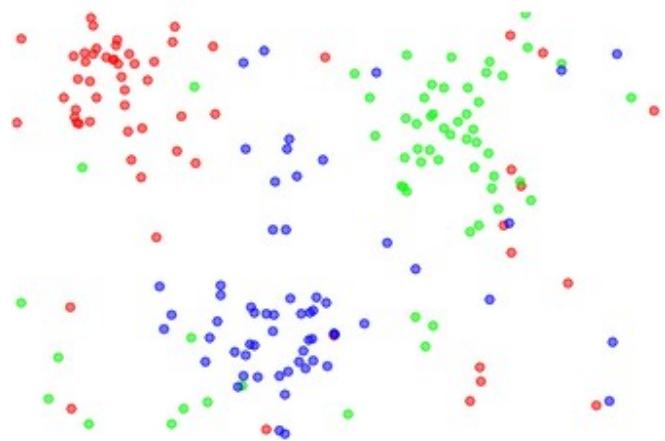
K-nearest Neighbor classifier

the data

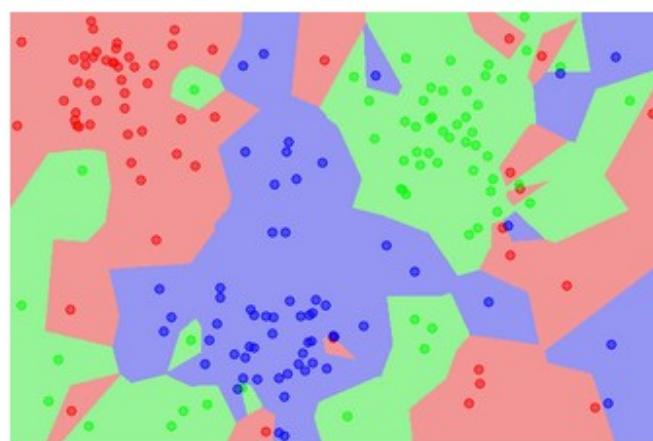


K-nearest Neighbor classifier

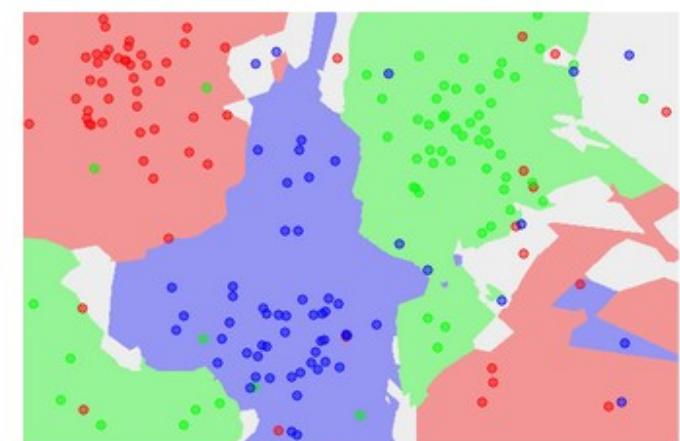
the data



NN classifier

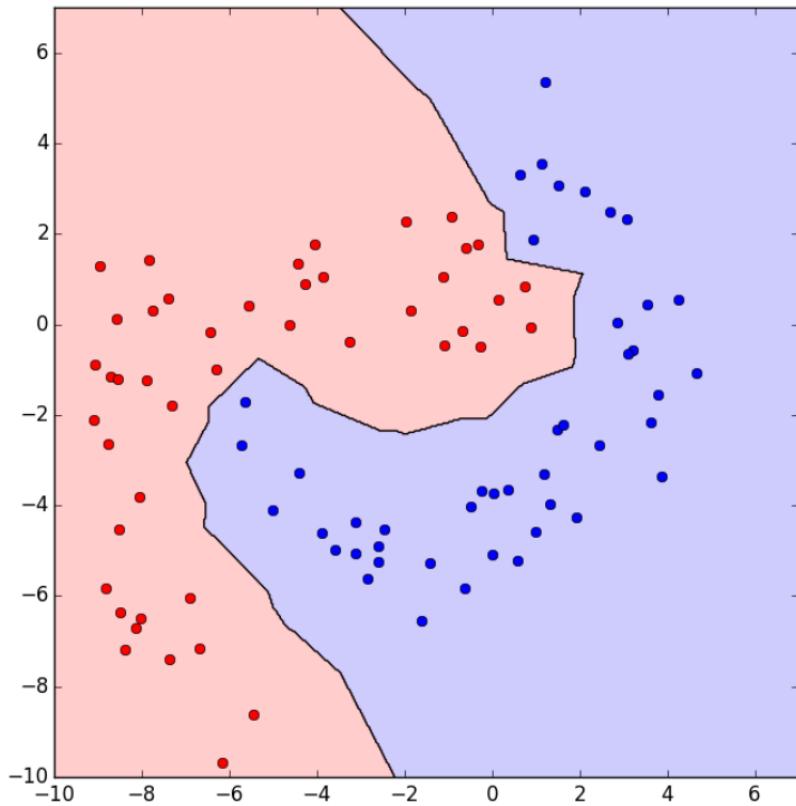


5-NN classifier



K-nearest Neighbor classifier

$K=1$



$K=25$

