

Introduction to Computer Vision

5. 2D Image Transformation

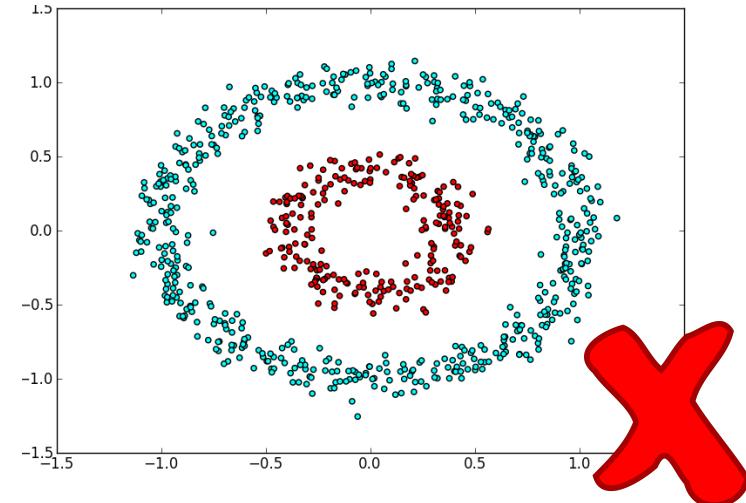
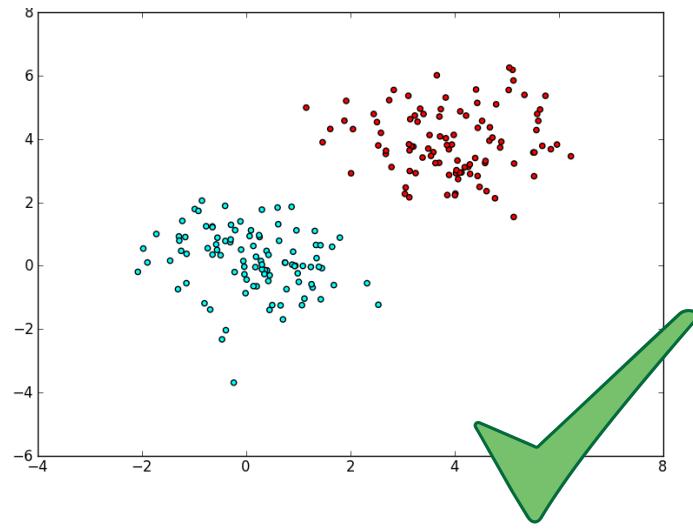
UCLA – CS 188 – Fall 2019

Fabien Scalzo, Ph.D.

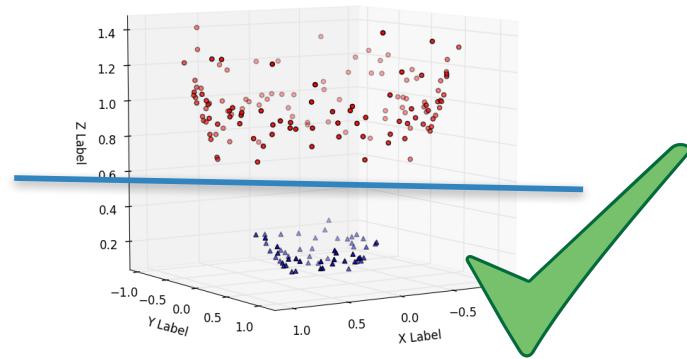
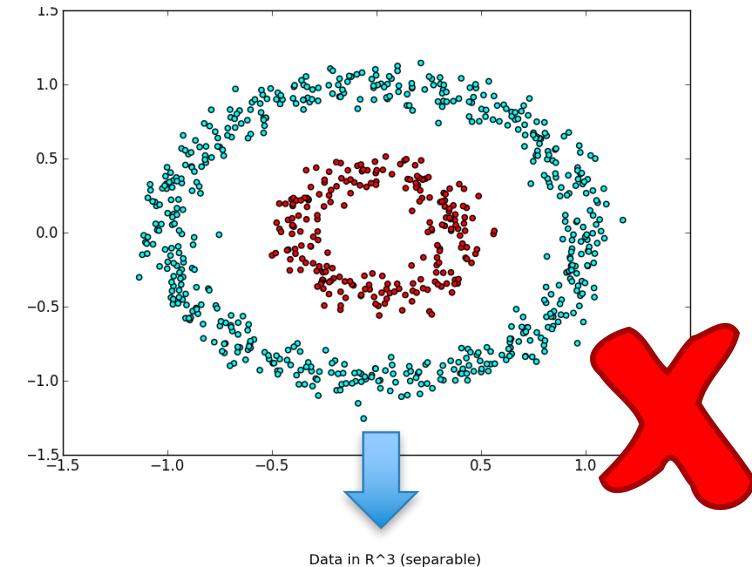
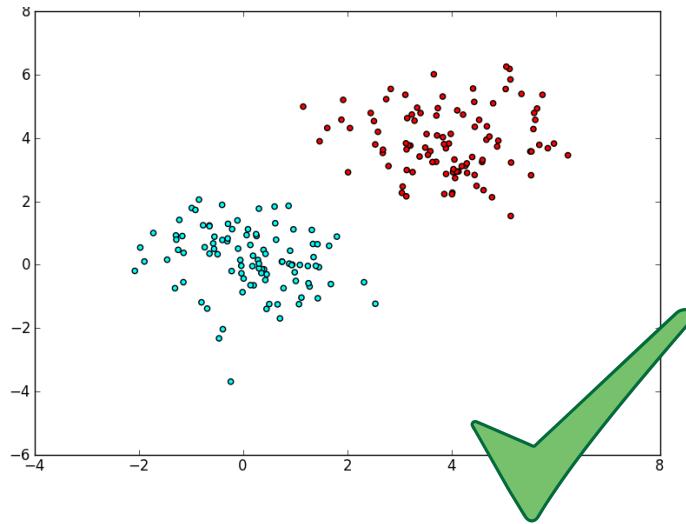
Today

| | | | | |
|----------------|--------|-------------------------------------------------|--------|-------------------------------------------------|
| Week 1 | | | 26-Sep | Introduction |
| Week 2 | 1-Oct | Basic Image Processing | 3-Oct | Feature Extraction and Classification |
| Week 3 | 8-Oct | Classification of Visual Features | 10-Oct | 2D Image transformation |
| Week 4 | 15-Oct | SVD, 2D camera model, projective plane | 17-Oct | Euclidean geometry, rigid body motion |
| Week 5 | 22-Oct | Epipolar Geometry | 24-Oct | 3D Cameras and processing |
| Week 6 | 29-Oct | Midterm | 31-Oct | Statistical decision theory/Pattern Recognition |
| Week 7 | 5-Nov | Deep Learning | 7-Nov | Deep Learning for Image Classification |
| Week 8 | 12-Nov | Object Detection | 14-Nov | Generative models |
| Week 9 | 19-Nov | Medical Imaging | 21-Nov | Autonomous Navigation |
| Week 10 | 26-Nov | Guest Lecture (Nikhil Naik) | 28-Nov | |
| Week 11 | 3-Dec | Recursive 3D reconstruction and pose estimation | 5-Dec | Recap |
| Week 12 | 10-Dec | | 12-Dec | Final |

Linear Classifier



Kernel Projection + Linear Classifier



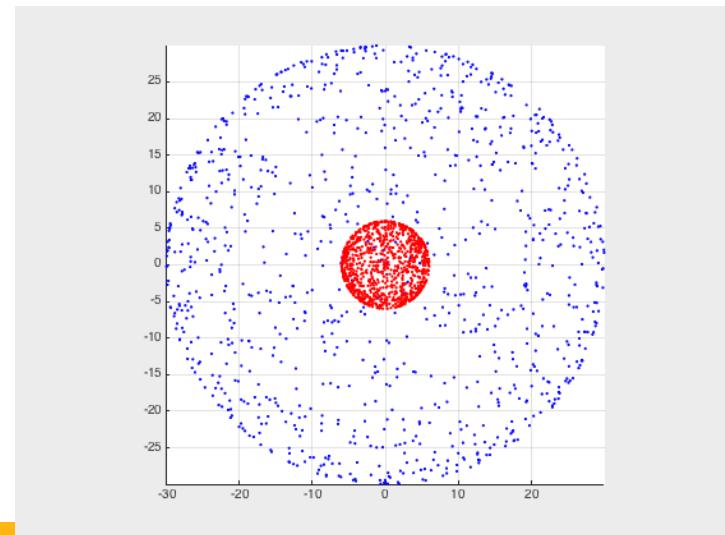
$$[x_1, x_2] = [x_1, x_2, x_1^2 + x_2^2]$$

Kernel Projection + Linear Classifier

Radial Basis Function (RBF) Kernel

$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$

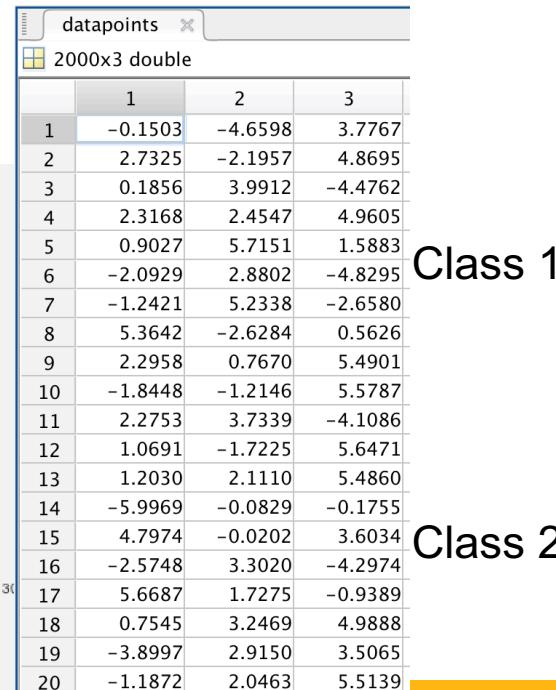
Trick: represents each datapoint as a vector of weights to all the other datapoints



Kernel Projection + Linear Classifier

Represents each datapoint as a vector of weights to all the datapoints

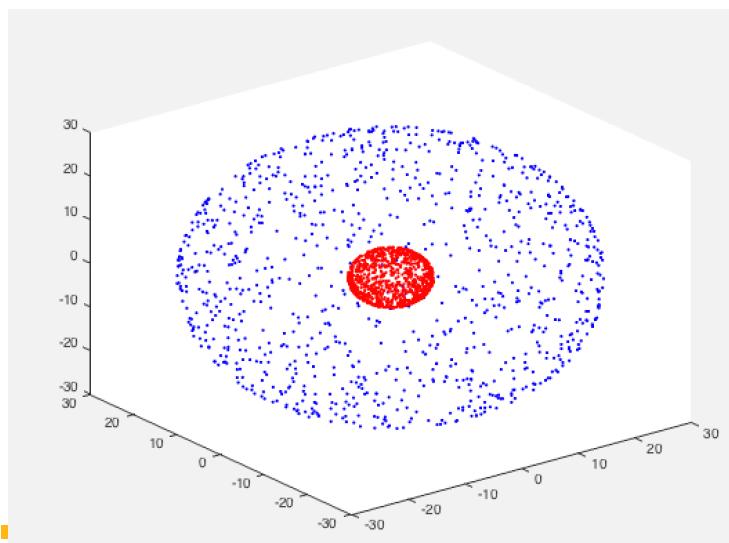
$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$



Kernel Projection + Linear Classifier

Represents each datapoint as a vector of weights to all the datapoints

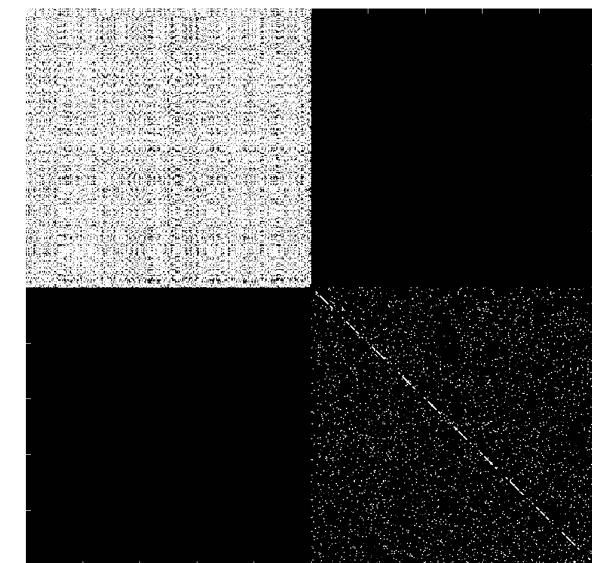
$$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$



| | 1 | 2 | 3 |
|----|---------|---------|---------|
| 1 | -0.1503 | -4.6598 | 3.7767 |
| 2 | 2.7325 | -2.1957 | 4.8695 |
| 3 | 0.1856 | 3.9912 | -4.4762 |
| 4 | 2.3168 | 2.4547 | 4.9605 |
| 5 | 0.9027 | 5.7151 | 1.5883 |
| 6 | -2.0929 | 2.8802 | -4.8295 |
| 7 | -1.2421 | 5.2338 | -2.6580 |
| 8 | 5.3642 | -2.6284 | 0.5626 |
| 9 | 2.2958 | 0.7670 | 5.4901 |
| 10 | -1.8448 | -1.2146 | 5.5787 |
| 11 | 2.2753 | 3.7339 | -4.1086 |
| 12 | 1.0691 | -1.7225 | 5.6471 |
| 13 | 1.2030 | 2.1110 | 5.4860 |
| 14 | -5.9969 | -0.0829 | -0.1755 |
| 15 | 4.7974 | -0.0202 | 3.6034 |
| 16 | -2.5748 | 3.3020 | -4.2974 |
| 17 | 5.6687 | 1.7275 | -0.9389 |
| 18 | 0.7545 | 3.2469 | 4.9888 |
| 19 | -3.8997 | 2.9150 | 3.5065 |
| 20 | -1.1872 | 2.0463 | 5.5139 |

Class 1

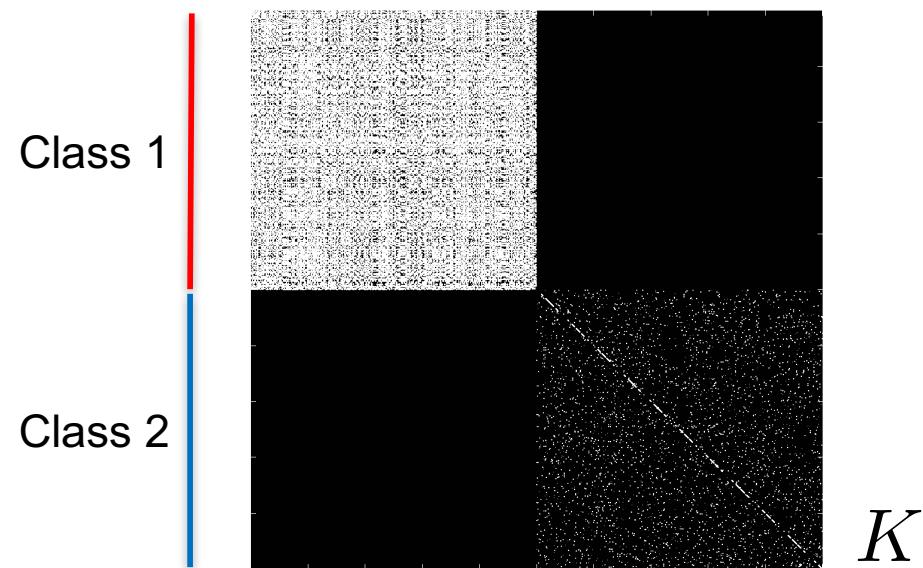
Class 2



K

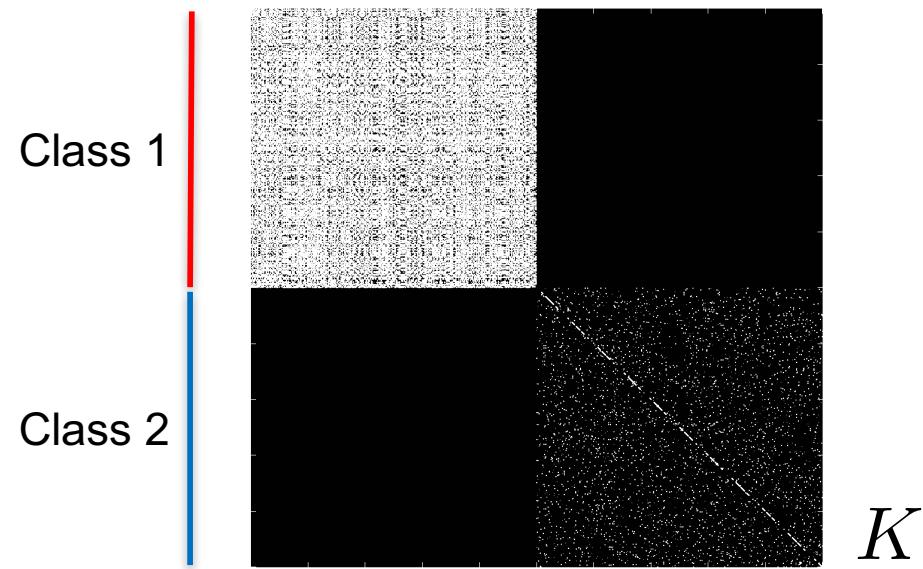
Kernel Projection + Linear Classifier

$$\operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^N L(y_i, \mathbf{w}^\top K(\mathbf{x}_i, \mathbf{x}))$$



Kernel Projection + Linear Classifier

$$\underset{\mathbf{w}}{\operatorname{argmin}} R(\mathbf{w}) + C \sum_{i=1}^N L(y_i, \mathbf{w}^\top K(\mathbf{x}_i, \mathbf{x}))$$







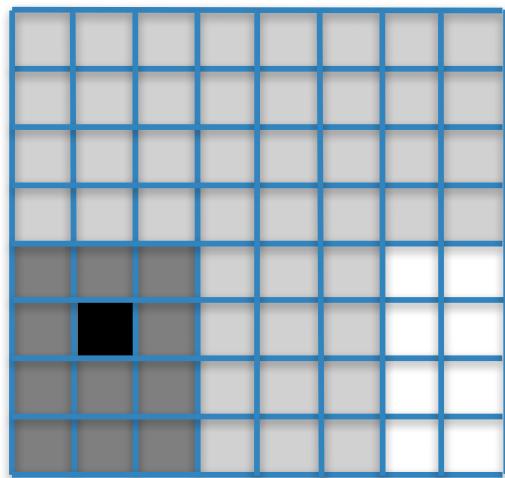


How?

- Rotation, Scaling

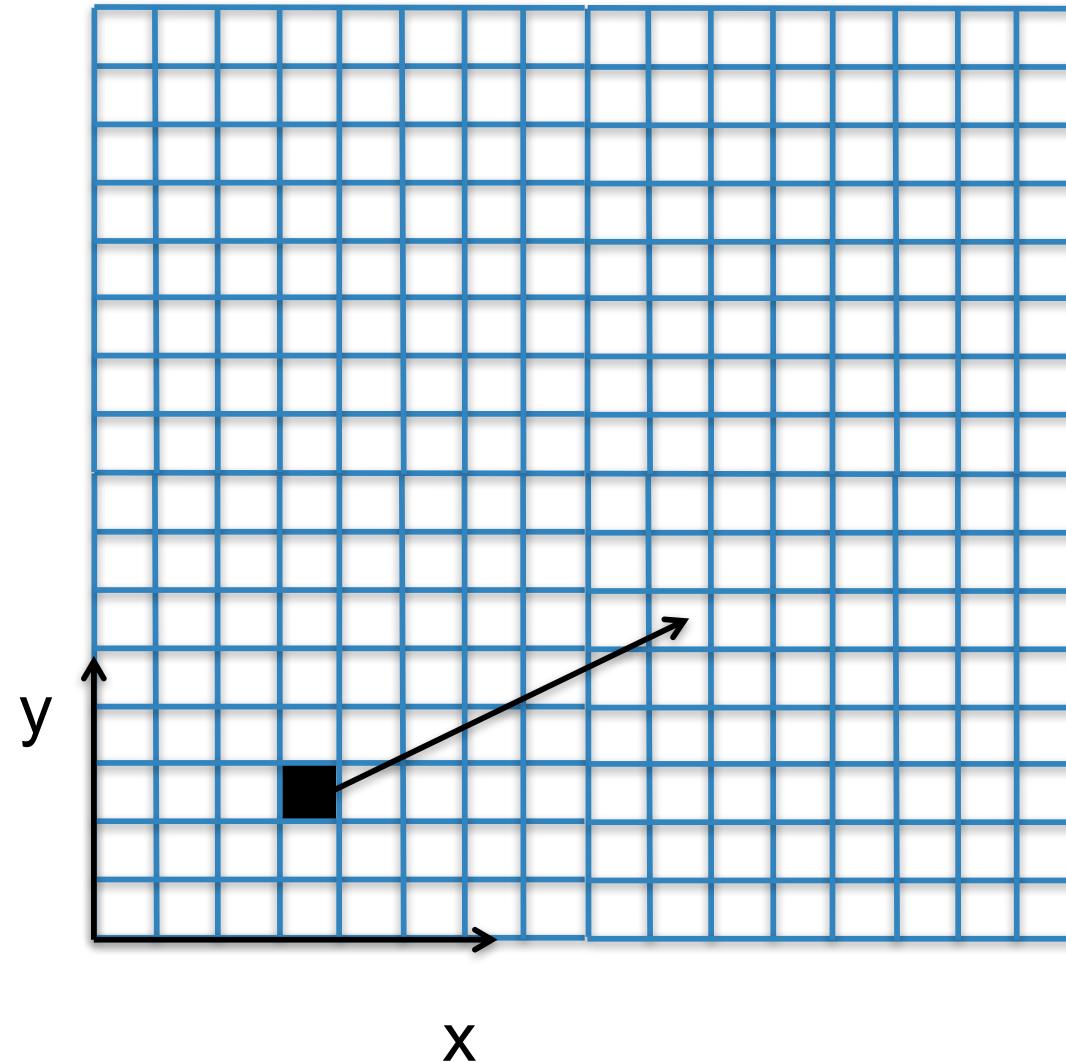


Image representation



Each pixel is described by a position (x,y) in the image and an intensity value $I(x,y)$

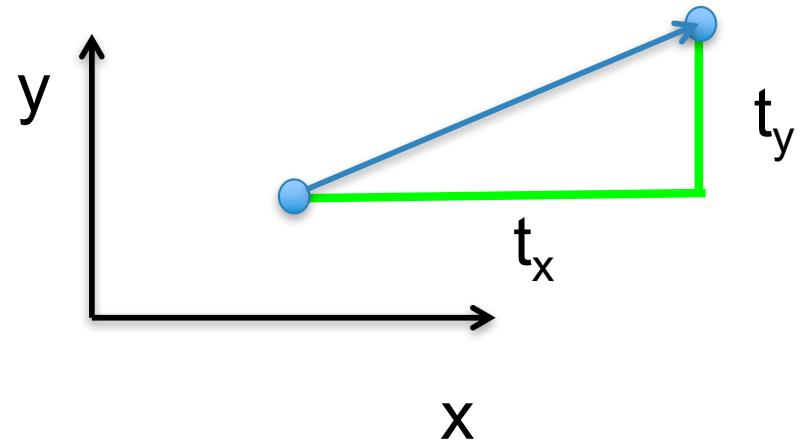
Translation



Translation

t_x : shift in x

t_y : shift in y



Translation



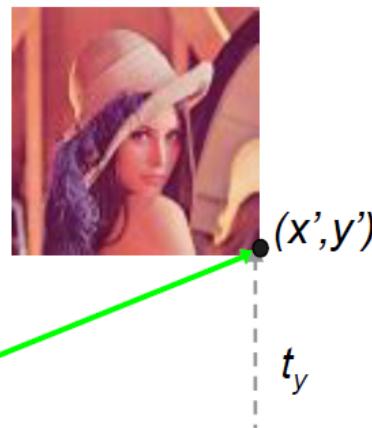
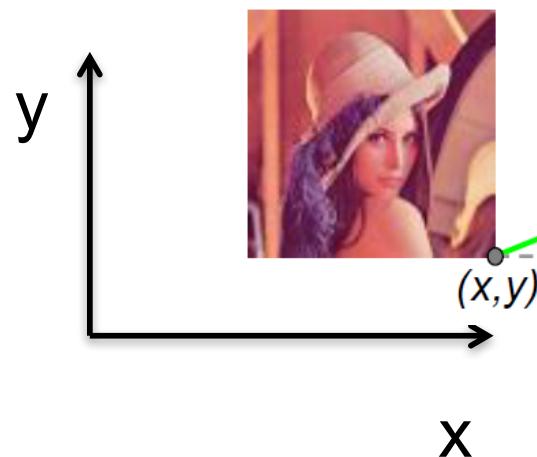
Translation

$$x' = x + t_x$$

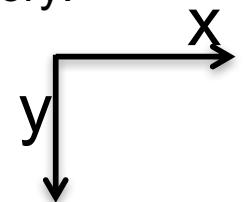
$$y' = y + t_y$$

t_x : shift in x

t_y : shift in y



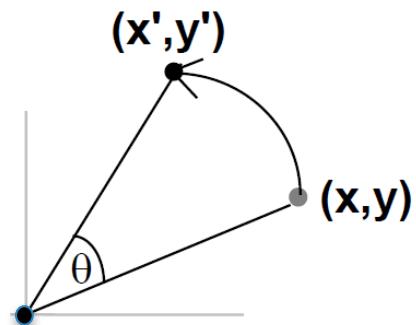
Array in memory:



Rotation

$$x' = x \cos(\theta) - y \sin(\theta)$$
$$y' = y \cos(\theta) + x \sin(\theta)$$

θ is the counter-clockwise rotation angle



The Unit Circle

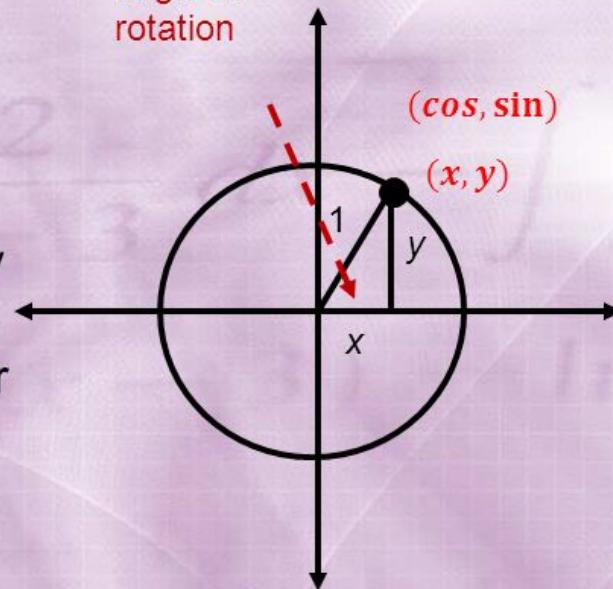
- The **coordinates** of the chosen point are the cosine and sine of the angle θ .

$$\cos \theta = x$$

$$\sin \theta = y$$

- This provides a way to define functions $\sin(\theta)$ and $\cos(\theta)$ for all real numbers θ .
- The other trigonometric functions can be defined from these.

θ is the angle of rotation



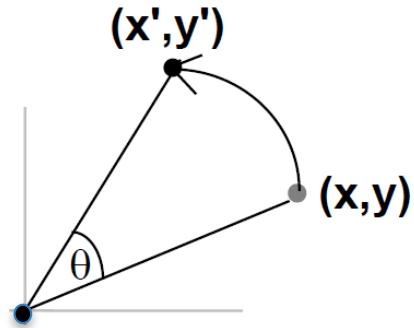
Rotation

Rotation

$$x' = x \cos(\theta) - y \sin(\theta)$$

$$y' = y \cos(\theta) + x \sin(\theta)$$

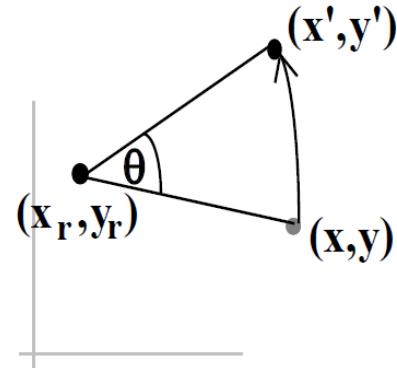
θ is the counter-clockwise rotation angle



Rotation about a point

$$x' = x_r + (x - x_r) \cos(\theta) - (y - y_r) \sin(\theta)$$

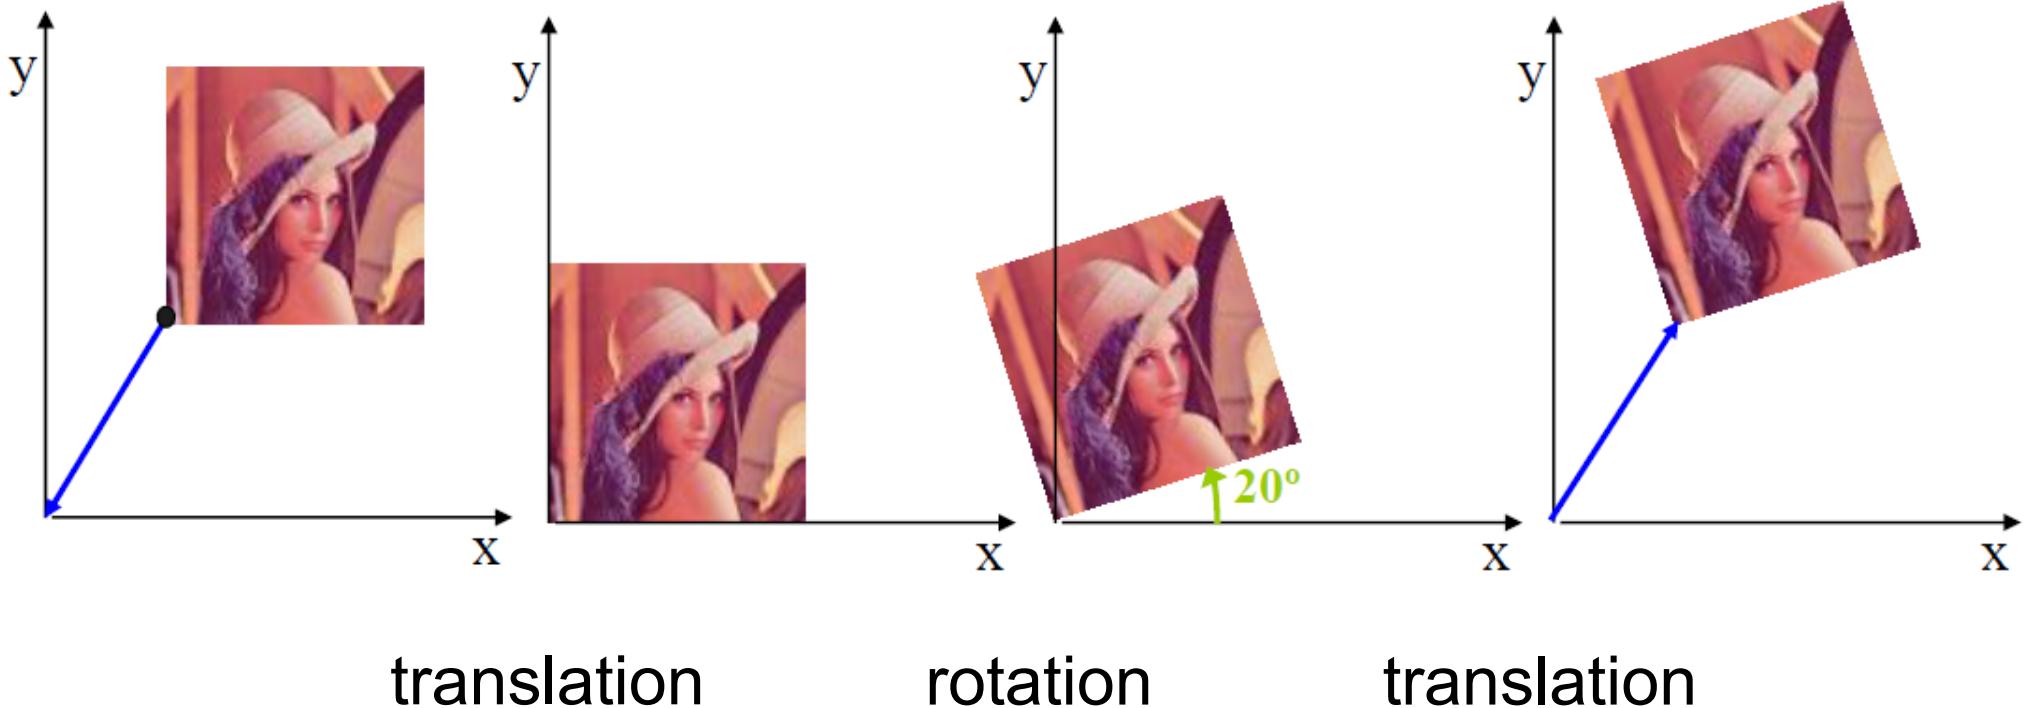
$$y' = y_r + (y - y_r) \cos(\theta) + (x - x_r) \sin(\theta)$$



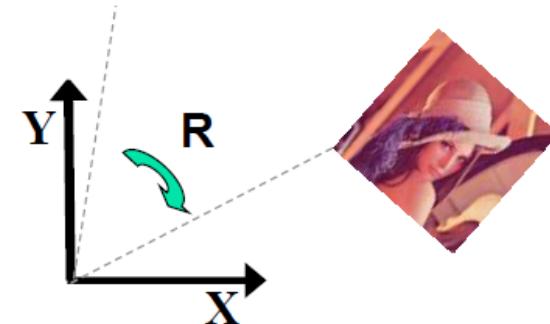
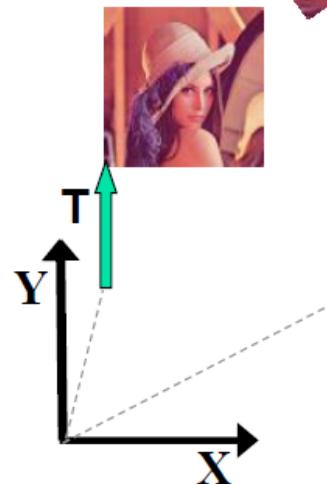
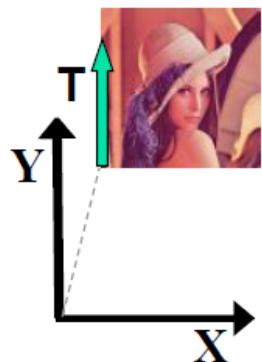
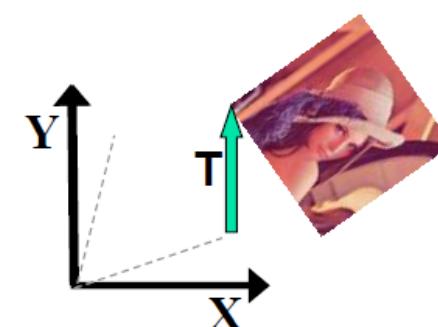
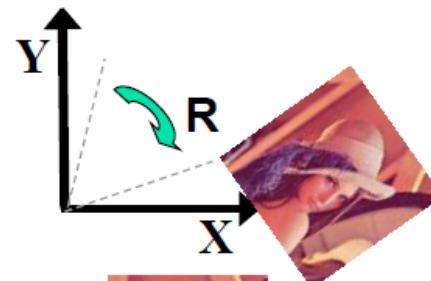
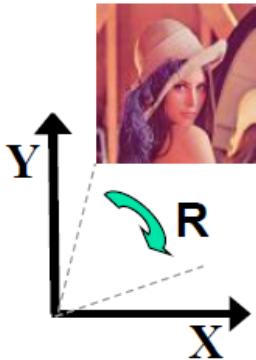
Example: rotation about a point

$$x' = x_r + (x - x_r) \cos(\theta) - (y - y_r) \sin(\theta)$$

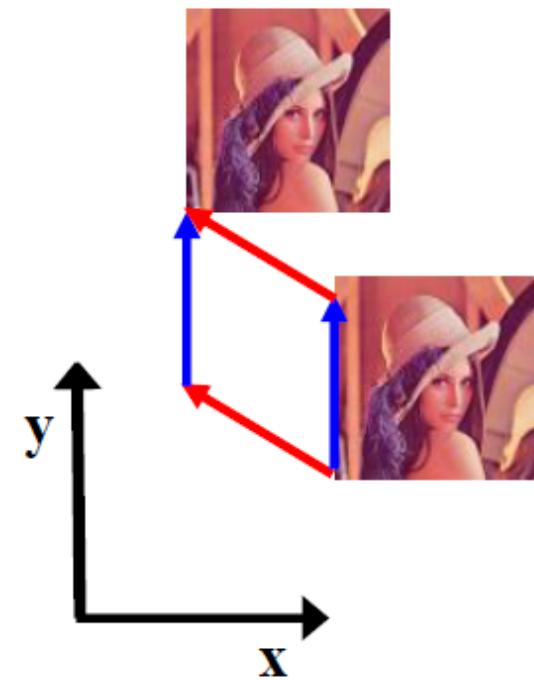
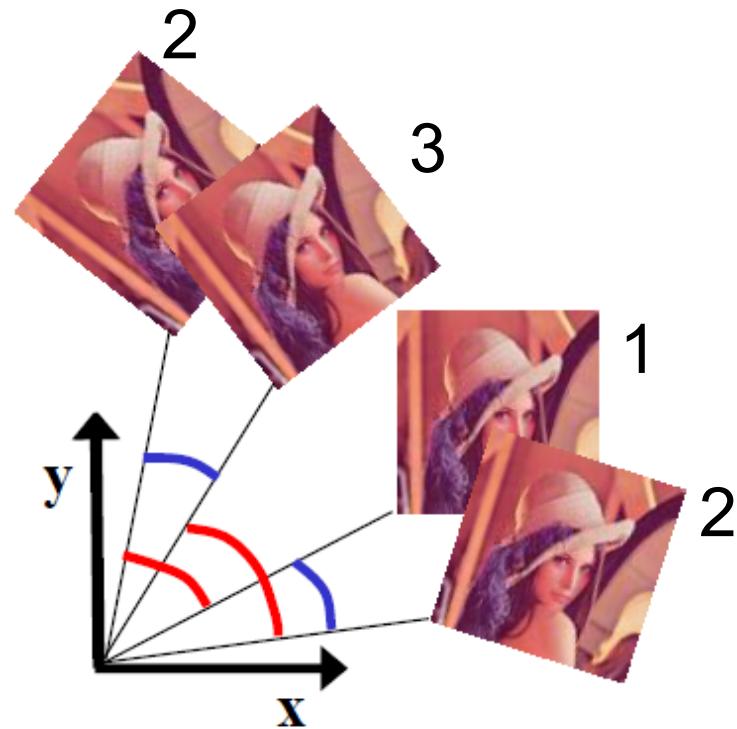
$$y' = y_r + (y - y_r) \cos(\theta) + (x - x_r) \sin(\theta)$$



Rotation x Translation \neq Translation x Rotation



Order

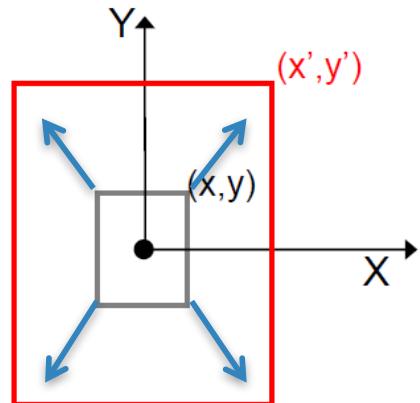


Scaling (w.r.t. the origin)

$$x' = s_x \cdot x \quad y' = s_y \cdot y$$

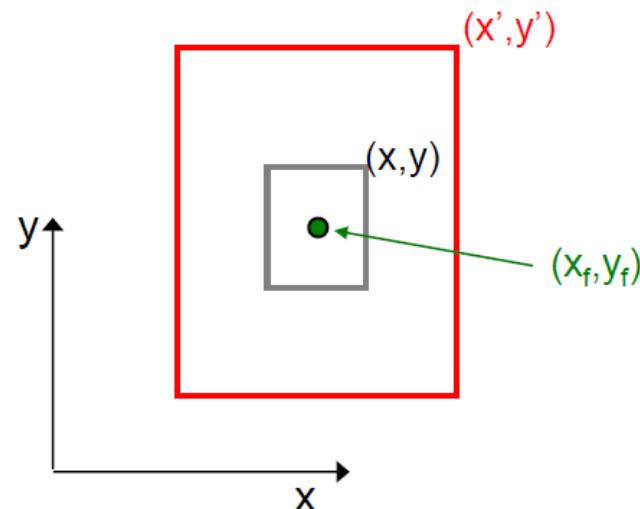
s_x = scaling factor in x

s_y = scaling factor in y

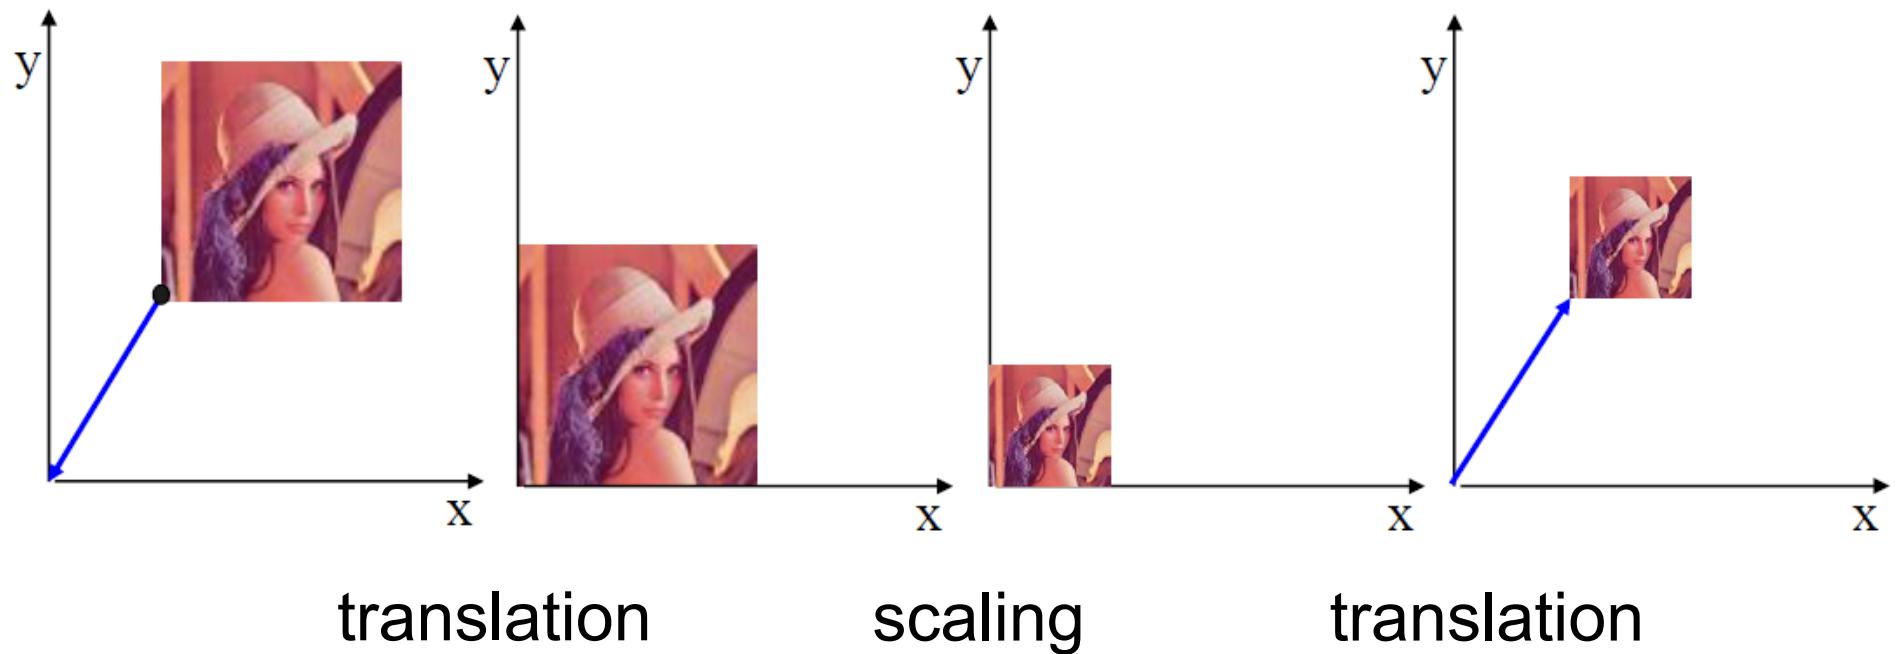


Scaling (about a point)

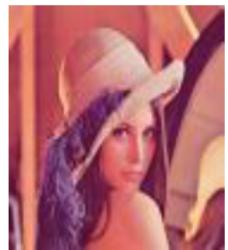
$$x' = x_f + (x - x_f) \cdot s_x \quad y' = y_f + (y - y_f) \cdot s_y$$



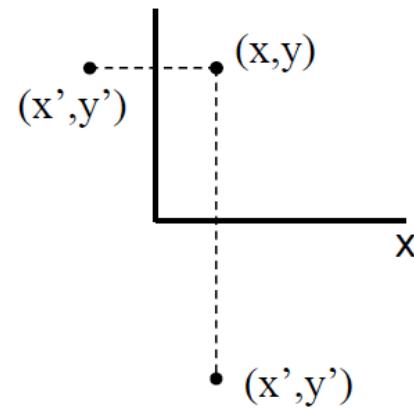
Example



Flip



$$\begin{bmatrix} x' = x \\ y' = -y \end{bmatrix} \text{ et } \begin{bmatrix} x' = -x \\ y' = y \end{bmatrix}$$



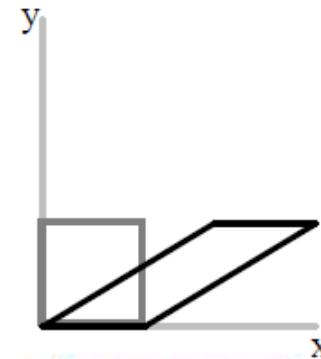
Shearing

$$\begin{cases} x' = x \\ y' = y + sh_y \cdot x \end{cases}$$

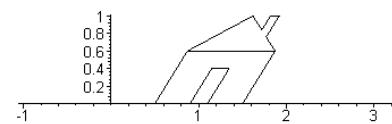


along Y-axis

$$\begin{cases} x' = x + sh_x \cdot y \\ y' = y \end{cases}$$



along X-axis



Inverse transforms

- Translation: $t_x, t_y \rightarrow -t_x, -t_y$
- Rotation: $\theta \rightarrow -\theta$
- Scaling: $s_x, s_y \rightarrow 1/s_x, 1/s_y$
- Shearing: $sh_x, sh_y \rightarrow -sh_x, -sh_y$

Matrix representation

- Rotation

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$
$$x' = x \cos(\theta) - y \sin(\theta)$$
$$y' = y \cos(\theta) + x \sin(\theta);$$

- Scaling

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$
$$x' = s_x \cdot x \quad y' = s_y \cdot y$$

- Shearing

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & sh_x \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$
$$x' = x + sh_x \cdot y$$
$$y' = y$$

- Translation

~~$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$~~

- For convenience, we use **Homogeneous Coordinates**
 - For combinations of transforms: easier to read

Translation:

$$\begin{aligned}x' &= 1 \cdot x + 0 \cdot y + t_x \cdot 1 \\x' &= x + t_x\end{aligned}$$

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

T

Homogeneous Coordinates

Scaling:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

S

Rotation:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 \\ \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

R

Homogeneous Coordinates

$$M_{ST} = S(s_x, s_y) \cdot T(t_x, t_y) = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & s_x \cdot t_x \\ 0 & s_y & s_y \cdot t_y \\ 0 & 0 & 1 \end{bmatrix}$$

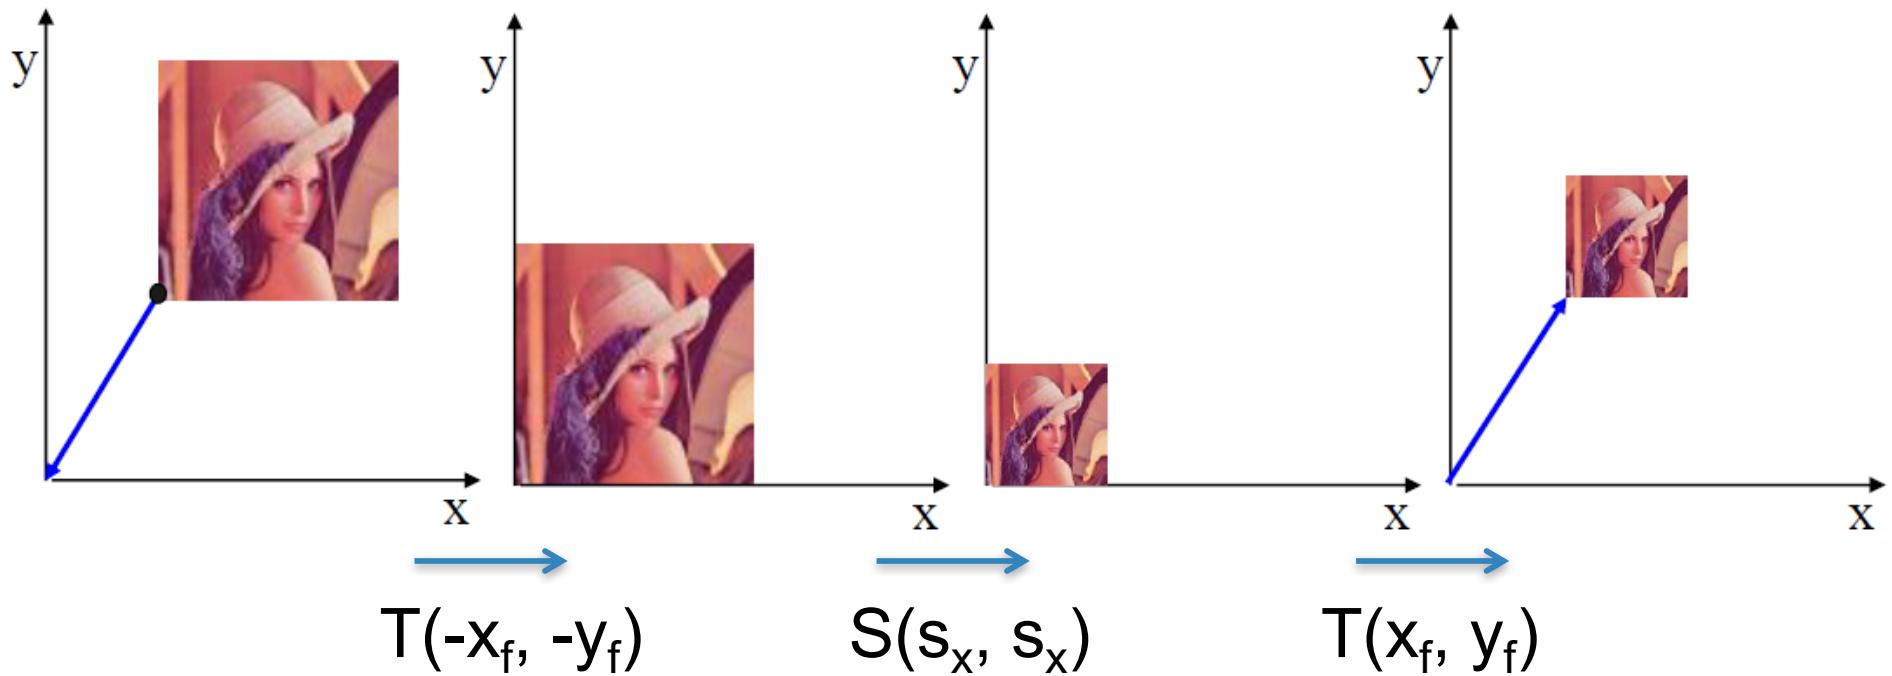
Homogeneous Coordinates

$$M_{ST} = S(s_x, s_y) \cdot T(t_x, t_y) = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & s_x \cdot t_x \\ 0 & s_y & s_y \cdot t_y \\ 0 & 0 & 1 \end{bmatrix}$$

$$\neq \begin{bmatrix} s_x & 0 & t_x \\ 0 & s_y & t_y \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} = T(t_x, t_y) \cdot S(s_x, s_y) = M_{TS}$$

Example: scaling

$$x' = x_f + (x - x_f) \cdot s_x \quad y' = y_f + (y - y_f) \cdot s_y$$



$$= \begin{bmatrix} 1 & 0 & x_f \\ 0 & 1 & y_f \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & -x_f \\ 0 & 1 & -y_f \\ 0 & 0 & 1 \end{bmatrix}$$

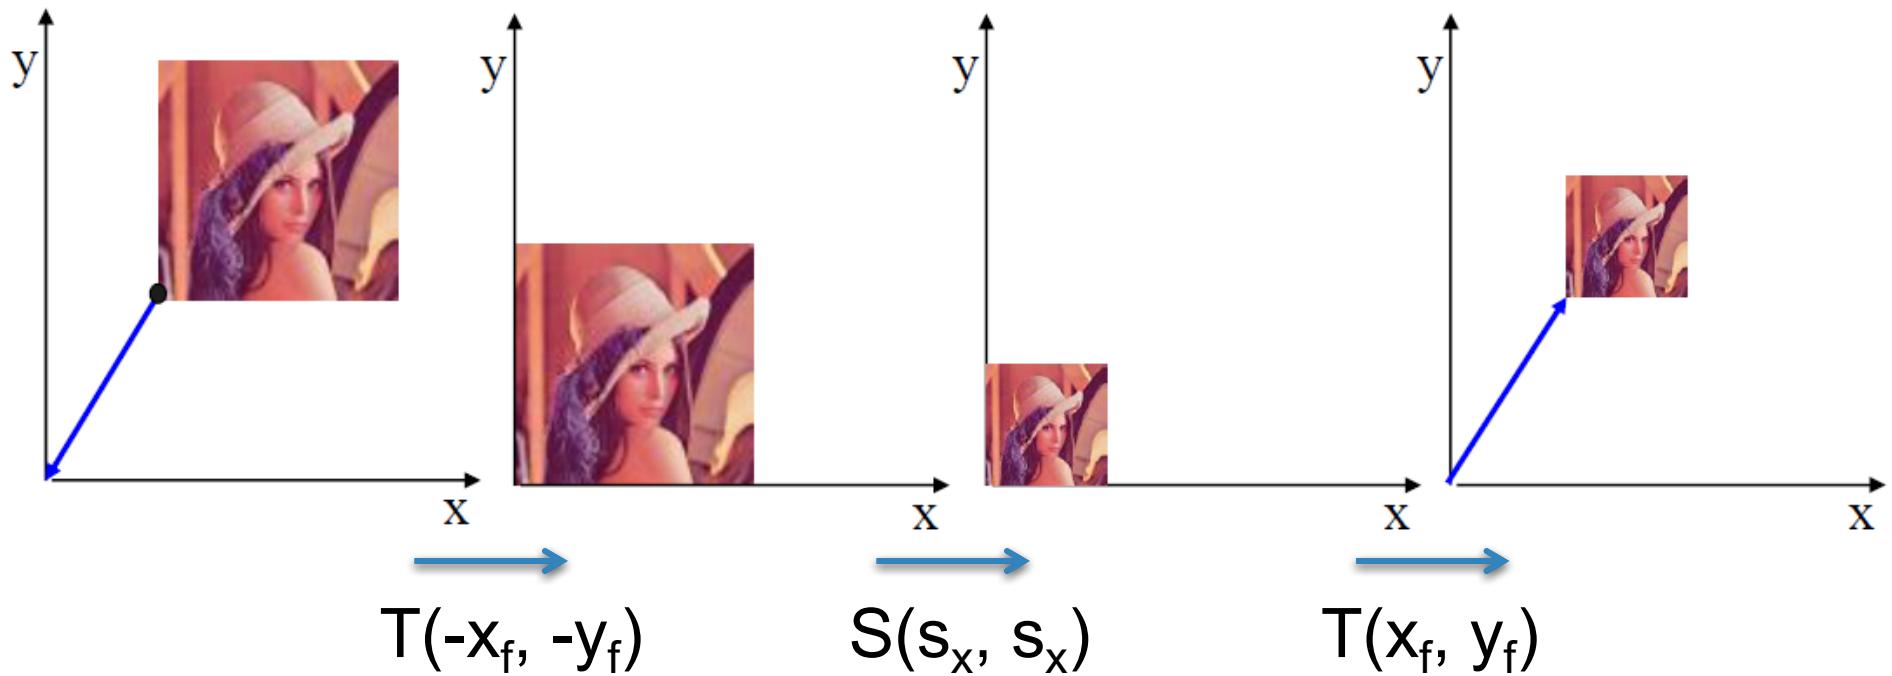
When we write transformations using standard math notation, the closest transformation to the point is applied first:

$$T \ R \ S \ p = T(R(Sp))$$

**first, the object is scaled,
then rotated,
then translated**

Example: scaling

$$x' = x_f + (x - x_f) \cdot s_x \quad y' = y_f + (y - y_f) \cdot s_y$$



$$= \begin{bmatrix} 1 & 0 & x_f \\ 0 & 1 & y_f \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 & 0 & -x_f \\ 0 & 1 & -y_f \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & x_f - s_x \cdot x_f \\ 0 & s_y & y_f - s_y \cdot y_f \\ 0 & 0 & 1 \end{bmatrix}$$

Homogeneous Coordinates

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = M_{TST^{-1}} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & x_f - s_x \cdot x_f \\ 0 & s_y & y_f - s_y \cdot y_f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \Leftrightarrow$$

$$x' = x_f + (x - x_f) \cdot s_x \quad y' = y_f + (y - y_f) \cdot s_y$$

Summary

Translation

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{pmatrix} * \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

Scaling

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{pmatrix} * \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

Rotation

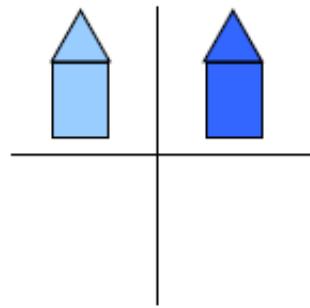
$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{pmatrix} * \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

Shearing

$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} * \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

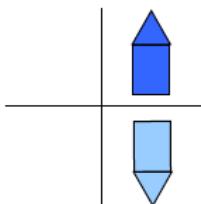
Other transformations

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

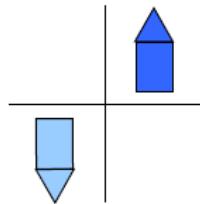


Flip along x

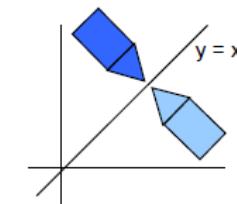
$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



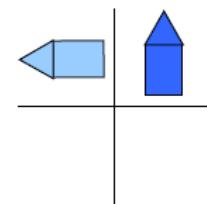
$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

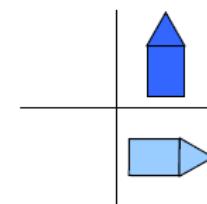


$$\begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



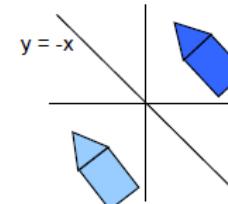
Rotation 90°

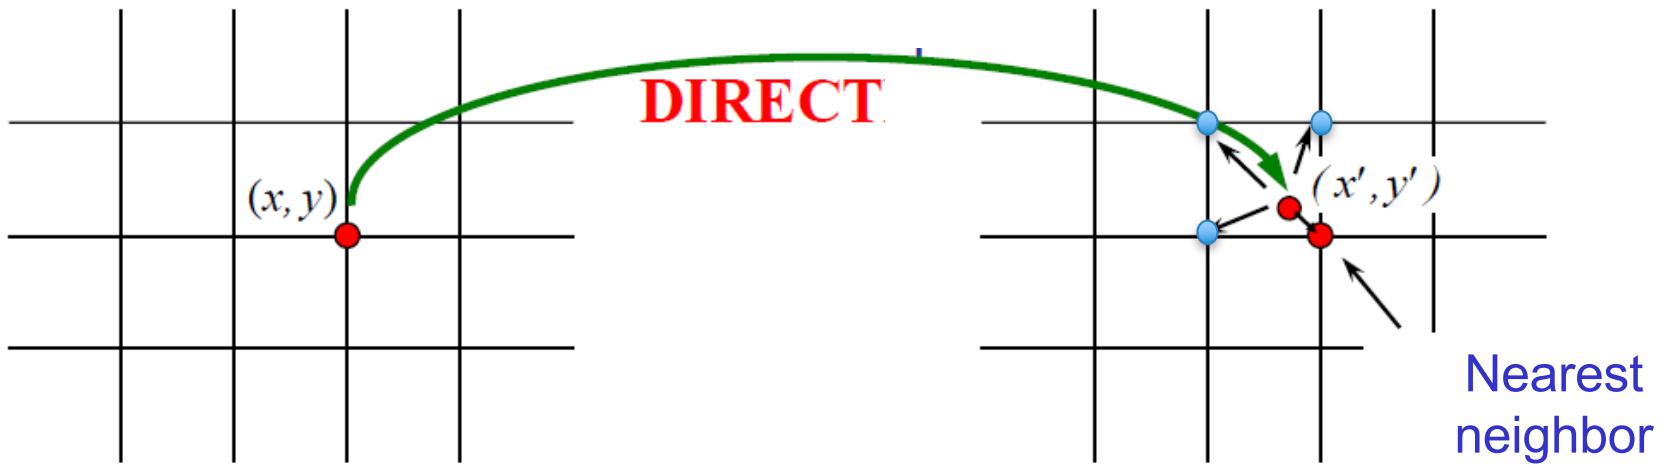
$$\begin{bmatrix} 0 & 1 & 0 \\ -1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



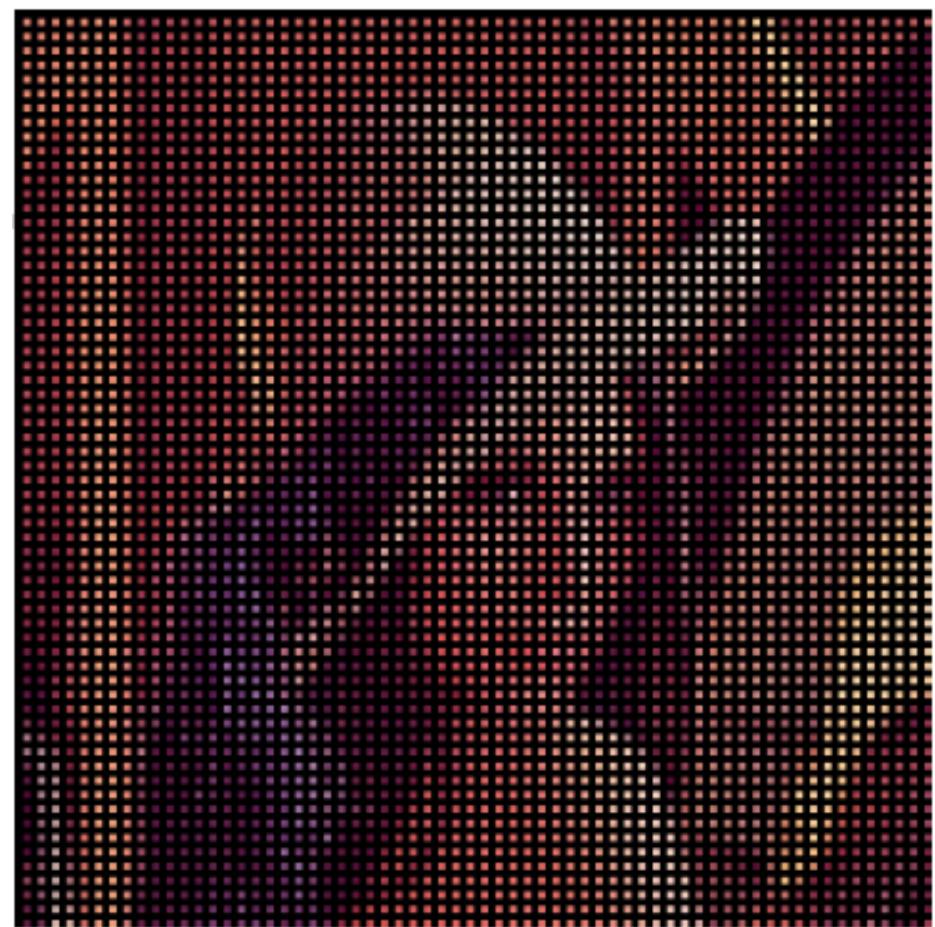
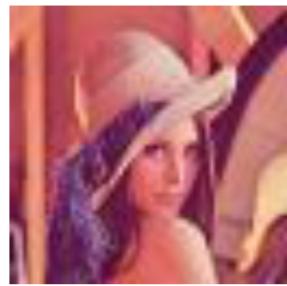
Rotation -90°

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

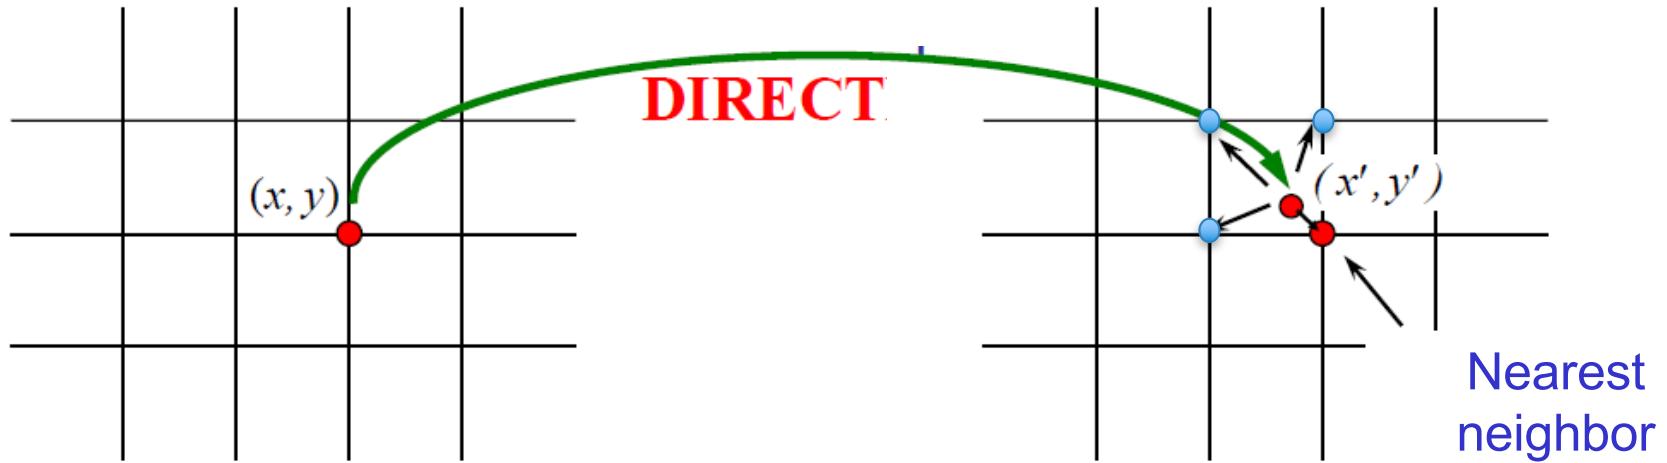




Direct

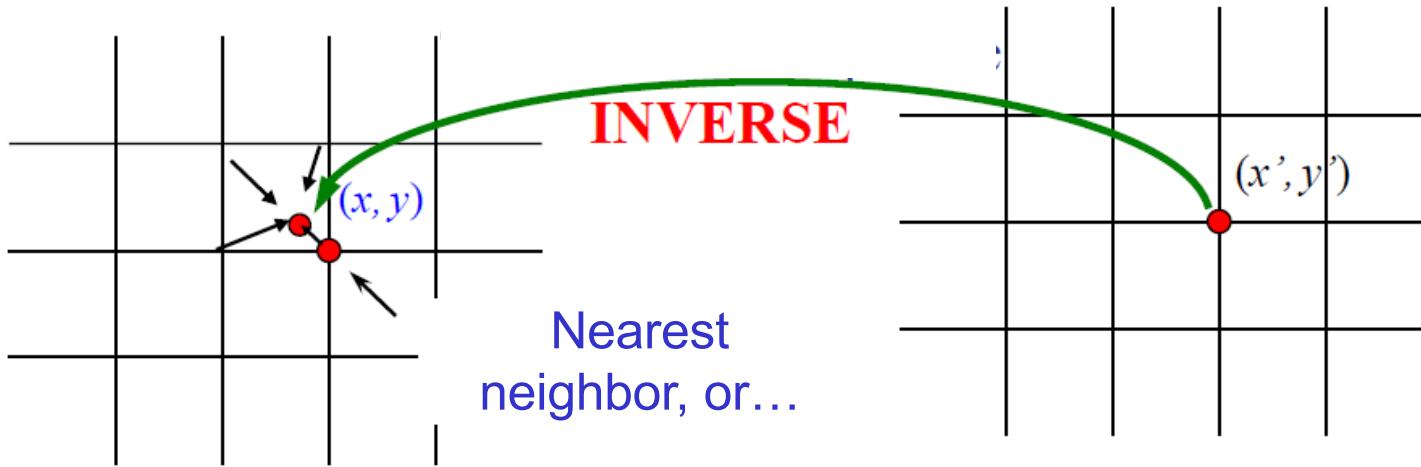


How not to do it



- leads to holes in the image, multiple values for some pixels

How to do it



Inverse

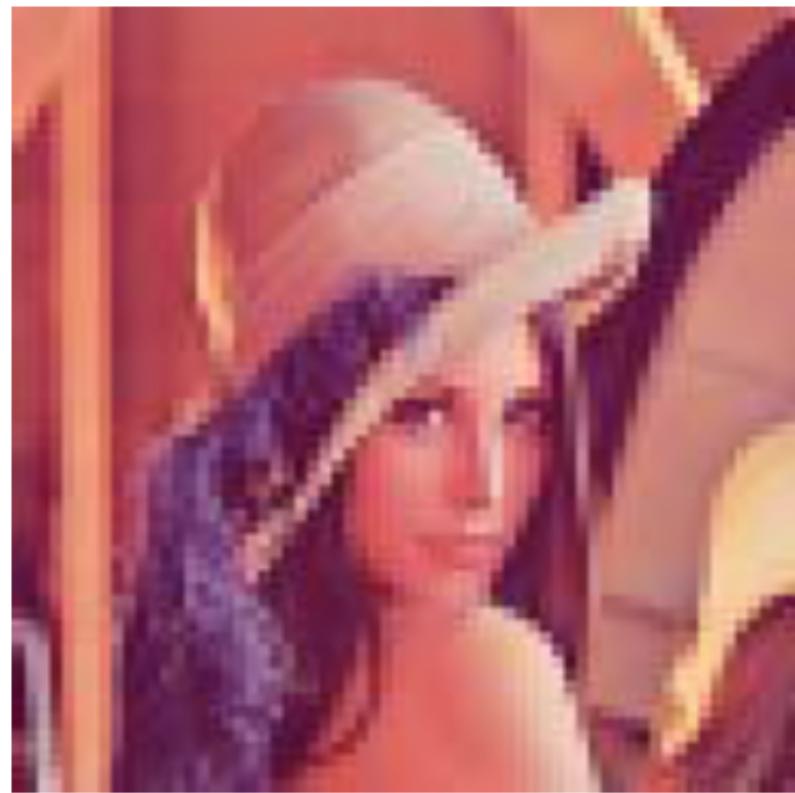
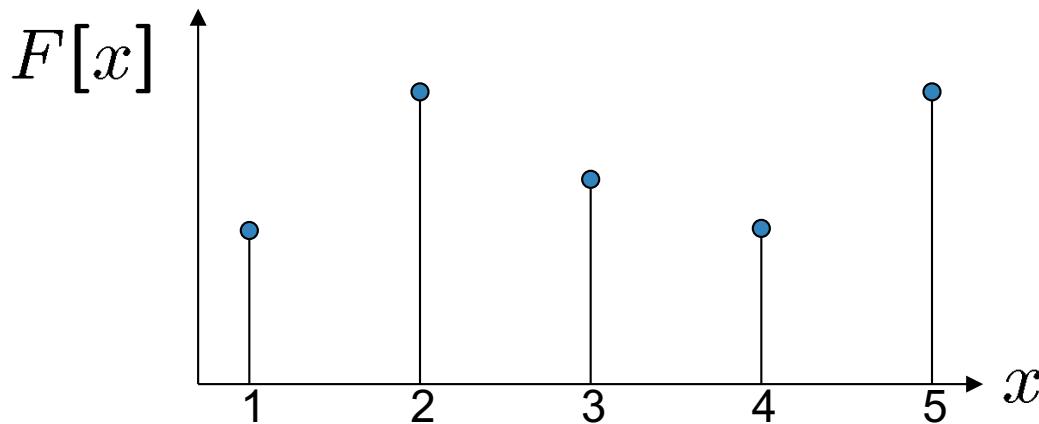
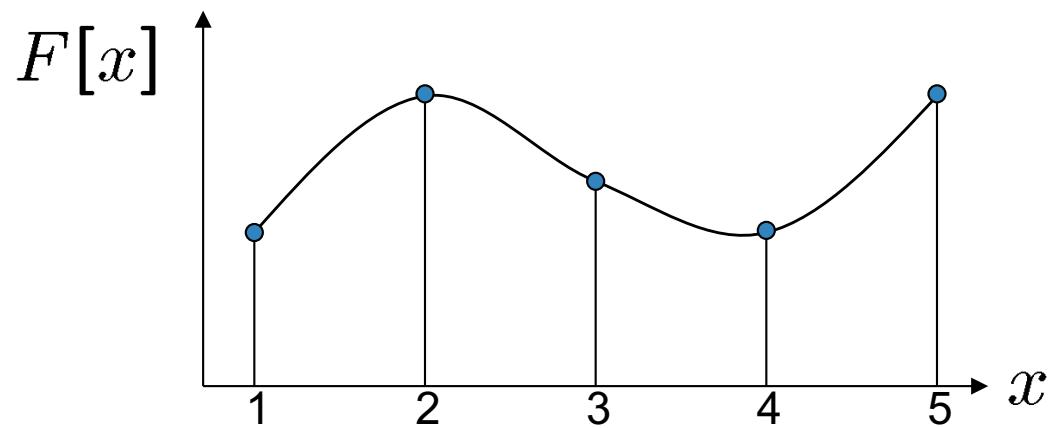


Image interpolation



- What if we don't know F ?
 - It is a discrete point-sampling of a continuous function
 - If we could somehow reconstruct the original function, any new image could be generated, at any resolution and scale

Image interpolation



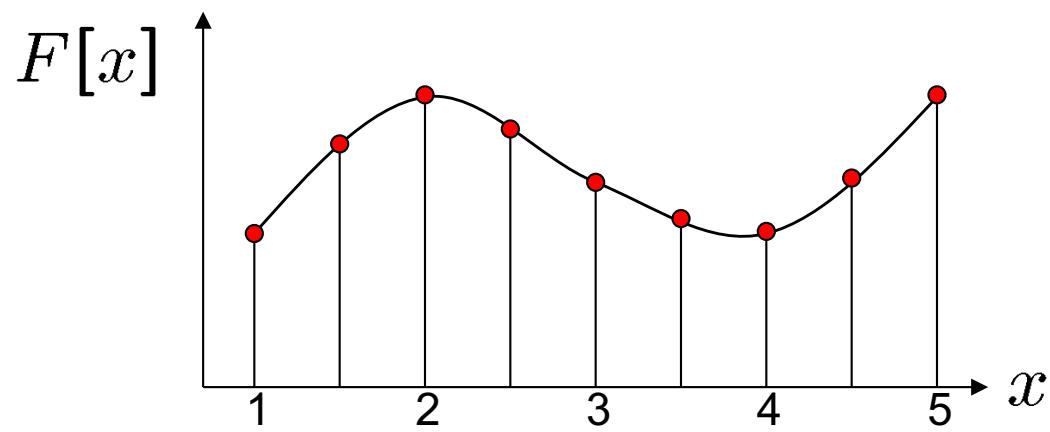


Image interpolation

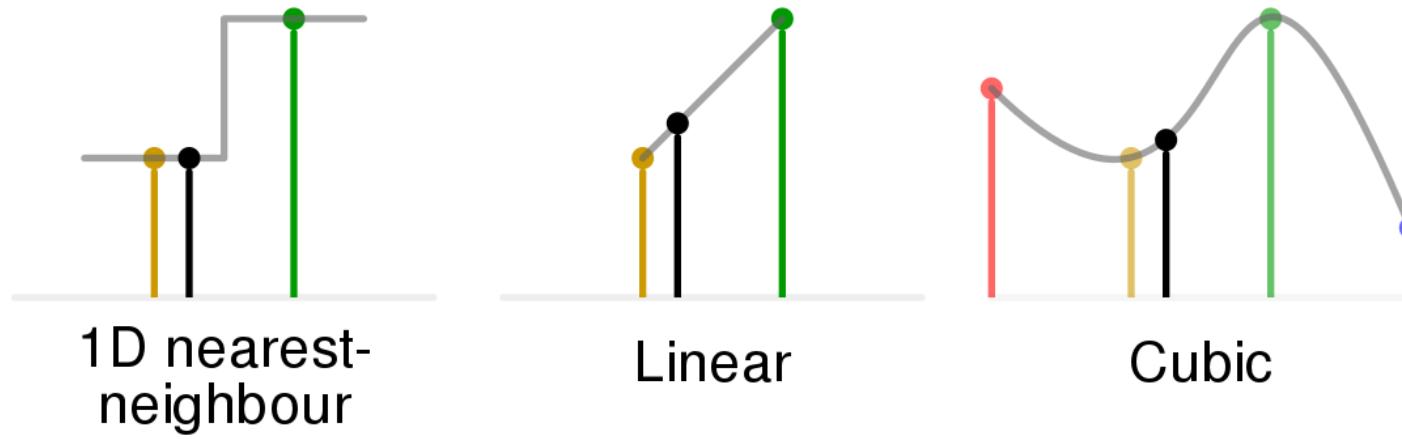
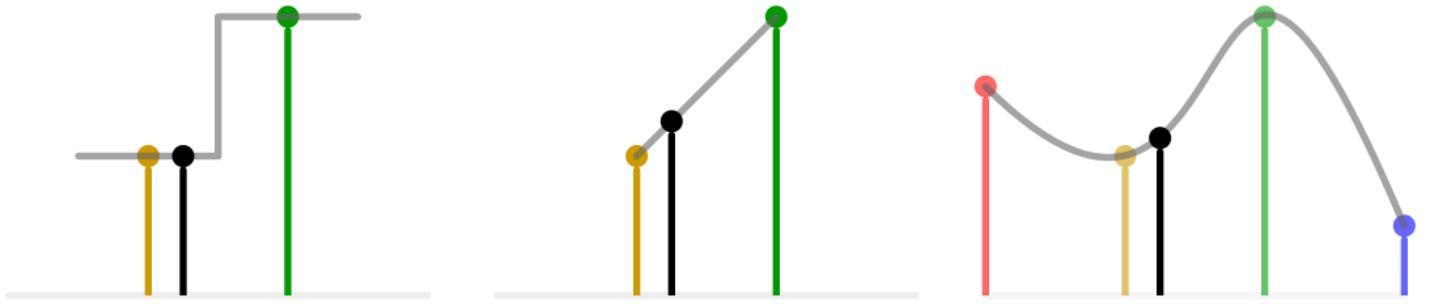


Image interpolation

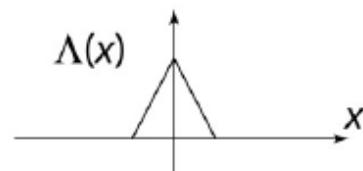
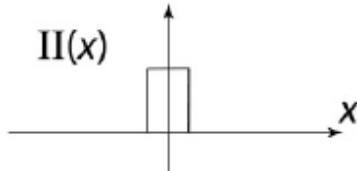


1D nearest-neighbour

Linear

Cubic

kernel



response

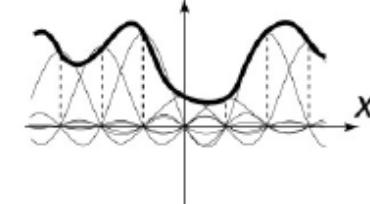
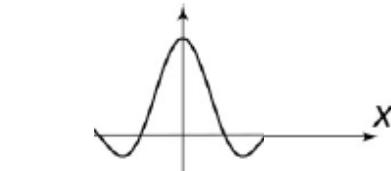
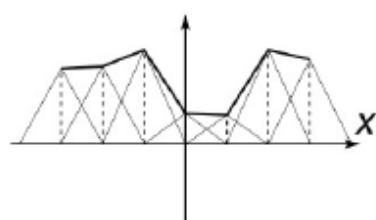
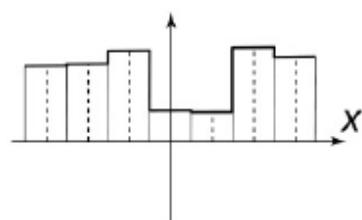
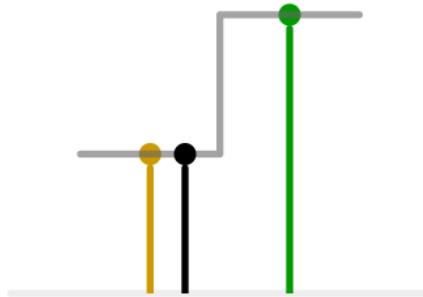
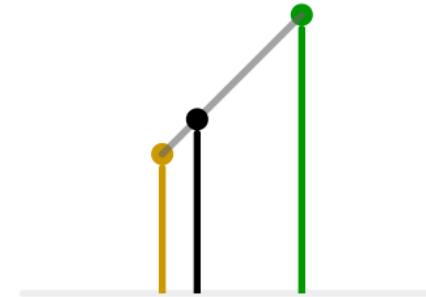


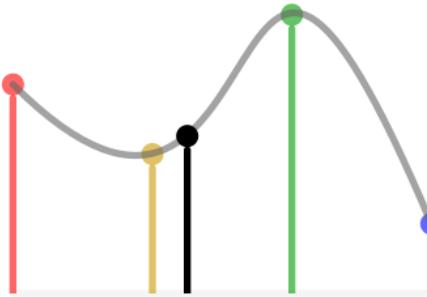
Image interpolation



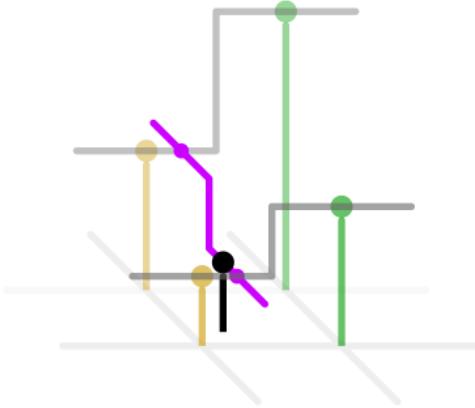
1D nearest-neighbour



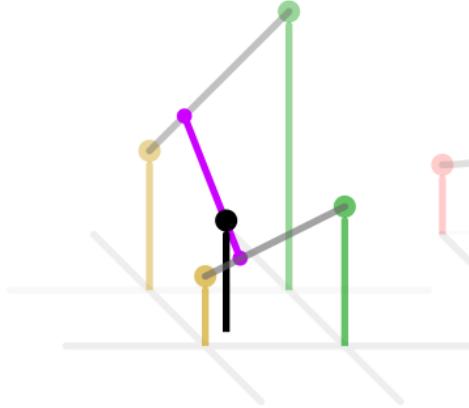
Linear



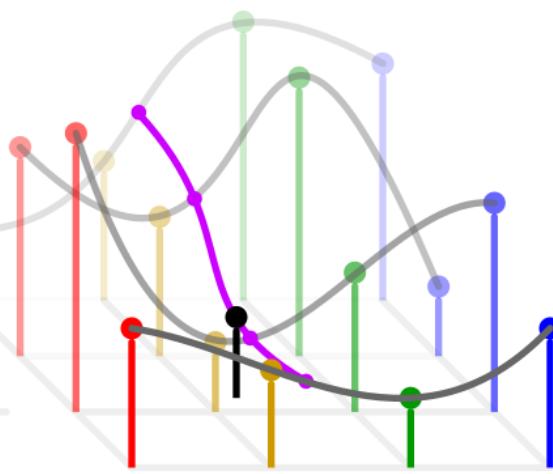
Cubic



2D nearest-neighbour

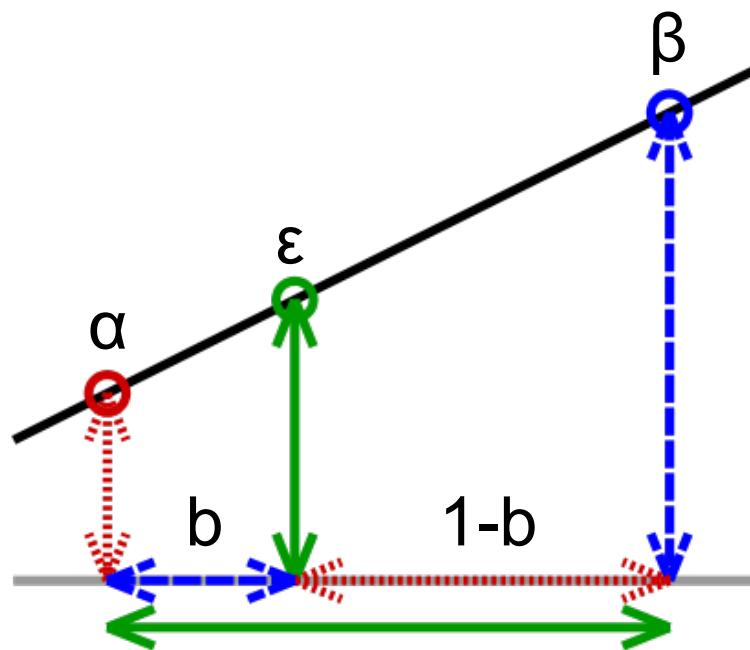


Bilinear



Bicubic

Linear Interpolation



Linear Interpolation

The value at the green circle multiplied by the distance between the red and blue circles

=

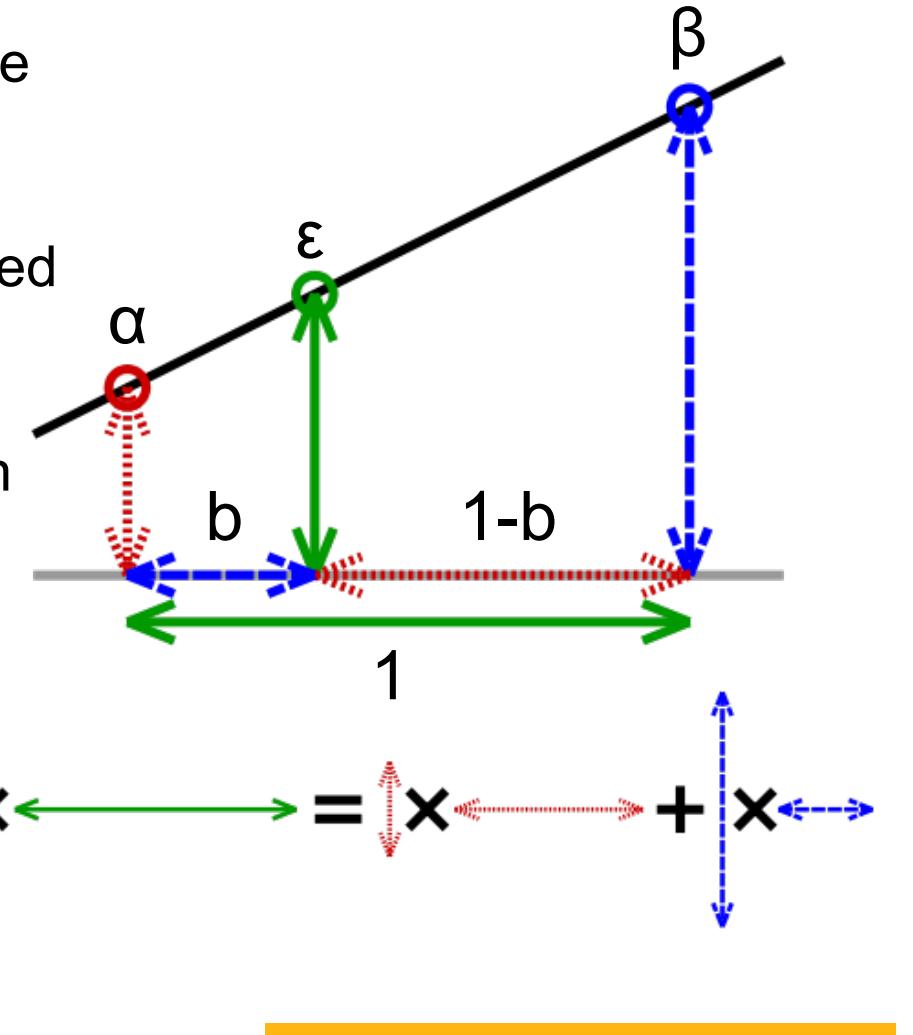
the sum of the value at the red circle multiplied by the distance between the green and blue circles, and the value at the blue circle multiplied by the distance between the green and red circles.

$$\varepsilon = (b\beta + (1 - b)\alpha)$$

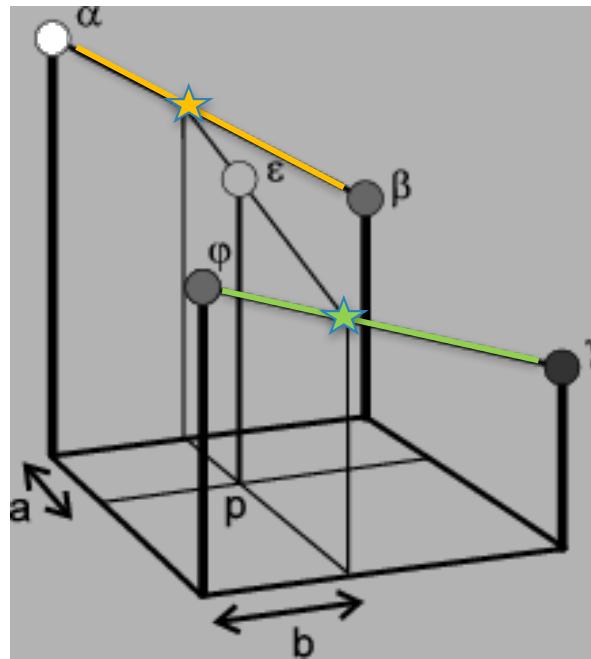
$$I(1) = 175$$

$$I(2) = 200$$

$$I(1.6) = ?$$



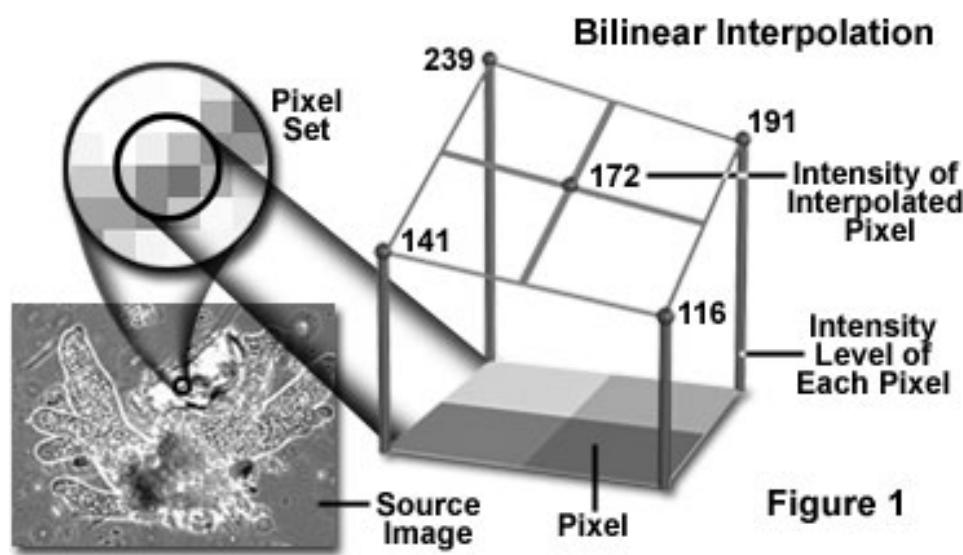
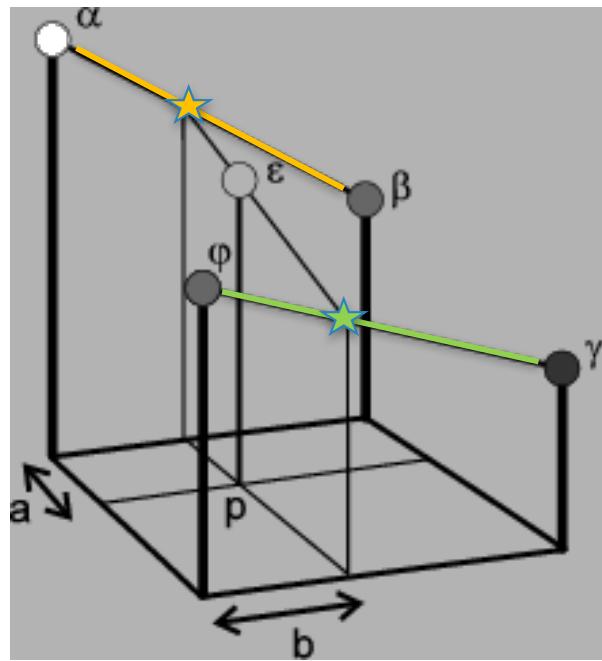
Bilinear interpolation



$$\varepsilon = a(b\gamma + (1 - b)\varphi) + (1 - a)(b\beta + (1 - b)\alpha)$$



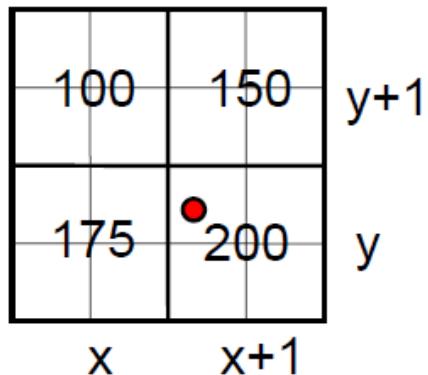
Bilinear interpolation



$$\varepsilon = a(b\gamma + (1 - b)\varphi) + (1 - a)(b\beta + (1 - b)\alpha)$$



Example



Interpolated value for

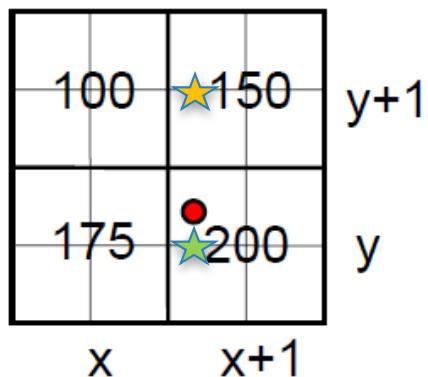
$(x+0.7, y+0.2)?$

x' y'

★ $I(x',y) = (175 * .3) + (200 * .7) = 192.5$

★ $I(x',y+1) = (100 * .3) + (150 * .7) = 135$

Example



Interpolated value for

$(x+0.7, y+0.2)?$

x' y'

★ $I(x',y) = (175 * .3) + (200 * .7) = 192.5$

★ $I(x',y+1) = (100 * .3) + (150 * .7) = 135$

$I(x',y') = (192.5 * .8) + (135 * .2) = 181$

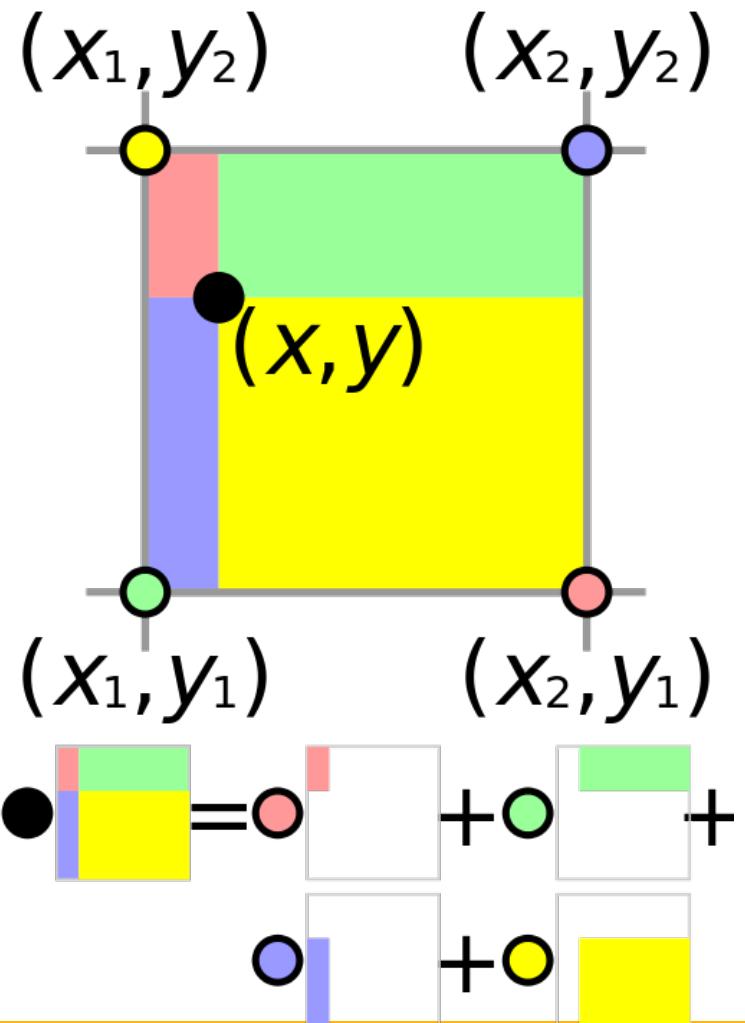
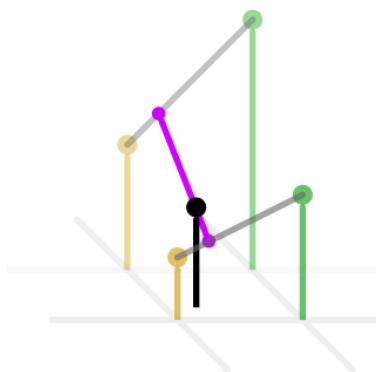
Example

| | | X | |
|---|---|----|----|
| | | 1 | 2 |
| | | 1 | 55 |
| Y | 1 | 55 | 45 |
| | 2 | 16 | 55 |

$X=1.2, Y=1.35$

Bilinear interpolation

The value at the black spot is the sum of the value at each colored spot multiplied by the area of the rectangle of the same color, divided by the total area of all four rectangles.



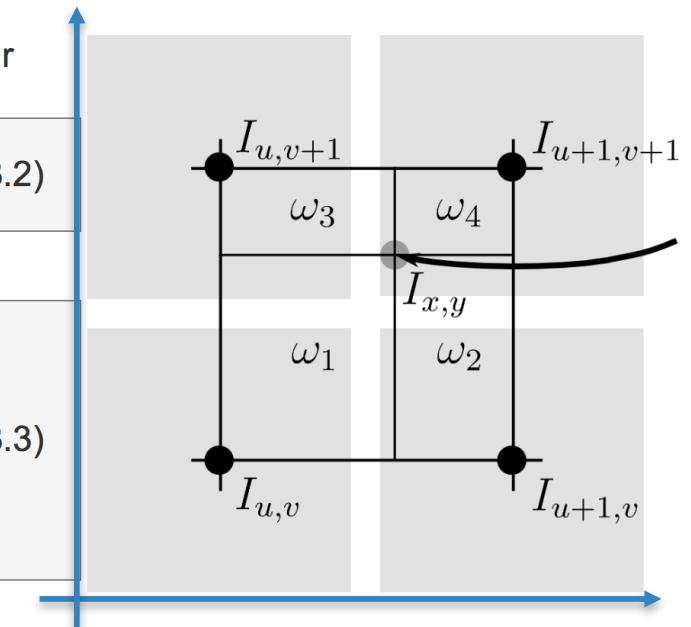
Bilinear interpolation

Bilinear interpolation determines the grey level from the weighted average of the four closest pixels and can be defined as

$$I_{x,y} = \omega_4 I_{u,v} + \omega_2 I_{u,v+1} + \omega_3 I_{u+1,v} + \omega_1 I_{u+1,v+1} \quad (3.2)$$

where

$$\begin{aligned}\omega_1 &= (x - u)(y - v) \\ \omega_2 &= (u + 1 - x)(y - v) \\ \omega_3 &= (x - u)(v + 1 - y) \\ \omega_4 &= (u + 1 - x)(v + 1 - y)\end{aligned}$$



Bicubic interpolation

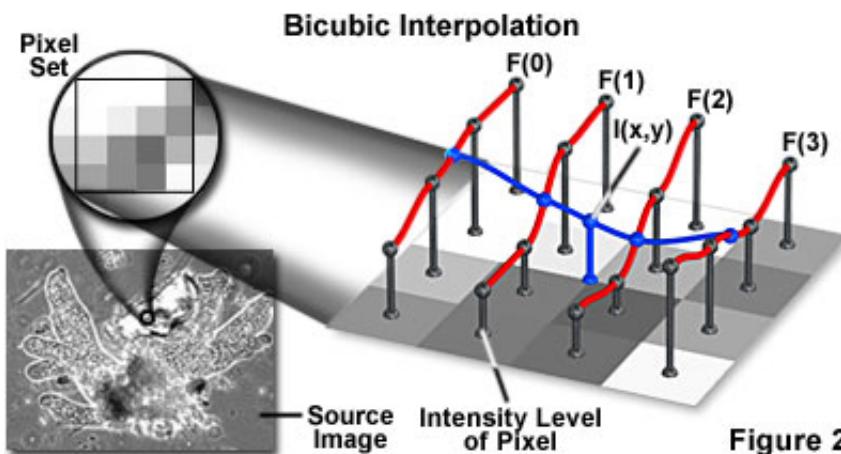
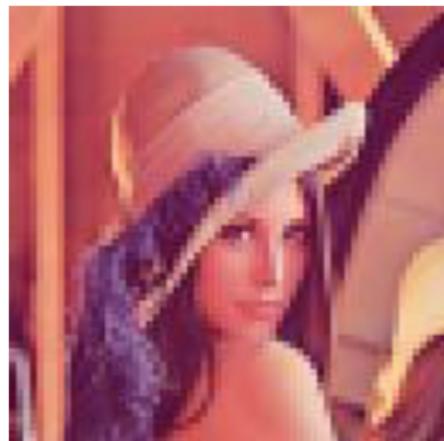
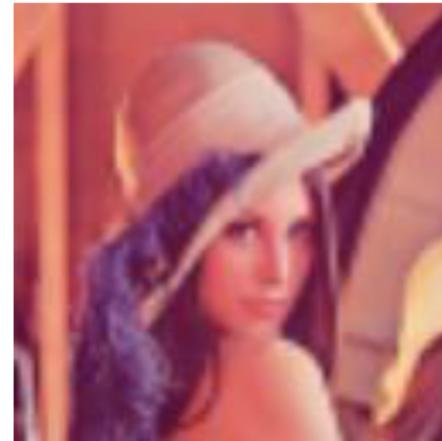


Figure 2

$$\varphi_{CC3}(x) = \begin{cases} \frac{3}{2}|x|^3 - \frac{5}{2}|x|^2 + 1 & \text{if } 0 \leq |x| \leq 1, \\ -\frac{1}{2}|x|^3 + \frac{5}{2}|x|^2 - 4|x| + 2 & \text{if } 1 \leq |x| \leq 2, \\ 0 & \text{if } 2 \leq |x|. \end{cases}$$



Nearest neighbor

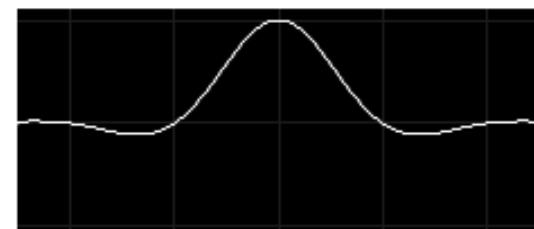
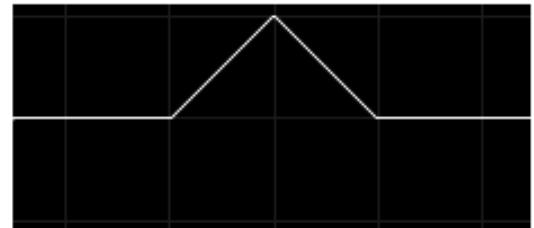
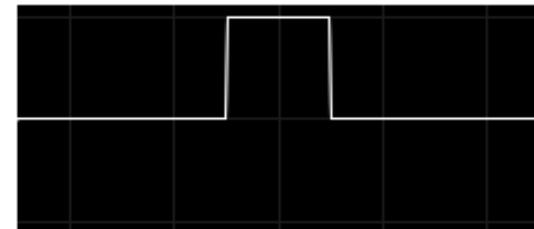
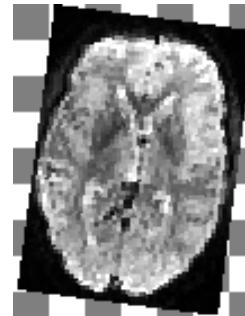


Bilinear



Bicubic

- Cubic interpolation is better at retaining the original grey values than NN and bilinear interpolation.
- Computational requirement is up to 10 times higher than NN interpolation.
- Computational requirement is up to 2 times higher than bi-linear interpolation.



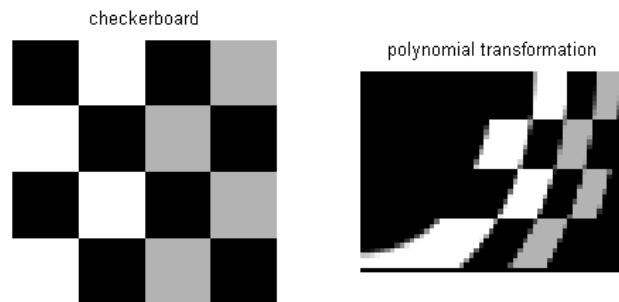
Matlab: how to infer a transformation?

- $t = cp2tform(input_points, base_points);$

```
pairs1 = [1 1; 5 21; 17 40; 28 1; 32 20; 45 40; 72 1; 77 20; 90 40];
pairs2 = [1 1; 1 21; 1 40; 20 1; 20 20; 20 40; 40 1; 40 20; 40 40];
t_poly = cp2tform(pairs1, pairs2, 'polynomial',2);
I = checkerboard(10, 2);
J = imtransform(I,t_poly);

subplot(1,2,1)
imshow(I), title('checkerboard')

subplot(1,2,2)
imshow(J), title('polynomial transformation')
```



How to find correspondence automatically?