

# Introduction to Computer Vision

## *14. Deep Learning*

UCLA – CS 188 – Fall 2019

Fabien Scalzo, Ph.D.

<b>Week 1</b>			26-Sep	<b>Introduction</b>
<b>Week 2</b>	1-Oct	<b>Basic Image Processing</b>	3-Oct	<b>Feature Extraction and Classification</b>
<b>Week 3</b>	8-Oct	<b>Feature Tracking/Optical Flow</b>	10-Oct	<b>SVD, 2D camera model, projective plane</b>
<b>Week 4</b>	15-Oct	<b>2D Image transformations, RANSAC</b>	17-Oct	<b>Euclidean geometry, rigid body motion</b>
<b>Week 5</b>	22-Oct	<b>Epipolar Geometry</b>	24-Oct	<b>3D Cameras and processing</b>
<b>Week 6</b>	29-Oct	<b>Midterm</b>	31-Oct	<b>Statistical decision theory/Pattern Recognition</b>
<b>Week 7</b>	5-Nov	<b>Learning from data</b>	7-Nov	<b>Neural Networks</b>
<b>Week 8</b>	12-Nov	<b>Deep Learning</b>	14-Nov	<b>Deep Learning</b>
<b>Week 9</b>	19-Nov	<b>Deep Learning (Large Scale and Object Detection)</b>	21-Nov	<b>Advanced Deep Learning</b>
<b>Week 10</b>	26-Nov	<b>Guest Lecture (Nikhil Naik)</b>	28-Nov	
<b>Week 11</b>	3-Dec	<b>Applications</b>	5-Dec	<b>Recap</b>
<b>Week 12</b>	10-Dec		12-Dec	<b>Final</b>

- 
- Batch-normalization
  - Evaluating Deep Learning models
  - Transposed Convolution (i.e. Deconv)
  - MNIST with Deep Learning
  - Beyond Deep Learning
-

# Batch, Epoch, and Iteration

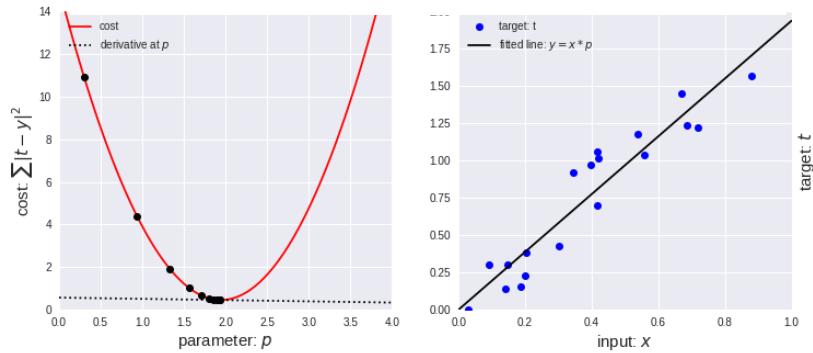
**epoch**: equivalent to the forward and backward processing of the entire training set

**Batch**: subset of training samples

**Batch size**: # of training samples in 1 batch

**Iteration**: # of batches to complete 1 epoch  
(or # of passes to complete 1 epoch.  
1 pass = 1 forward + 1 backward pass)

Given a dataset of 2,000 examples divided into batches of 500 then it will take 4 iterations to complete 1 epoch.



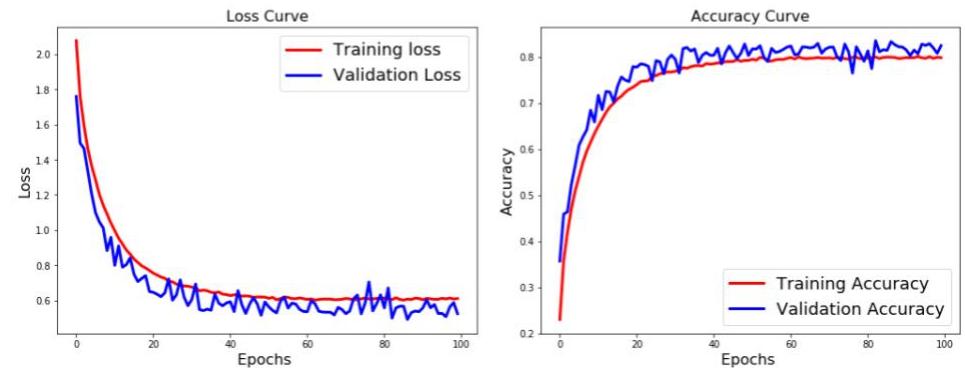
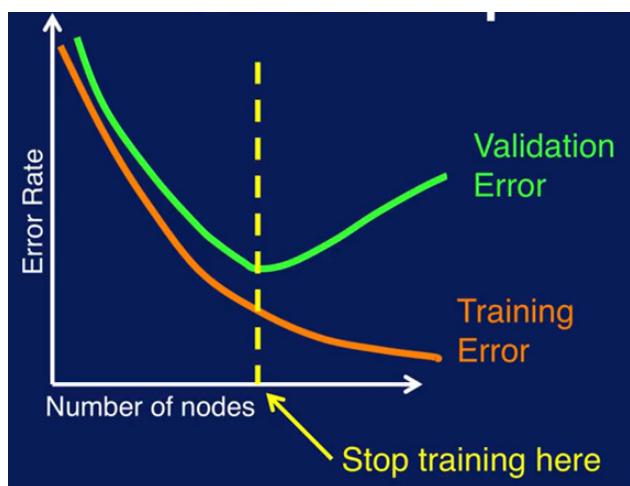
Training

Testing  
(holdout sample)

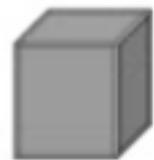
Training

Validation  
(validation holdout sample)

Testing  
(testing holdout sample)



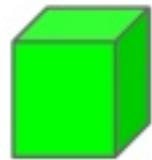
# Batch, Epoch, and Iteration



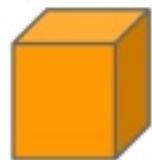
Dataset = 6 batches



- train batch



- validation batch



- test batch



←One Epoch's data

# Batch normalization (BN)

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

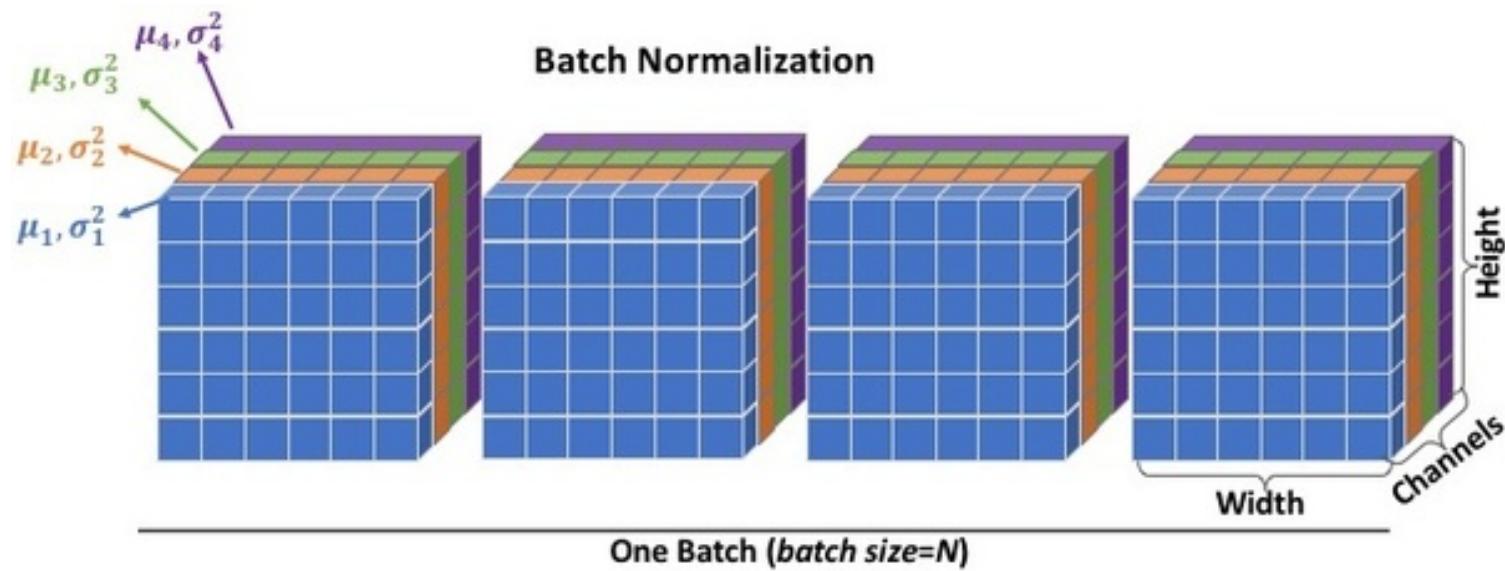
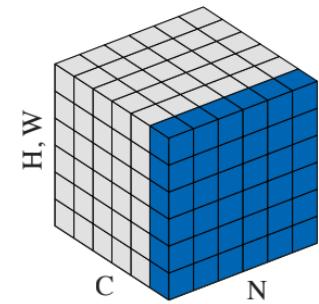
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

# Batch normalization (BN)



# Group Normalization

Yuxin Wu

Kaiming He

Facebook AI Research (FAIR)

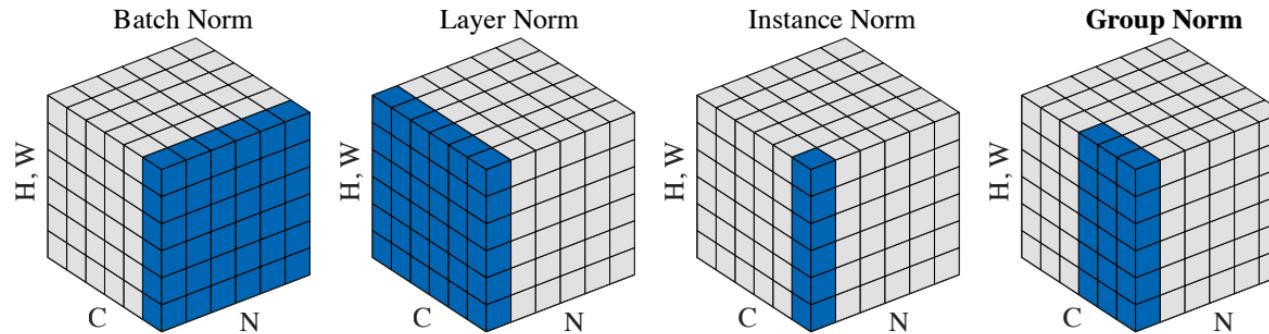
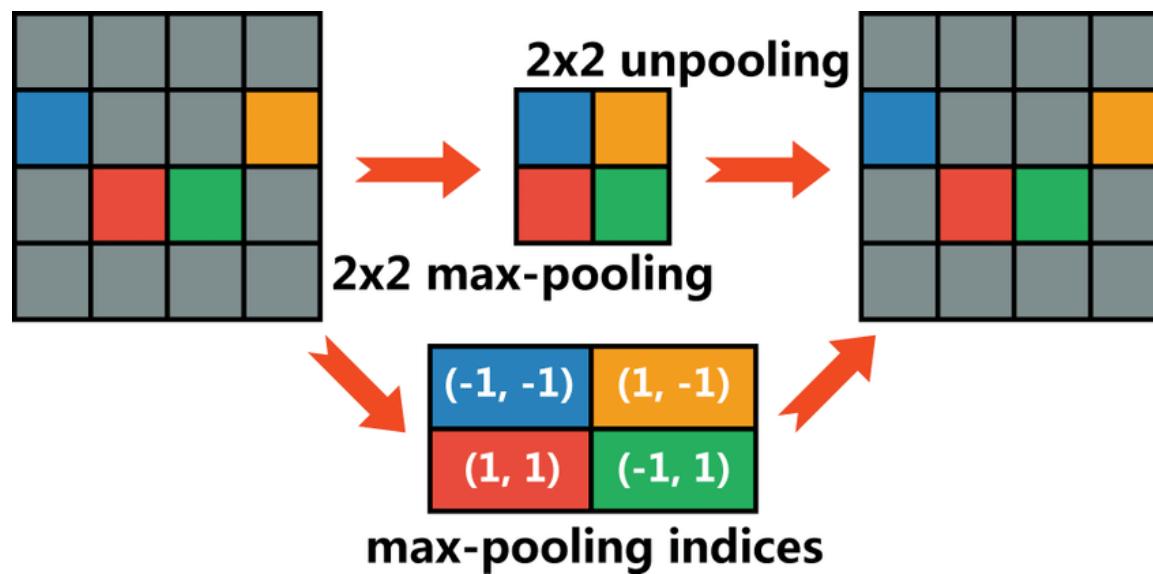


Figure 2. **Normalization methods.** Each subplot shows a feature map tensor, with  $N$  as the batch axis,  $C$  as the channel axis, and  $(H, W)$  as the spatial axes. The pixels in blue are normalized by the same mean and variance, computed by aggregating the values of these pixels.

# Max-Unpooling



# Transposed Convolution (i.e. Deconv)



Ground Truth



$\frac{1}{4}$  Sized  
Input

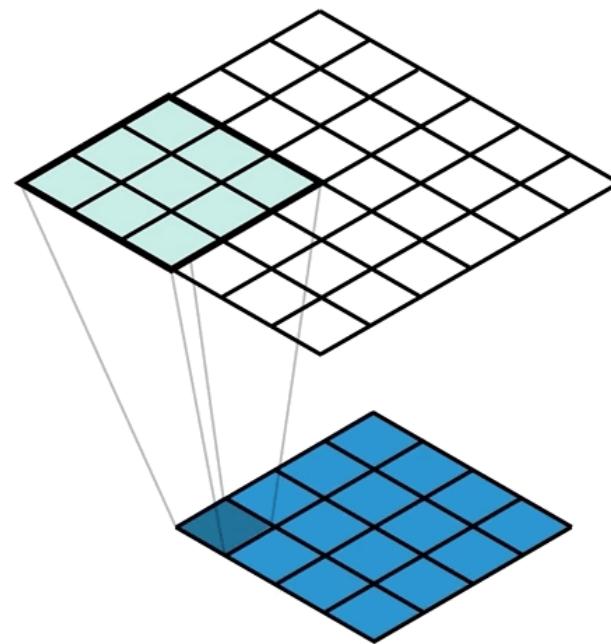


Bicubic



Super Resolution  
Network

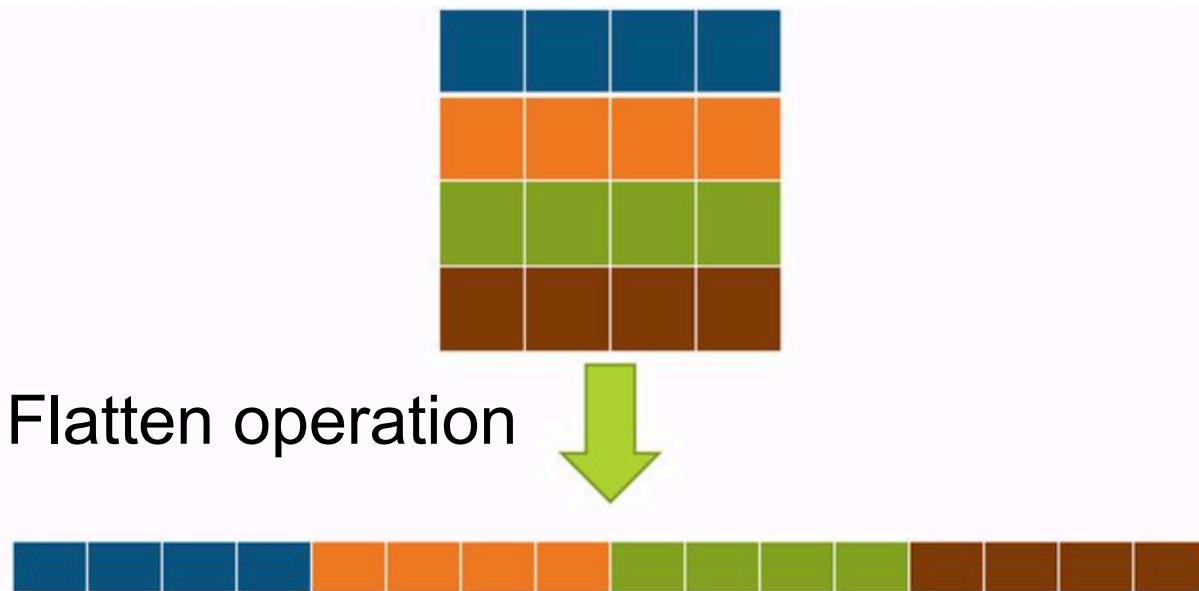
# Transposed Convolution (i.e. Deconv)



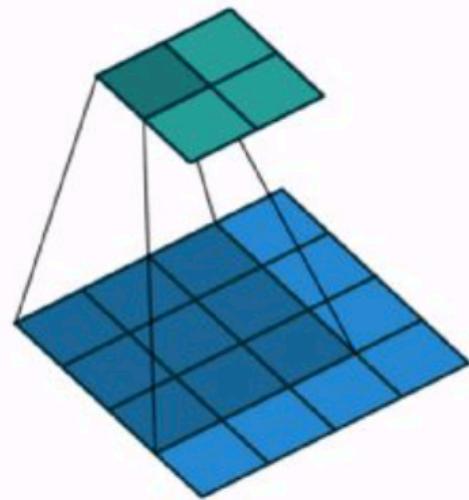
# Convolution as matrix multiplication

- ▶ Convolution can be expressed as a matrix operation
- ▶ The convolution matrix “**C**” is inferred from the kernel
- ▶ The output is obtained by multiplying the “C” with the **flattened** input column vector and then reshaped

# Convolution as matrix multiplication



# Convolution as matrix multiplication



Slide it over the input

kernel

$w_{0,0}$	$w_{0,1}$	$w_{0,2}$
$w_{1,0}$	$w_{1,1}$	$w_{1,2}$
$w_{2,0}$	$w_{2,1}$	$w_{2,2}$

$$\begin{pmatrix} w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} \end{pmatrix}$$

# Convolution as matrix multiplication

$$\text{output} = C \cdot i^T$$

Where  $i$  is the flattened input feature map.

From the previous two slides:

Assume the convolution on  $i = 4$  with  $k = 3, s = 1$ :

We get the output as 4x1 vector



**Reshape the  
output**

1
2
3
4



1	2
3	4

# Transposed convolution as matrix multiplication

- ▶ Output =  $C^T i'$
- ▶  $C^T$  is 16x4 matrix
- ▶  $i'$  is 4x1 column vector
- ▶ Output = 16x1 vector



**Reshape it to the  
shape of input**

# Transposed convolution as matrix multiplication

0	1	2
1	4	1
1	4	3
3	3	1

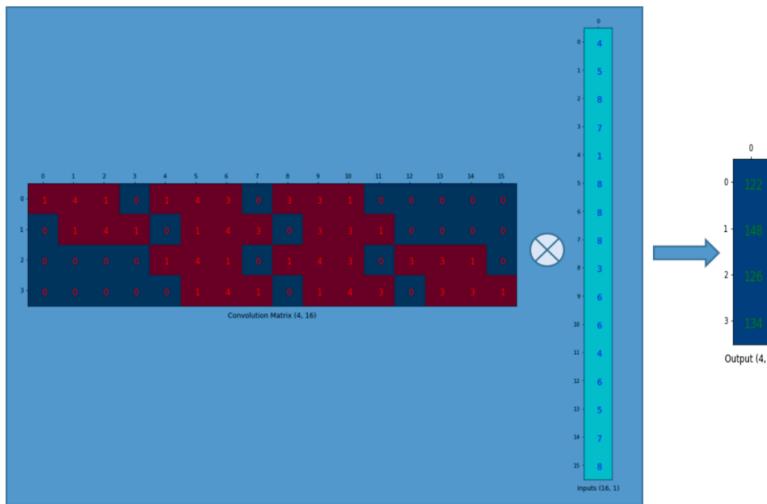
Kernel (3, 3)



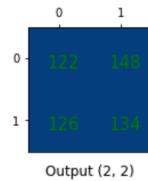
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0
1	0	1	4	1	0	1	4	3	0	3	3	1	0	0	0
2	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1
3	0	0	0	0	0	1	4	1	0	1	4	3	0	3	3

Convolution Matrix (4, 16)

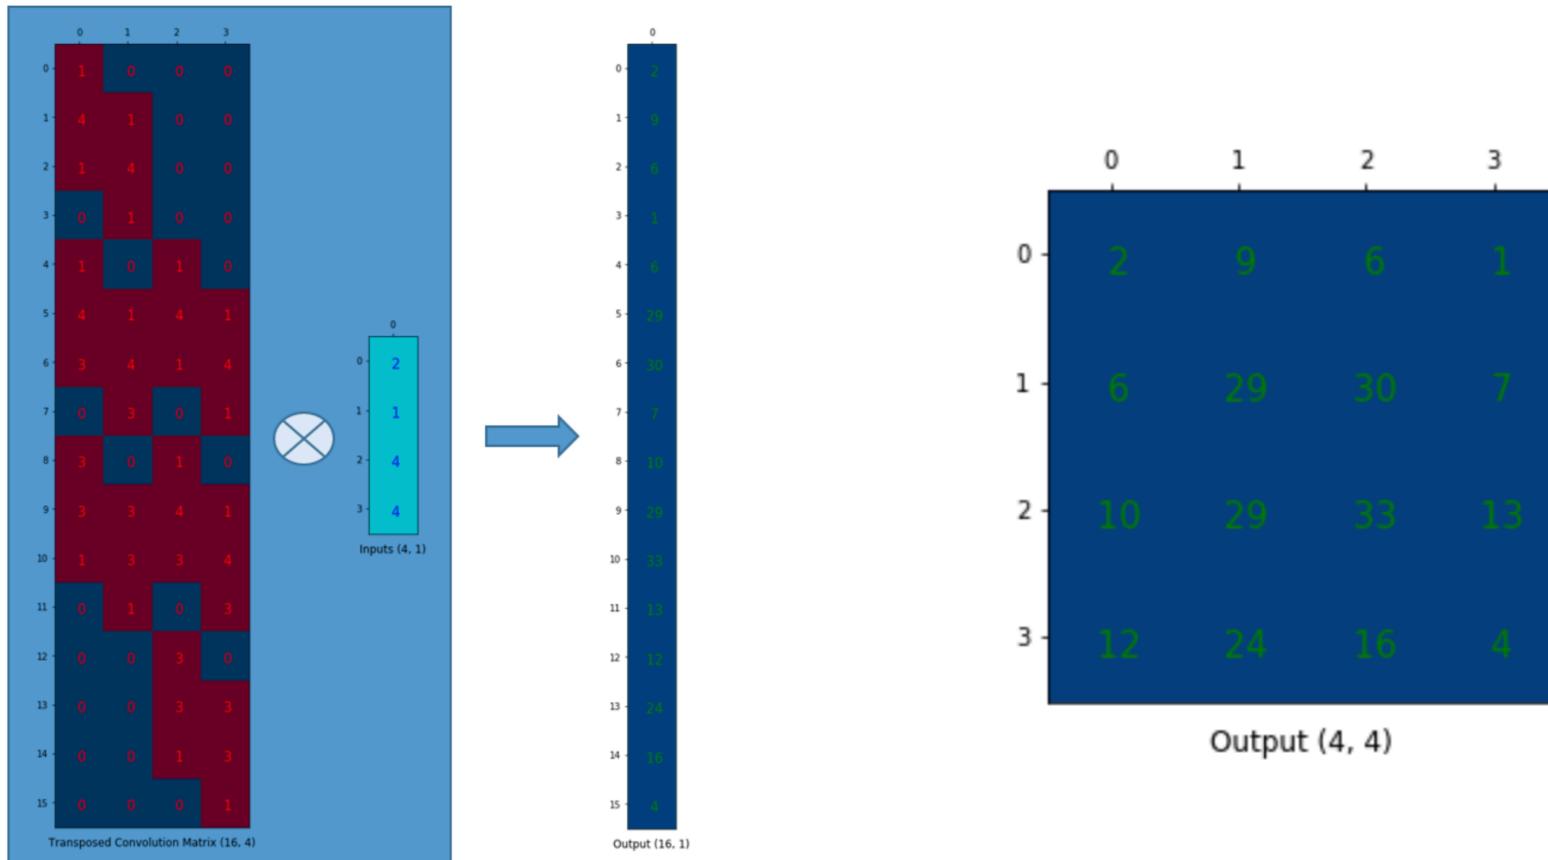
# Transposed convolution as matrix multiplication



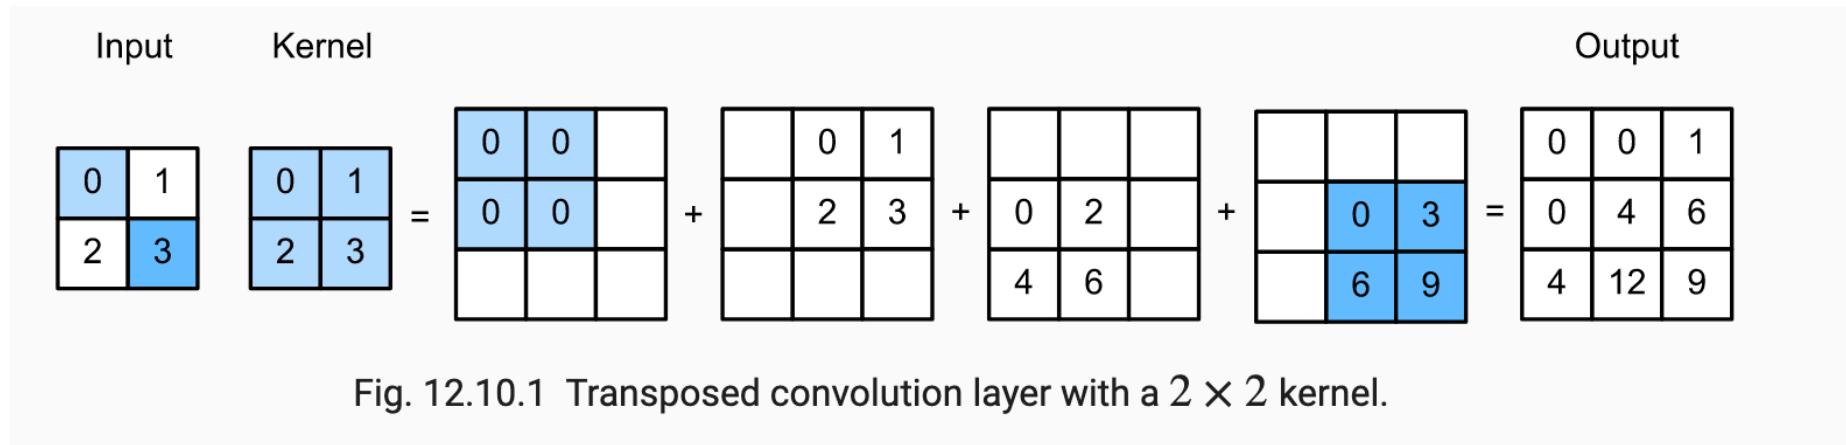
The output  $4 \times 1$  matrix can be reshaped into a  $2 \times 2$  matrix which gives us the same result as before.



# Transposed convolution as matrix multiplication



# Transposed Convolution (i.e. Deconv)

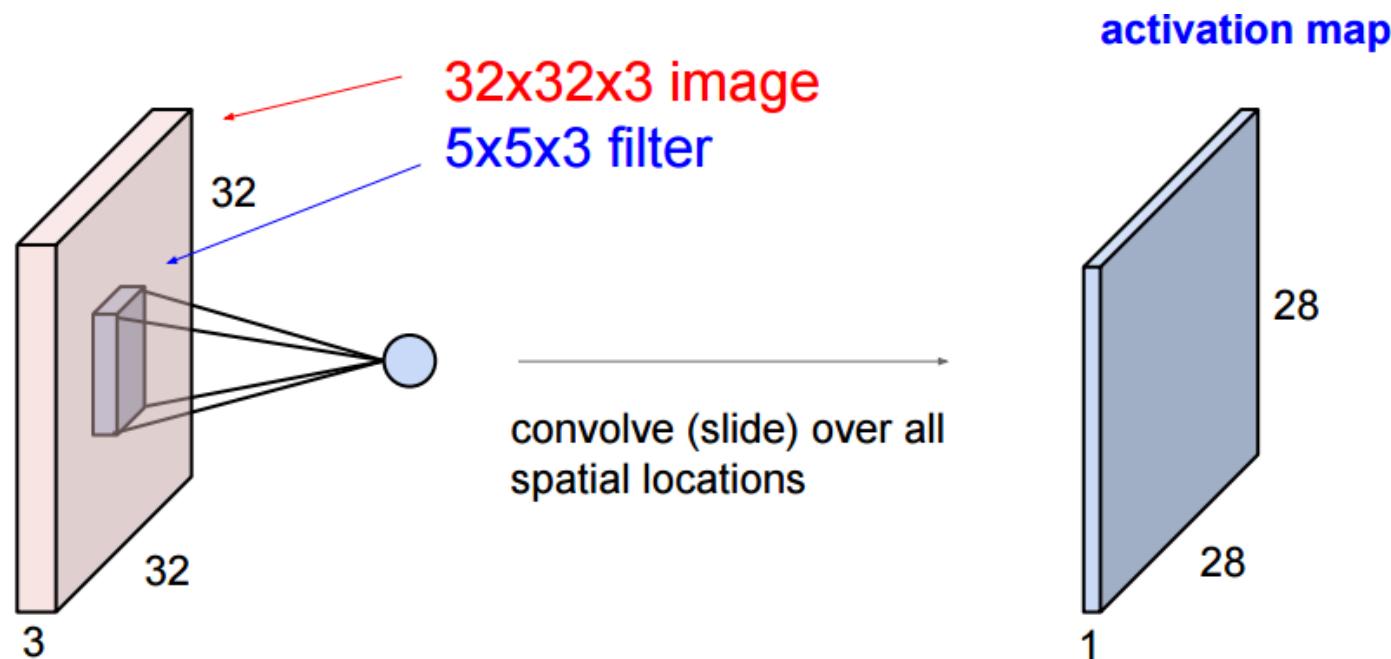


# Transposed Convolution (i.e. Deconv)

<u>Input</u>	<u>Kernel</u>	<u>Output</u>
0 0 0 0 0 0	1 2 3 0 1 0 2 1 2	3 6 12 6 9 0 3 0 3 0 7 5 16 5 9 0 1 0 1 0 2 1 4 1 2

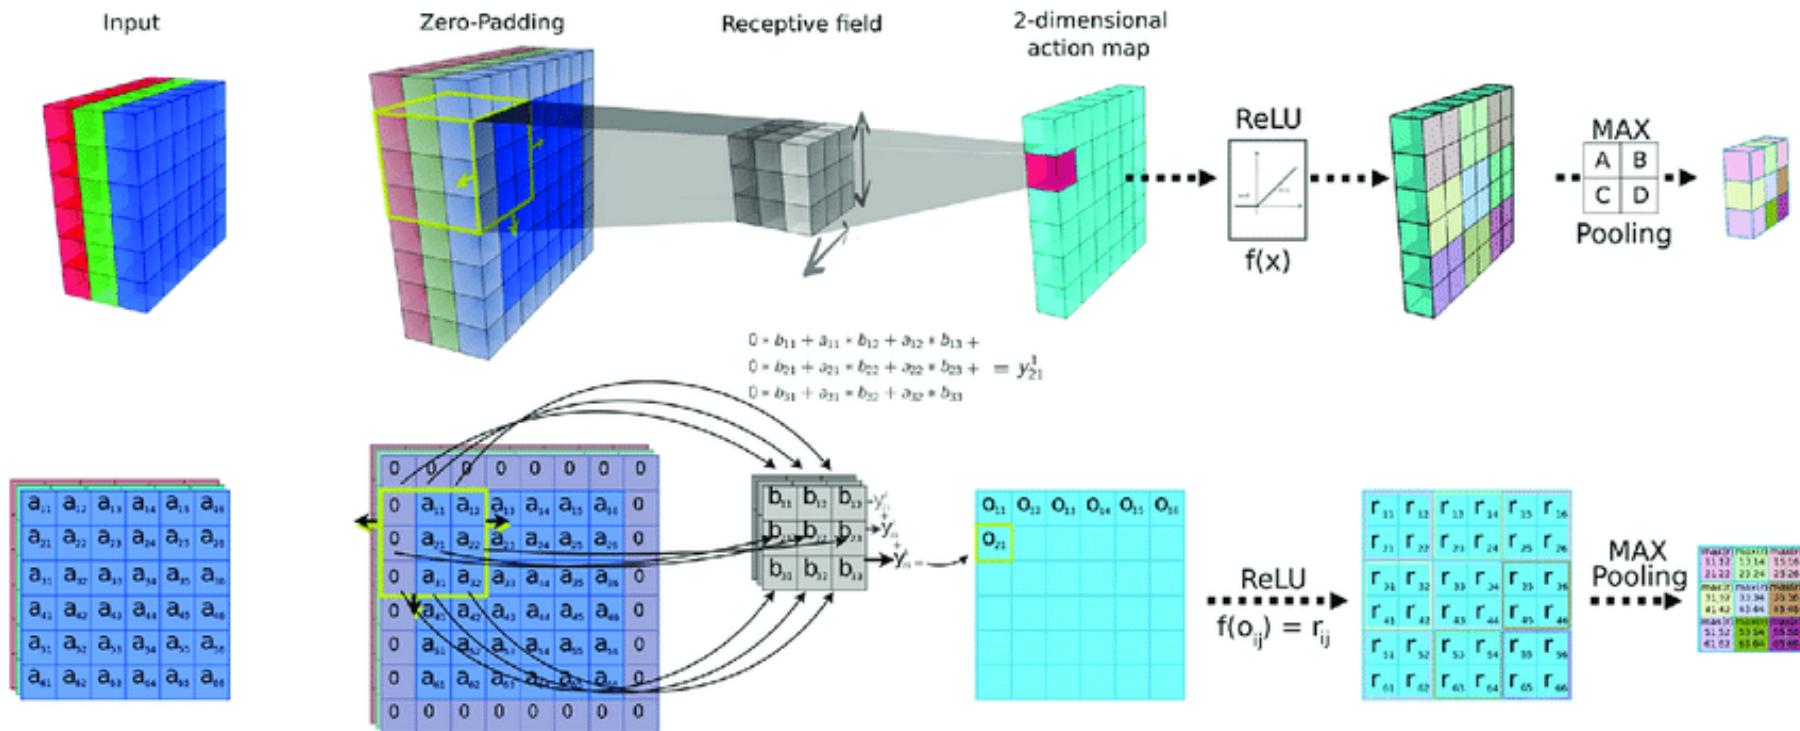
[https://docs.google.com/spreadsheets/d/1a0HUXcc\\_75IBZyPmfNfNd31wx9QCqHKI3u4JJes9BRs/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1a0HUXcc_75IBZyPmfNfNd31wx9QCqHKI3u4JJes9BRs/edit?usp=sharing)

# Convolution Layer

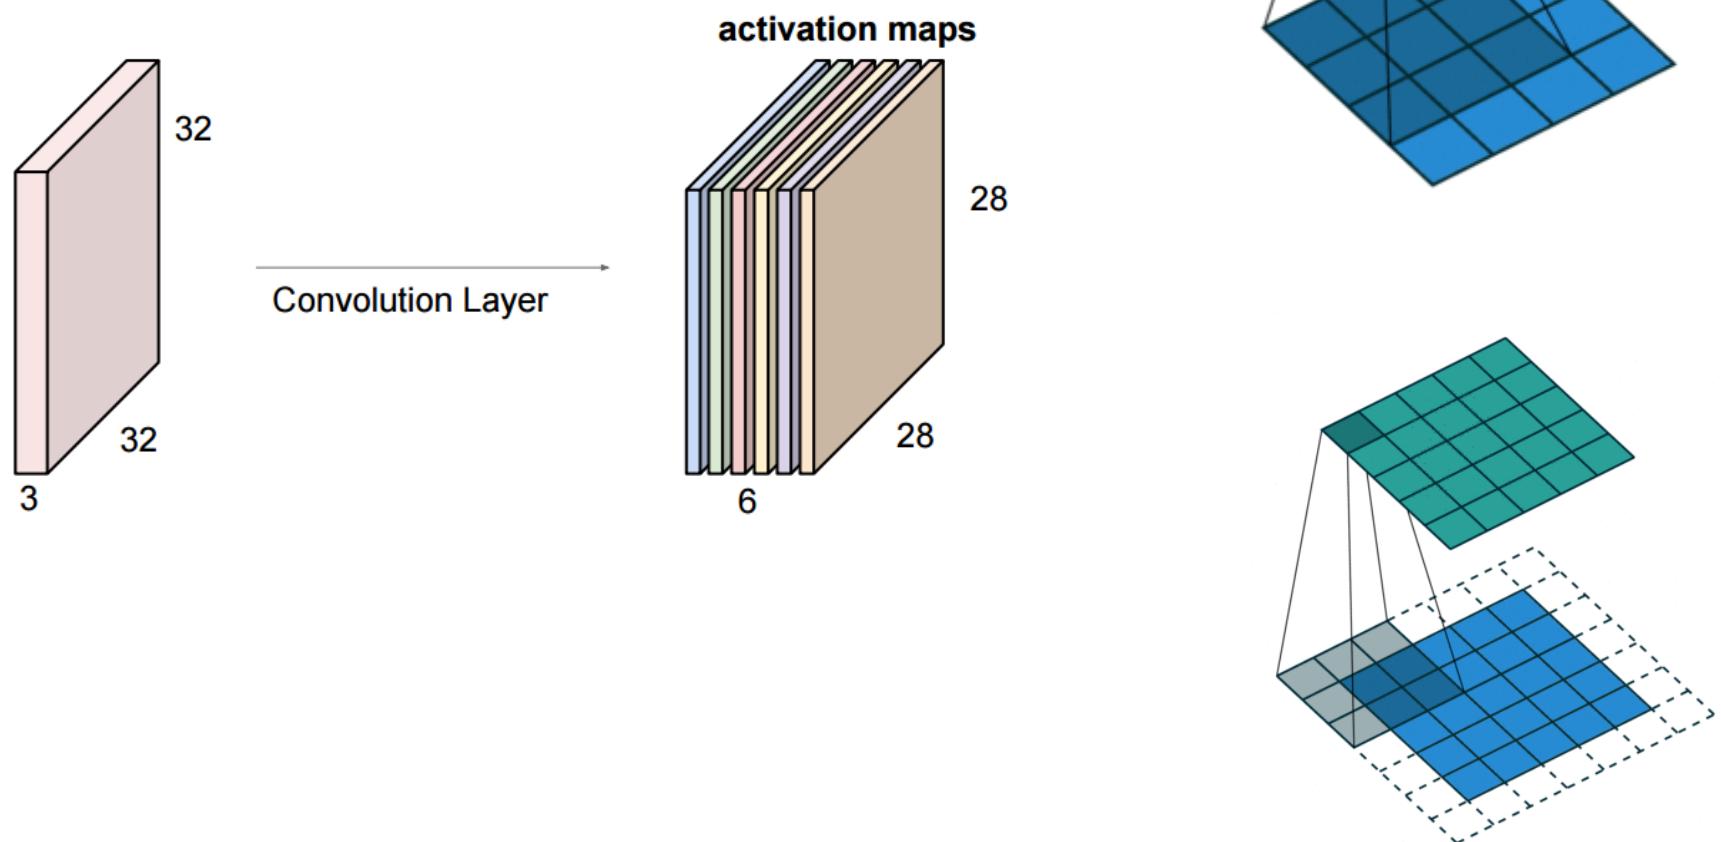


# Convolution layer

- Parametrized by: height, width, depth, stride, padding, number of filters, type of activation function



# Convolution layer



# MNIST

- (Modified National Institute of Standards and Technology database)

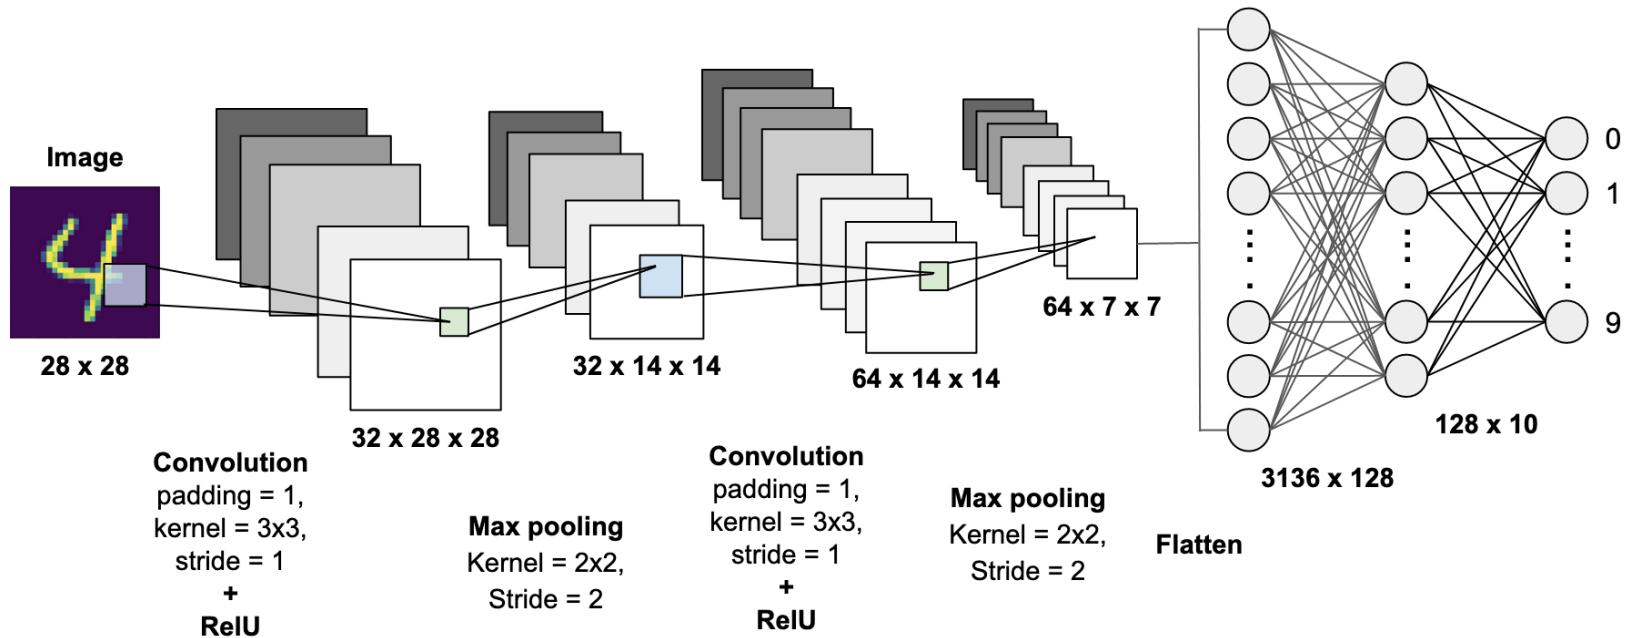
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

10 classes

60,000 training images and 10,000 testing images

Each image is 28x28 pixel

# MNIST



```

from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28

# the data, split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

```
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

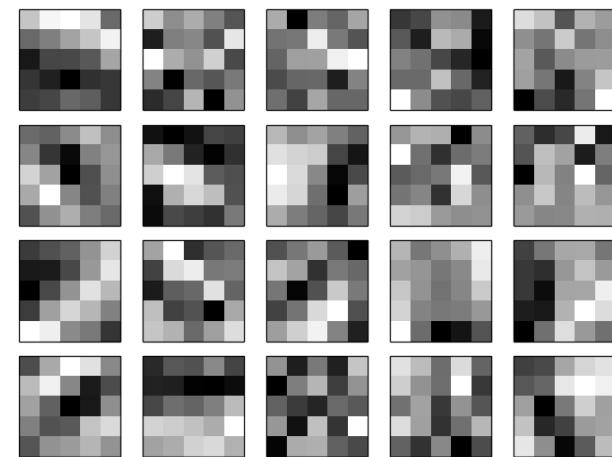
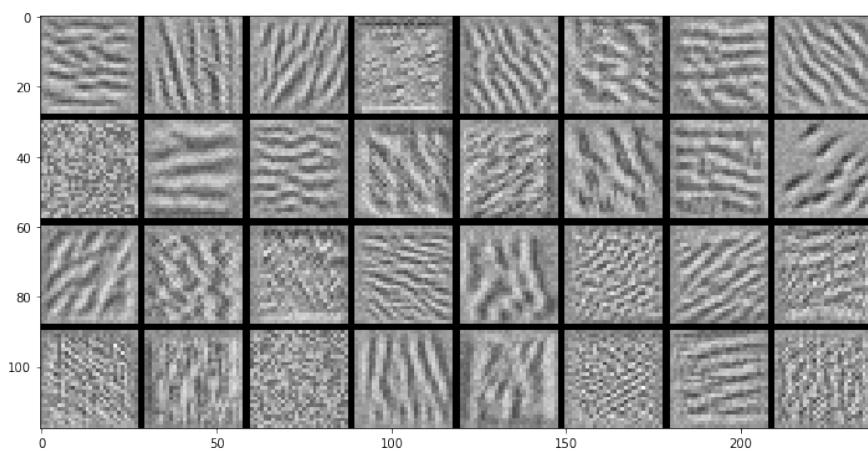
# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
```

```
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

99.25% test accuracy after 12 epochs



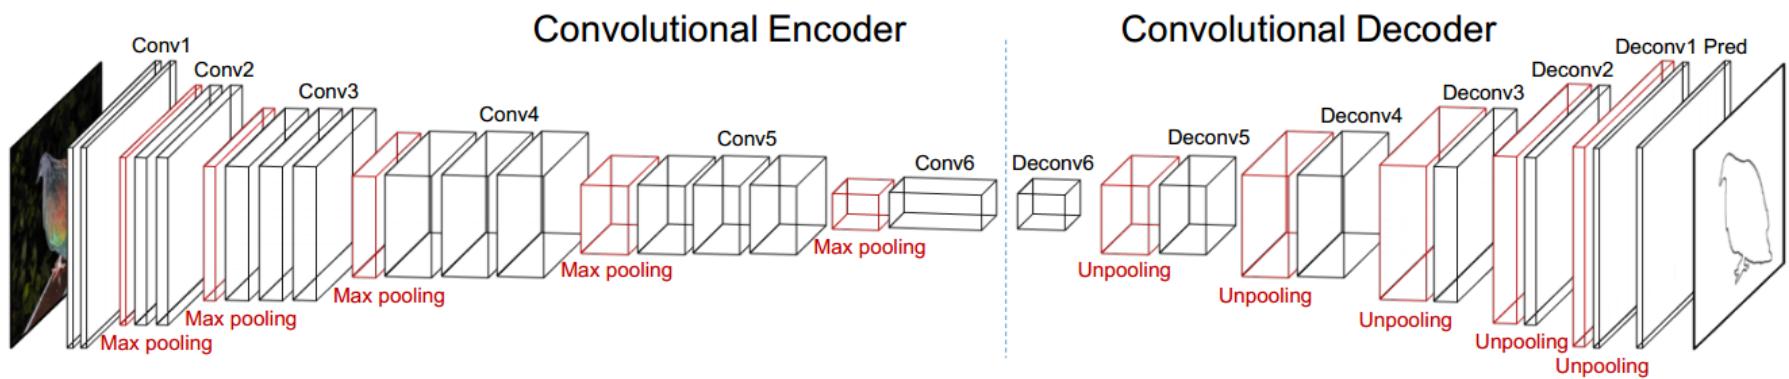
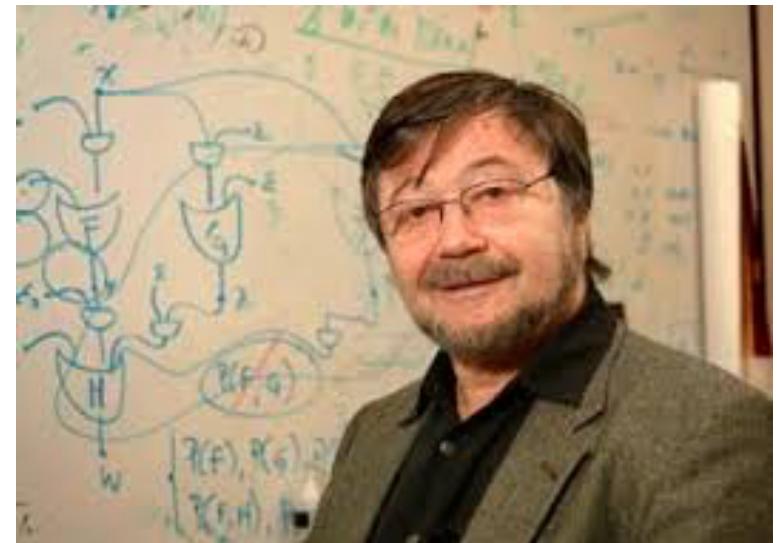
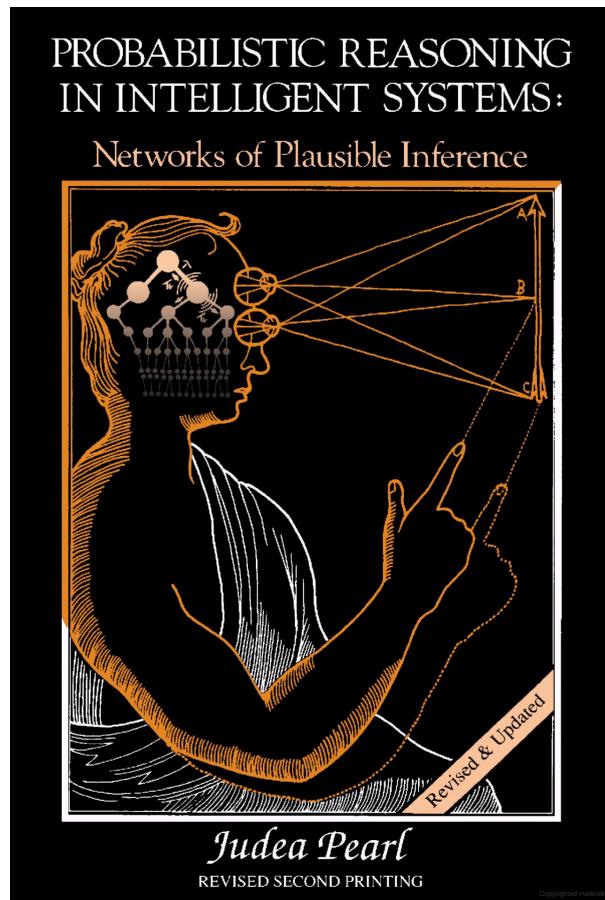


Figure 2. Architecture of the proposed fully convolutional encoder-decoder network.

# Graphical Models



Prof. Judea Pearl

# Graphical models

- Nodes: random variables
  - Discrete, Parametric, Nonparametric

